
Adversarial Networks for Spatial Context-Aware Spectral Image Reconstruction from RGB

Aitor Alvarez-Gila, Joost van de Weijer, Estibaliz Garrote
aitorshuffle.github.io

Venice, 2017/10/23

The task

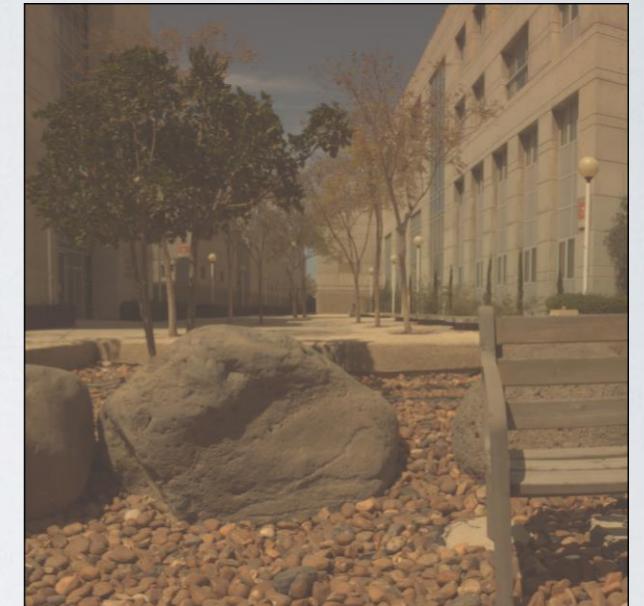
Adversarial Networks for Spatial
Context-Aware Spectral Image
Reconstruction from RGB

The task

Adversarial Networks for Spatial
Context-Aware Spectral Image
Reconstruction from RGB

The task

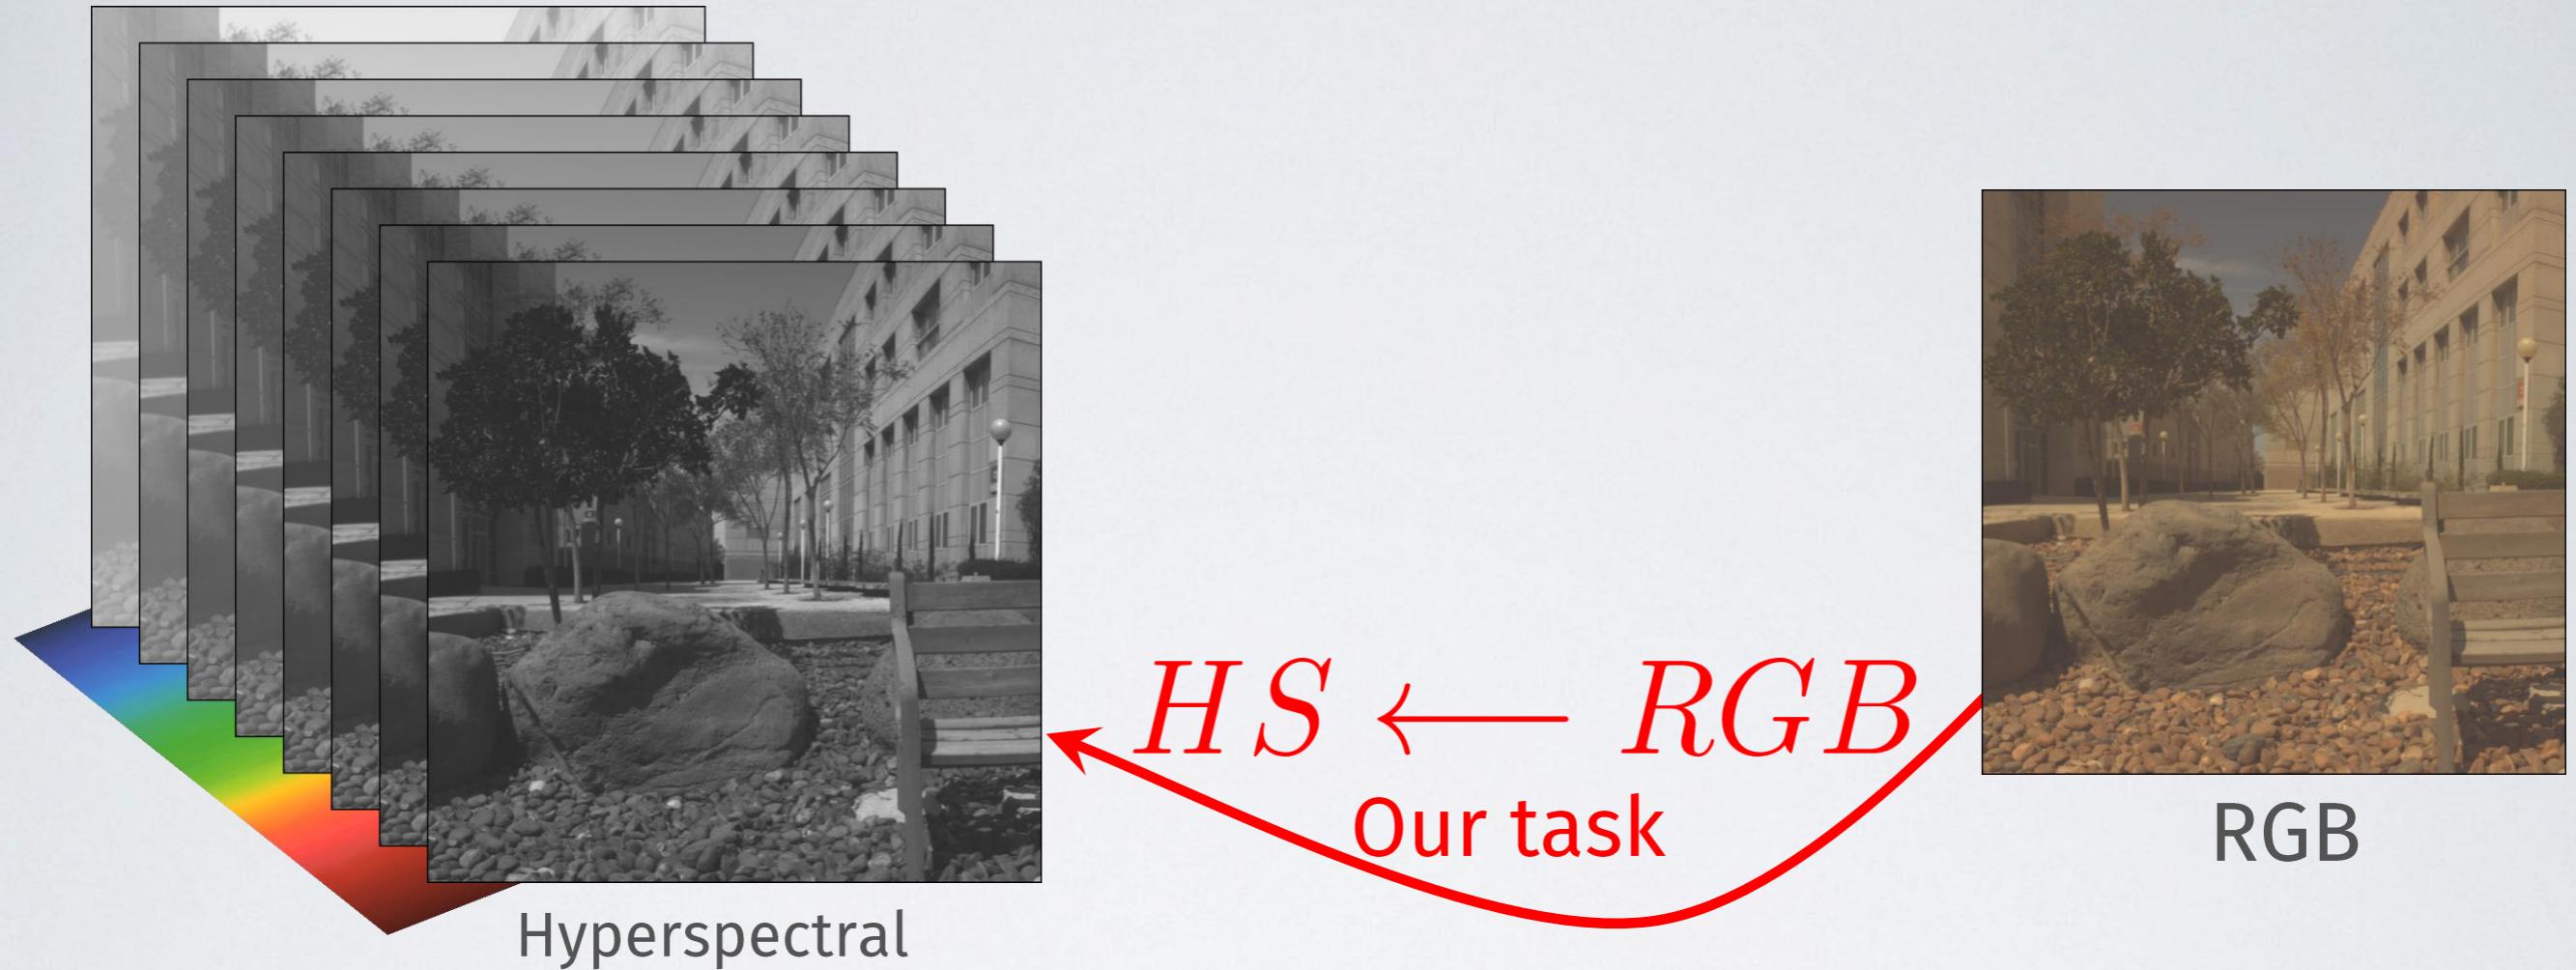
(Hyper)-Spectral Image reconstruction from RGB



RGB

- Learn a image to image mapping s.t., for each pixel:
 $\mathbb{R}^3 \rightarrow \mathbb{R}^c$ with $c \gg 3$
- In particular: $c = 31(400 : 10 : 700nm)$
- Heavily underconstrained (e.g. metamers), non-linear problem

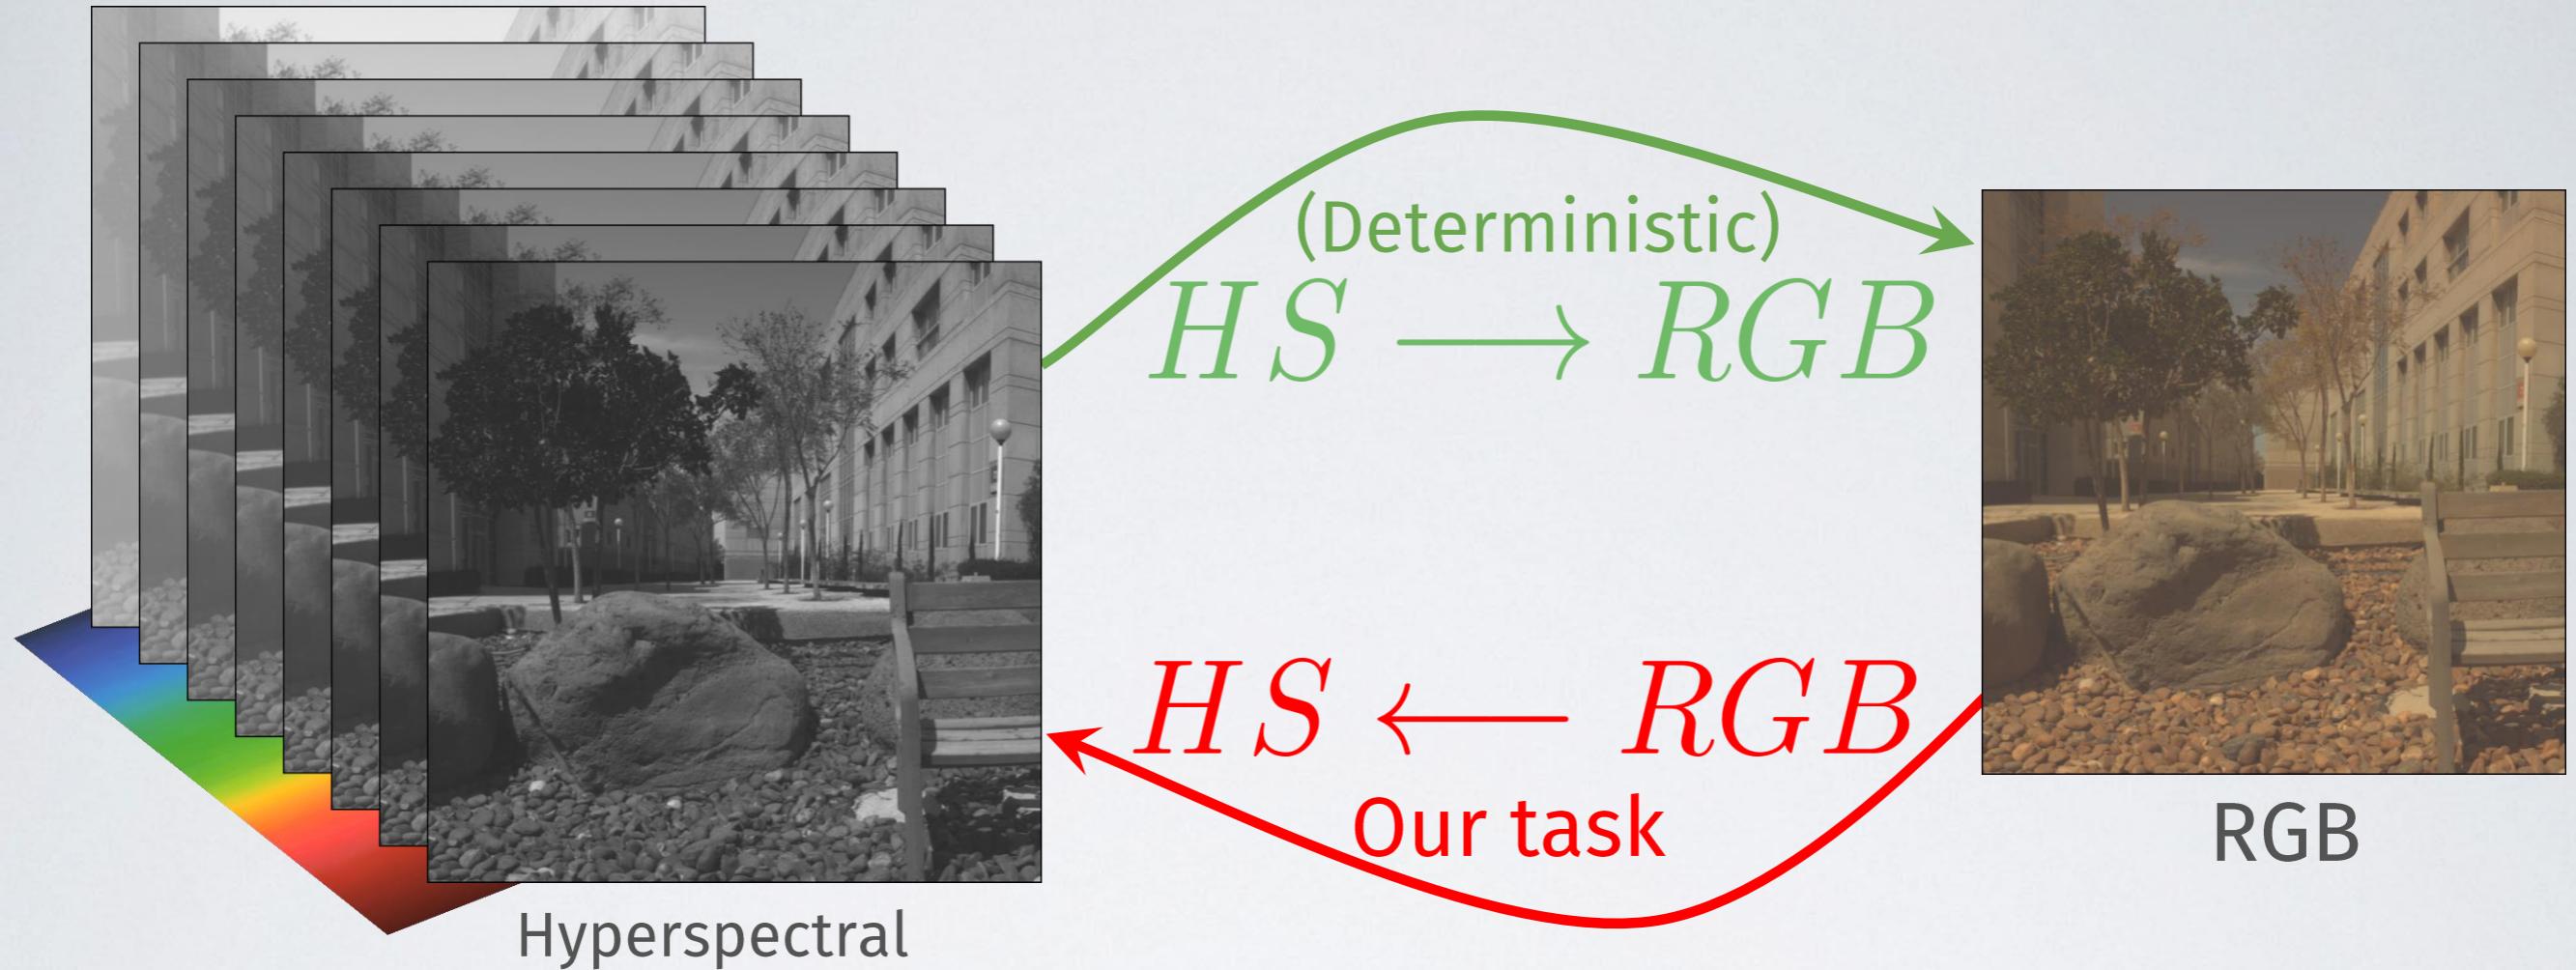
The task (Hyper)-Spectral Image reconstruction from RGB



- Learn a image to image mapping s.t., for each pixel:
 $\mathbb{R}^3 \rightarrow \mathbb{R}^c$ with $c \gg 3$
- In particular: $c = 31(400 : 10 : 700nm)$
- Heavily underconstrained (e.g. metamers), non-linear problem

The task

(Hyper)-Spectral Image reconstruction from RGB



- Learn a image to image mapping s.t., for each pixel:
 $\mathbb{R}^3 \rightarrow \mathbb{R}^c$ with $c \gg 3$
- In particular: $c = 31(400 : 10 : 700nm)$
- Heavily underconstrained (e.g. metamers), non-linear problem

Color Camera Image formation

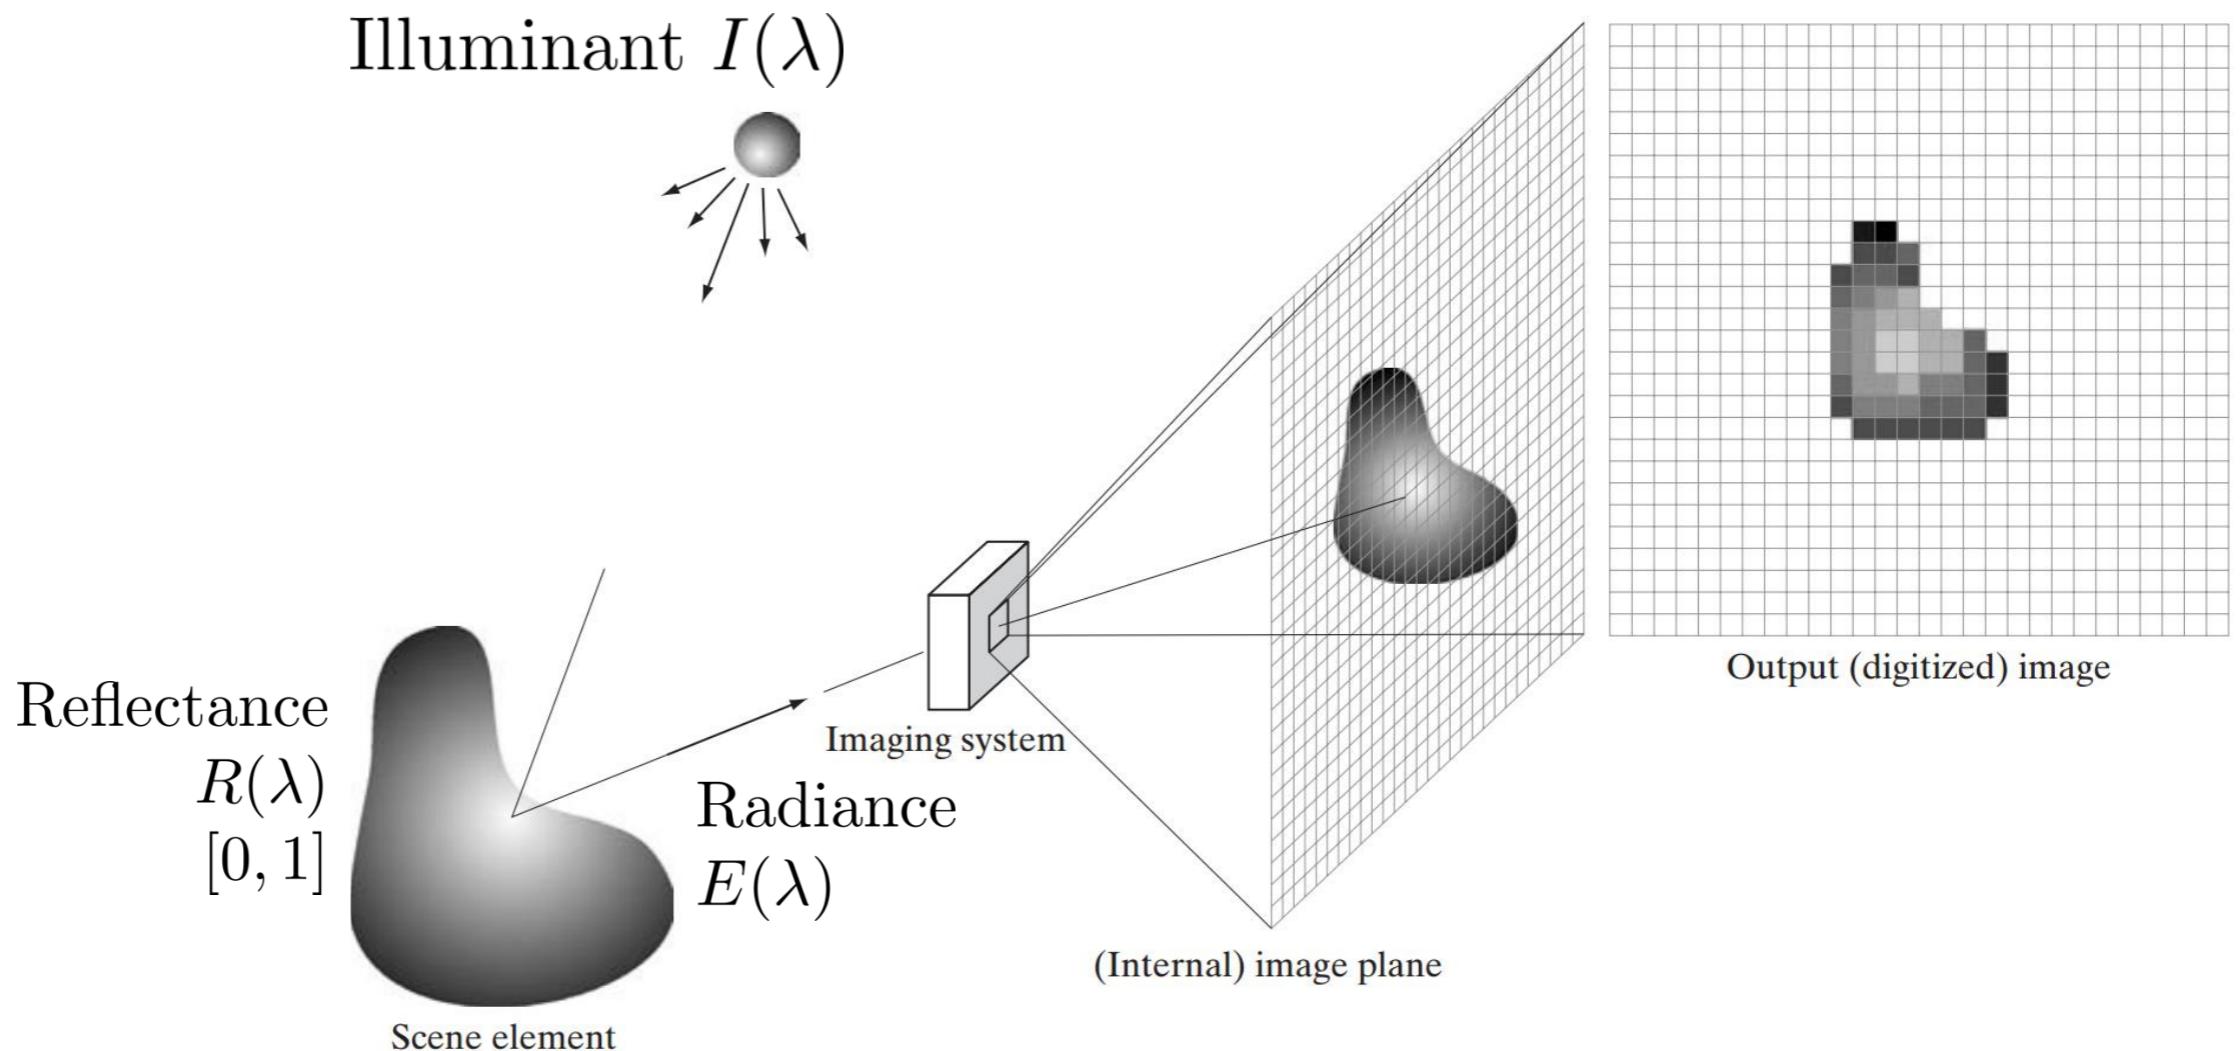


Image:

R. C. Gonzalez and R. E. Woods, Digital Image Processing, 3 edition. Upper Saddle River, N.J: Pearson, 2007.

M. S. Brown, "Understanding the In-Camera Image Processing Pipeline for Computer Vision," IEEE CVPR - Tutorial, 2016.

Color Camera Image formation

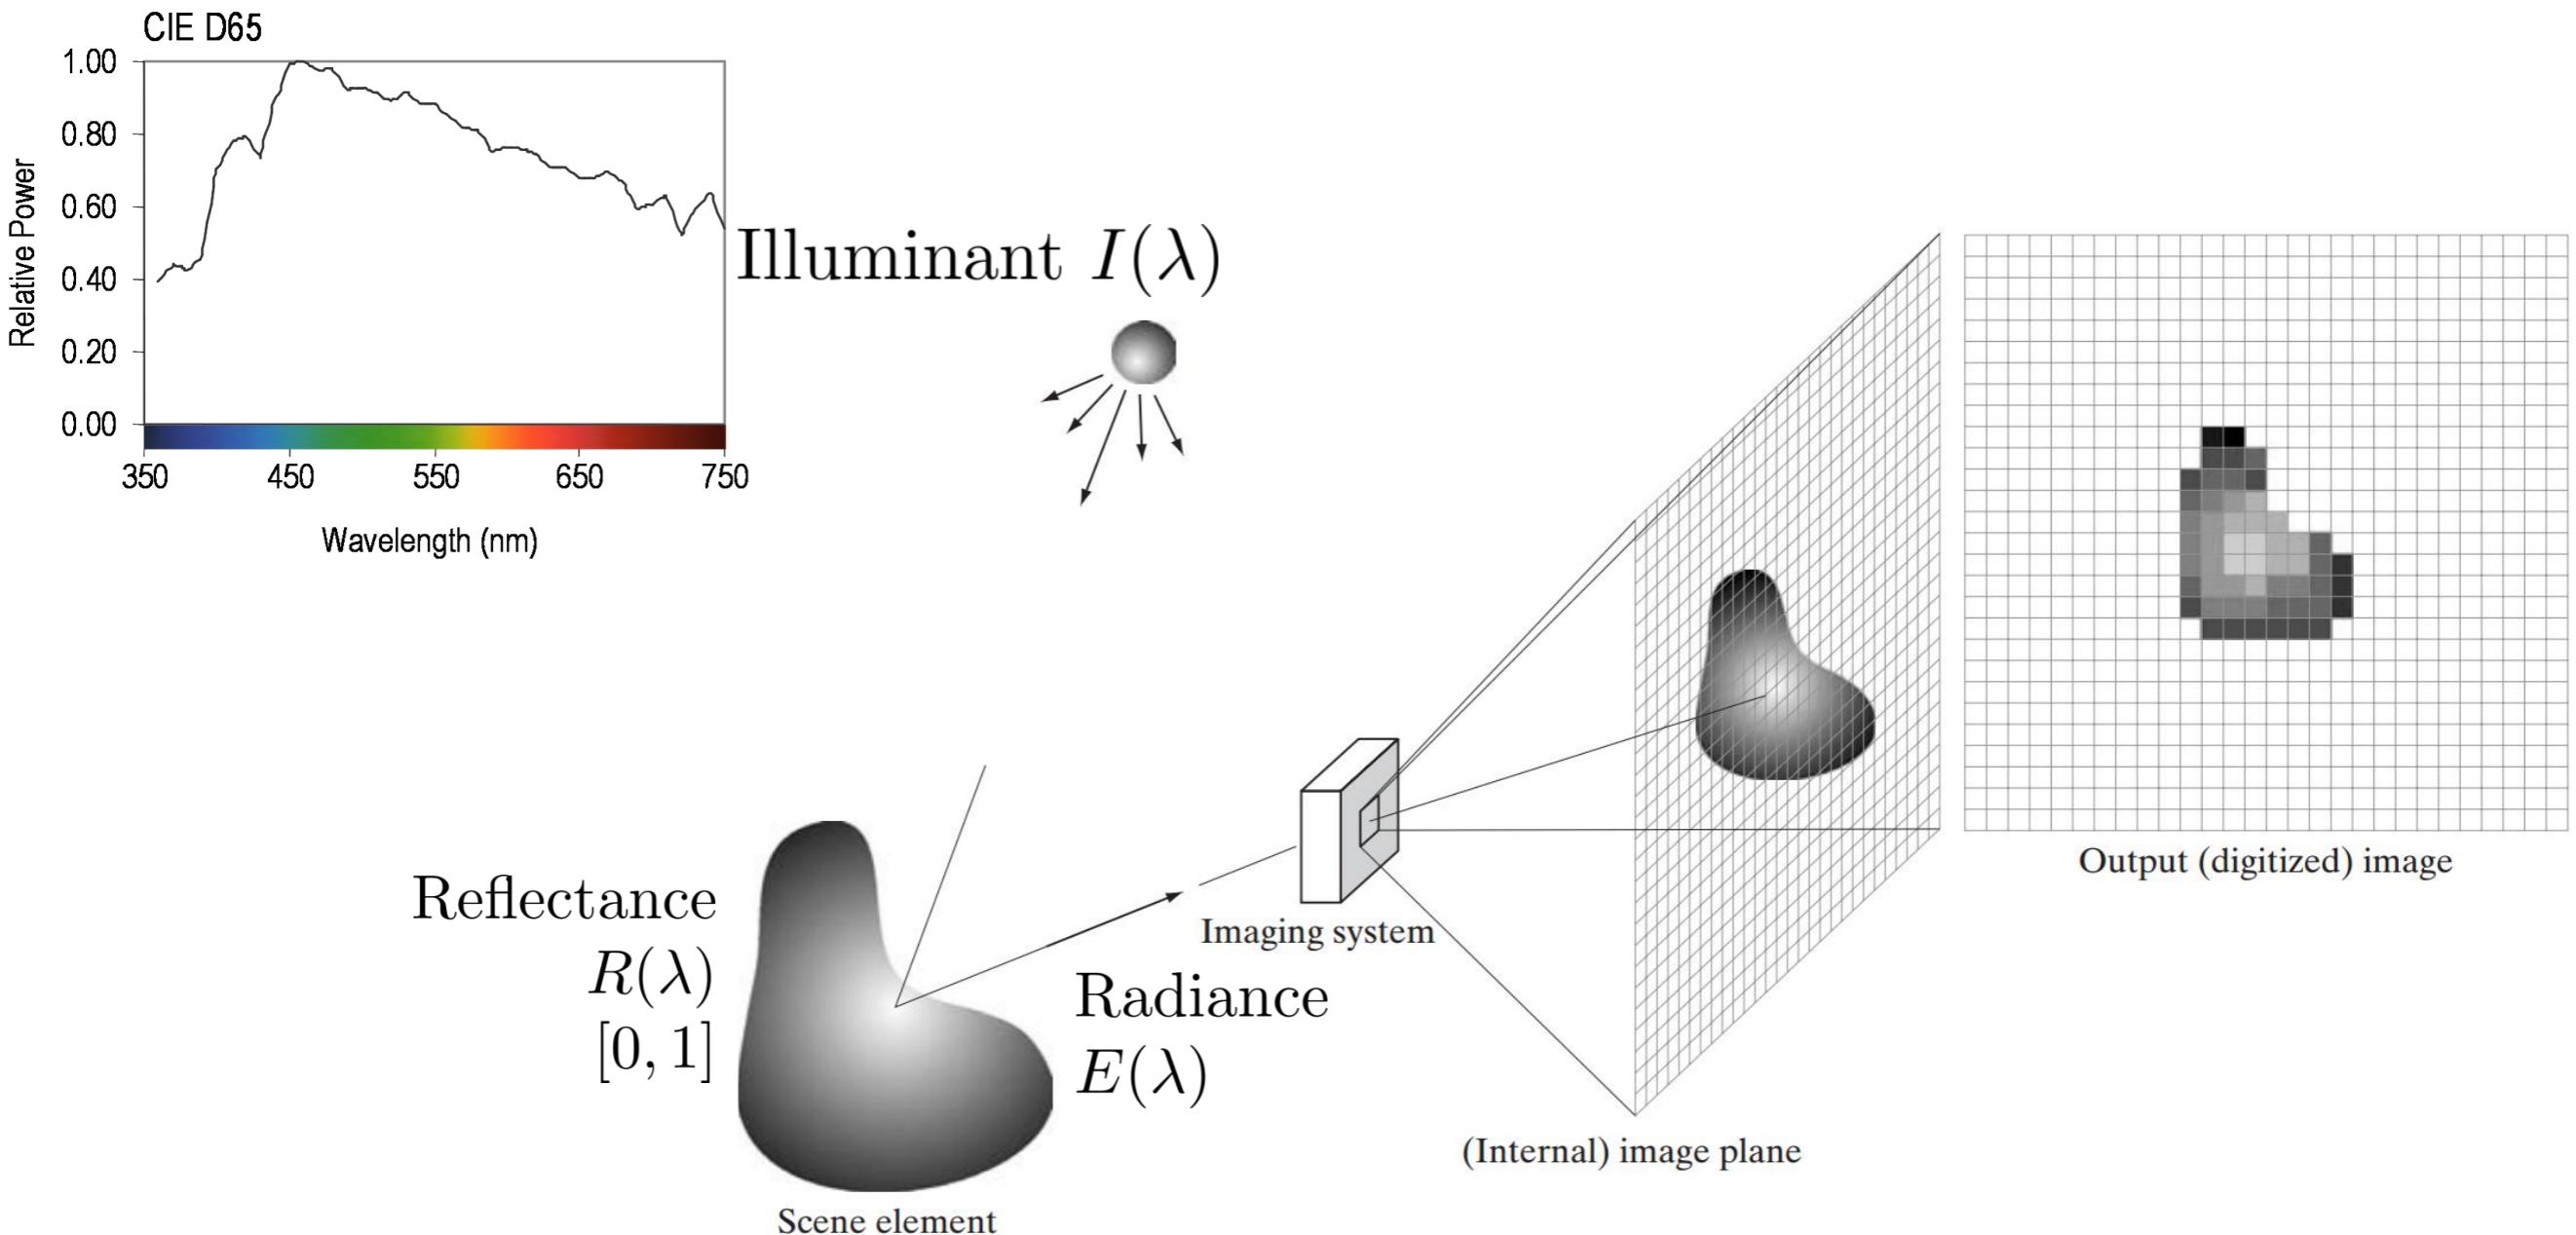


Image:

R. C. Gonzalez and R. E. Woods, Digital Image Processing, 3 edition. Upper Saddle River, N.J: Pearson, 2007.

M. S. Brown, "Understanding the In-Camera Image Processing Pipeline for Computer Vision," IEEE CVPR - Tutorial, 2016.

Color Camera Image formation

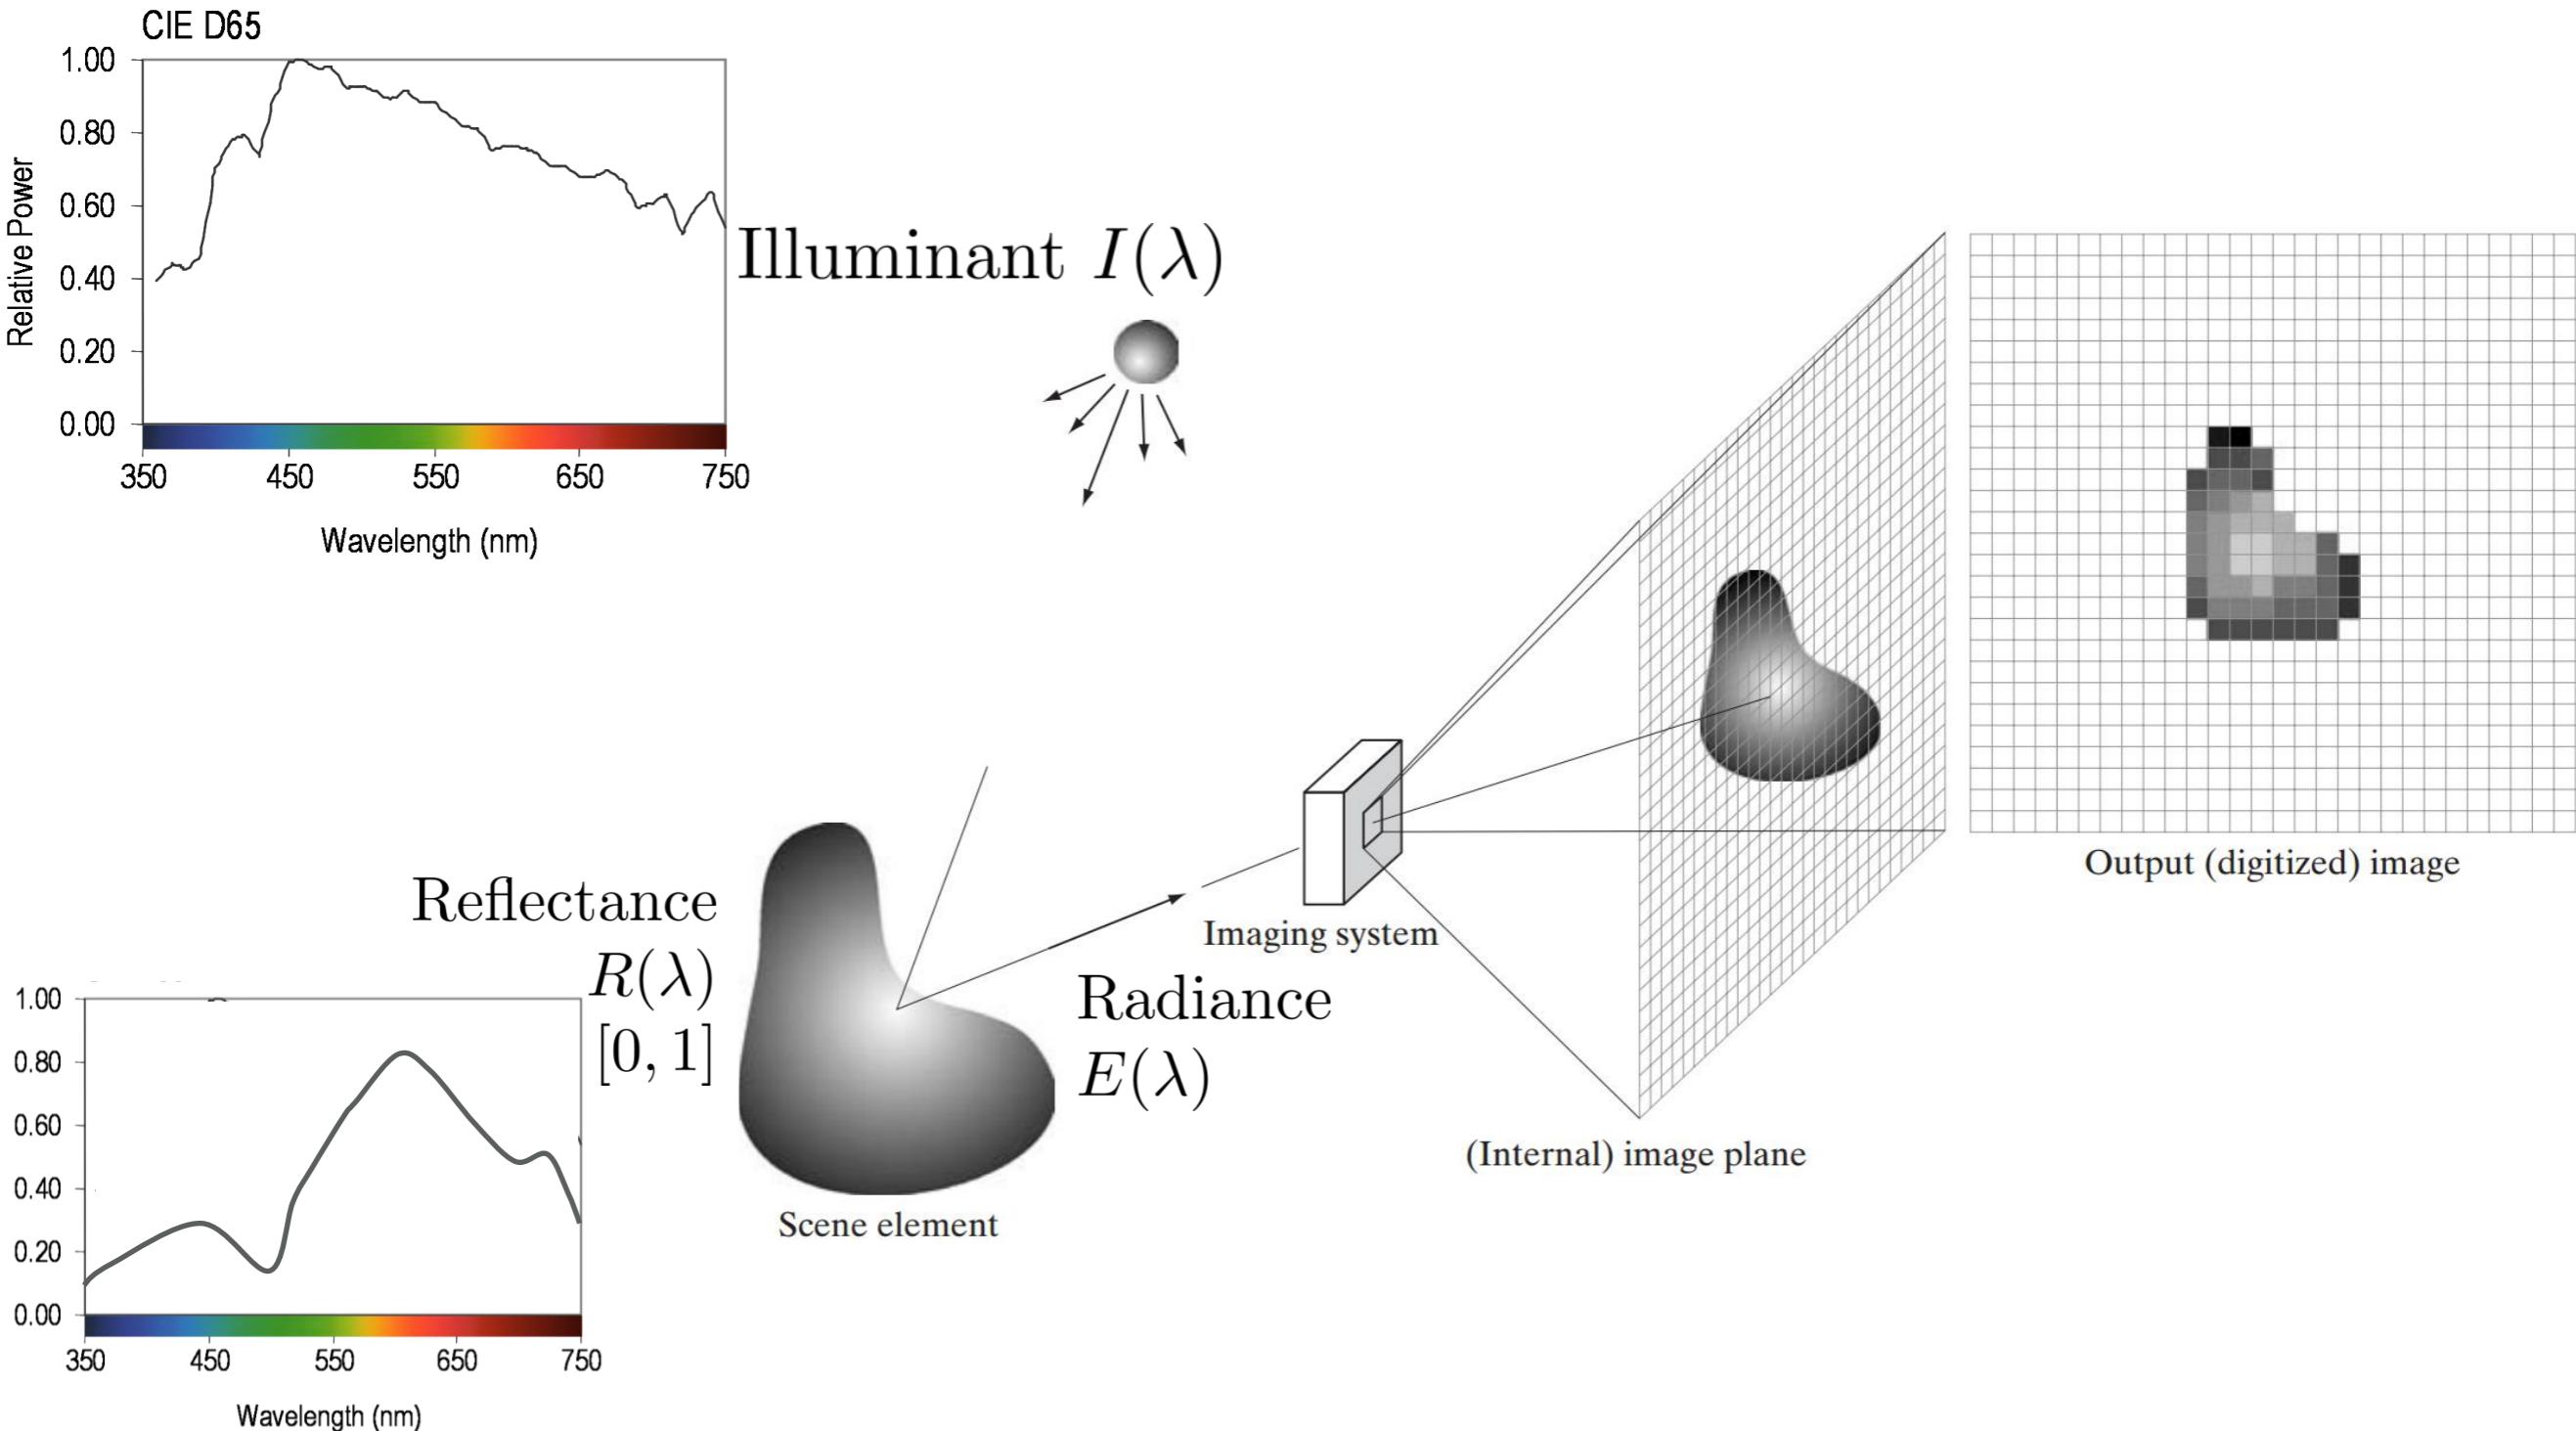


Image:

R. C. Gonzalez and R. E. Woods, Digital Image Processing, 3 edition. Upper Saddle River, N.J: Pearson, 2007.

M. S. Brown, "Understanding the In-Camera Image Processing Pipeline for Computer Vision," IEEE CVPR - Tutorial, 2016.

Color Camera Image formation

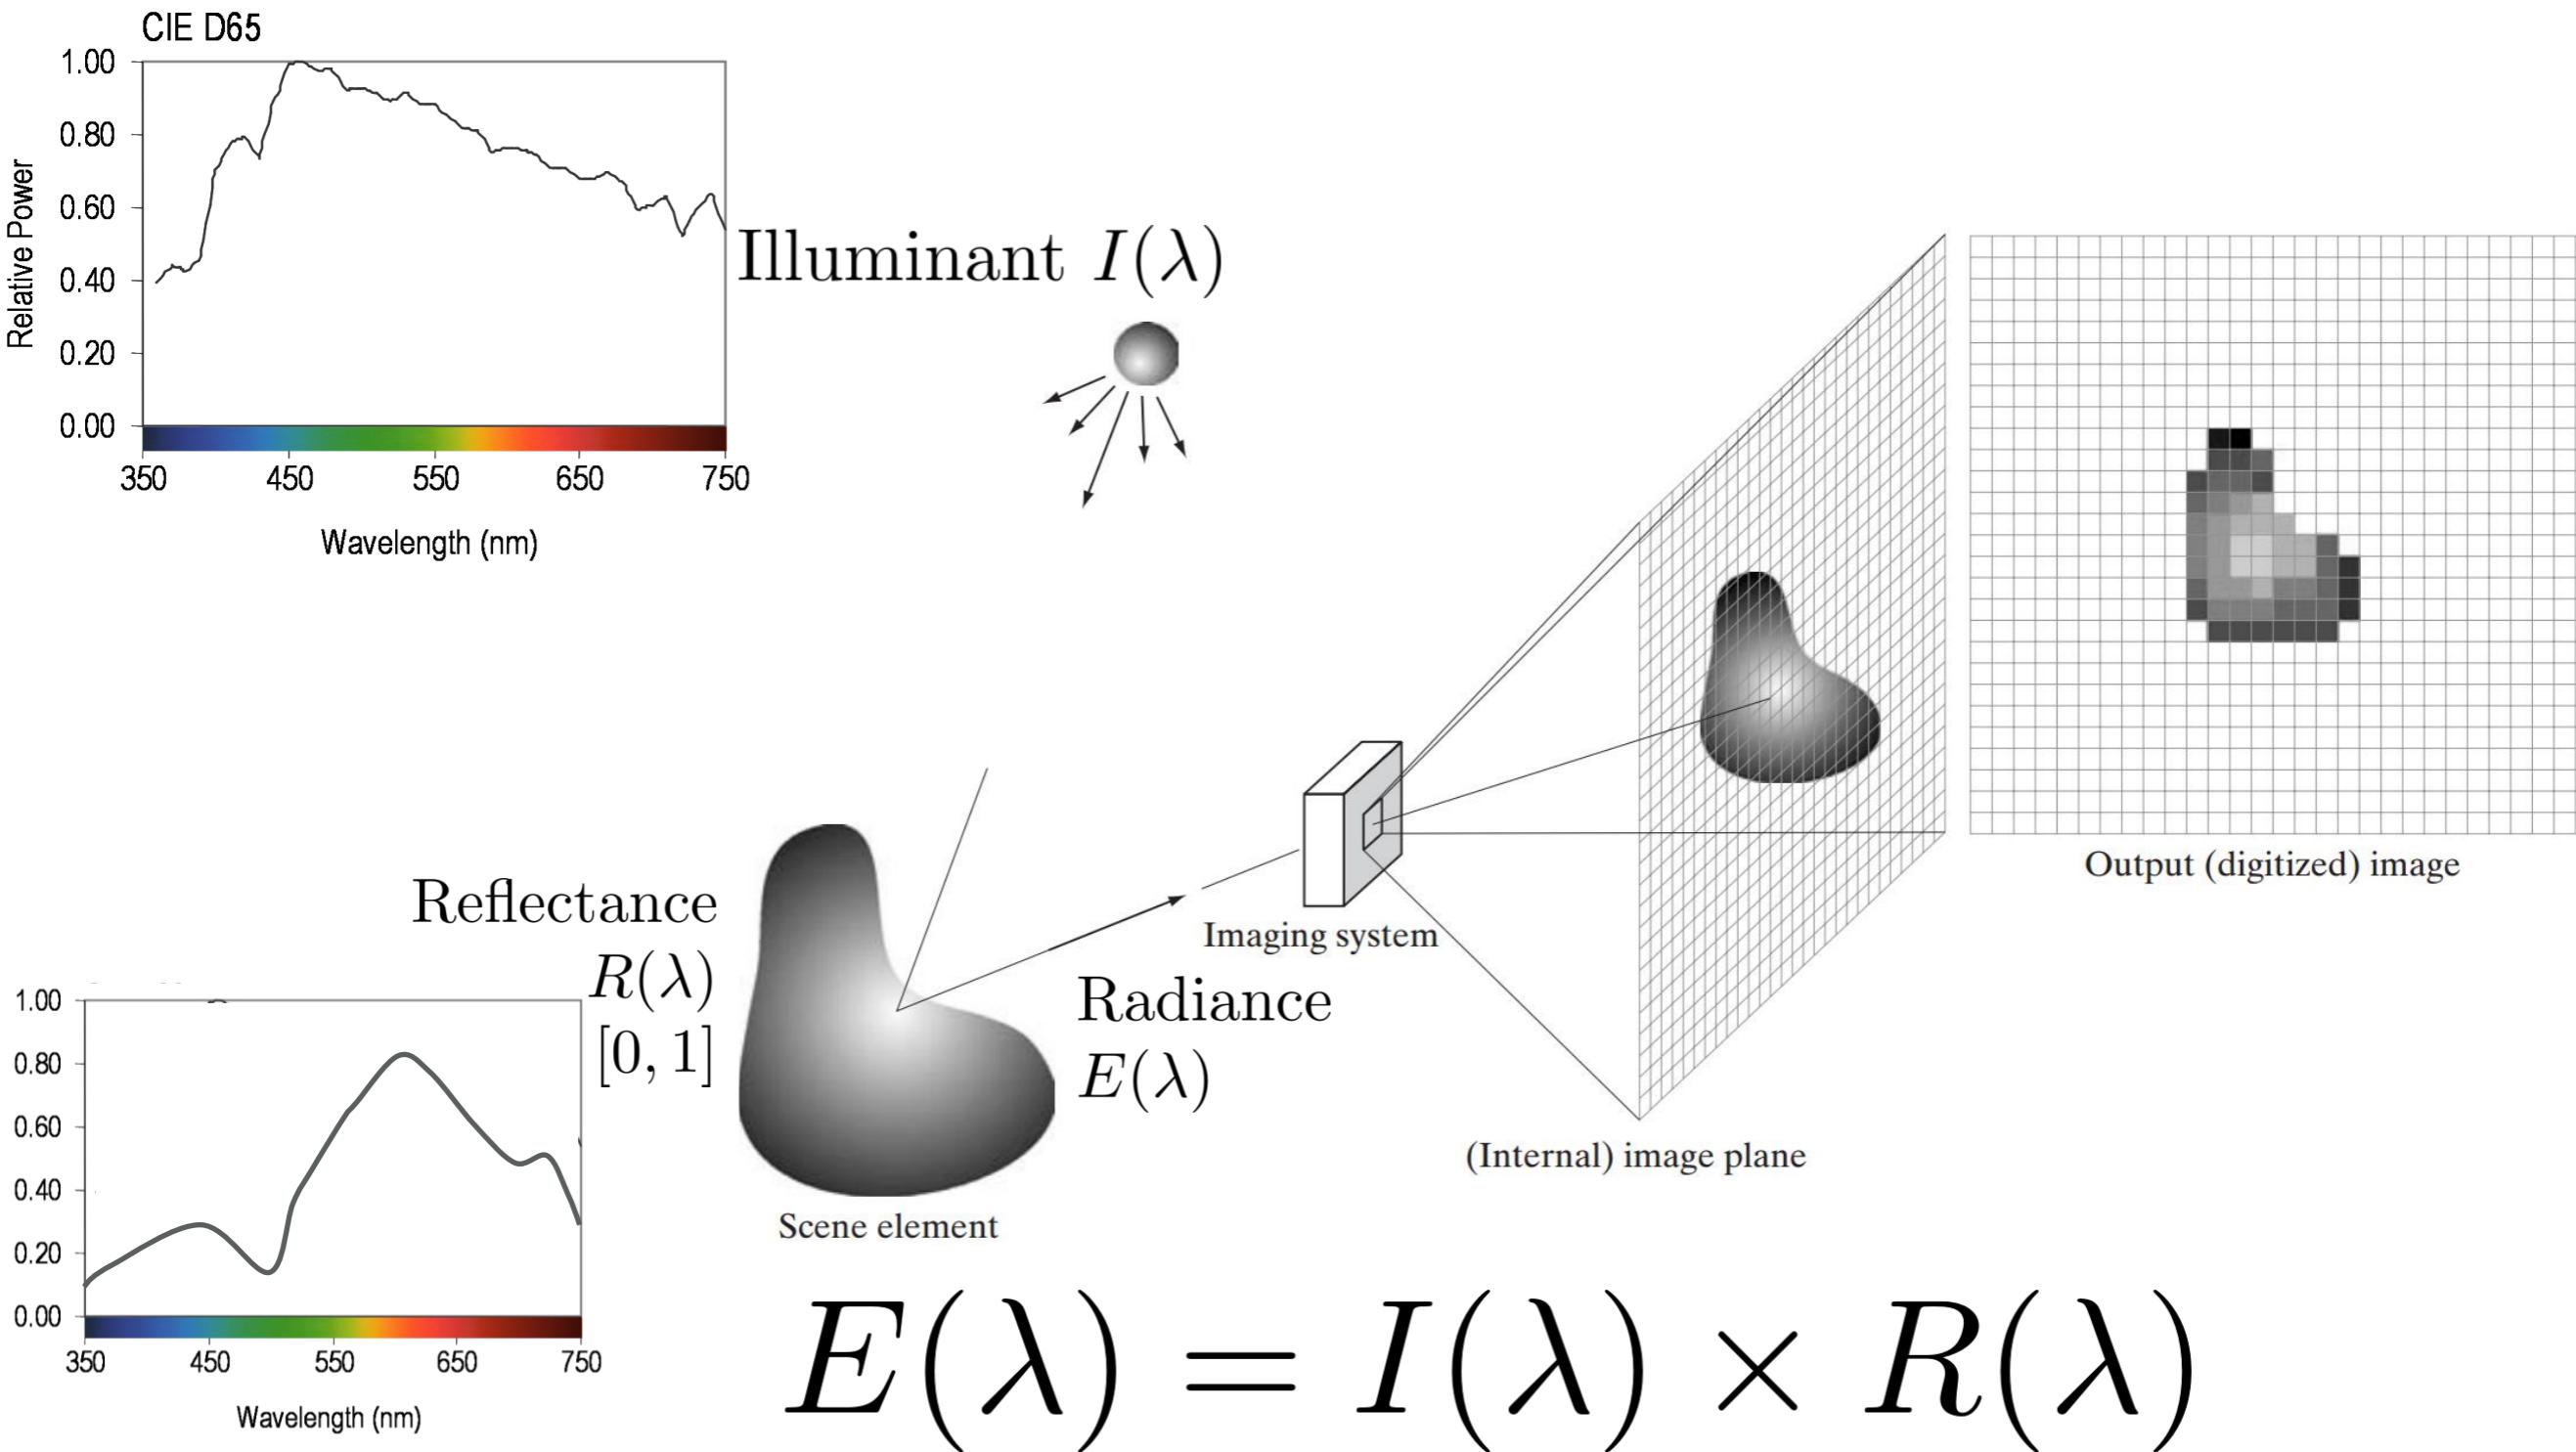
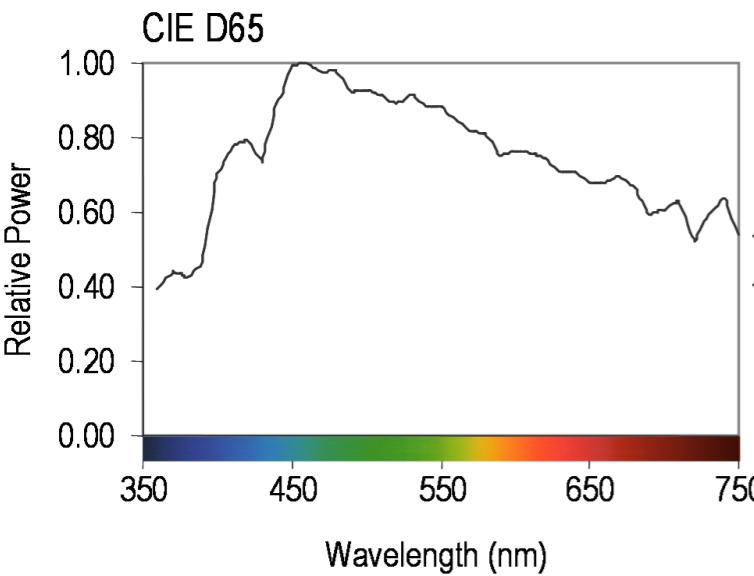


Image:

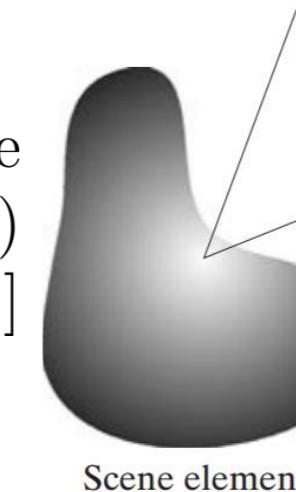
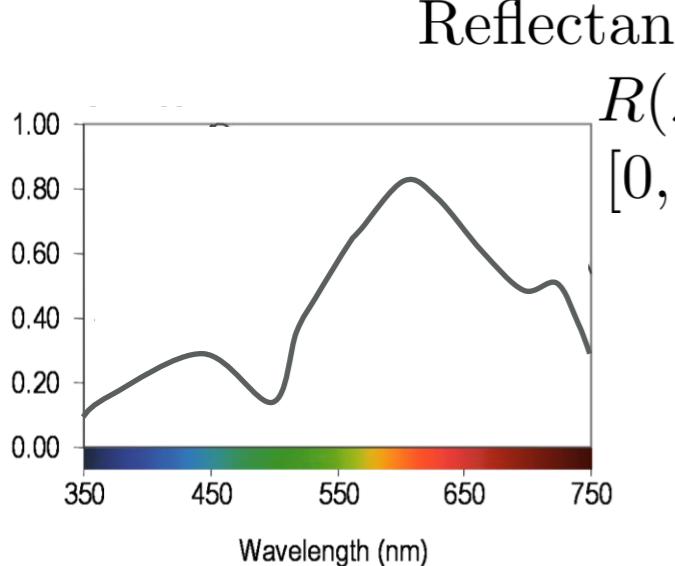
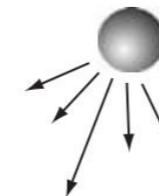
R. C. Gonzalez and R. E. Woods, Digital Image Processing, 3 edition. Upper Saddle River, N.J: Pearson, 2007.

M. S. Brown, "Understanding the In-Camera Image Processing Pipeline for Computer Vision," IEEE CVPR - Tutorial, 2016.

Color Camera Image formation



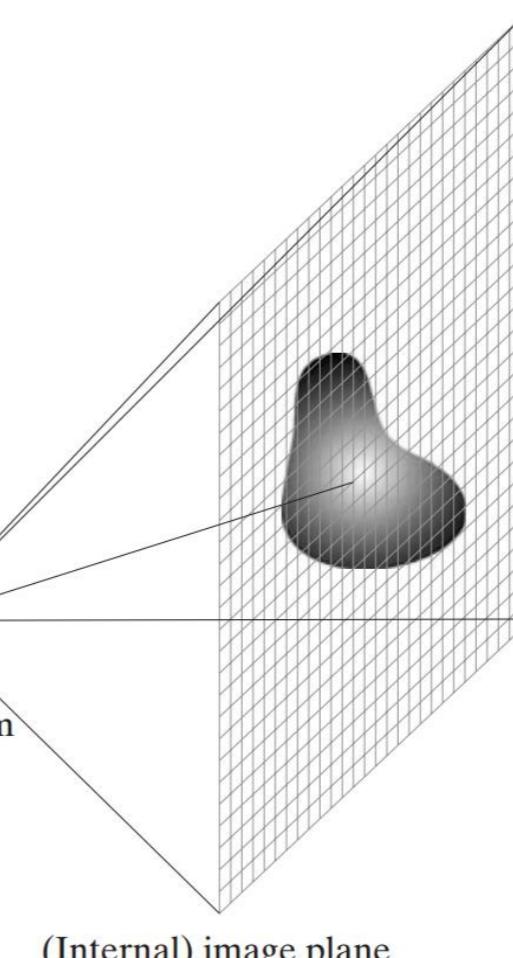
Illuminant $I(\lambda)$



Radiance
 $E(\lambda)$

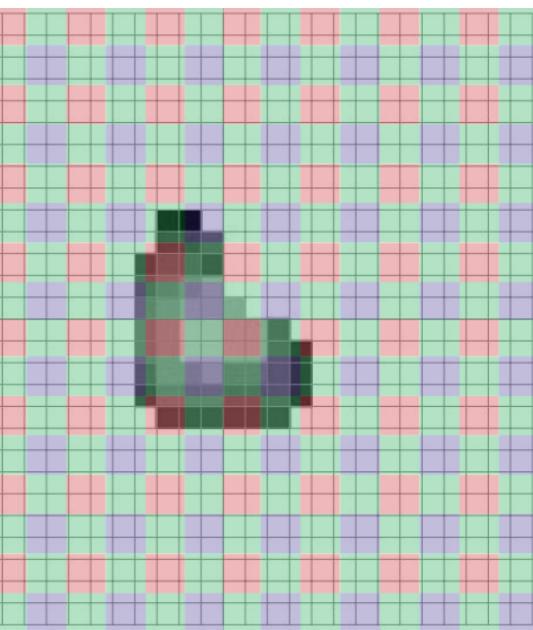
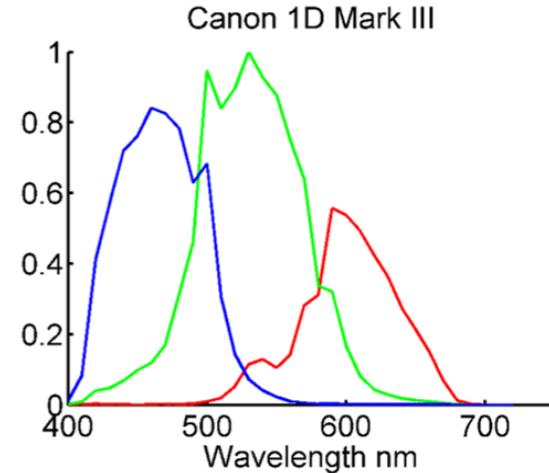


Imaging system



(Internal) image plane

RGB
Spectral
Sensitivity
Curves



Color Filter Array

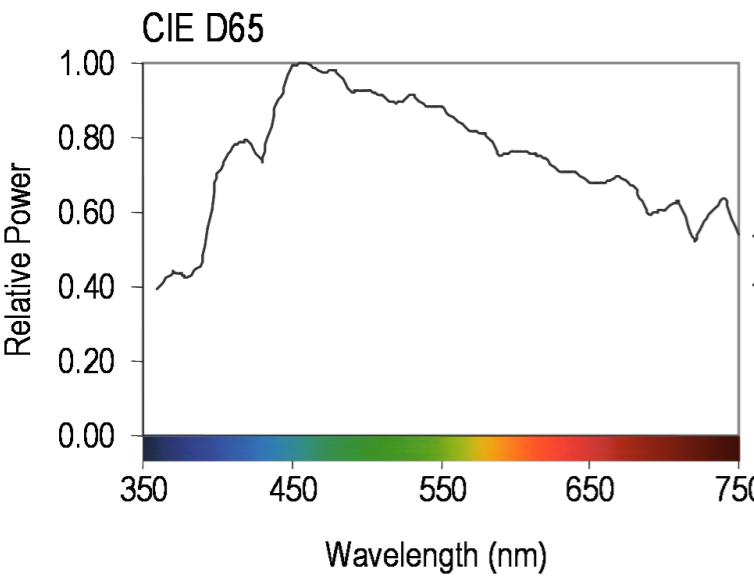
$$E(\lambda) = I(\lambda) \times R(\lambda)$$

Image:

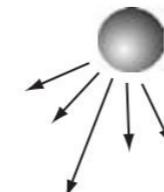
R. C. Gonzalez and R. E. Woods, Digital Image Processing, 3 edition. Upper Saddle River, N.J: Pearson, 2007.

M. S. Brown, "Understanding the In-Camera Image Processing Pipeline for Computer Vision," IEEE CVPR - Tutorial, 2016.

Color Camera Image formation

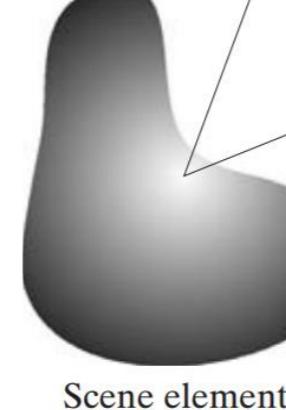
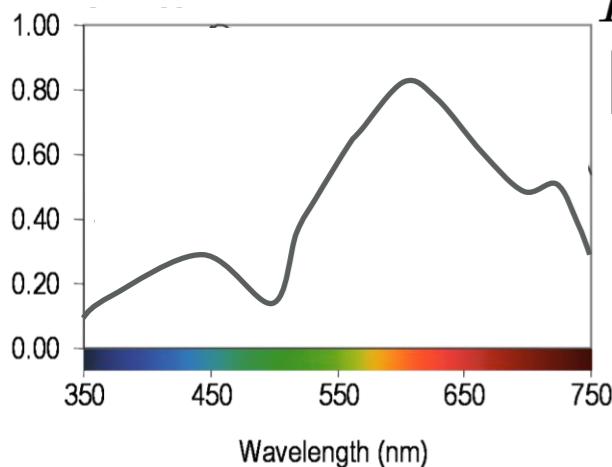


Illuminant $I(\lambda)$



Reflectance

$$R(\lambda) [0, 1]$$

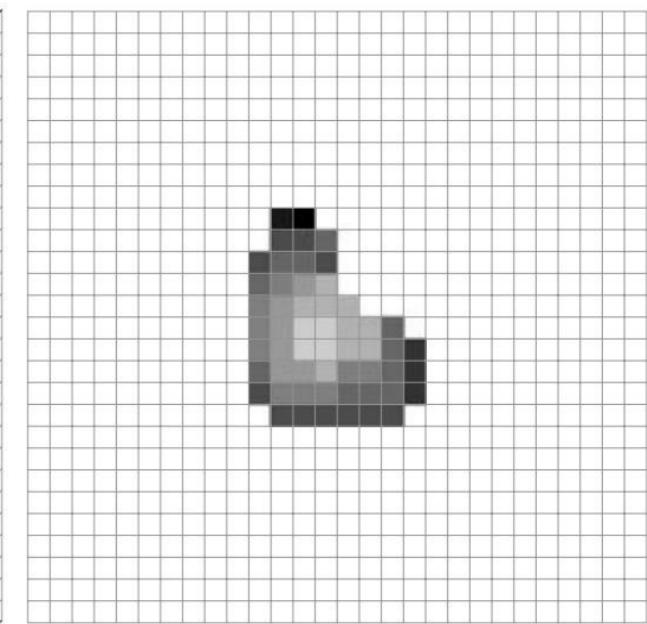
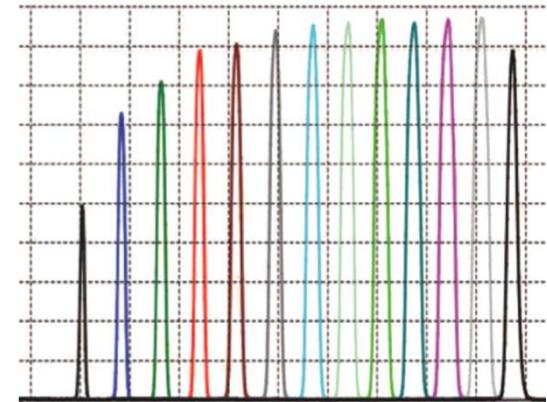


Radiance
 $E(\lambda)$

$$E(\lambda) = I(\lambda) \times R(\lambda)$$

Hyperspectral

Spectral
Sensitivity
Curves



Color Filter Array

Image:

R. C. Gonzalez and R. E. Woods, Digital Image Processing, 3 edition. Upper Saddle River, N.J: Pearson, 2007.

M. S. Brown, "Understanding the In-Camera Image Processing Pipeline for Computer Vision," IEEE CVPR - Tutorial, 2016.

Related work

Spectral reconstruction

- From RGB + extra help:

- Low-res HS image

[Cao2011] X. Cao, X. Tong, Q. Dai, and S. Lin, "High resolution multispectral video capture with a hybrid camera system," in CVPR 2011, 2011, pp. 297-304.

- Multiplexed light

[Park2007] J. I. Park, M. H. Lee, M. D. Grossberg, and S. K. Nayar, "Multispectral Imaging Using Multiplexed Illumination," in 2007 IEEE 11th International Conference on Computer Vision, 2007, pp. 1-8.

[Parmar2008] M. Parmar, S. Lansel, and B. A. Wandell, "Spatio-spectral reconstruction of the multispectral datacube using sparse recovery," in 2008 15th IEEE International Conference on Image Processing, 2008, pp. 473-476.

[Goel2015] M. Goel et al., "HyperCam: Hyperspectral Imaging for Ubiquitous Computing Applications," in Proceedings of the 2015 ACM International Joint Conference on Pervasive and Ubiquitous Computing, New York, NY, USA, 2015, pp. 145-156.

- No spatial info considered

[Nguyen2014] R. M. H. Nguyen, D. K. Prasad, and M. S. Brown, "Training-Based Spectral Reconstruction from a Single RGB Image," in Computer Vision ECCV 2014, D. Fleet, T. Pajdla, B. Schiele, and T. Tuytelaars, Eds. Springer International Publishing, 2014, pp. 186-201.

[Arad2016] B. Arad and O. Ben-Shahar, "Sparse Recovery of Hyperspectral Signal from Natural RGB Images," in ECCV, 2016, pp. 19-34.

- Leveraging spatial context:

[Robles-Kelly2015] A. Robles-Kelly, "Single Image Spectral Reconstruction for Multimedia Applications," in ACM Multimedia, 2015.

[Galliani2017] S. Galliani, C. Lanaras, D. Marmanis, E. Baltasvias, and K. Schindler, "Learned Spectral Super-Resolution," arXiv:1703.09470 [cs], Mar. 2017.

Related work

Spectral reconstruction

- From RGB + extra help:

- Low-res HS image

[Cao2011] X. Cao, X. Tong, Q. Dai, and S. Lin, "High resolution multispectral video capture with a hybrid camera system," in CVPR 2011, 2011, pp. 297-304.

- Multiplexed light

[Park2007] J. I. Park, M. H. Lee, M. D. Grossberg, and S. K. Nayar, "Multispectral Imaging Using Multiplexed Illumination," in 2007 IEEE 11th International Conference on Computer Vision, 2007, pp. 1-8.

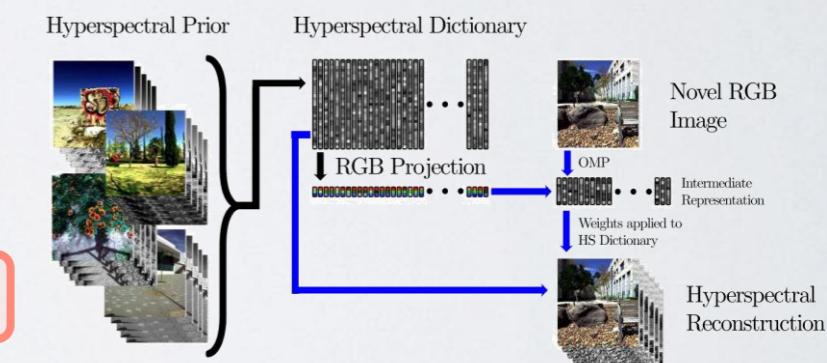
[Parmar2008] M. Parmar, S. Lansel, and B. A. Wandell, "Spatio-spectral reconstruction of the multispectral datacube using sparse recovery," in 2008 15th IEEE International Conference on Image Processing, 2008, pp. 473-476.

[Goel2015] M. Goel et al., "HyperCam: Hyperspectral Imaging for Ubiquitous Computing Applications," in Proceedings of the 2015 ACM International Joint Conference on Pervasive and Ubiquitous Computing, New York, NY, USA, 2015, pp. 145-156.

- No spatial info considered

[Nguyen2014] R. M. H. Nguyen, D. K. Prasad, and M. S. Brown, "Training-Based Spectral Reconstruction from a Single RGB Image," in Computer Vision ECCV 2014, D. Fleet, T. Pajdla, B. Schiele, and T. Tuytelaars, Eds. Springer International Publishing, 2014, pp. 186-201.

[Arad2016] B. Arad and O. Ben-Shahar, "Sparse Recovery of Hyperspectral Signal from Natural RGB Images," in ECCV, 2016, pp. 19-34.



- Leveraging spatial context:

[Robles-Kelly2015] A. Robles-Kelly, "Single Image Spectral Reconstruction for Multimedia Applications," in ACM Multimedia, 2015.

[Galliani2017] S. Galliani, C. Lanaras, D. Marmanis, E. Baltasvias, and K. Schindler, "Learned Spectral Super-Resolution," arXiv:1703.09470 [cs], Mar. 2017.

Related work

Spectral reconstruction

- From RGB + extra help:

- Low-res HS image

[Cao2011] X. Cao, X. Tong, Q. Dai, and S. Lin, "High resolution multispectral video capture with a hybrid camera system," in CVPR 2011, 2011, pp. 297-304.

- Multiplexed light

[Park2007] J. I. Park, M. H. Lee, M. D. Grossberg, and S. K. Nayar, "Multispectral Imaging Using Multiplexed Illumination," in 2007 IEEE 11th International Conference on Computer Vision, 2007, pp. 1-8.

[Parmar2008] M. Parmar, S. Lansel, and B. A. Wandell, "Spatio-spectral reconstruction of the multispectral datacube using sparse recovery," in 2008 15th IEEE International Conference on Image Processing, 2008, pp. 473-476.

[Goel2015] M. Goel et al., "HyperCam: Hyperspectral Imaging for Ubiquitous Computing Applications," in Proceedings of the 2015 ACM International Joint Conference on Pervasive and Ubiquitous Computing, New York, NY, USA, 2015, pp. 145-156.

- No spatial info considered

[Nguyen2014] R. M. H. Nguyen, D. K. Prasad, and M. S. Brown, "Training-Based Spectral Reconstruction from a Single RGB Image," in Computer Vision ECCV 2014, D. Fleet, T. Pajdla, B. Schiele, and T. Tuytelaars, Eds. Springer International Publishing, 2014, pp. 186-201.

[Arad2016] B. Arad and O. Ben-Shahar, "Sparse Recovery of Hyperspectral Signal from Natural RGB Images," in ECCV, 2016, pp. 19-34.

- Leveraging spatial context:

[Robles-Kelly2015] A. Robles-Kelly, "Single Image Spectral Reconstruction for Multimedia Applications," in ACM Multimedia, 2015.

[Galliani2017] S. Galliani, C. Lanaras, D. Marmanis, E. Baltasvias, and K. Schindler, "Learned Spectral Super-Resolution," arXiv:1703.09470 [cs], Mar. 2017.

Related work

Spectral reconstruction

- From RGB + extra help:

- Low-res HS image

[Cao2011] X. Cao, X. Tong, Q. Dai, and S. Lin, "High resolution multispectral video capture with a hybrid camera system," in CVPR 2011, 2011, pp. 297-304.

- Multiplexed light

[Park2007] J. I. Park, M. H. Lee, M. D. Grossberg, and S. K. Nayar, "Multispectral Imaging Using Multiplexed Illumination," in 2007 IEEE 11th International Conference on Computer Vision, 2007, pp. 1-8.

[Parmar2008] M. Parmar, S. Lansel, and B. A. Wandell, "Spatio-spectral reconstruction of the multispectral datacube using sparse recovery," in 2008 15th IEEE International Conference on Image Processing, 2008, pp. 473-476.

[Goel2015] M. Goel et al., "HyperCam: Hyperspectral Imaging for Ubiquitous Computing Applications," in Proceedings of the 2015 ACM International Joint Conference on Pervasive and Ubiquitous Computing, New York, NY, USA, 2015, pp. 145-156.

- No spatial info considered

[Nguyen2014] R. M. H. Nguyen, D. K. Prasad, and M. S. Brown, "Training-Based Spectral Reconstruction from a Single RGB Image," in Computer Vision ECCV 2014, D. Fleet, T. Pajdla, B. Schiele, and T. Tuytelaars, Eds. Springer International Publishing, 2014, pp. 186-201.

[Arad2016] B. Arad and O. Ben-Shahar, "Sparse Recovery of Hyperspectral Signal from Natural RGB Images," in ECCV, 2016, pp. 19-34.

- Leveraging spatial context:

[Robles-Kelly2015] A. Robles-Kelly, "Single Image Spectral Reconstruction for Multimedia Applications," in ACM Multimedia, 2015.

[Galliani2017] S. Galliani, C. Lanaras, D. Marmanis, E. Baltsavias, and K. Schindler, "Learned Spectral Super-Resolution," arXiv:1703.09470 [cs], Mar. 2017.

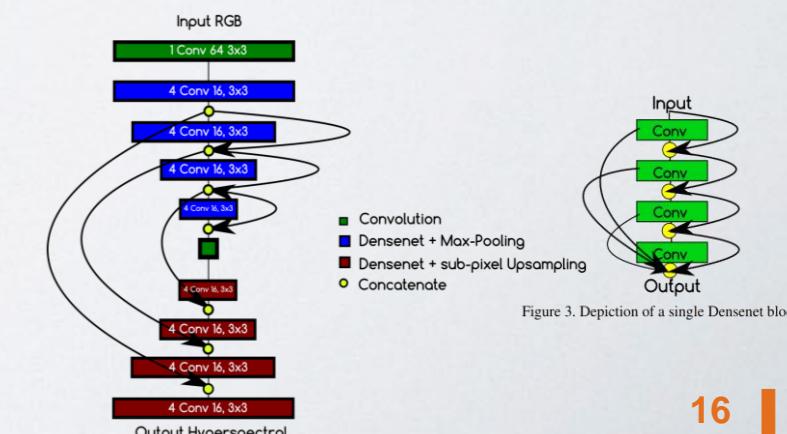
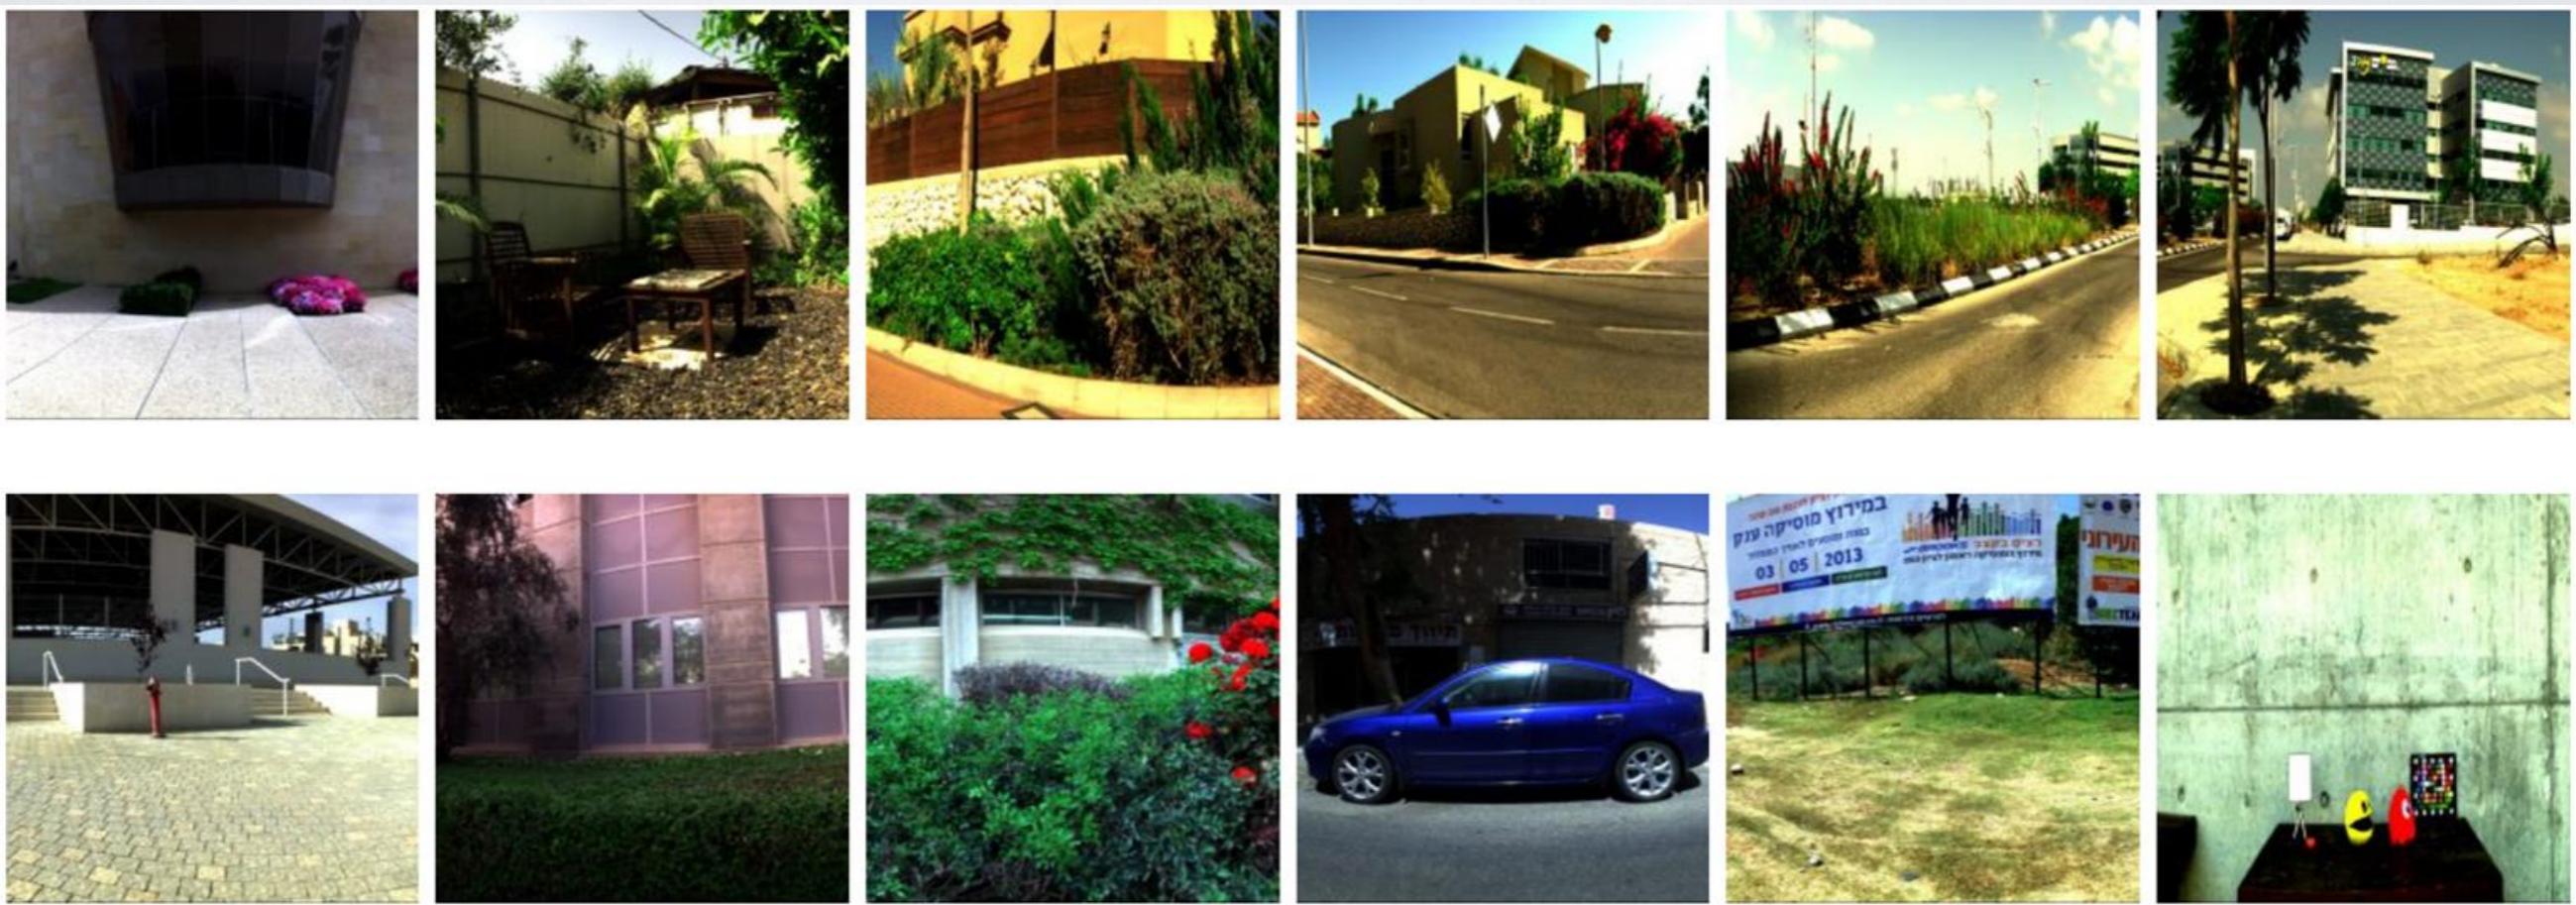


Figure 3. Depiction of a single DenseNet block

Dataset

ICVL dataset [Arad2016]

- ~200 images
- Specim PS Kappa DX4 + rotary
- 1392×1300 spatial resolution over 519 spectral bands (400-1,000nm at roughly 1.25nm increments).
- Downsampled to 31 spectral channels from 400nm to 700nm at 10nm increments

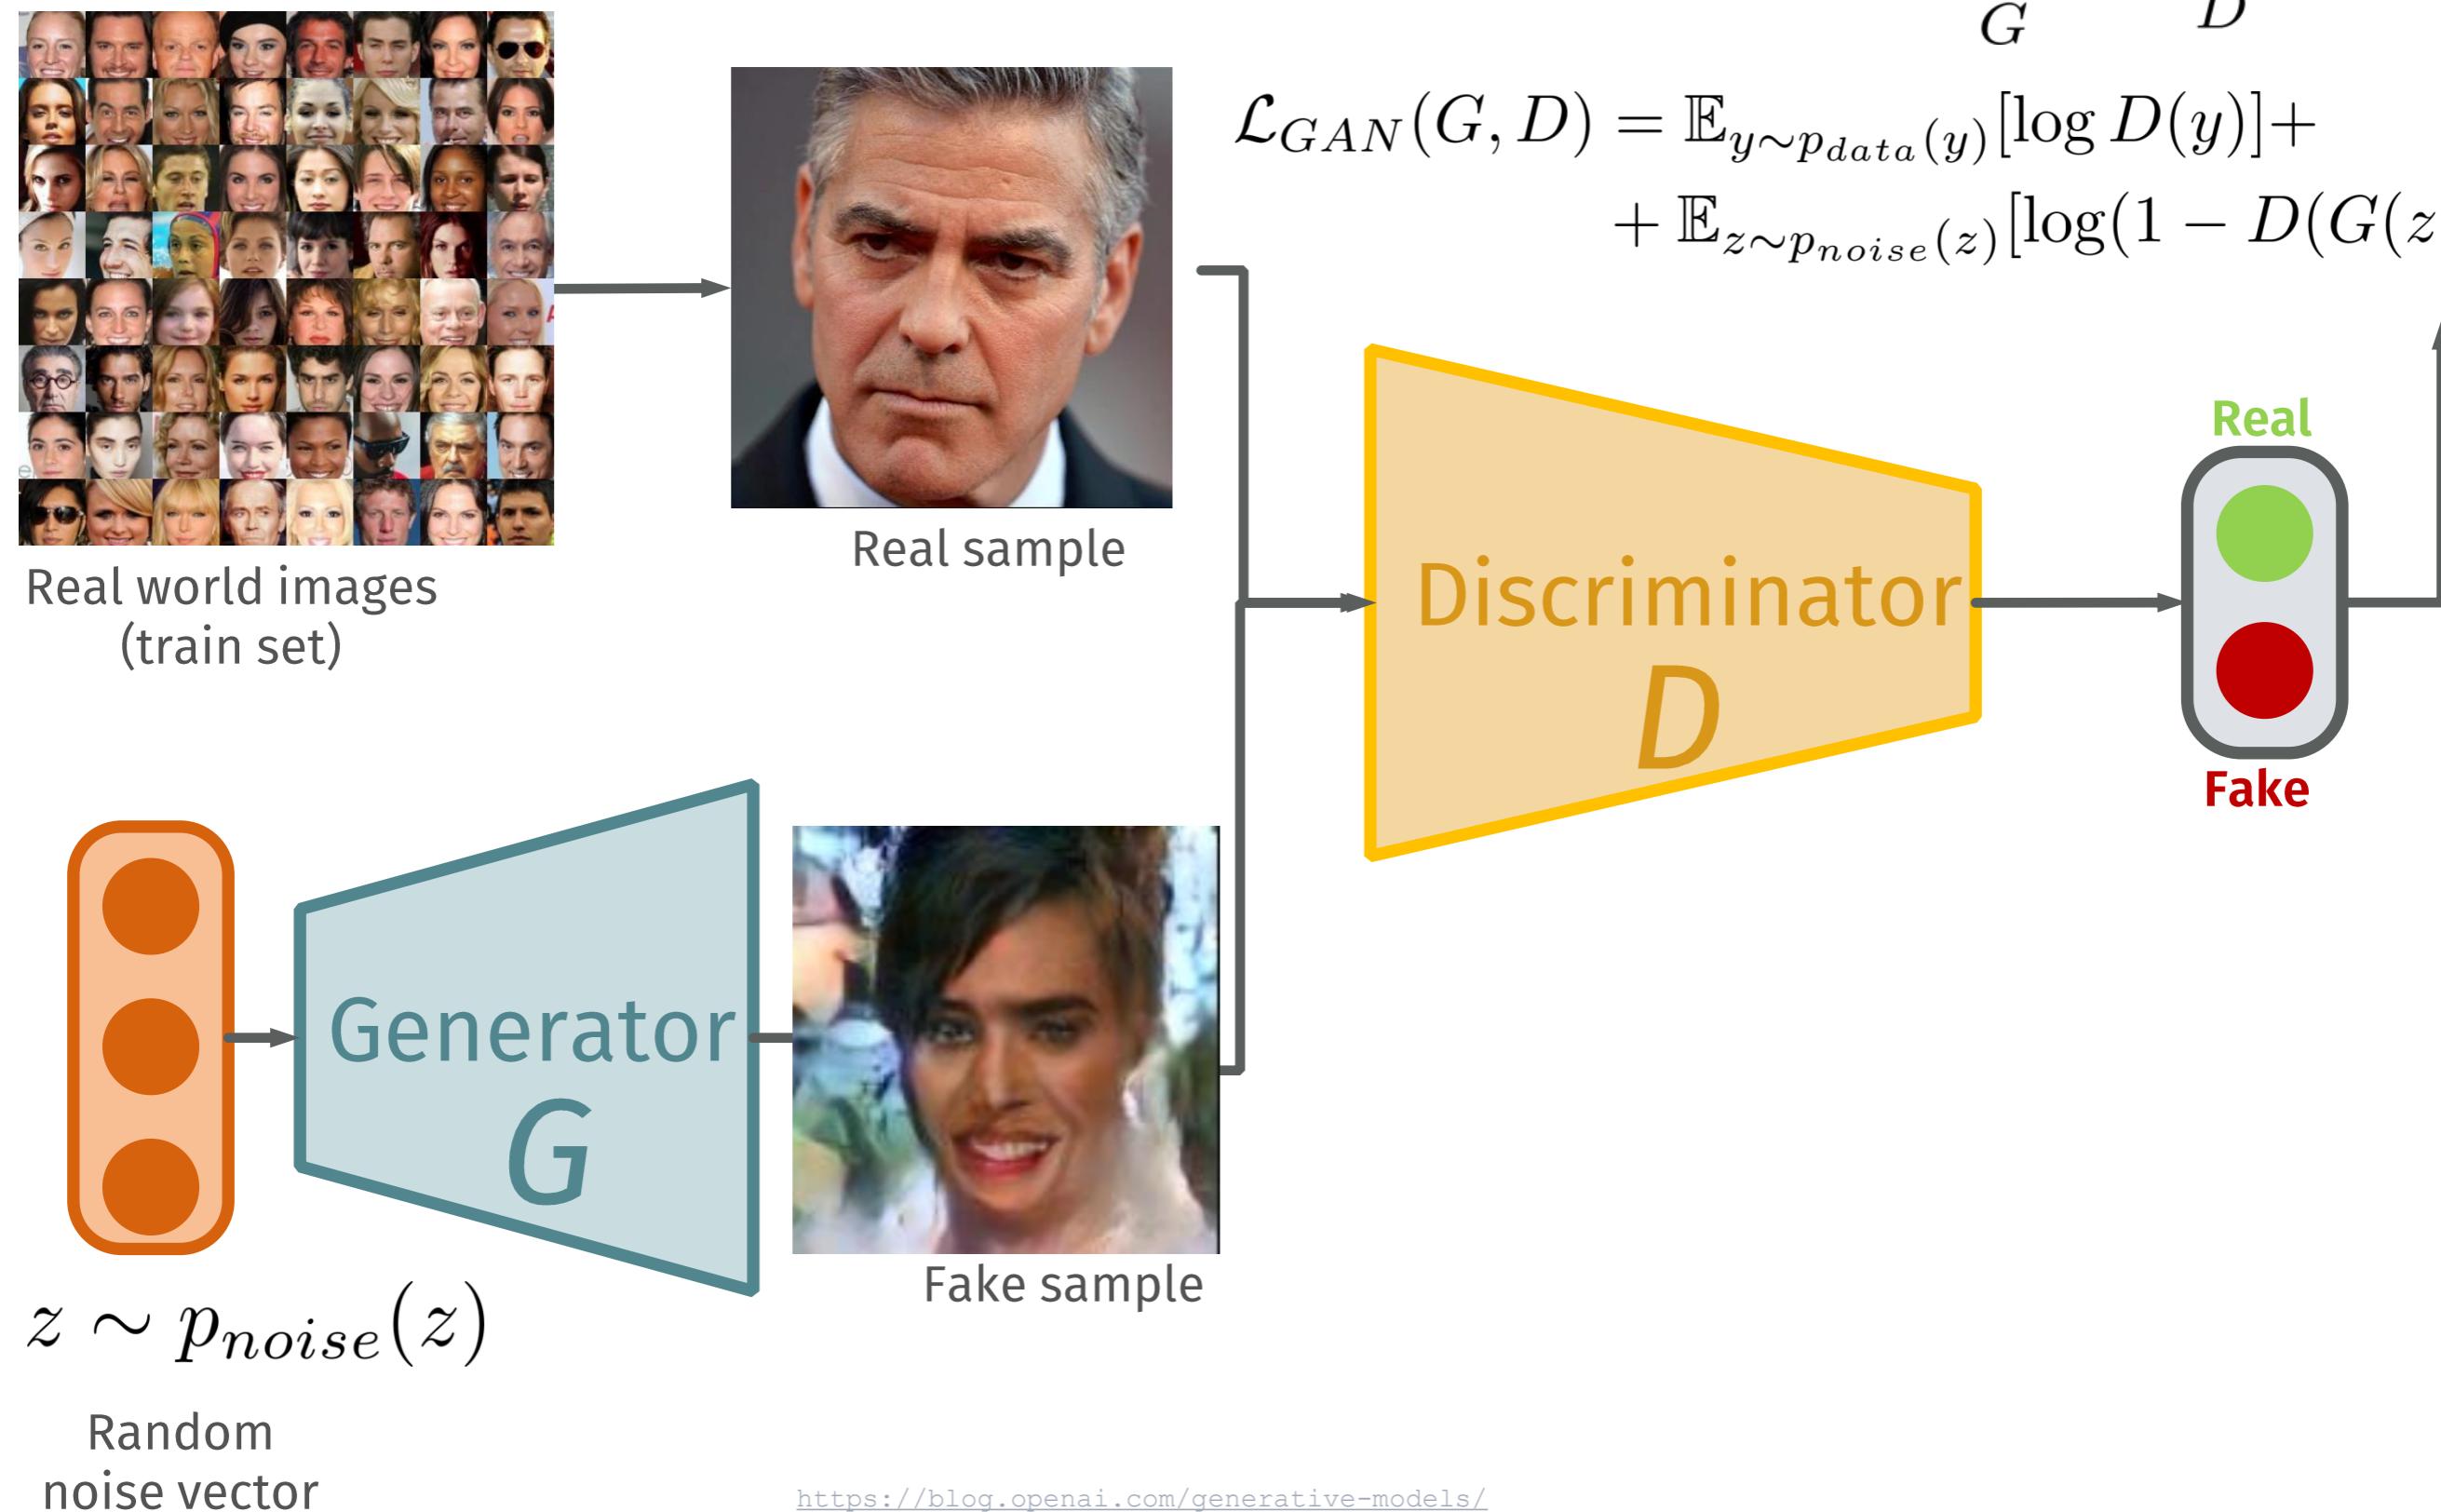


Adversarial Networks for Spatial Context-Aware Spectral Image Reconstruction from RGB

Adversarial Networks for Spatial Context-Aware Spectral Image Reconstruction from RGB

Generative Adversarial Networks (GANs)

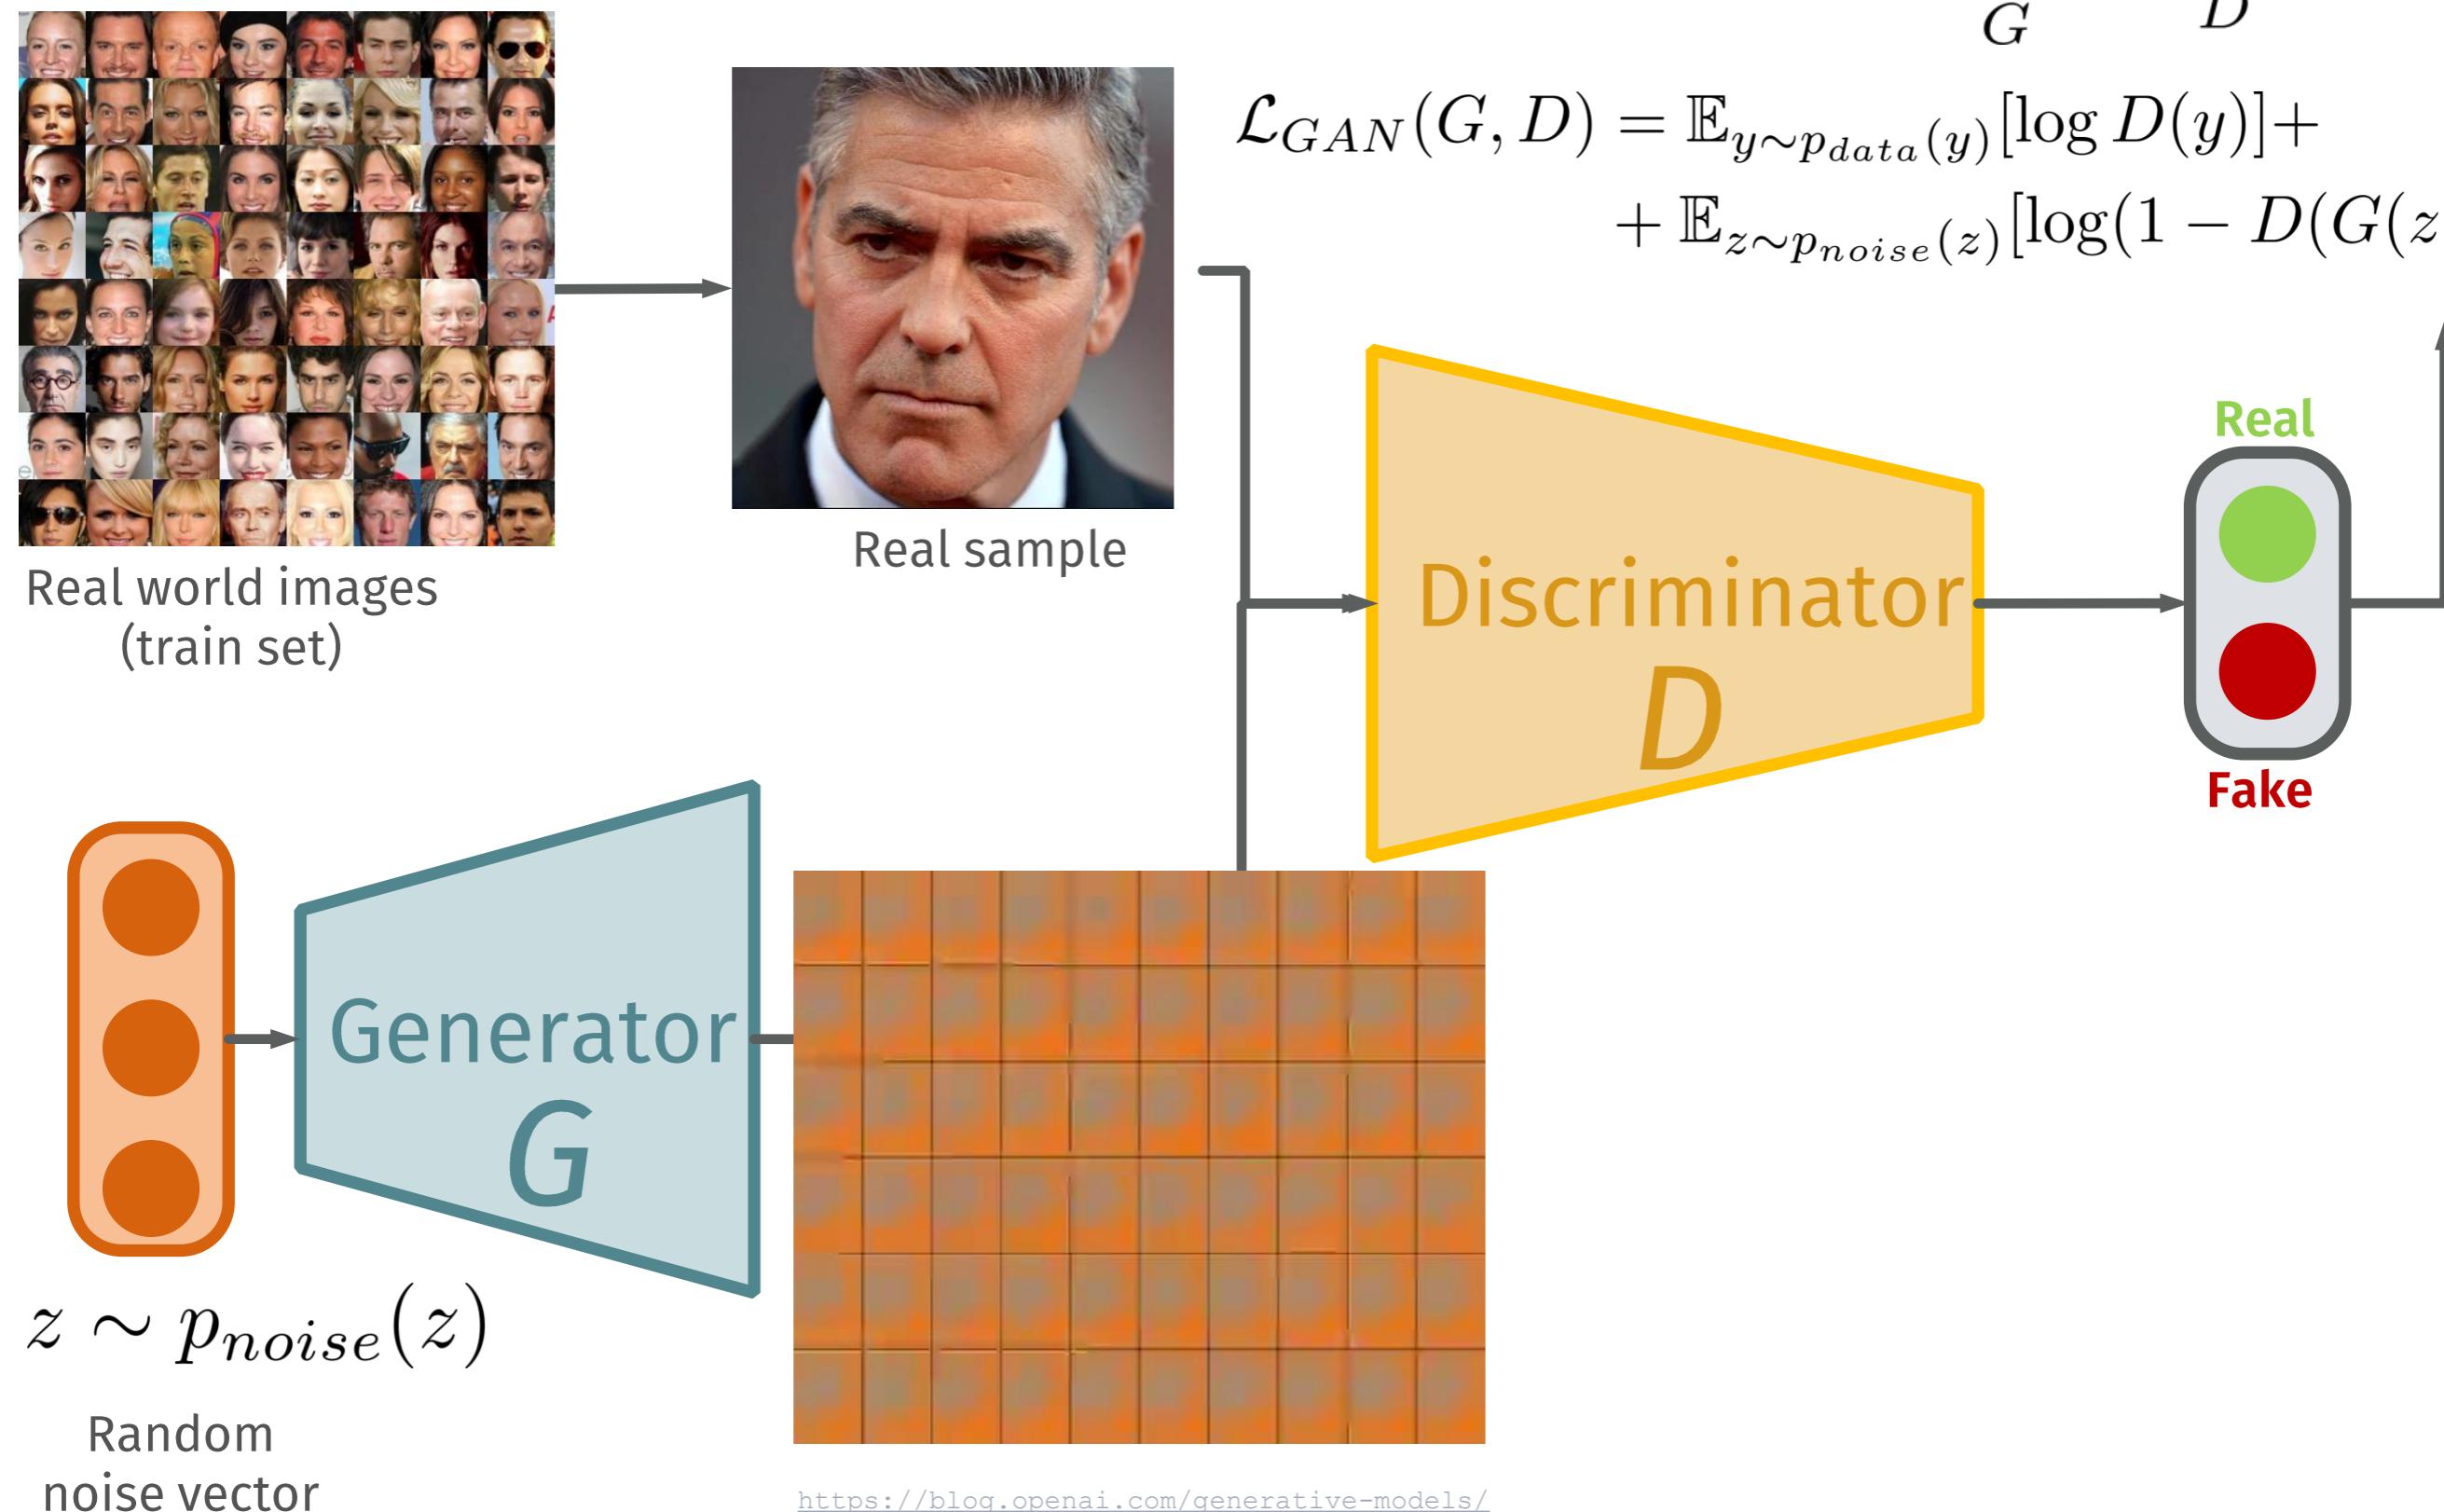
$$G^* = \operatorname{argmin}_G \max_D \mathcal{L}_{GAN}$$



<https://blog.openai.com/generative-models/>
<https://www.youtube.com/watch?v=J0o6LhaUSSc&vl=en>

Generative Adversarial Networks (GANs)

$$G^* = \operatorname{argmin}_G \max_D \mathcal{L}_{GAN}$$



<https://blog.openai.com/generative-models/>
<https://www.youtube.com/watch?v=J0o6LhaUSSc&vl=en>

GANs

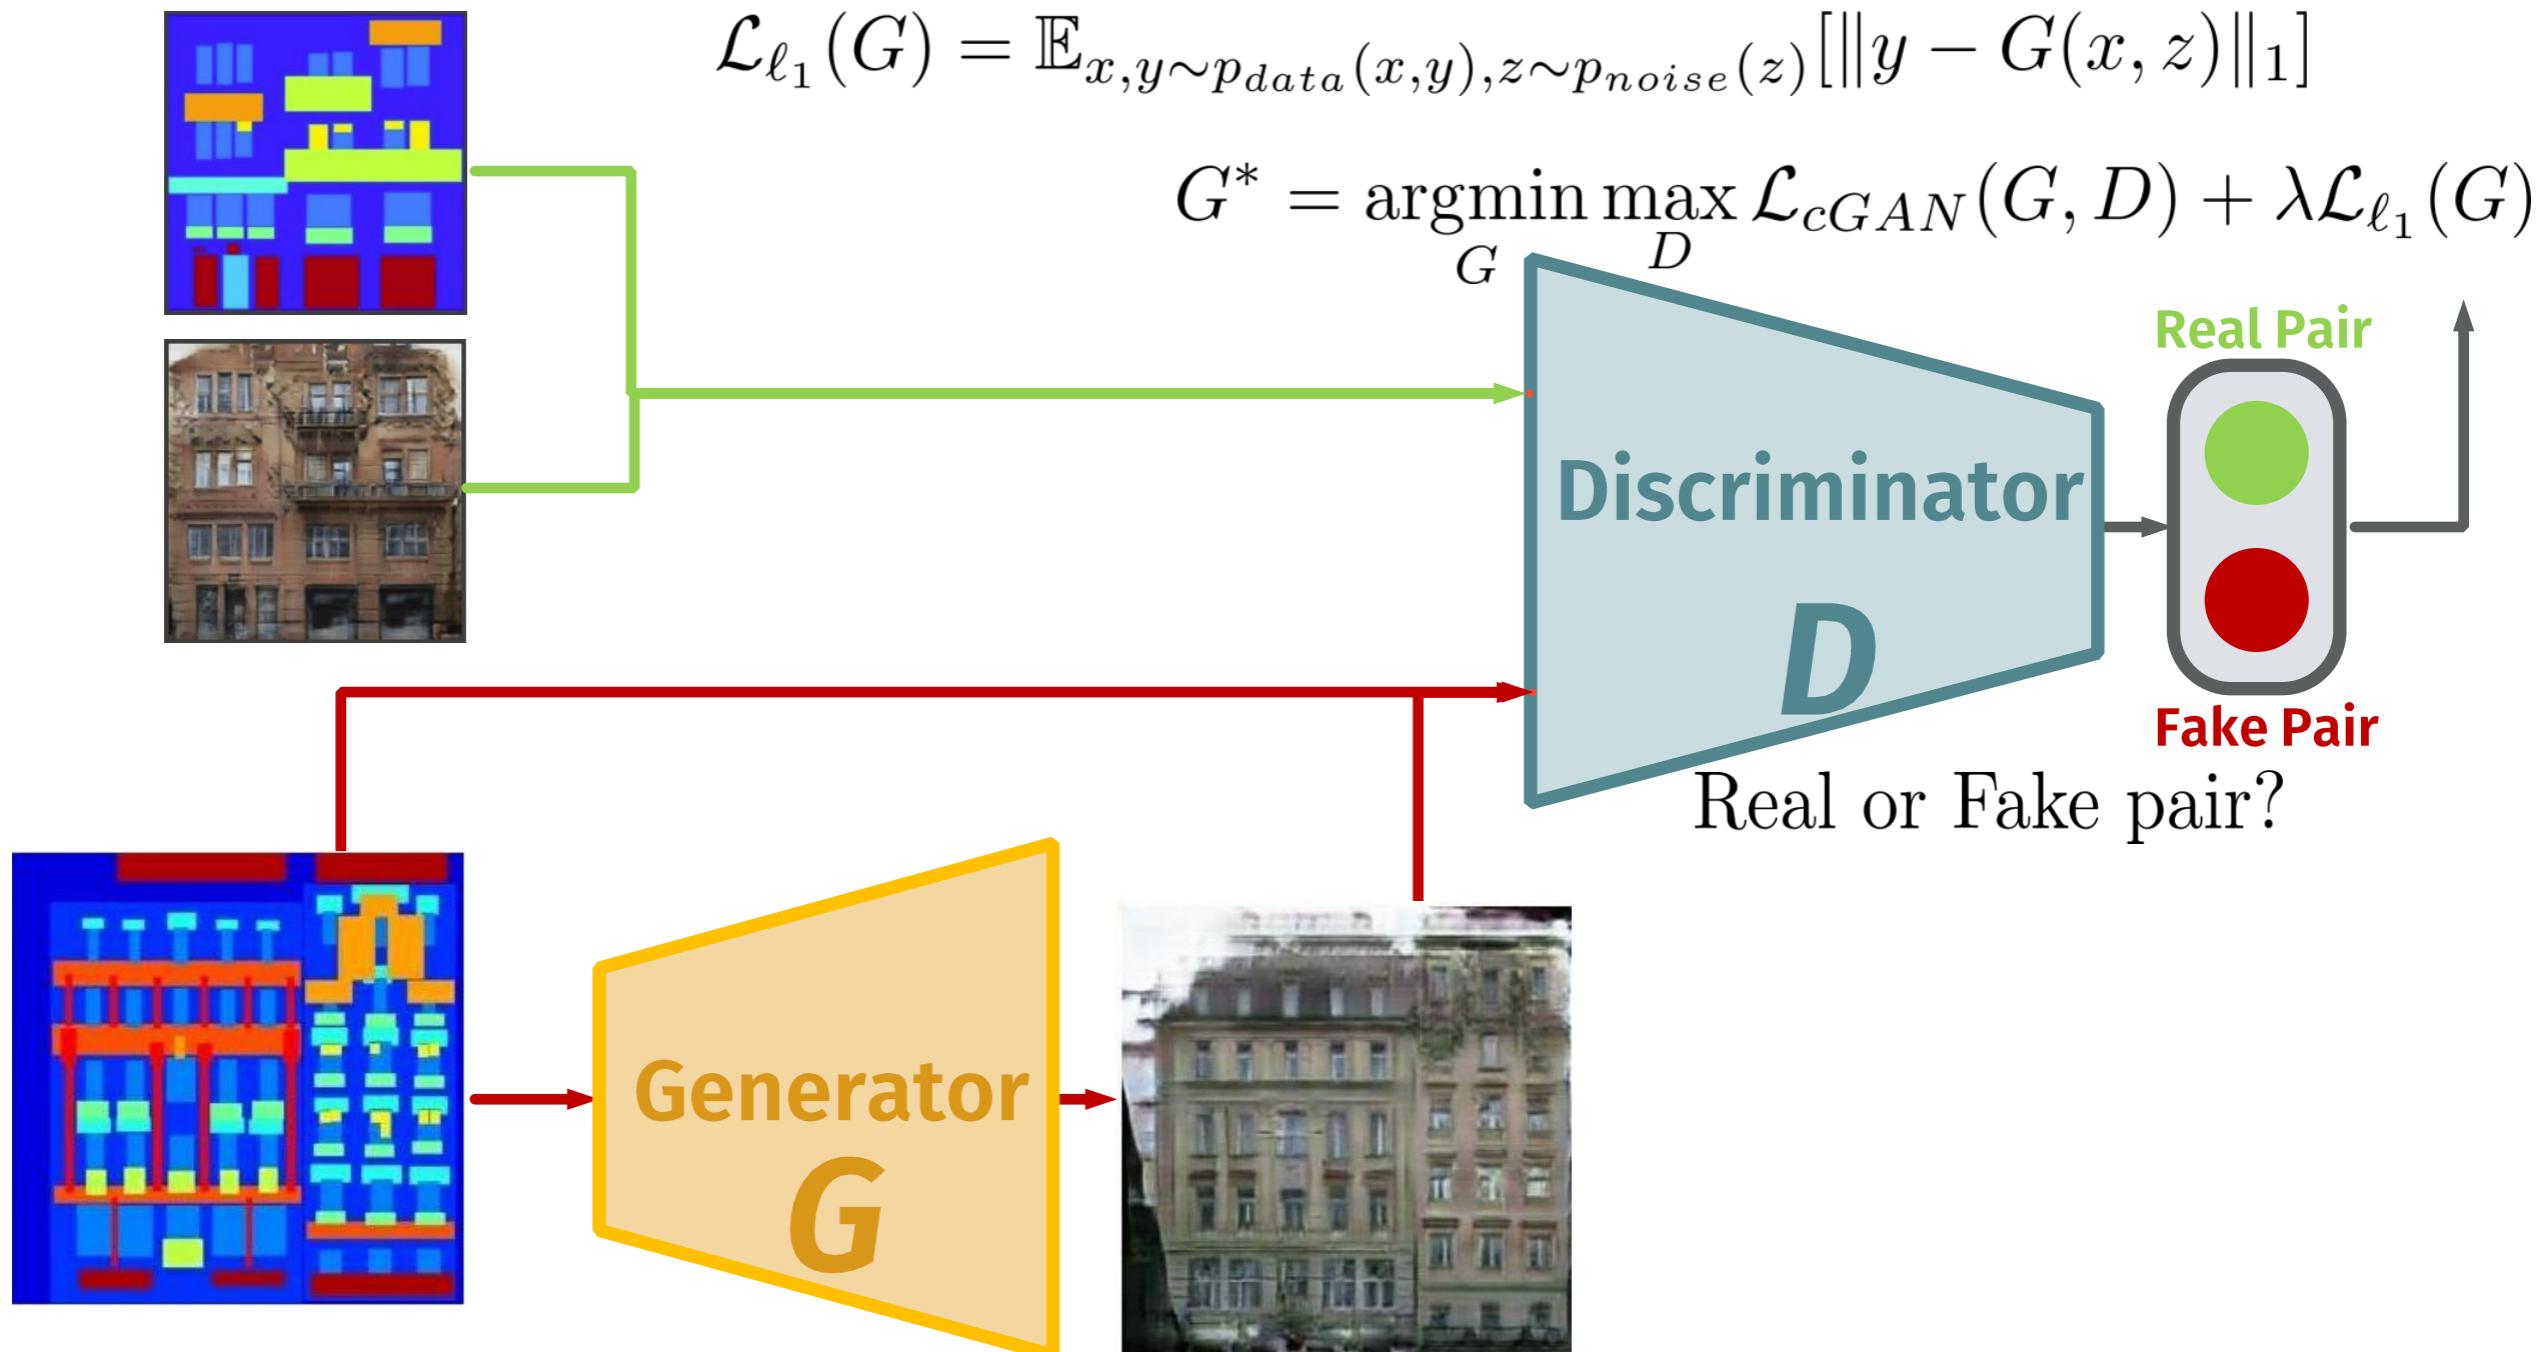
Generated fake samples



Conditional GANs: pix2pix [Isola2017]

Sketch-real image
aligned pairs

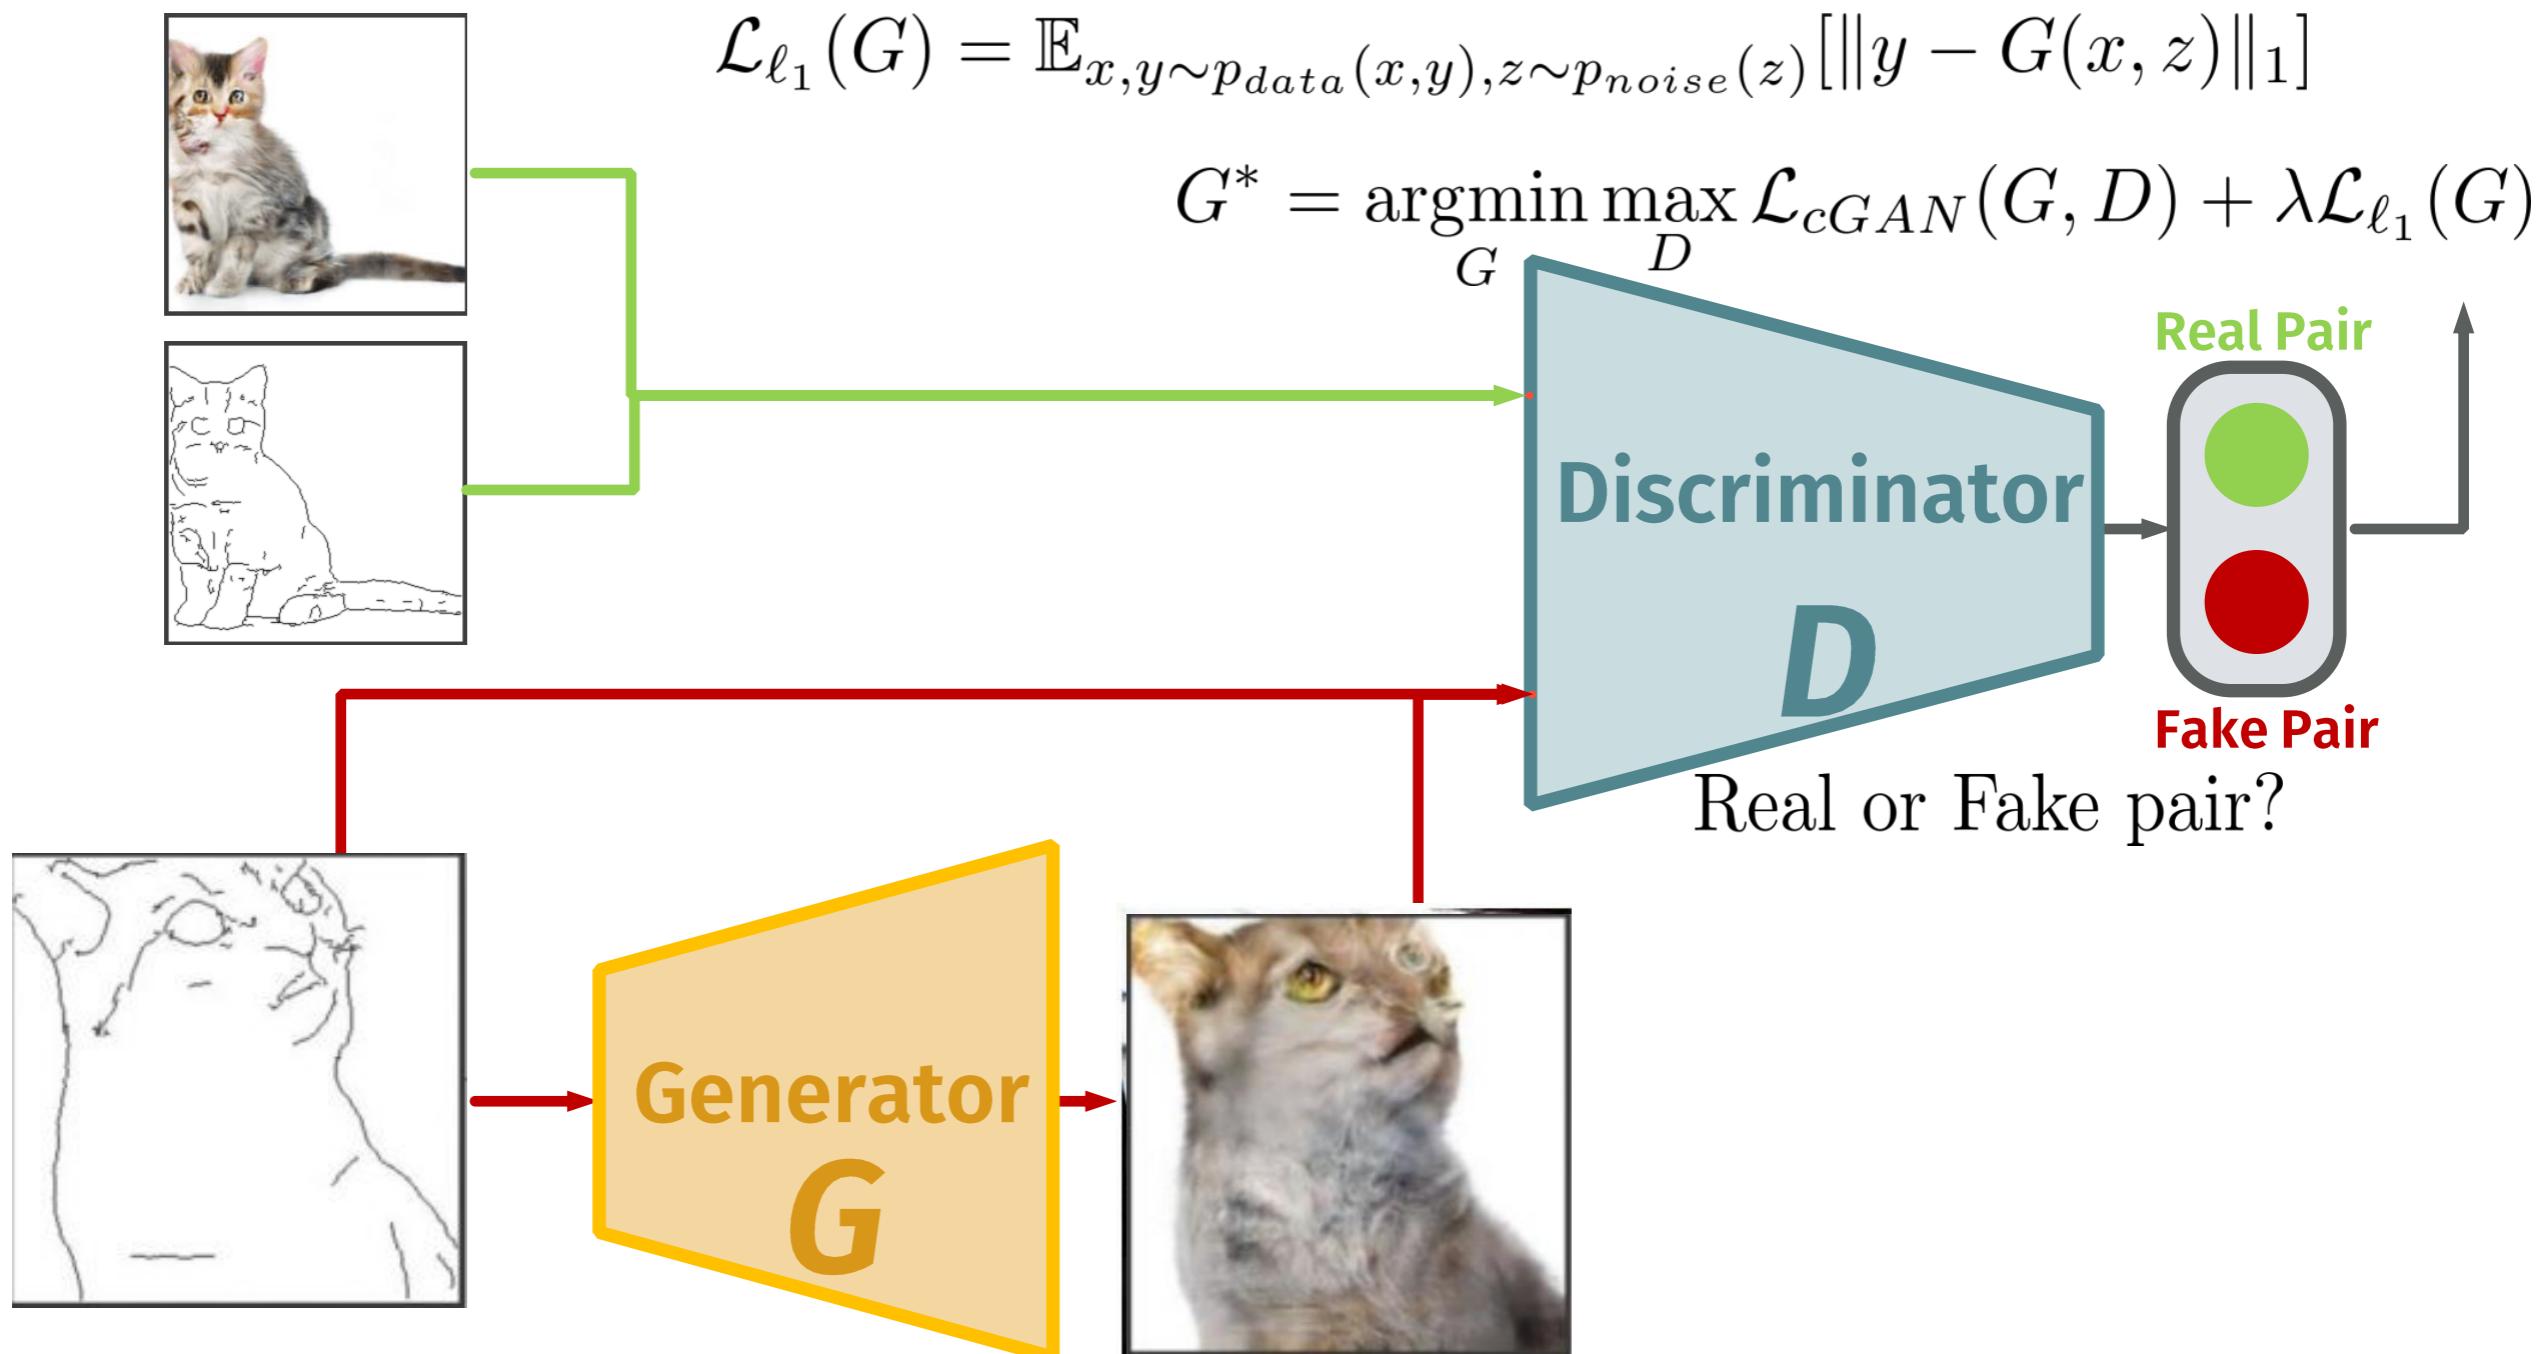
$$\mathcal{L}_{cGAN}(G, D) = \mathbb{E}_{x, y \sim p_{data}(x, y)} [\log D(x, y)] + \mathbb{E}_{x \sim p_{data}(x), z \sim p_{noise}(z)} [\log(1 - D(x, G(x, z)))]$$



Conditional GANs: pix2pix [Isola2017]

Sketch-real image
aligned pairs

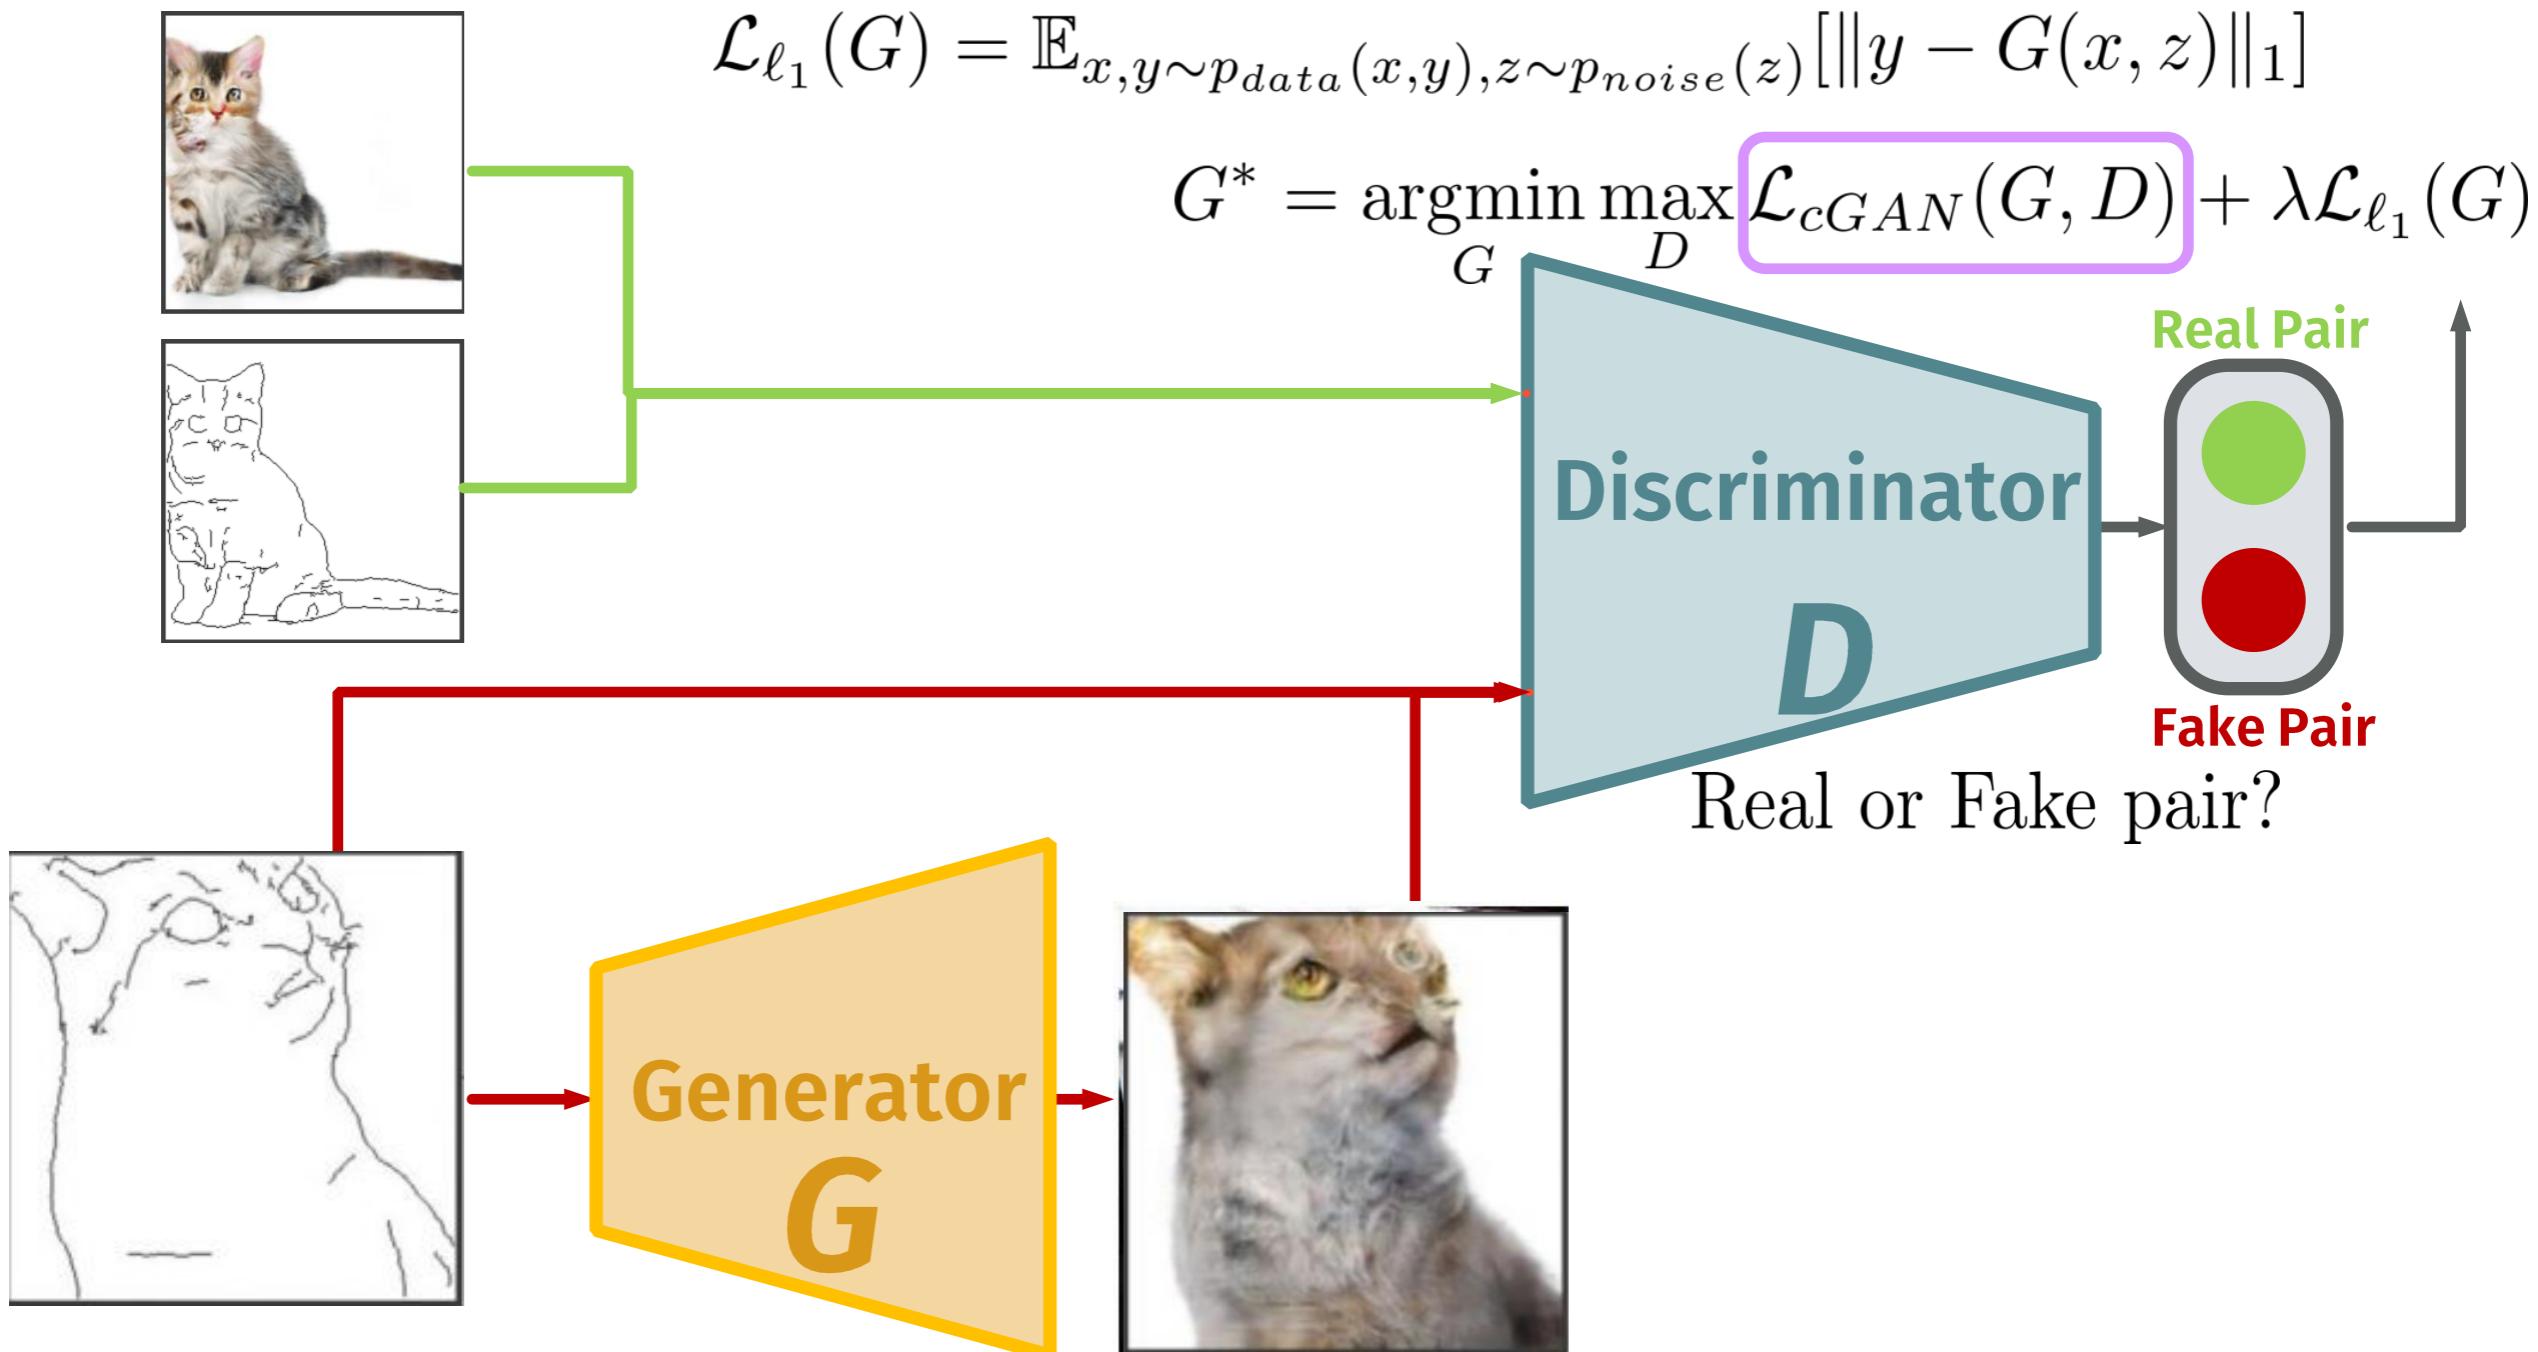
$$\mathcal{L}_{cGAN}(G, D) = \mathbb{E}_{x, y \sim p_{data}(x, y)} [\log D(x, y)] + \mathbb{E}_{x \sim p_{data}(x), z \sim p_{noise}(z)} [\log(1 - D(x, G(x, z)))]$$



Conditional GANs: pix2pix [Isola2017]

Sketch-real image aligned pairs

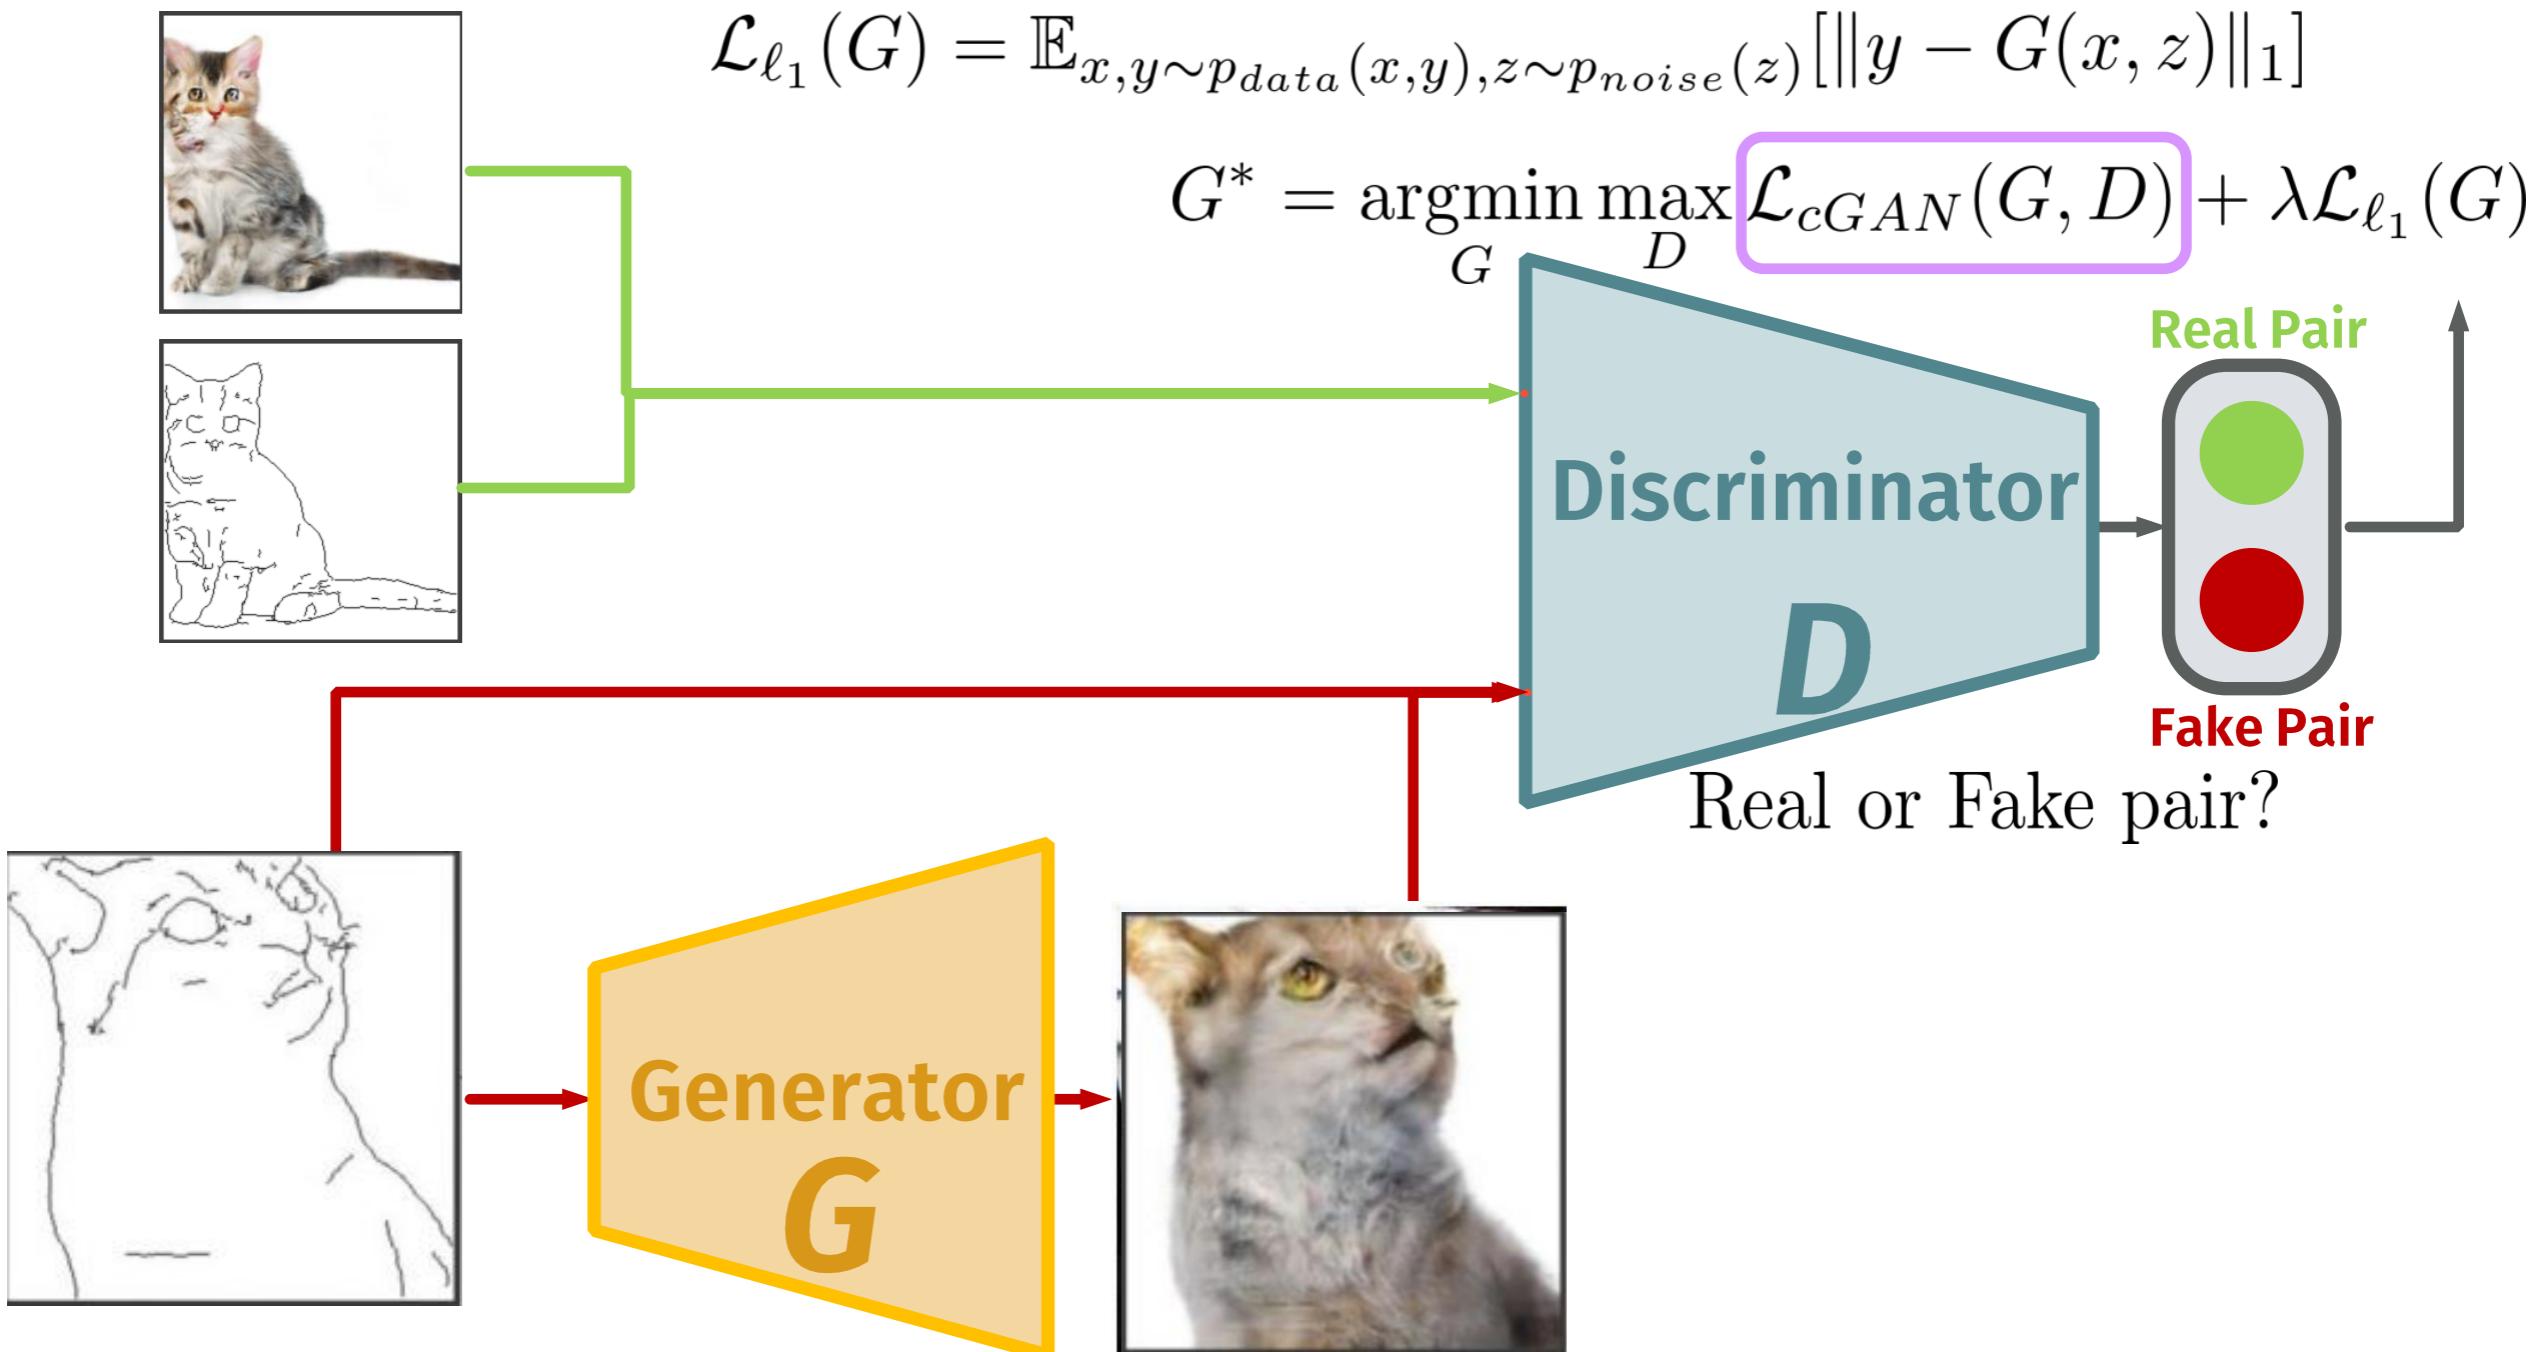
$$\mathcal{L}_{cGAN}(G, D) = \mathbb{E}_{x, y \sim p_{data}(x, y)} [\log D(x, y)] + \mathbb{E}_{x \sim p_{data}(x), z \sim p_{noise}(z)} [\log(1 - D(x, G(x, z)))]$$



Conditional GANs: pix2pix [Isola2017]

Sketch-real image aligned pairs

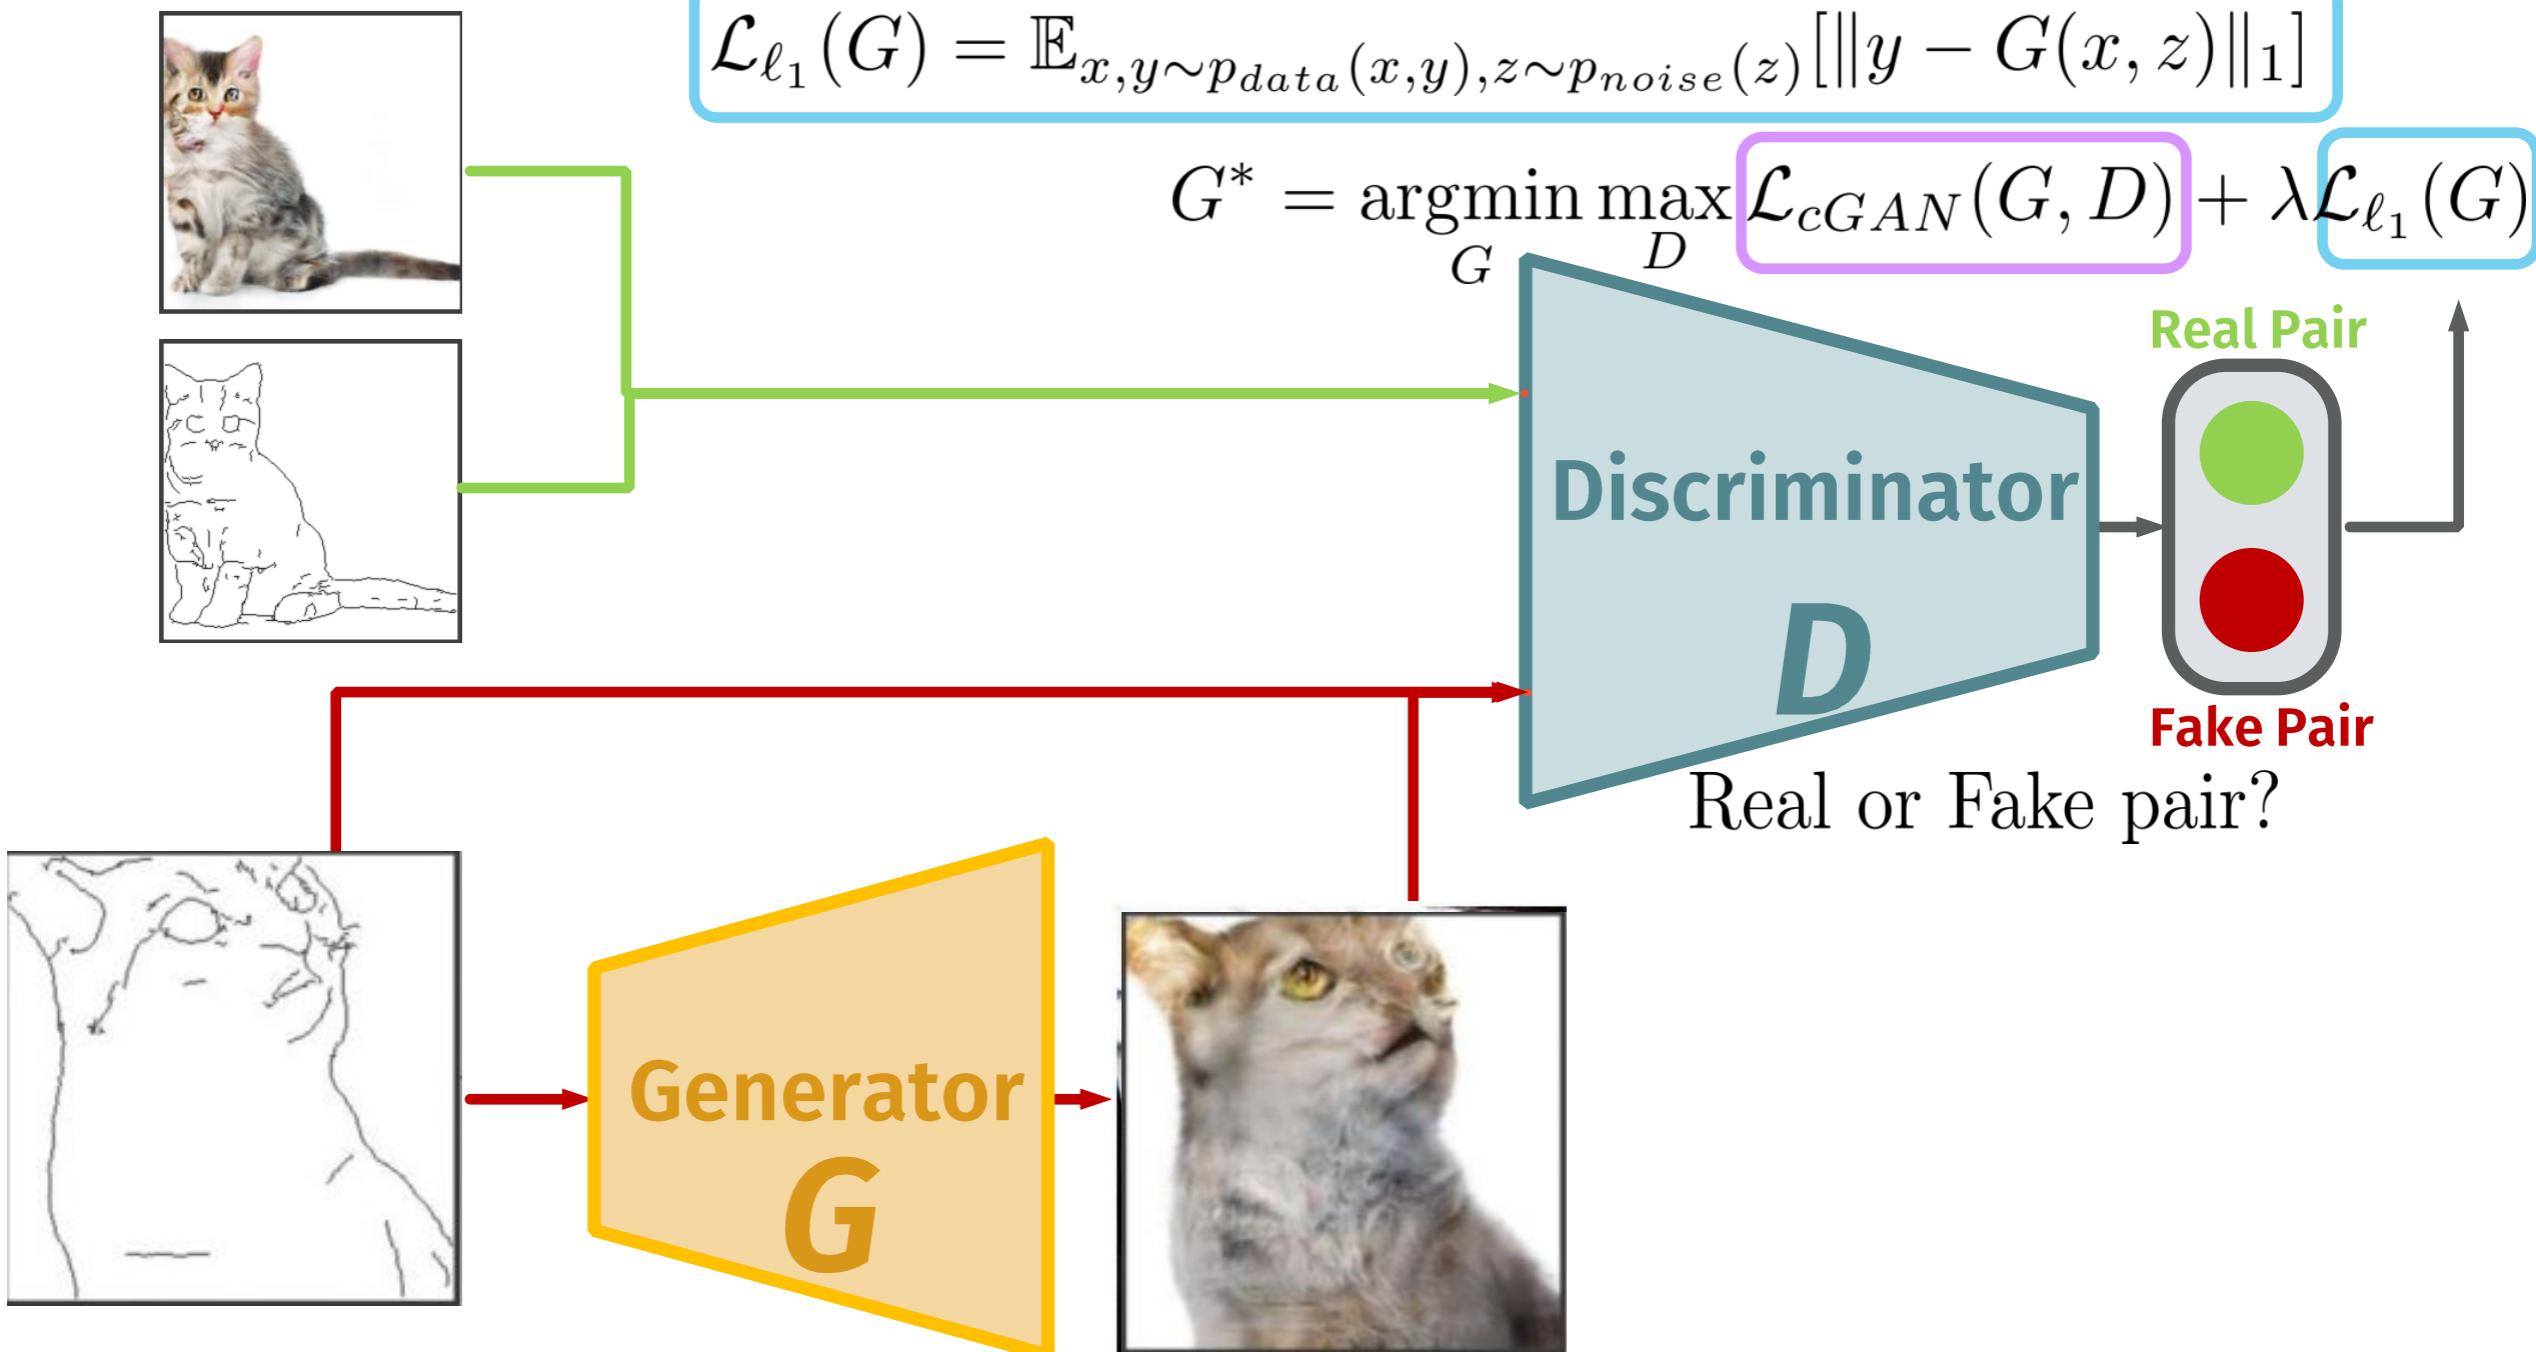
$$\mathcal{L}_{cGAN}(G, D) = \mathbb{E}_{x,y \sim p_{data}(x,y)} [\log D(x,y)] + \mathbb{E}_{x \sim p_{data}(x), z \sim p_{noise}(z)} [\log(1 - D(x, G(x, z)))]$$



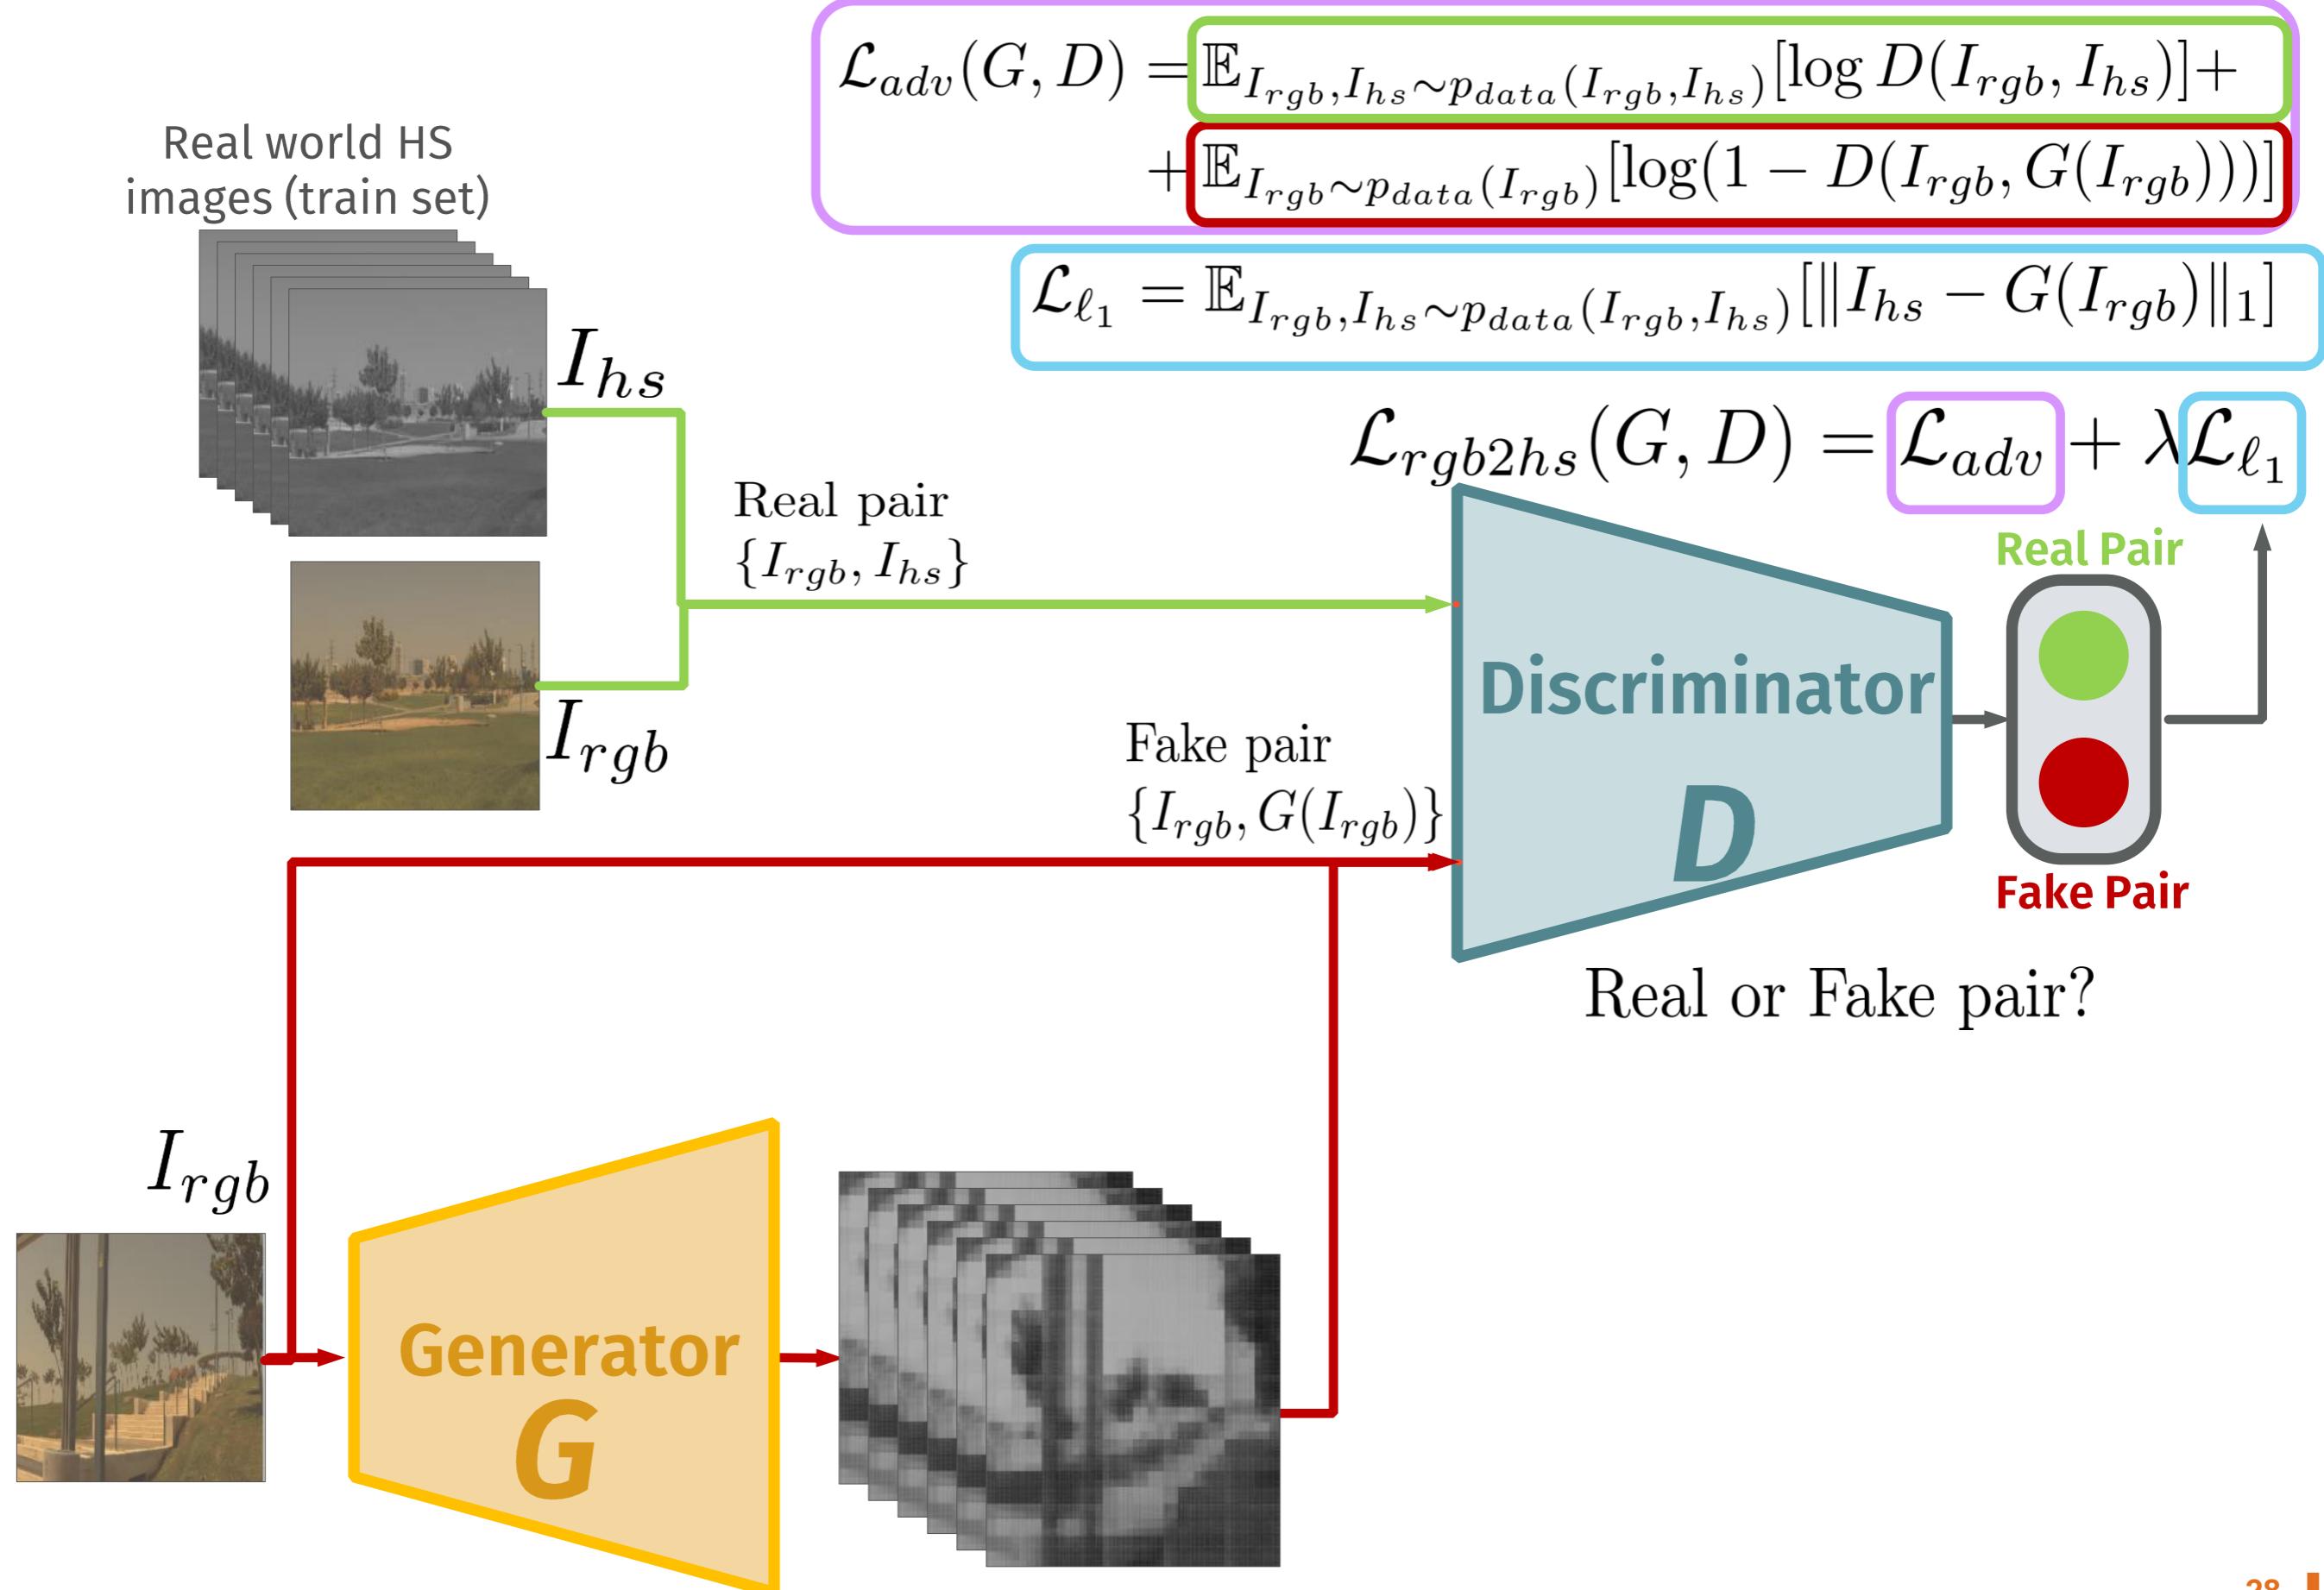
Conditional GANs: pix2pix [Isola2017]

Sketch-real image aligned pairs

$$\mathcal{L}_{cGAN}(G, D) = \mathbb{E}_{x,y \sim p_{data}(x,y)} [\log D(x,y)] + \mathbb{E}_{x \sim p_{data}(x), z \sim p_{noise}(z)} [\log(1 - D(x, G(x, z)))]$$



Adversarial spectral reconstruction networks



Adversarial spectral reconstruction networks

Real world HS
images (train set)

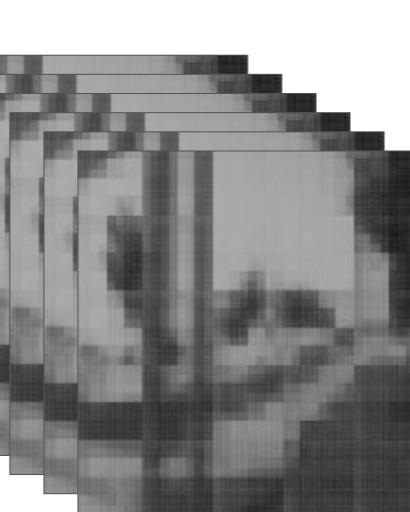


I_{hs}

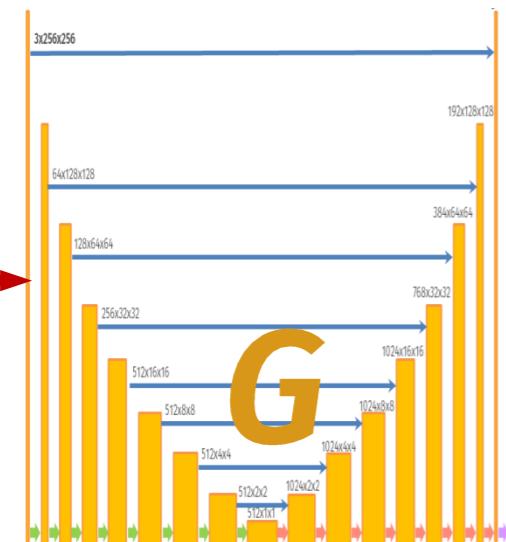


I_{rgb}

Real pair
 $\{I_{rgb}, I_{hs}\}$



I_{rgb}

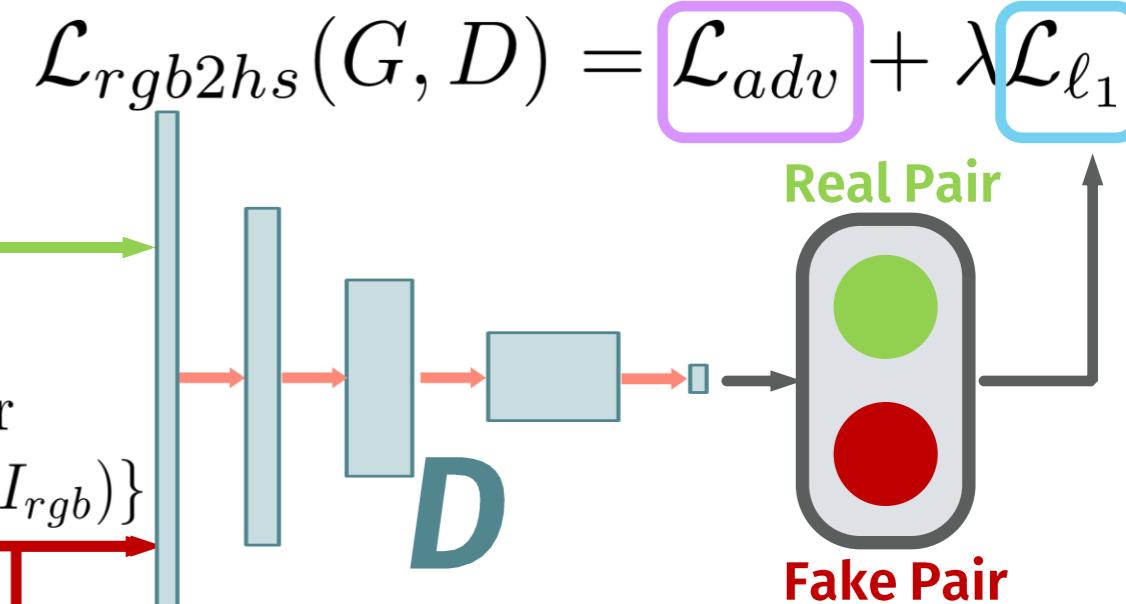


$$\mathcal{L}_{adv}(G, D) = \mathbb{E}_{I_{rgb}, I_{hs} \sim p_{data}(I_{rgb}, I_{hs})} [\log D(I_{rgb}, I_{hs})] + \\ + \mathbb{E}_{I_{rgb} \sim p_{data}(I_{rgb})} [\log(1 - D(I_{rgb}, G(I_{rgb})))]$$

$$\mathcal{L}_{\ell_1} = \mathbb{E}_{I_{rgb}, I_{hs} \sim p_{data}(I_{rgb}, I_{hs})} [\|I_{hs} - G(I_{rgb})\|_1]$$

$$\mathcal{L}_{rgb2hs}(G, D) = \mathcal{L}_{adv} + \lambda \mathcal{L}_{\ell_1}$$

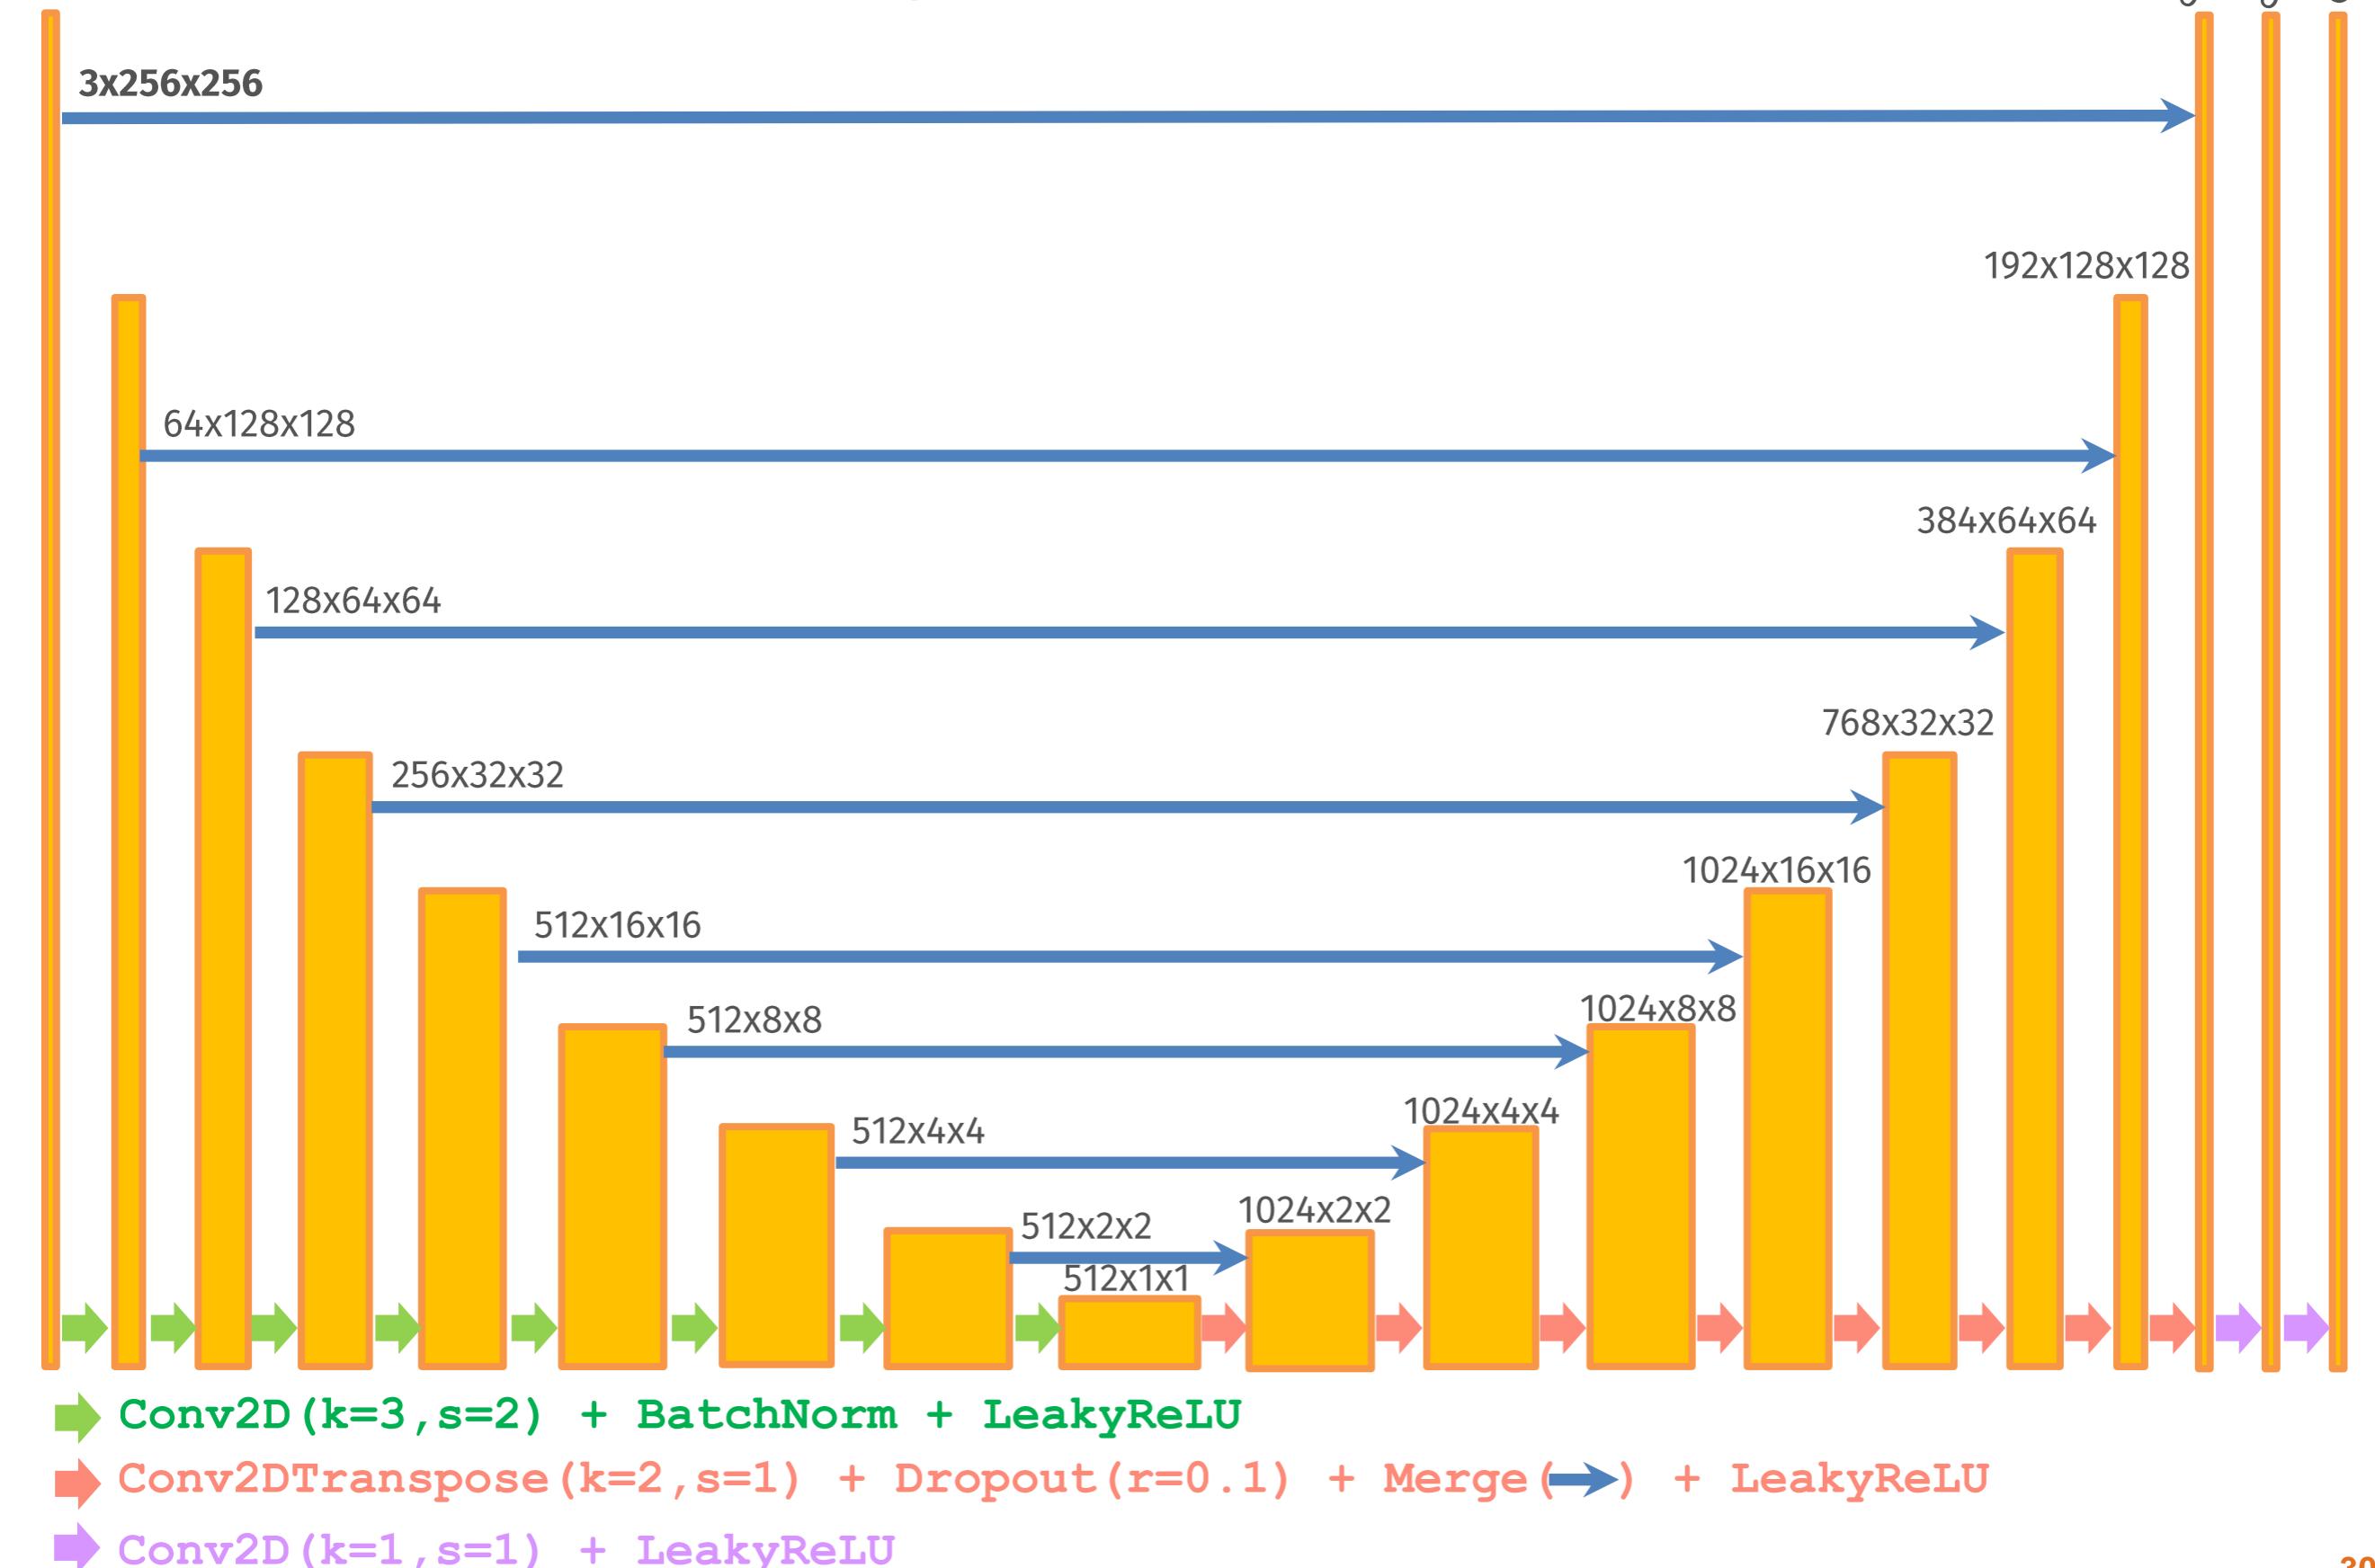
Fake pair
 $\{I_{rgb}, G(I_{rgb})\}$



Real or Fake pair?

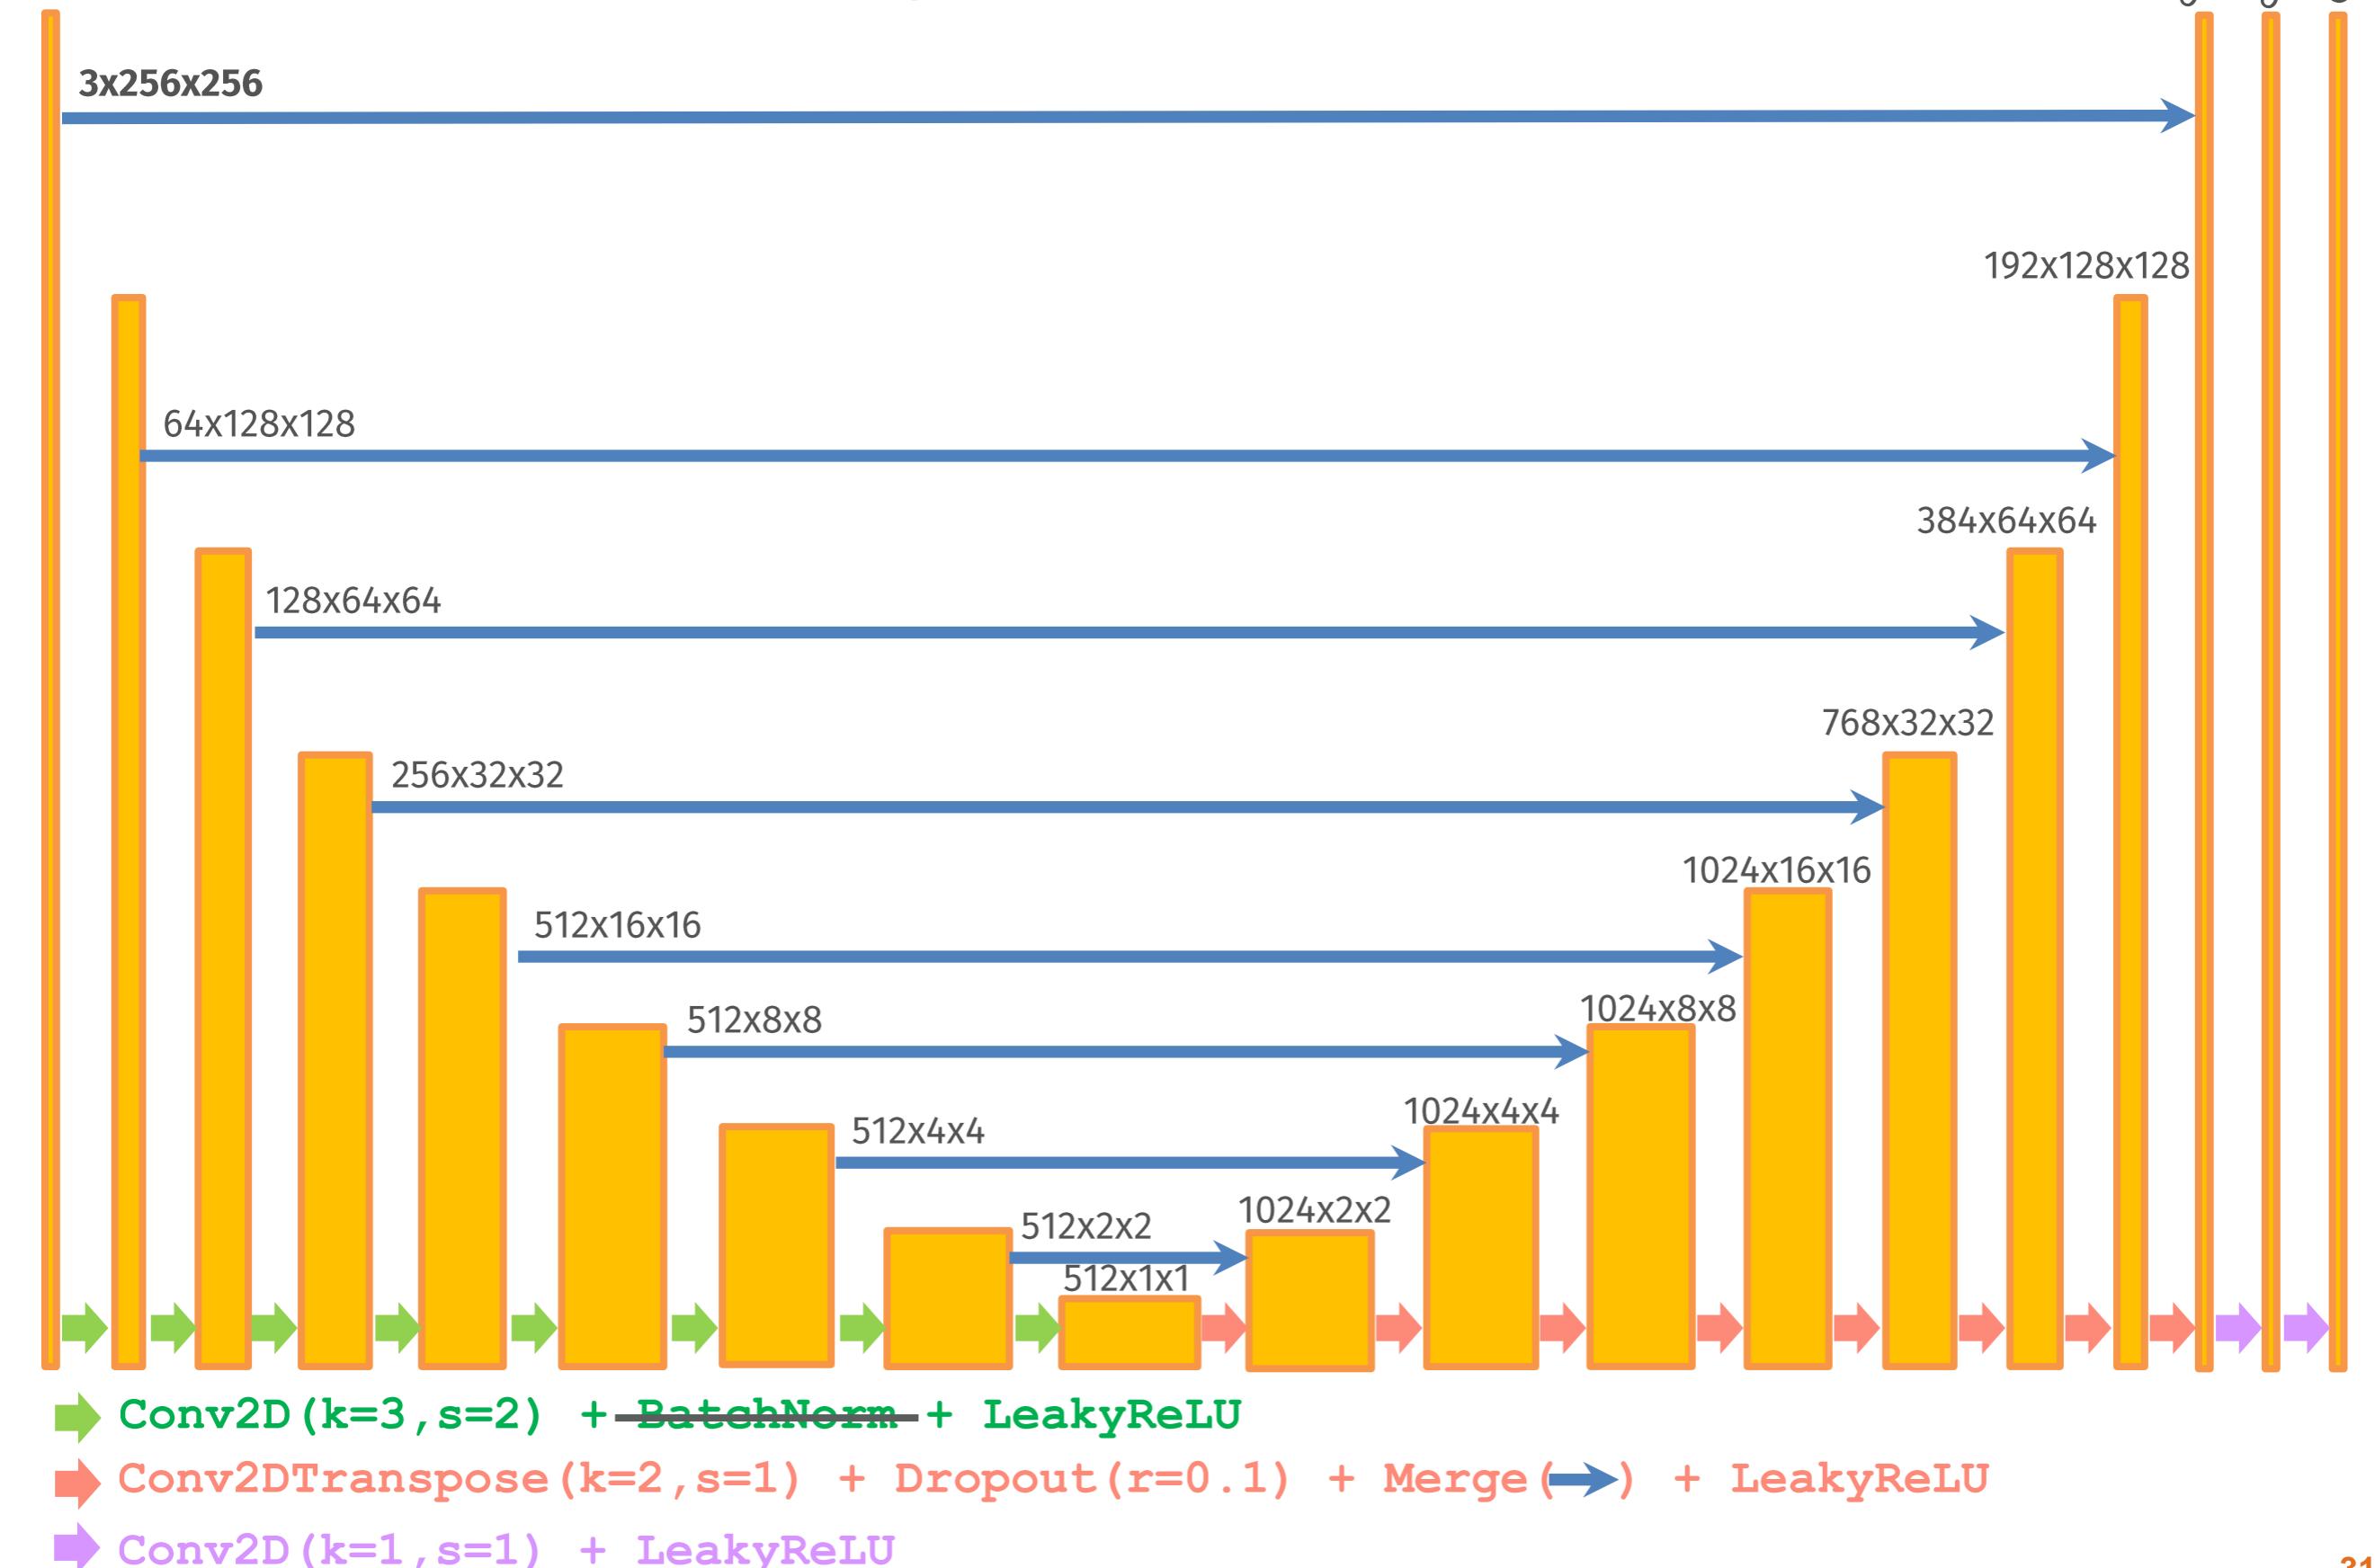
Architecture

G is a U-Net [Ronneberger2015]



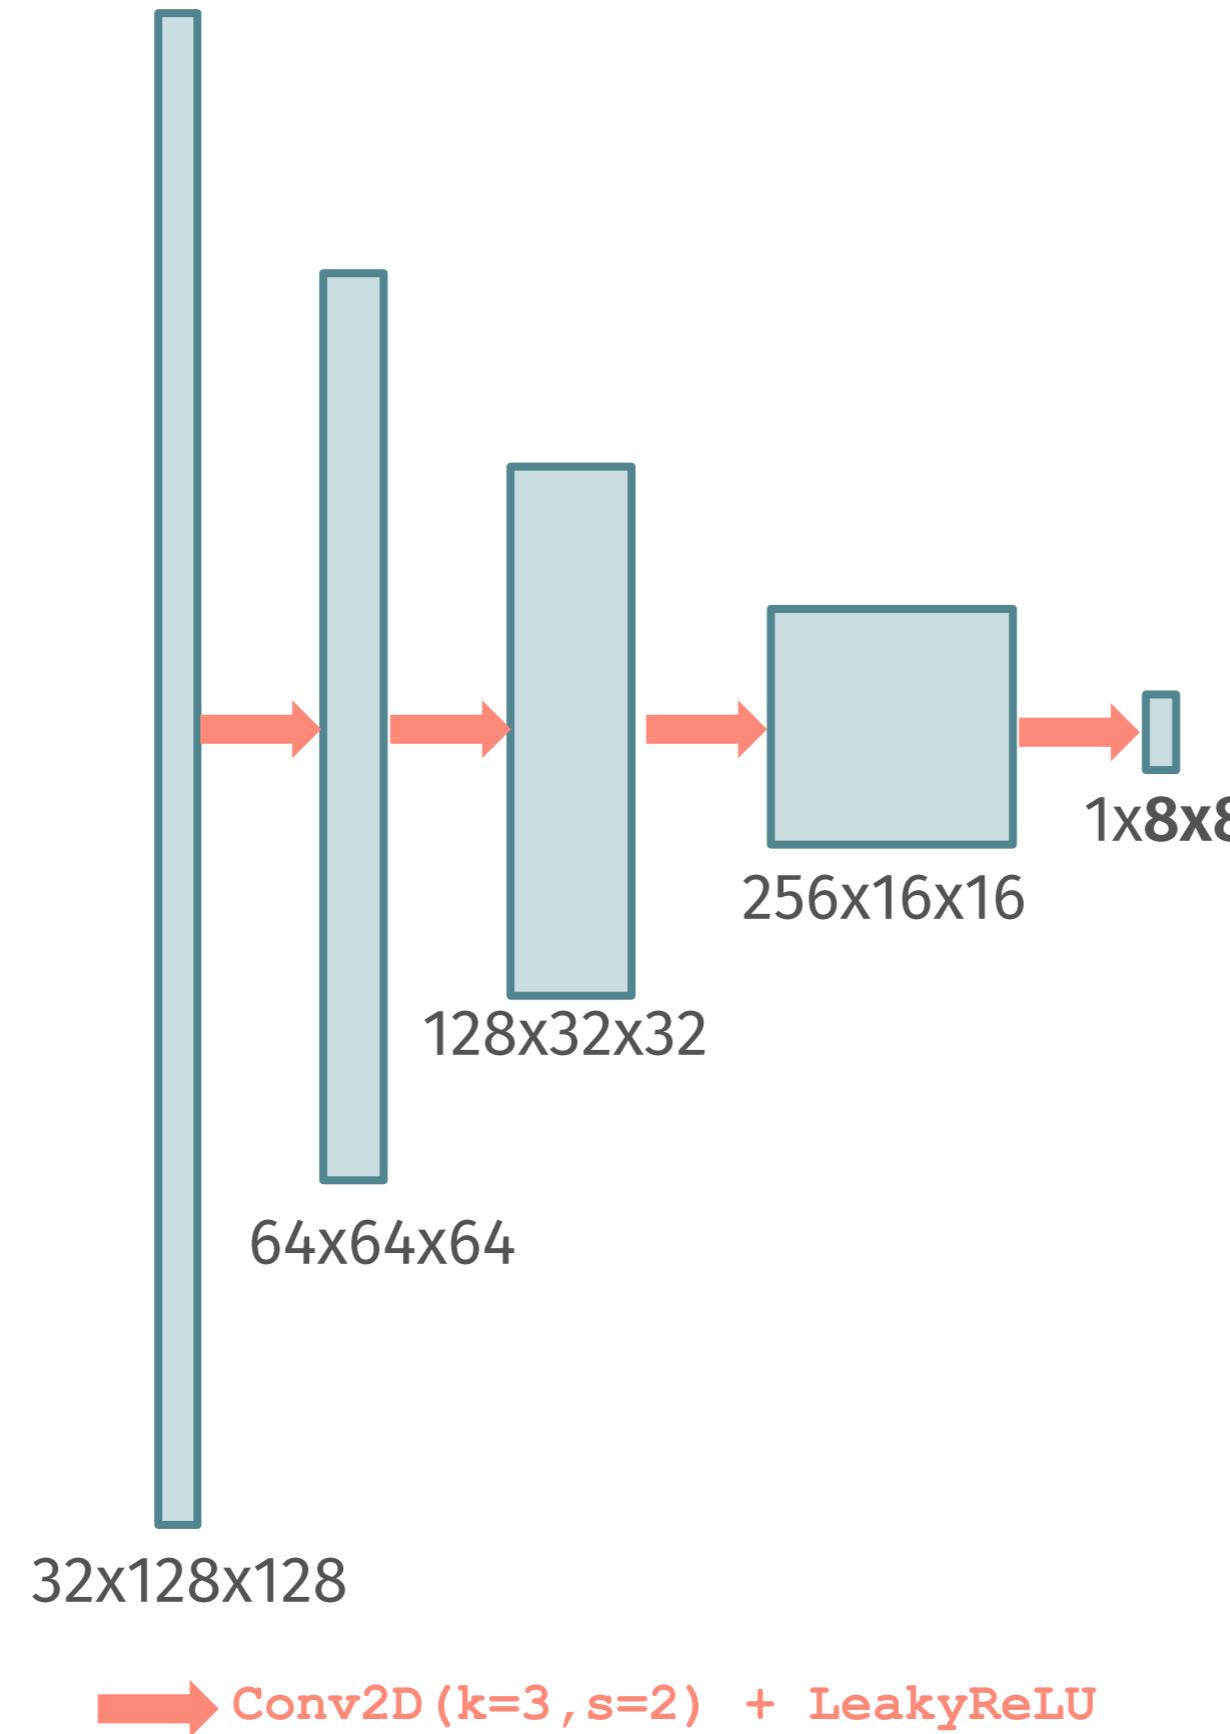
Architecture

G is a U-Net [Ronneberger2015]



Architecture

D is a Patch-CNN

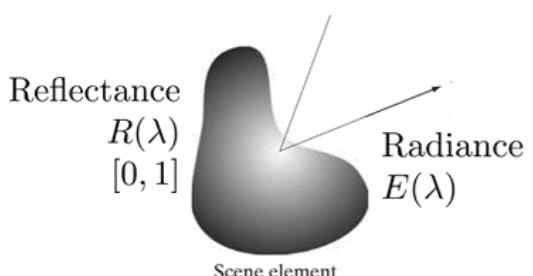


Experiments

Dataset preparation

- 2 fold cross-validation (2x100)
- From hyperspectral to sRGB:

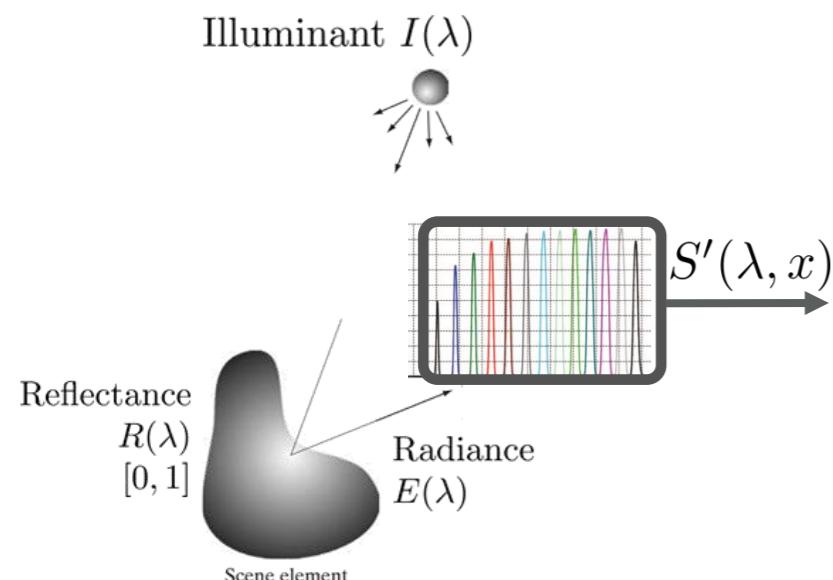
Illuminant $I(\lambda)$



Experiments

Dataset preparation

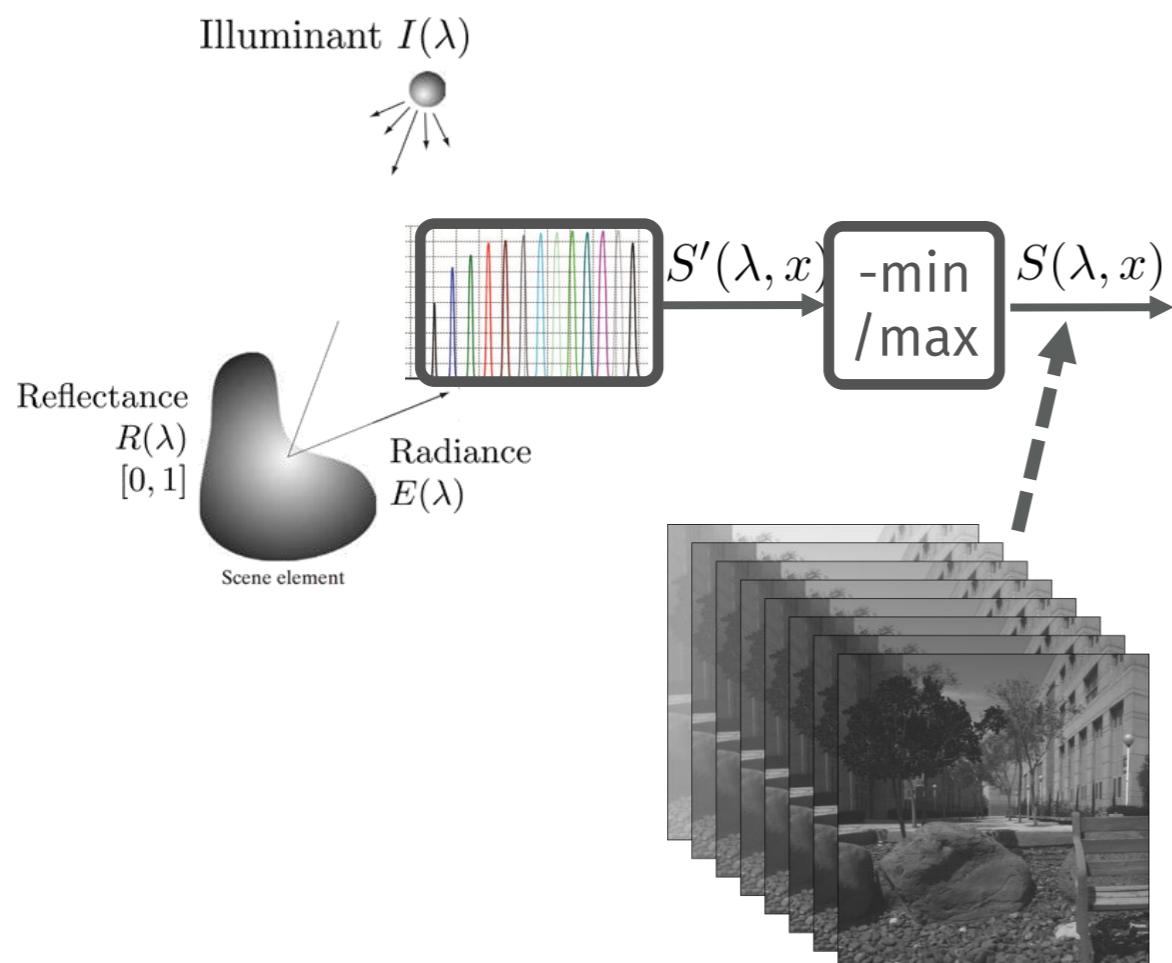
- 2 fold cross-validation (2x100)
- From hyperspectral to sRGB:



Experiments

Dataset preparation

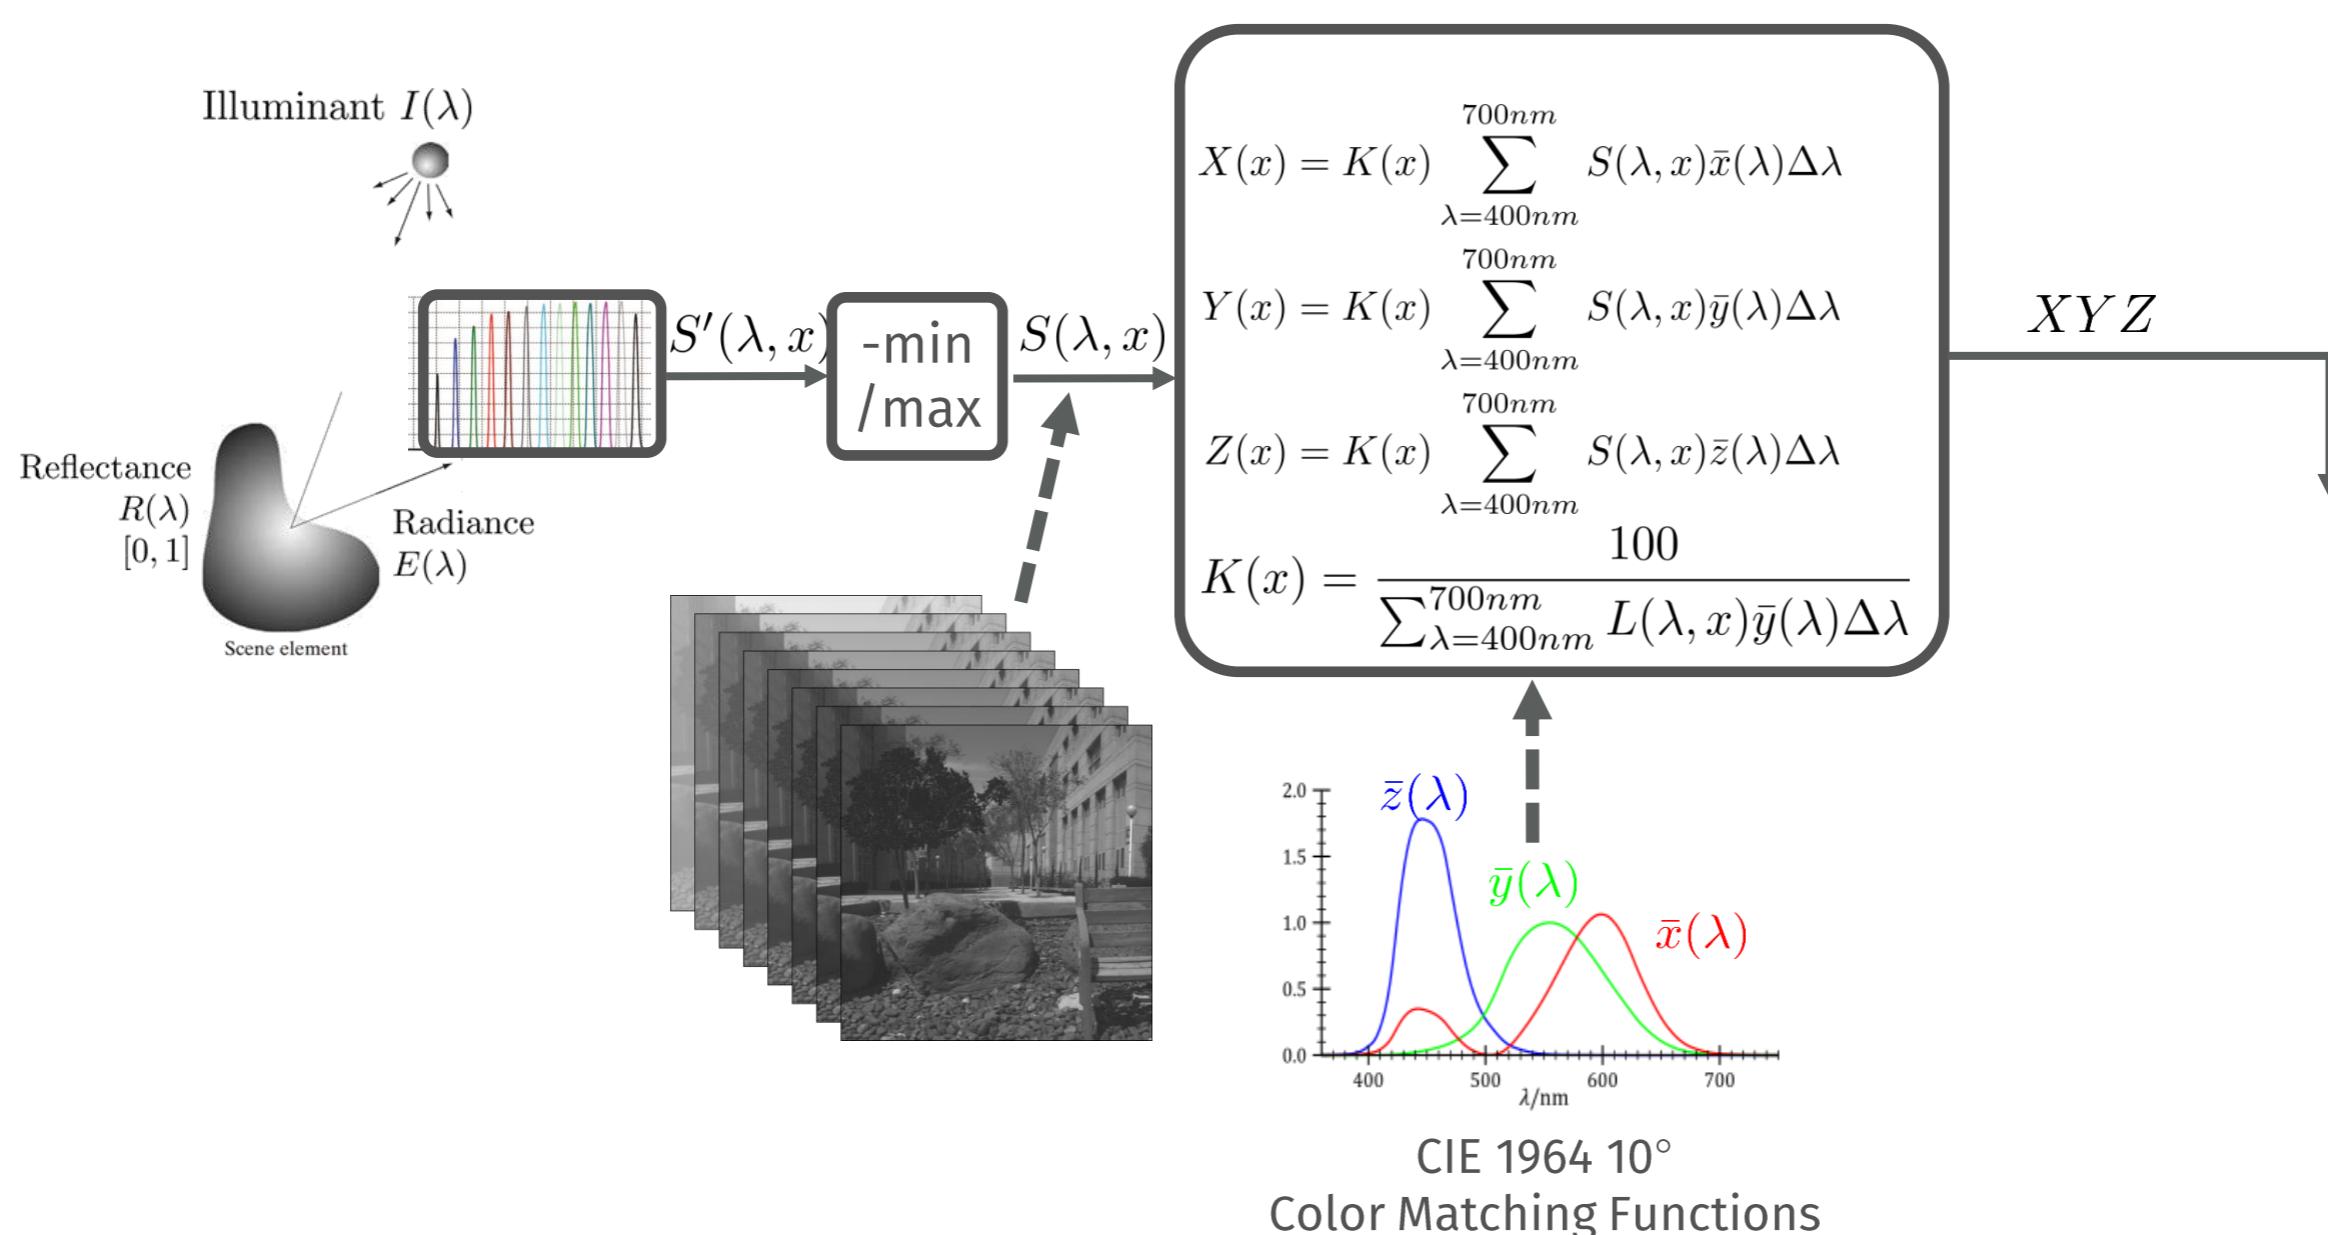
- 2 fold cross-validation (2x100)
- From hyperspectral to sRGB:



Experiments

Dataset preparation

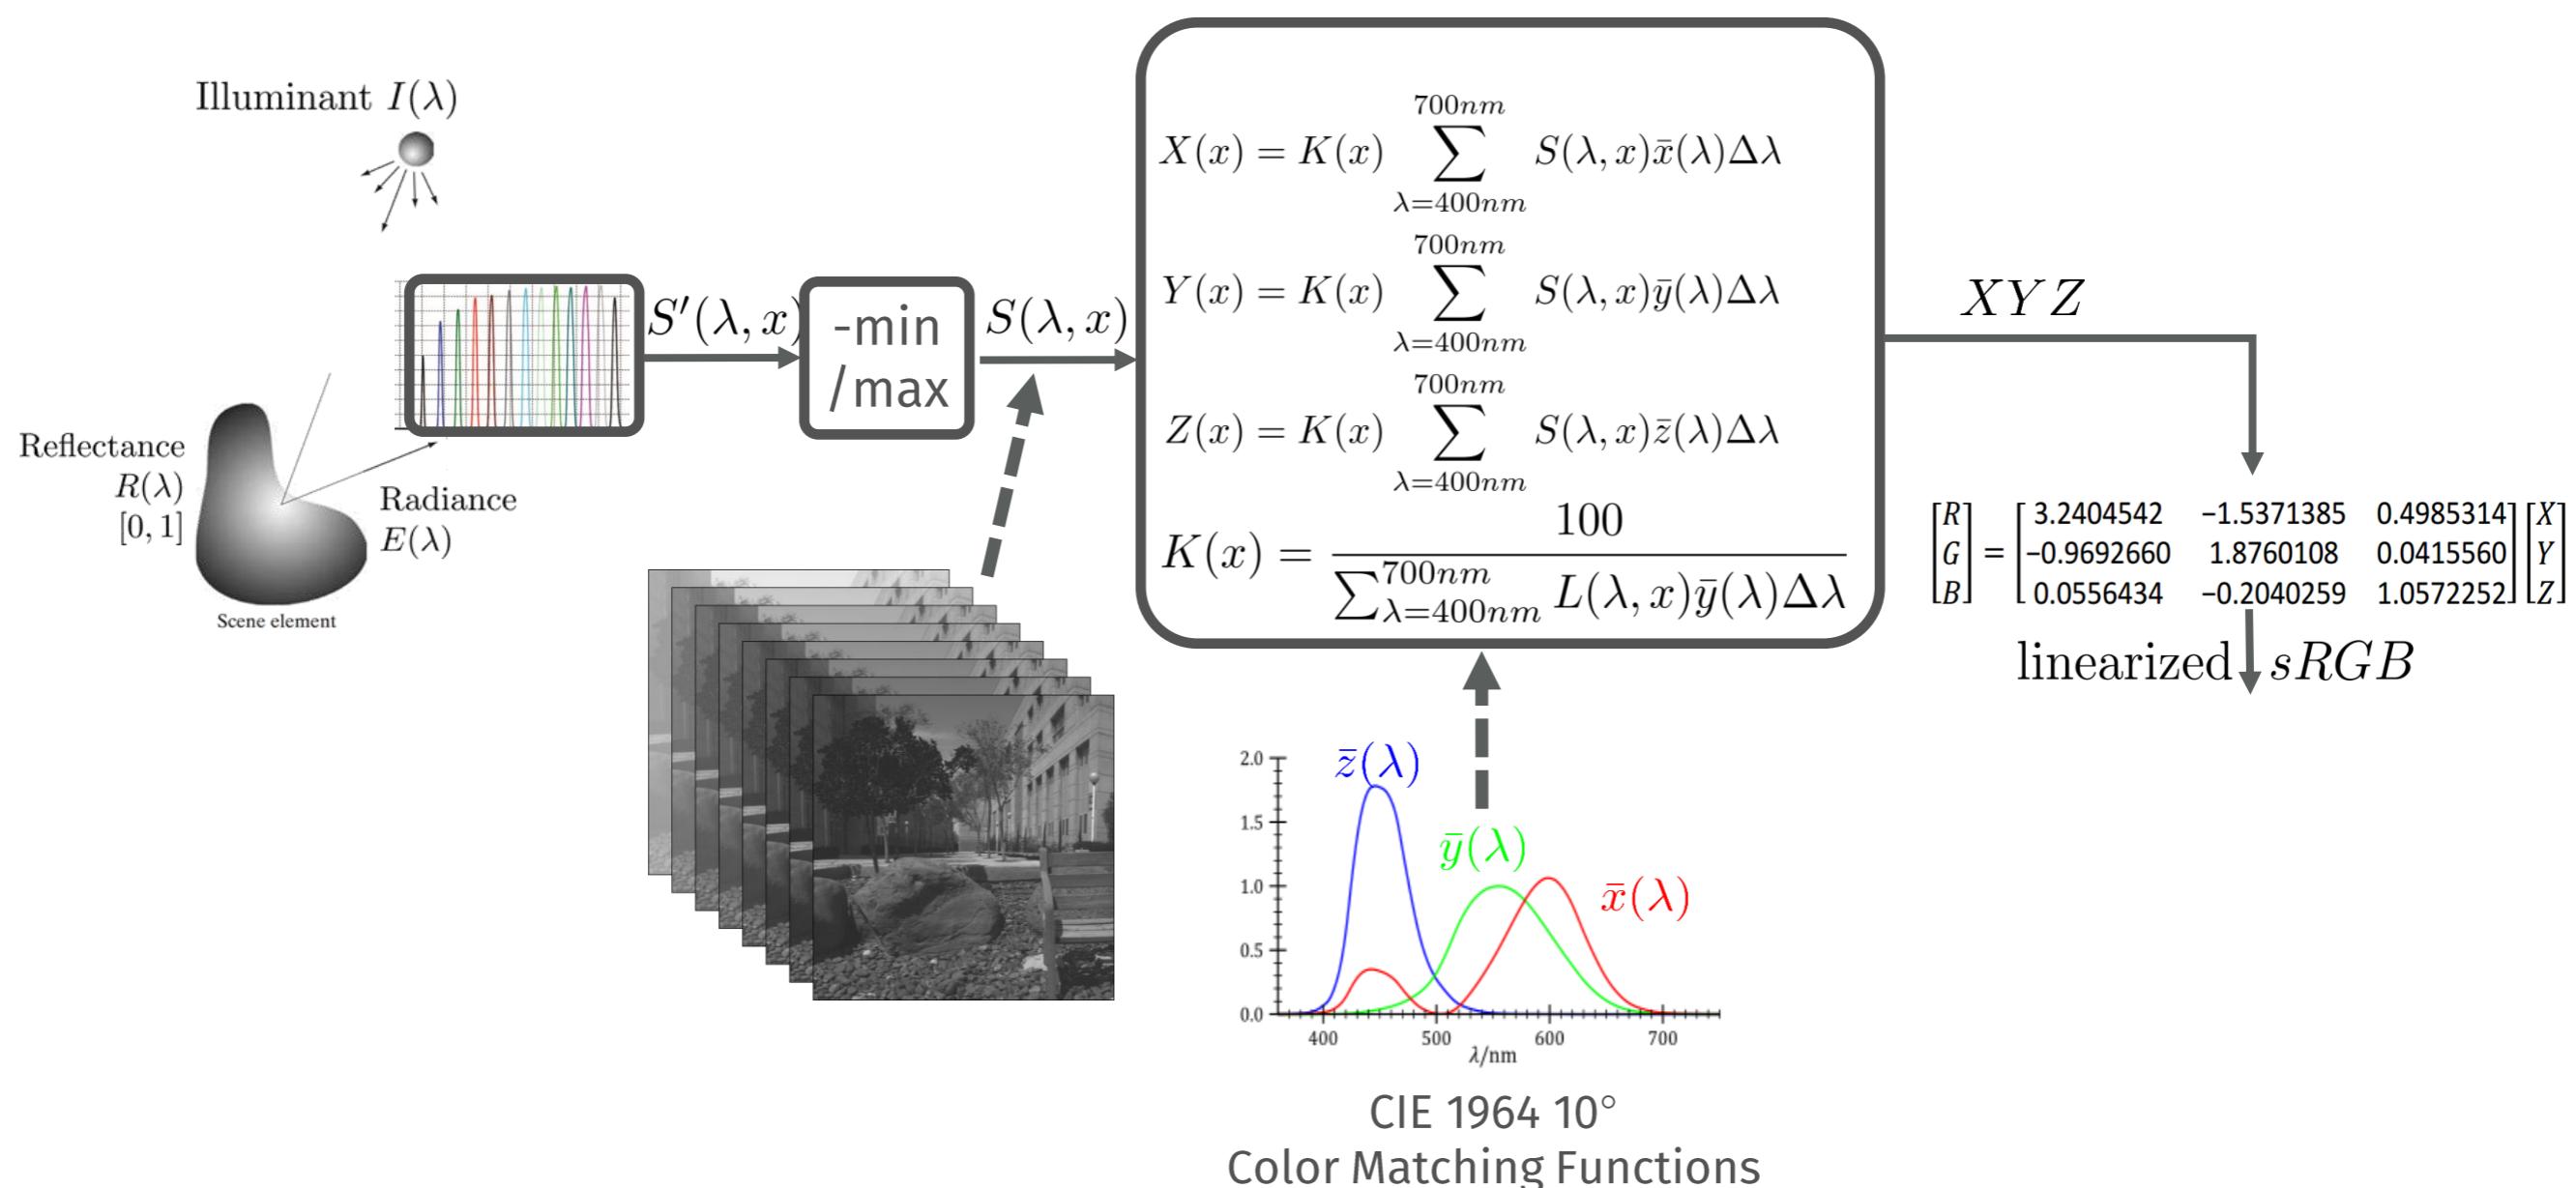
- 2 fold cross-validation (2x100)
- From hyperspectral to sRGB:



Experiments

Dataset preparation

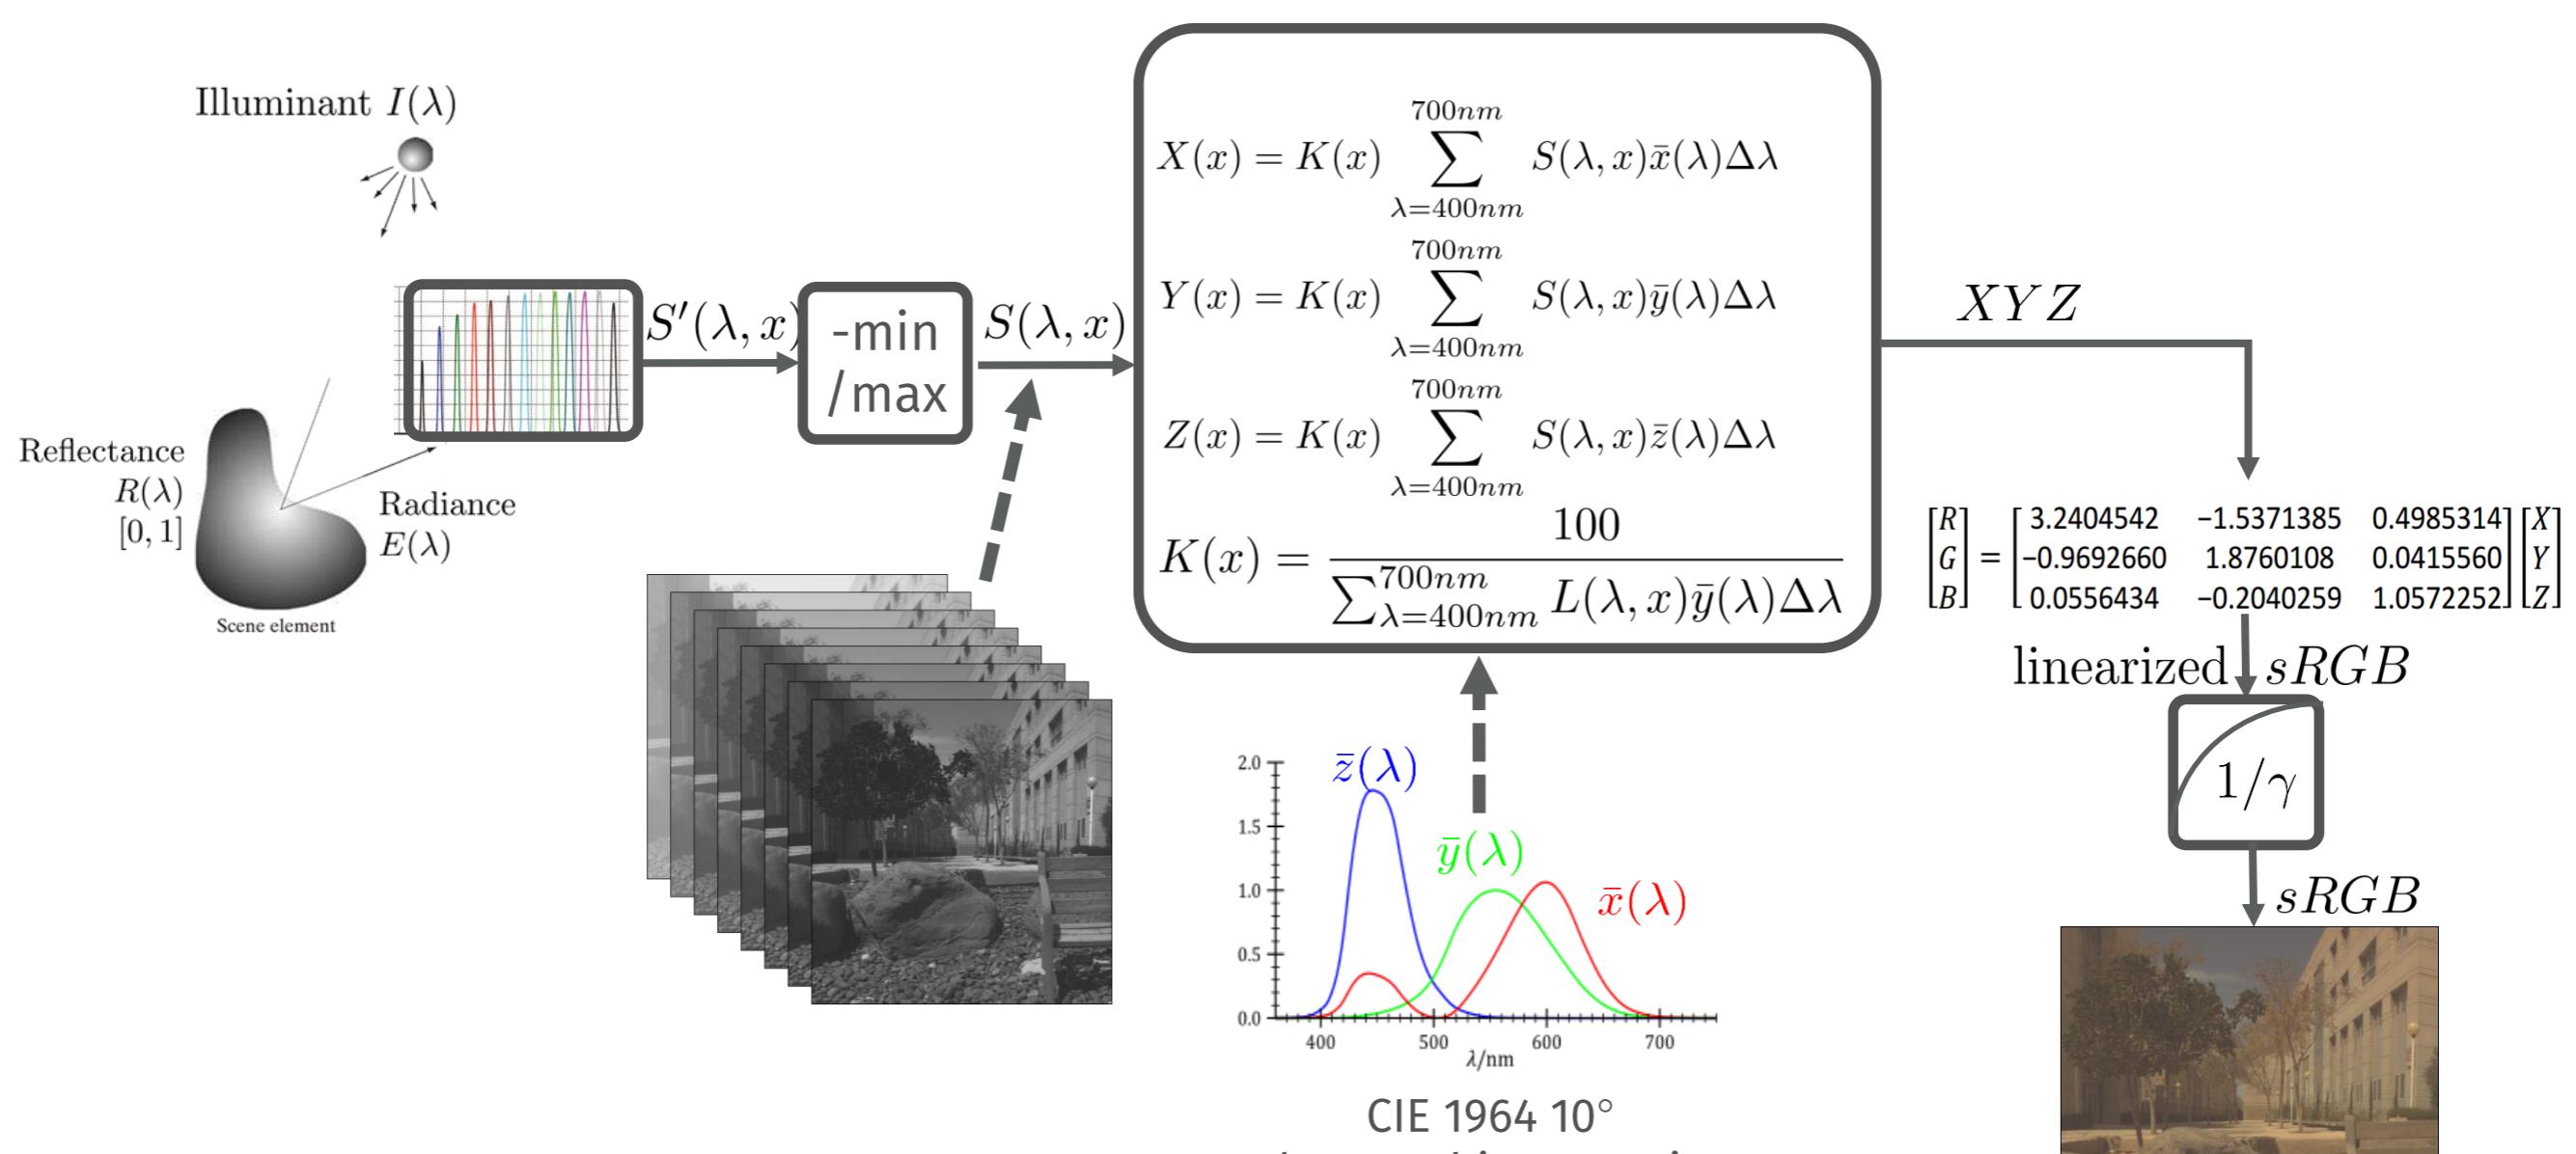
- 2 fold cross-validation (2x100)
- From hyperspectral to sRGB:



Experiments

Dataset preparation

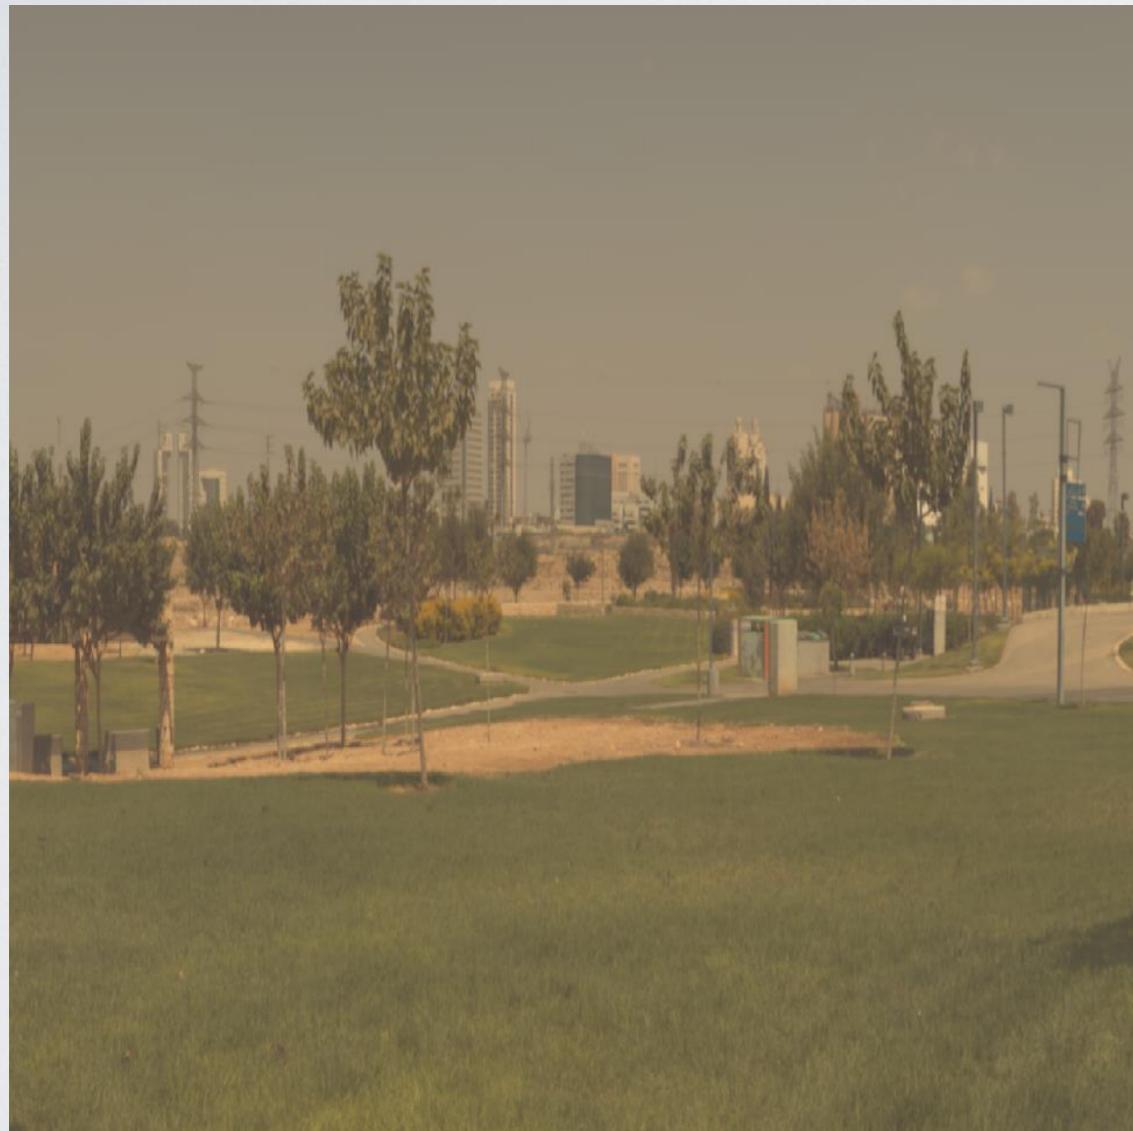
- 2 fold cross-validation (2x100)
- From hyperspectral to sRGB:



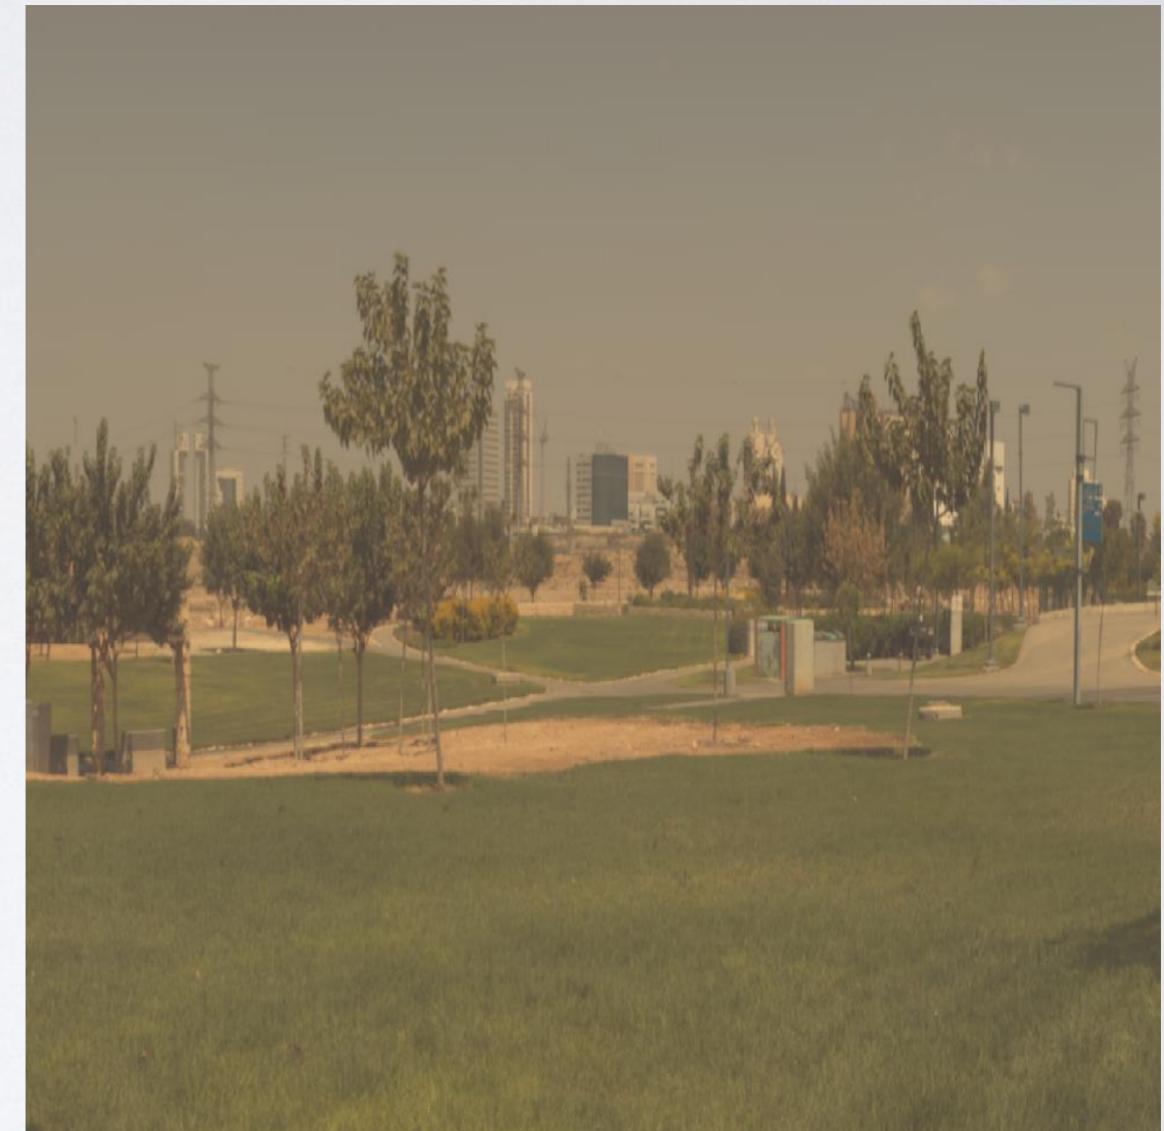
Experiments

Training and testing with large images

- Training: random crops
- Testing: tile decompositon + reconstruction



1280

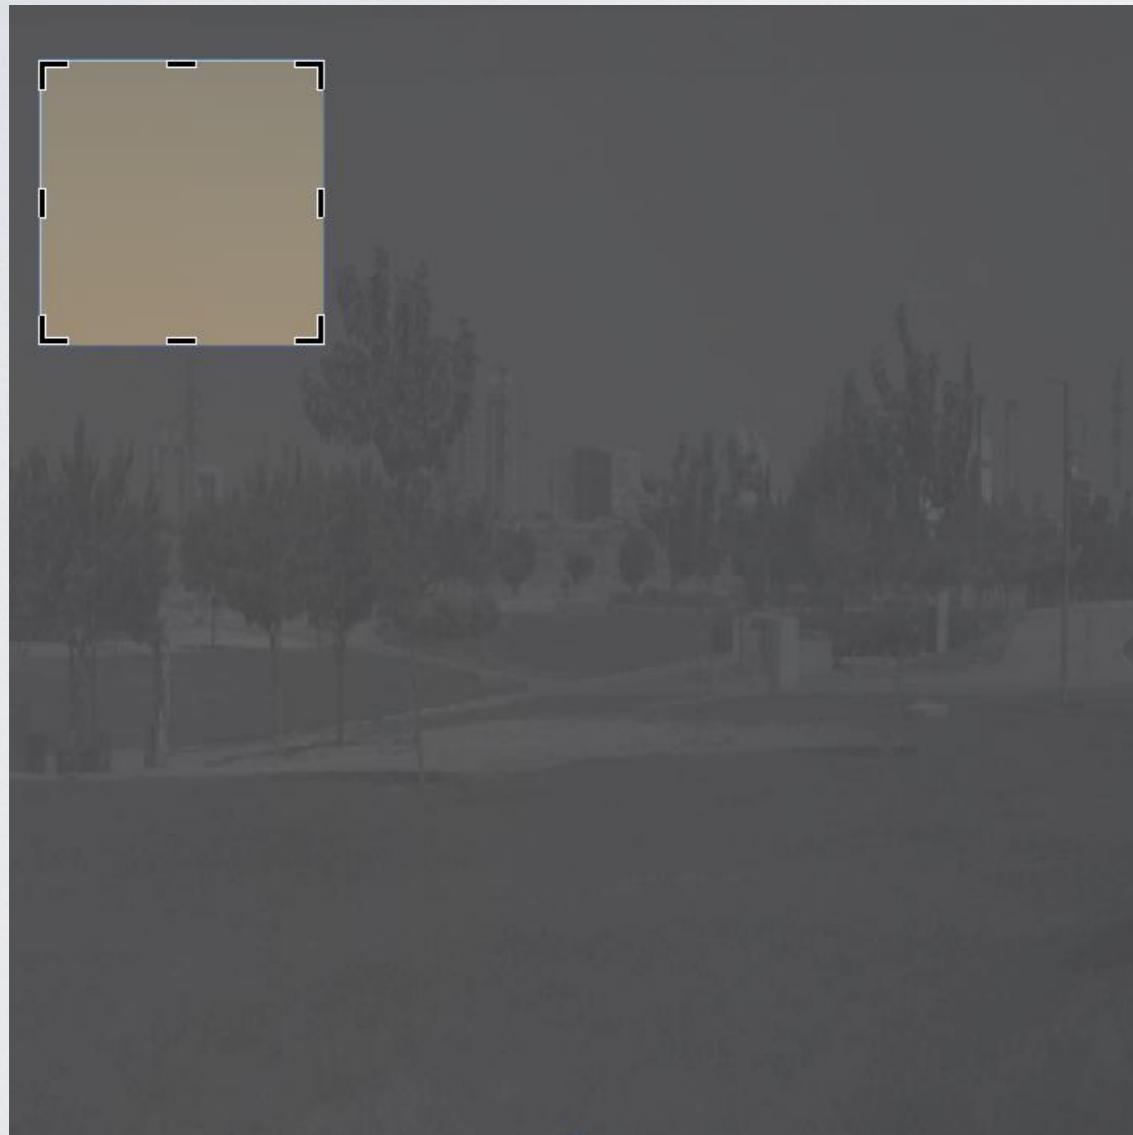


1280

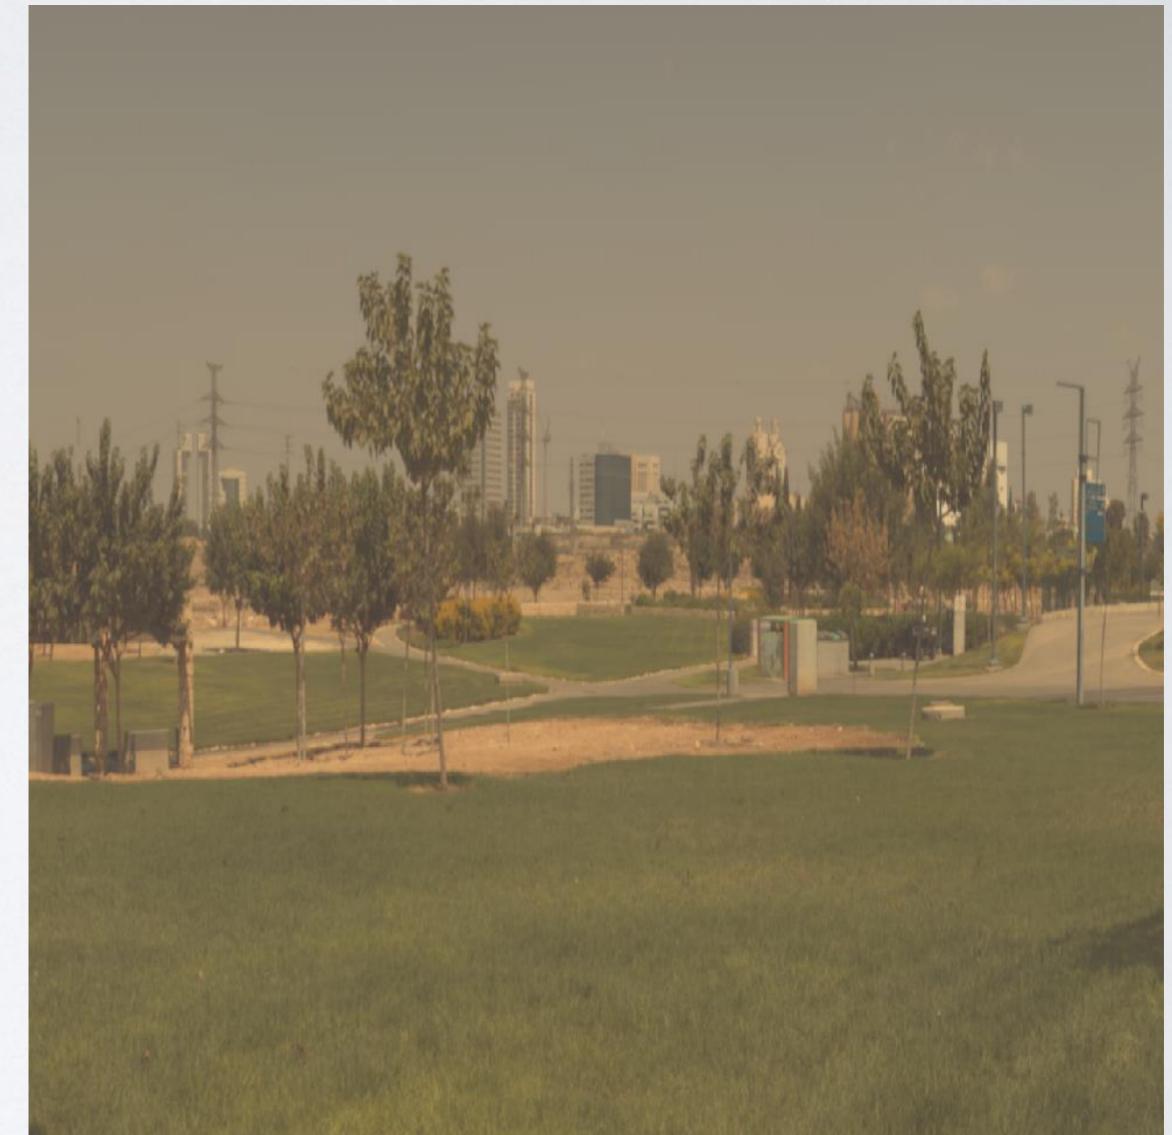
Experiments

Training and testing with large images

- Training: random crops
- Testing: tile decompositon + reconstruction



1280

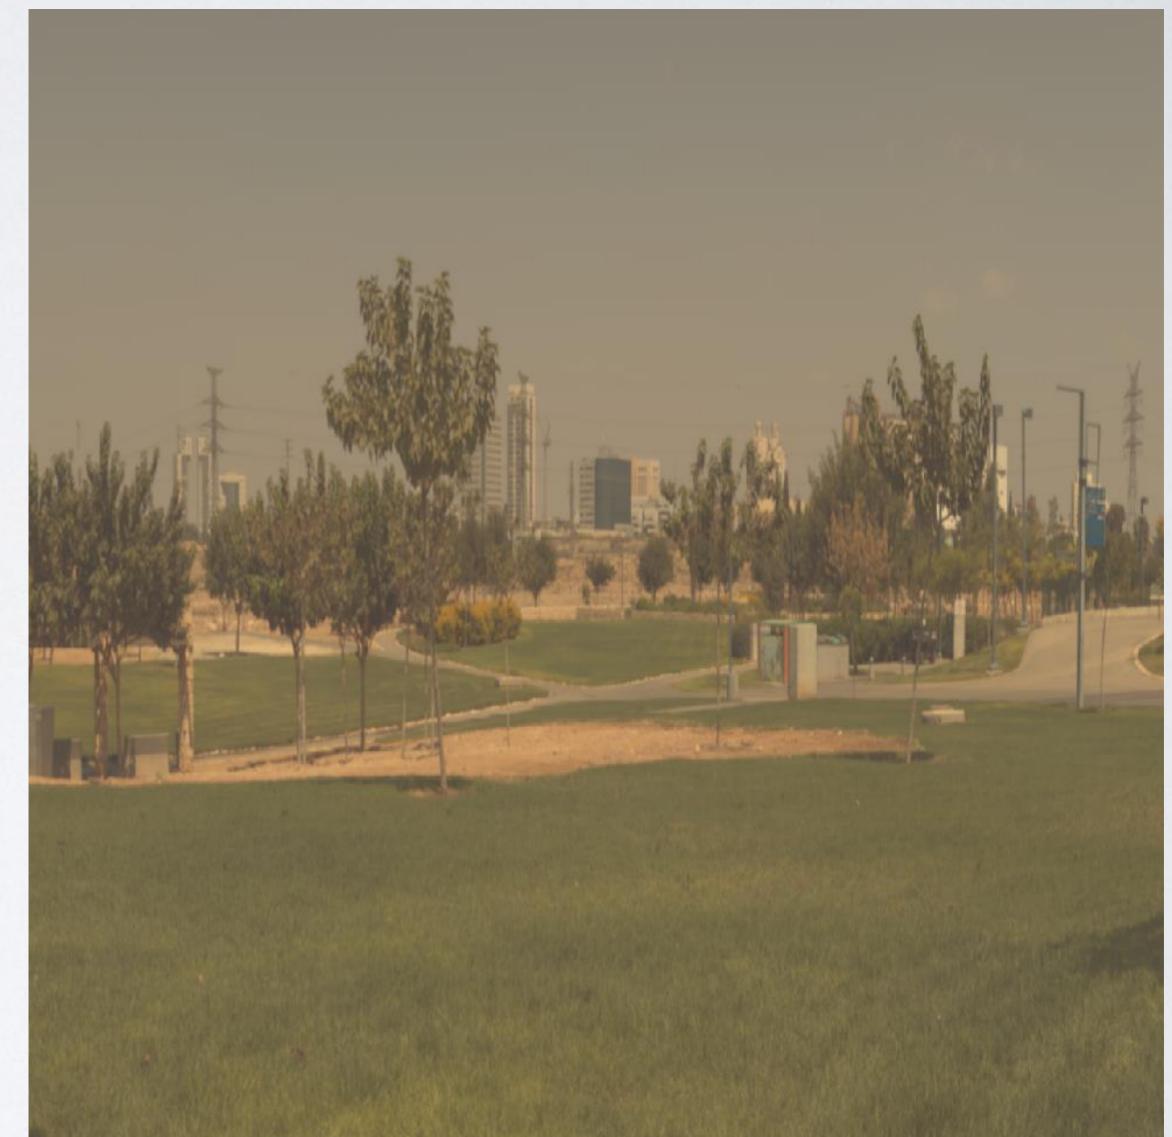
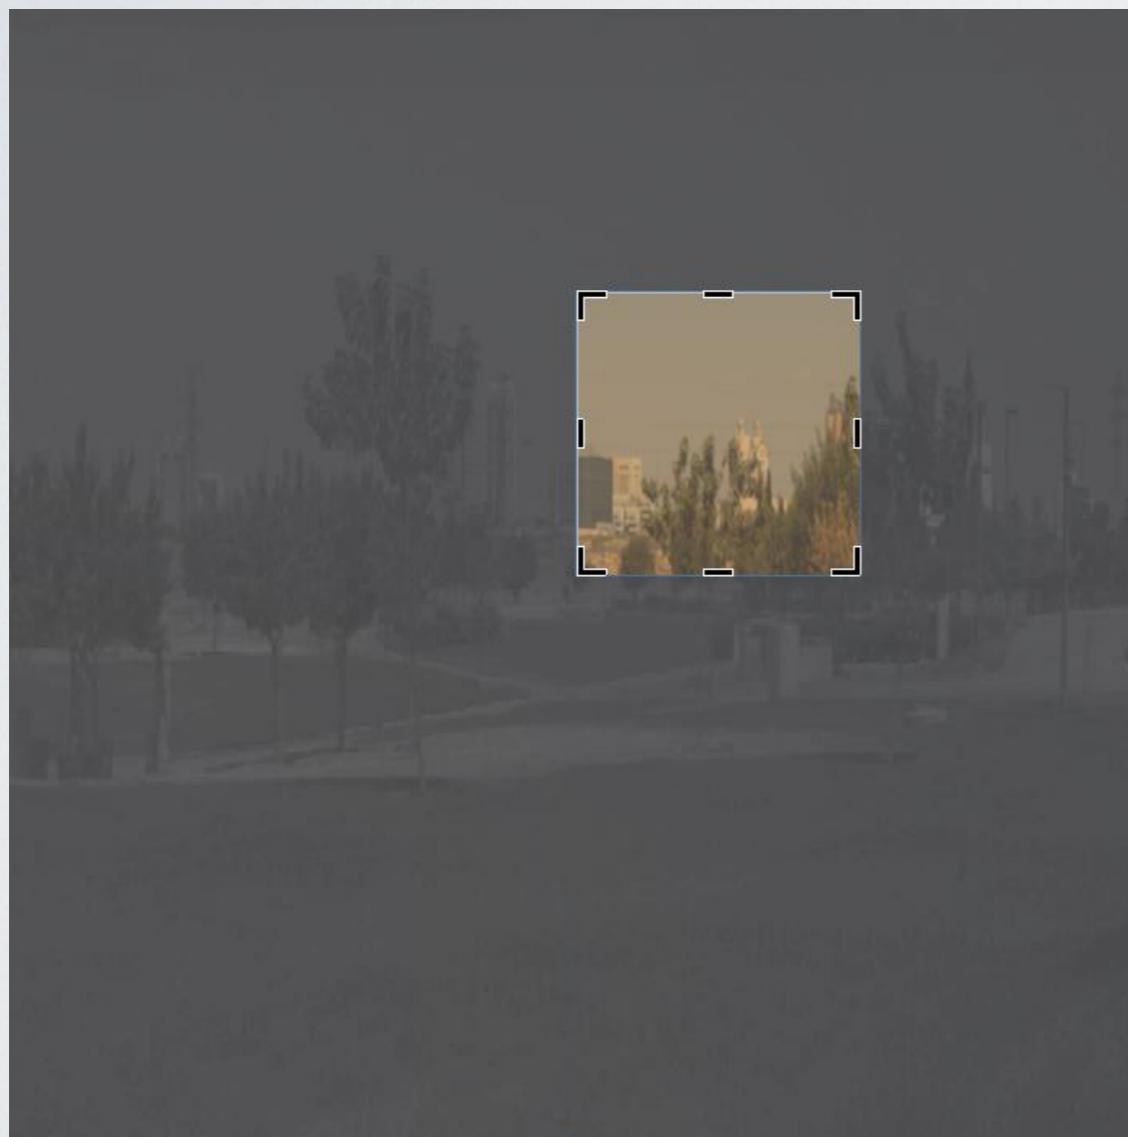


1280

Experiments

Training and testing with large images

- Training: random crops
- Testing: tile decompositon + reconstruction



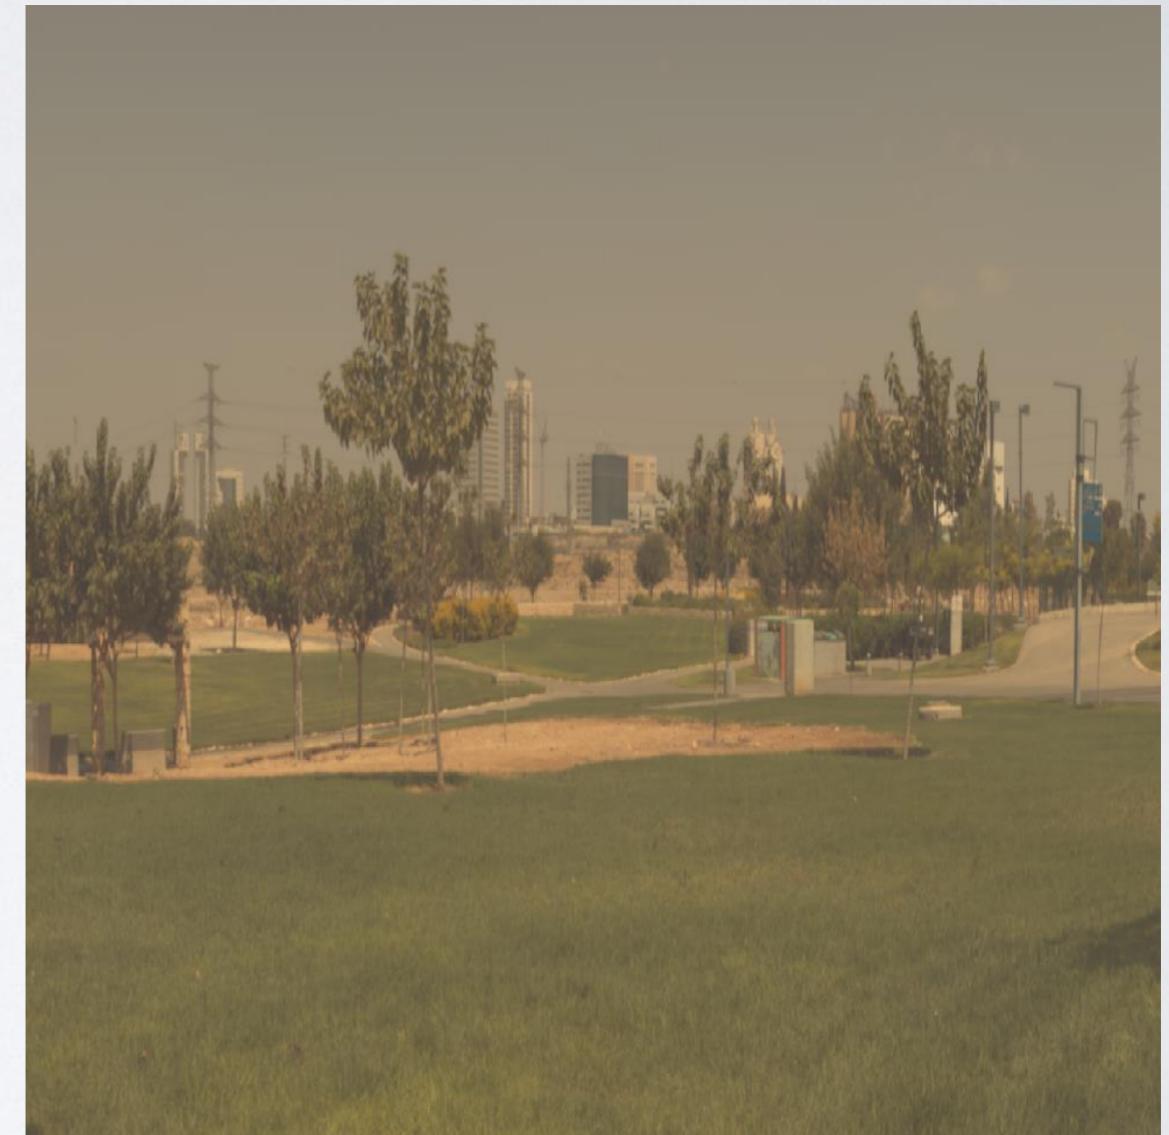
Experiments

Training and testing with large images

- Training: random crops
- Testing: tile decompositon + reconstruction



1280

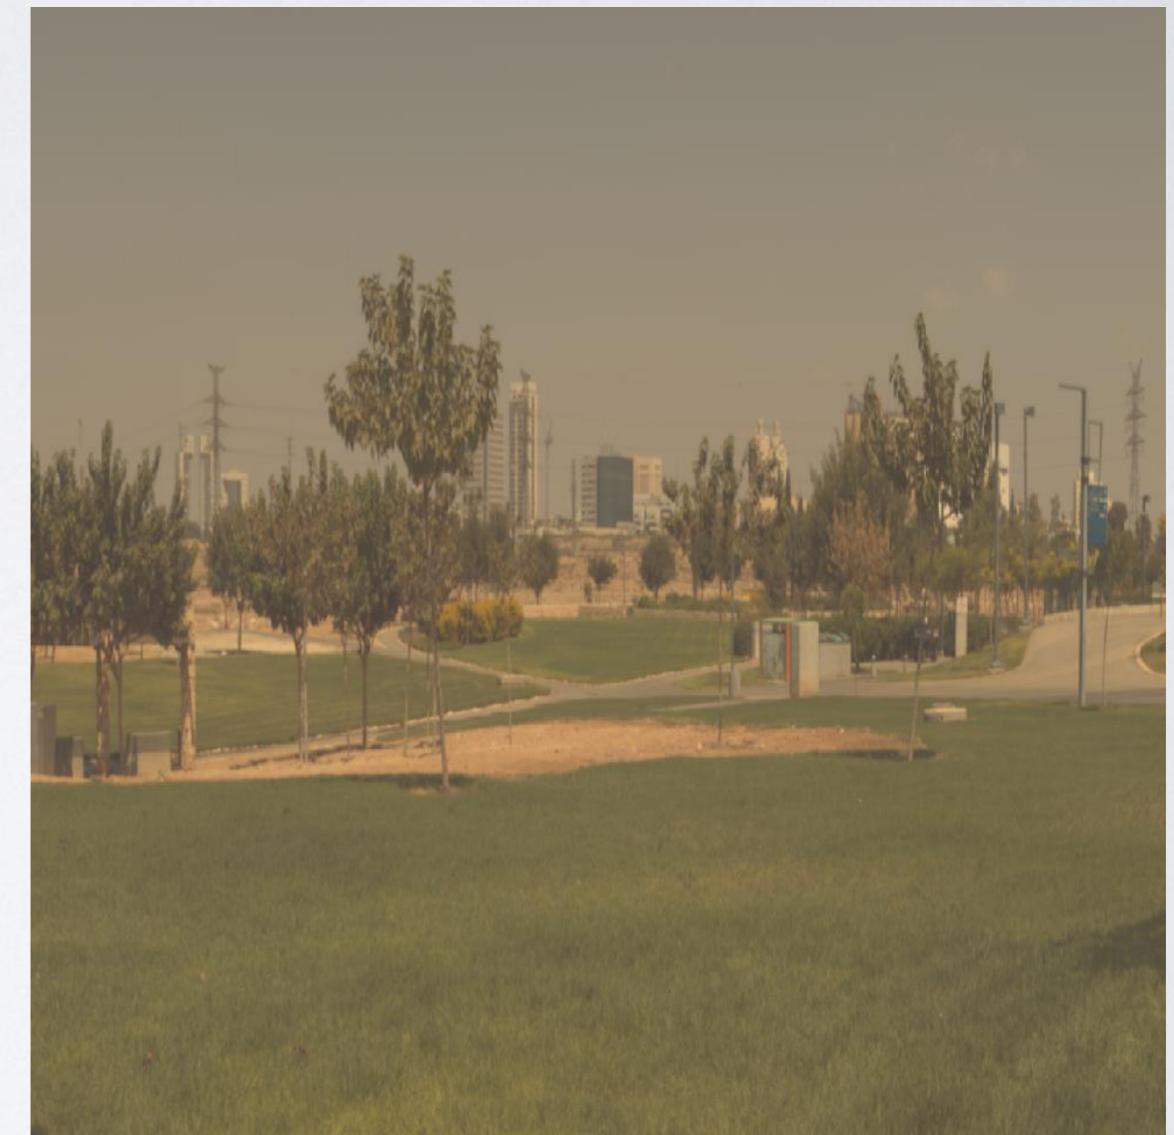


1280

Experiments

Training and testing with large images

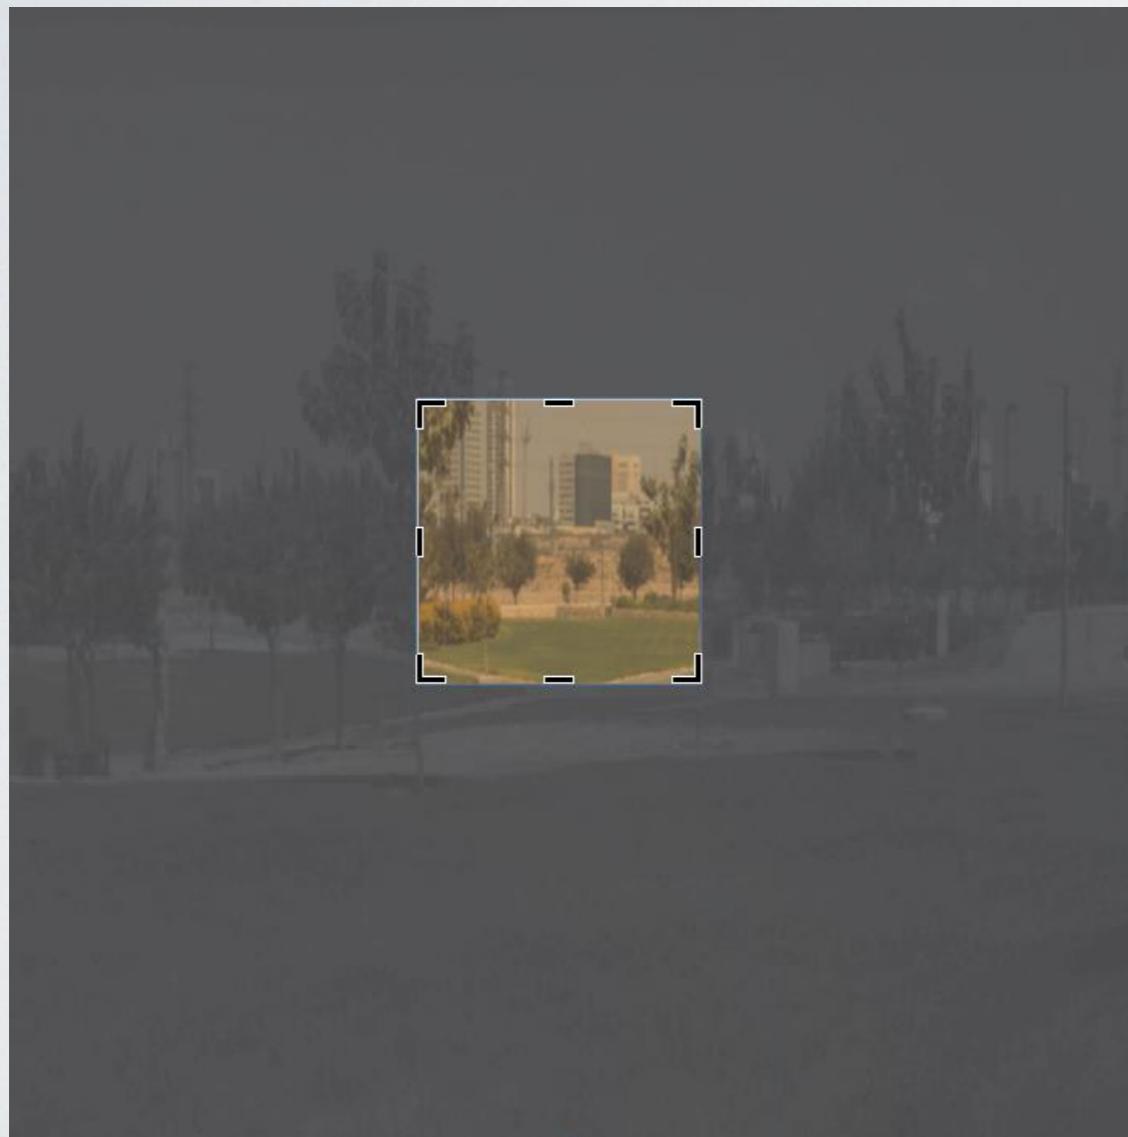
- Training: random crops
- Testing: tile decompositon + reconstruction



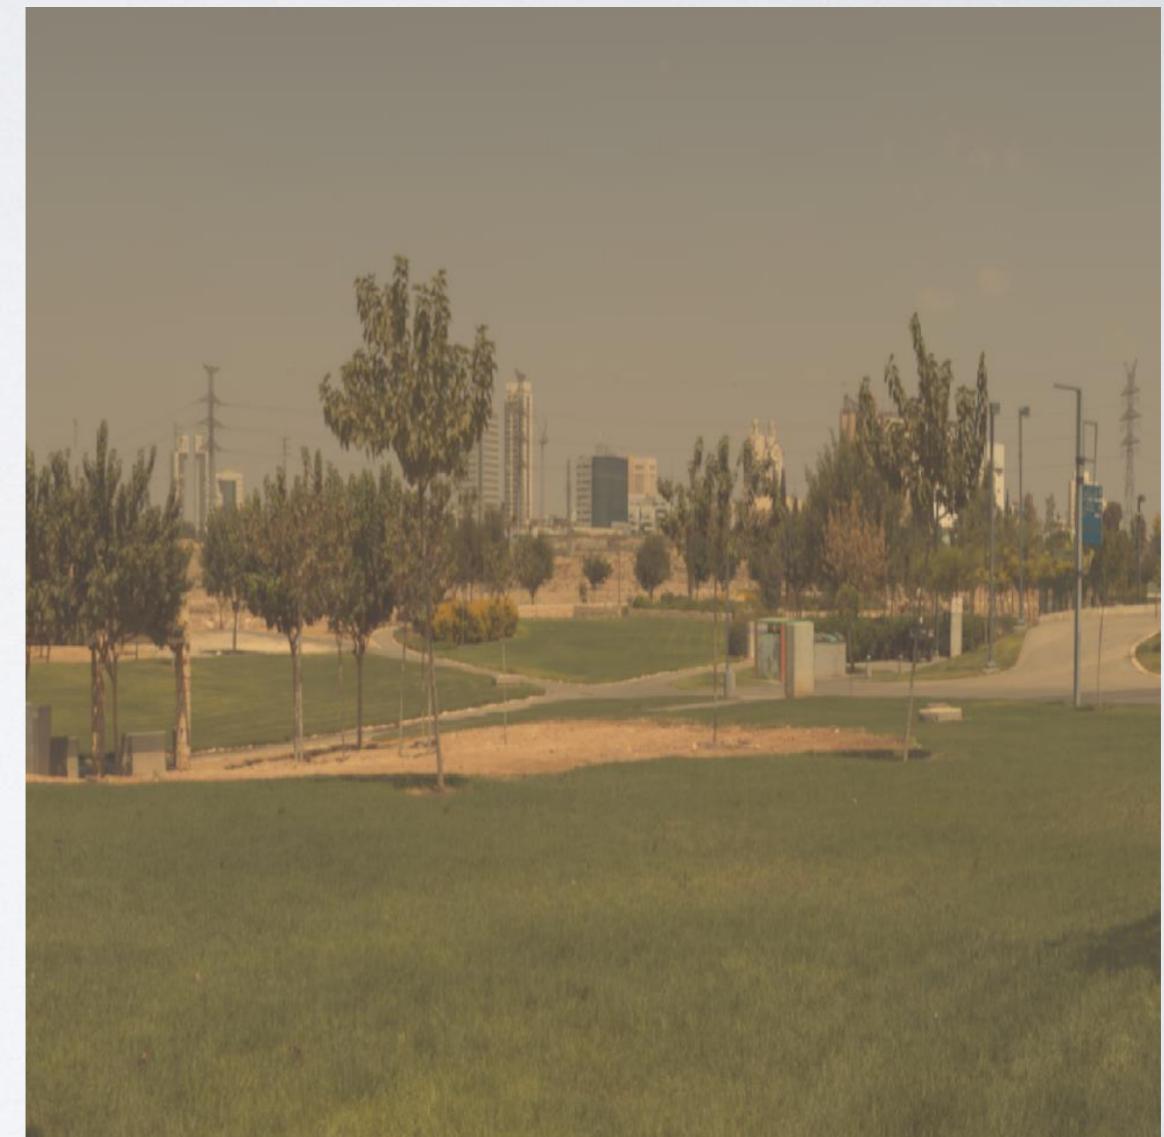
Experiments

Training and testing with large images

- Training: random crops
- Testing: tile decompositon + reconstruction



1280

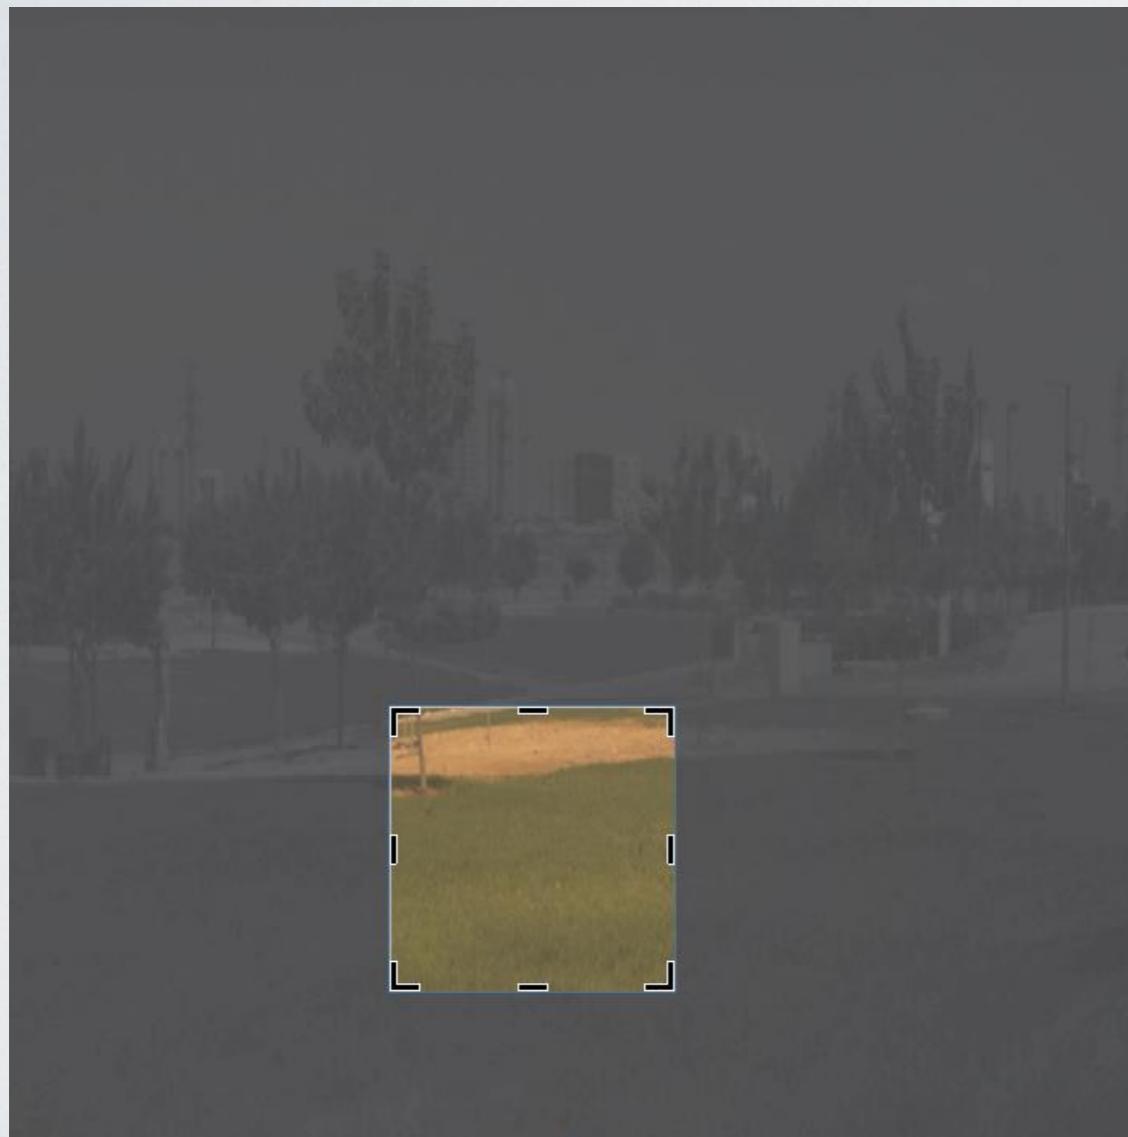


1280

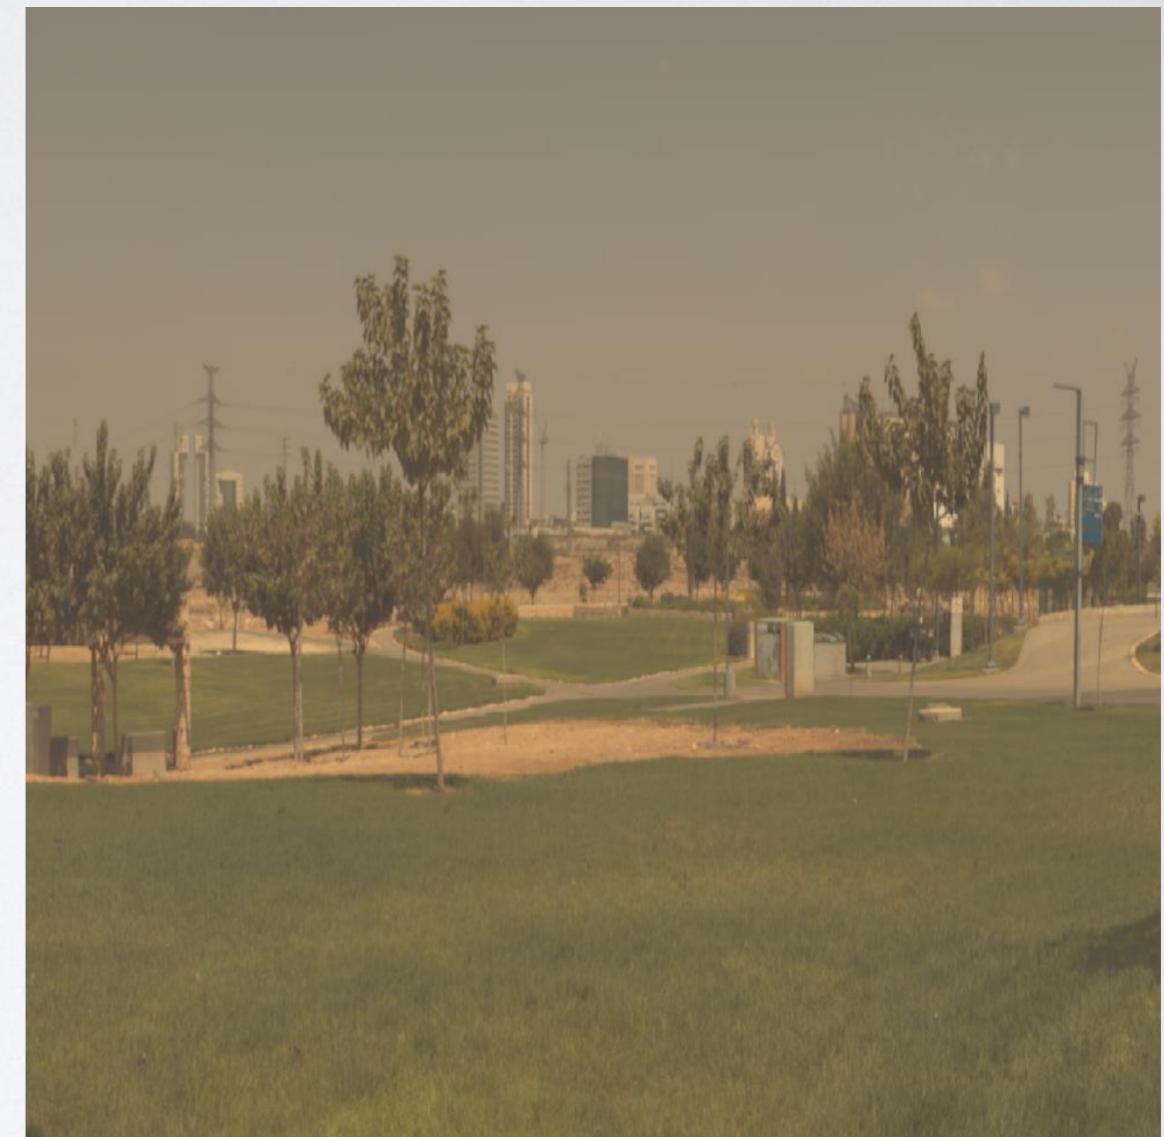
Experiments

Training and testing with large images

- Training: random crops
- Testing: tile decompositon + reconstruction



1280

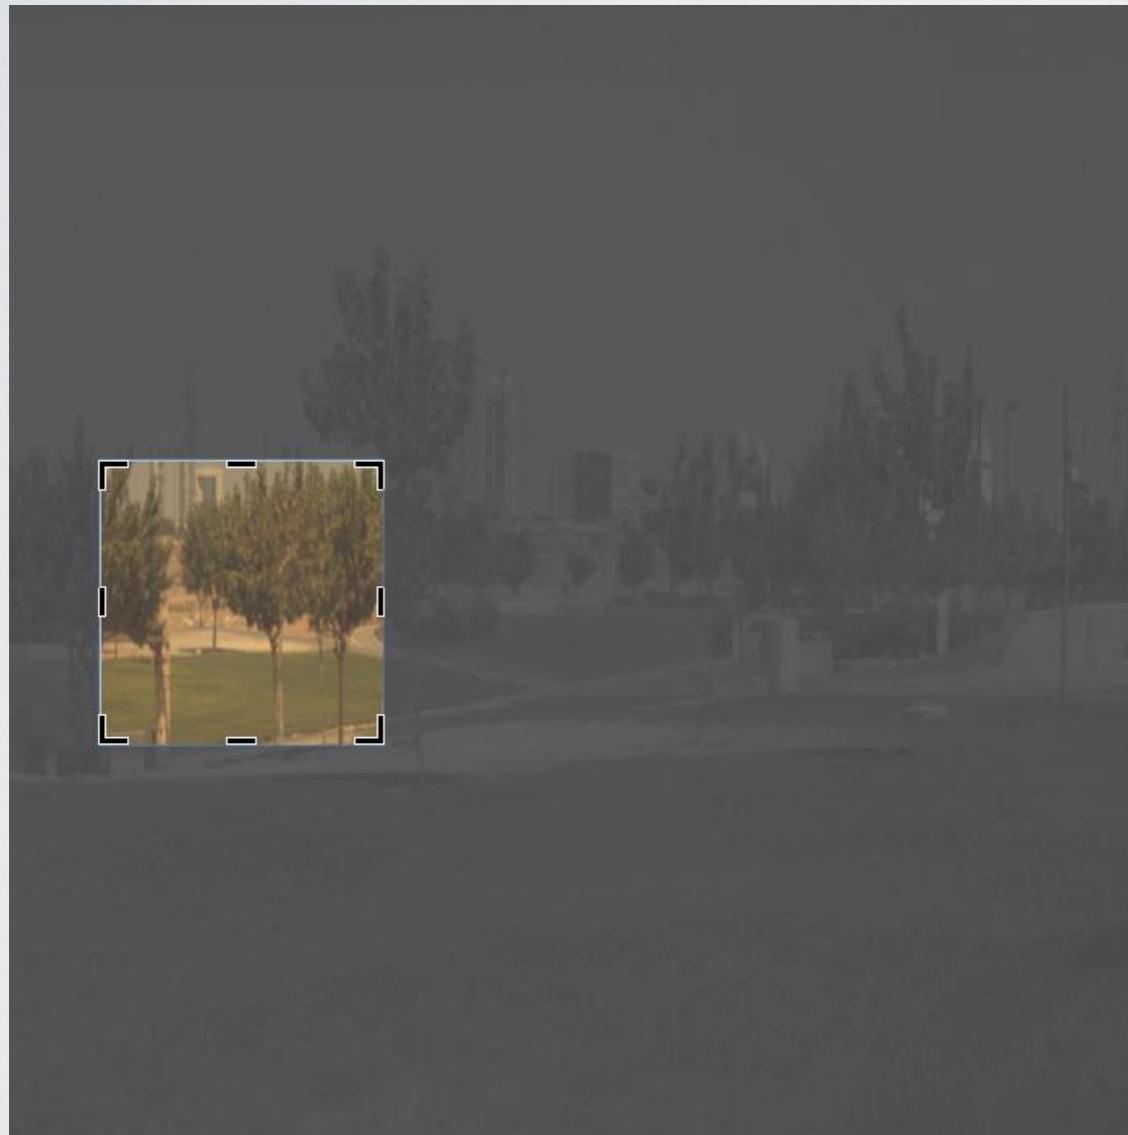


1280

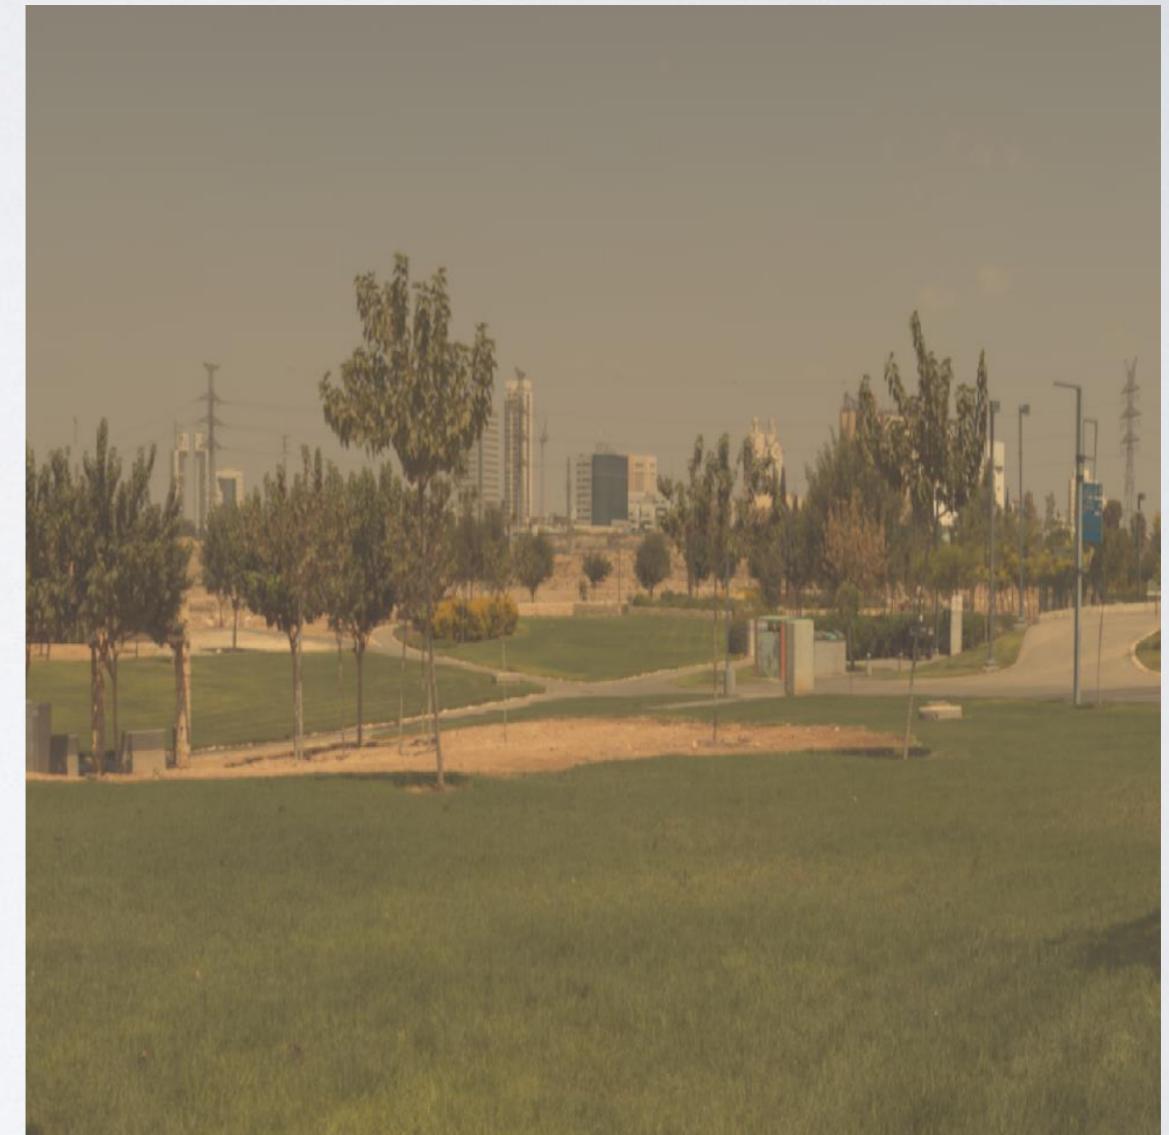
Experiments

Training and testing with large images

- Training: random crops
- Testing: tile decompositon + reconstruction



1280



1280

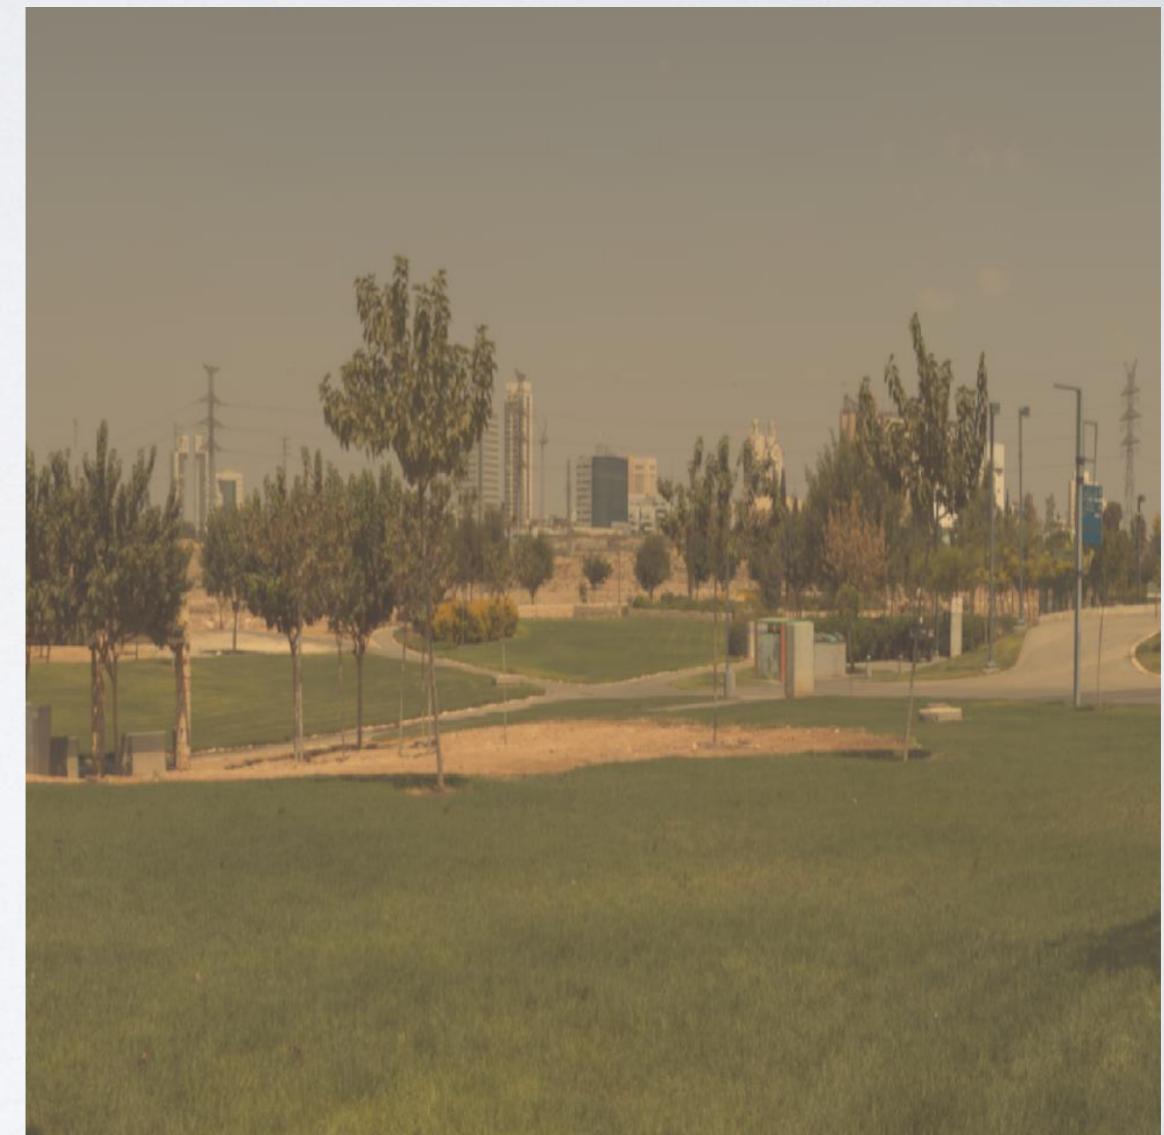
Experiments

Training and testing with large images

- Training: random crops
- Testing: tile decompositon + reconstruction



1280

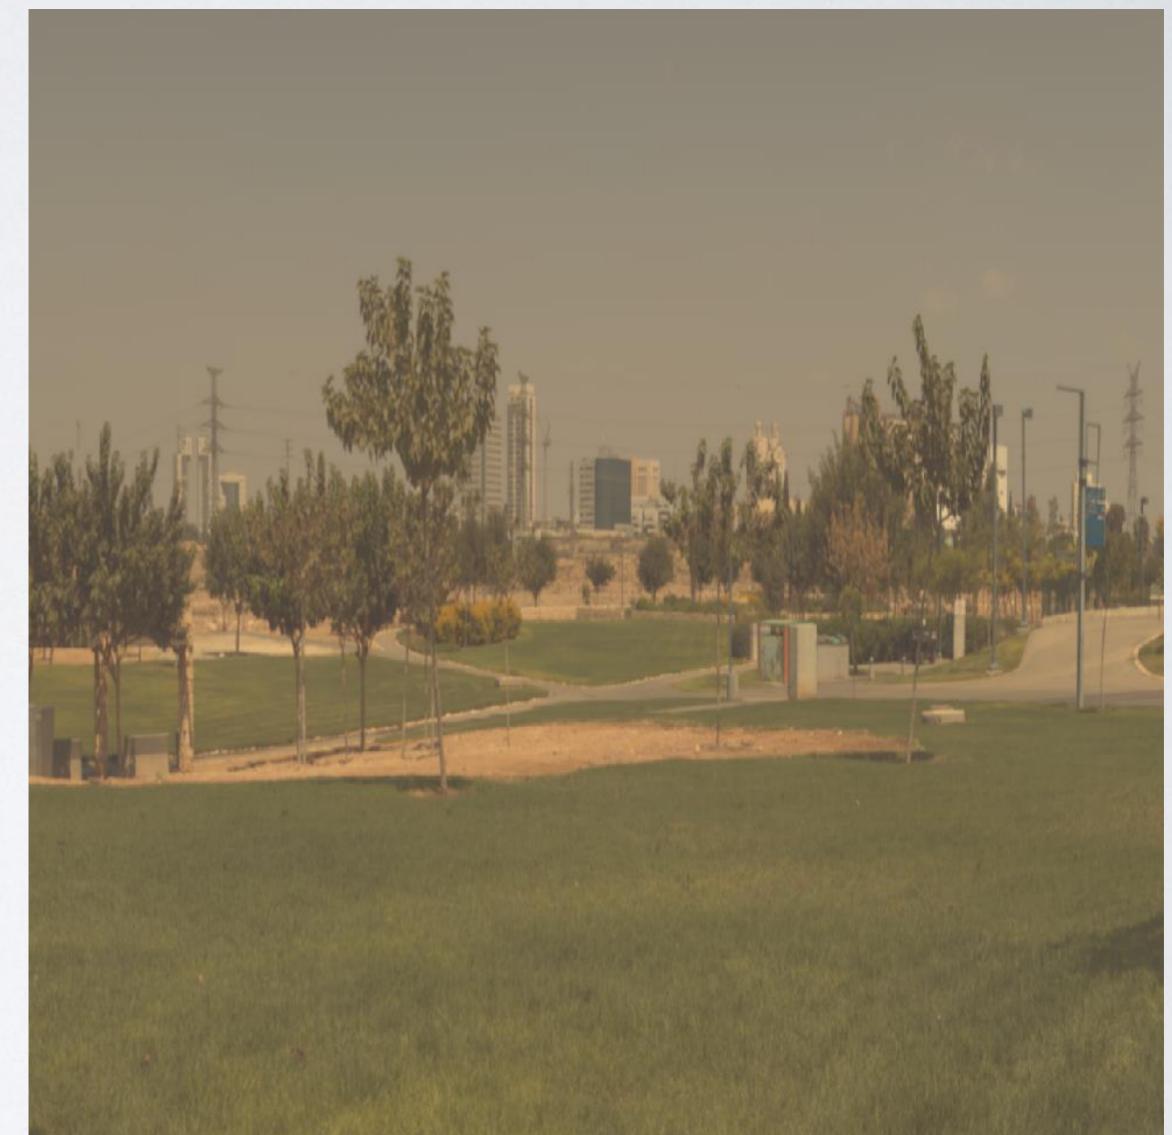


1280

Experiments

Training and testing with large images

- Training: random crops
- Testing: tile decompositon + reconstruction



Experiments

Training and testing with large images

- Training: random crops
- Testing: tile decompositon + reconstruction



1280



1280

Experiments

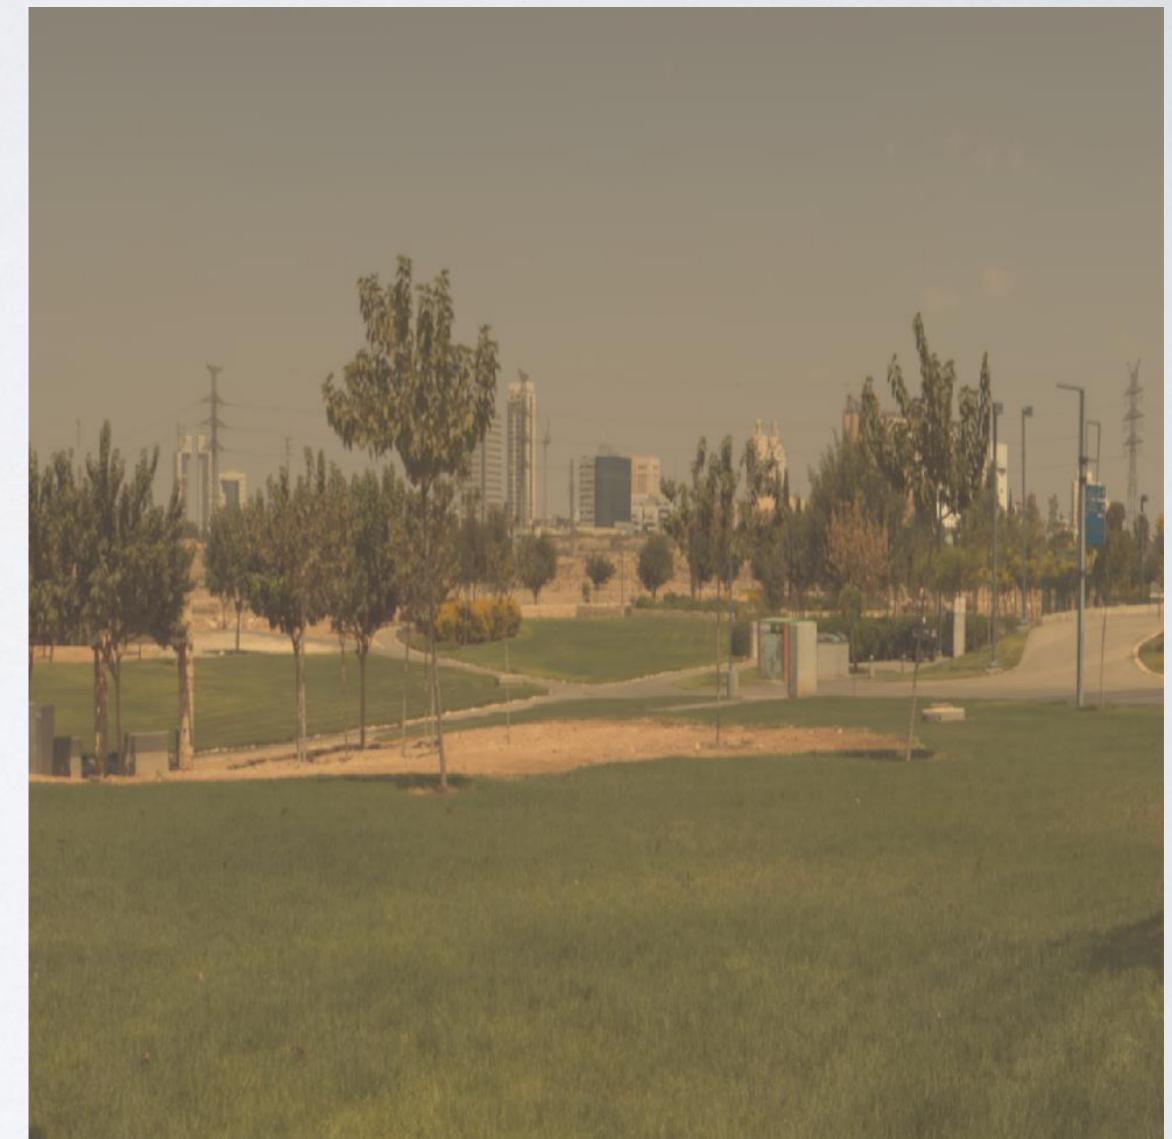
Training and testing with large images

- Training: random crops
- Testing: tile decompositon + reconstruction



1280

1280



1280

1280

Experiments

Training and testing with large images

- Training: random crops
- Testing: tile decompositon + reconstruction



1280



1280

Results

Global

- Metrics (per pixel, pixel-count weighted):
 - **RMSE**: per-pixel Root Mean SE accross the spectrum dimension
 - **Relative RMSE**: Normalized for radiance level to account for samples with low luminance
 - **GFC**: Goodness of fit coefficient
 - **CIE ΔE_{00}** : Perceptual color difference

Results

Global

- Metrics (per pixel, pixel-count weighted):
 - **RMSE**: per-pixel Root Mean SE accross the spectrum dimension
 - **Relative RMSE**: Normalized for radiance level to account for samples with low luminance
 - **GFC**: Goodness of fit coefficient
 - **CIE ΔE_{00}** : Perceptual color difference

Method	RMSE	RMSERel	GFC	ΔE_{00}
[Arad2016] reported	2.633	0.0756	-	-
[Arad2016] default params	2.440 ± 0.066	0.1003 ± 0.022	—	—
[Arad2016] optimized	2.184 ± 0.064	0.0872 ± 0.004	—	—
[Galliani2017] reported	1.980	0.0587	—	—
Ours	1.457 ± 0.040	0.0401 ± 0.0024	0.99921 ± 0.00012	2.044 ± 0.341
fold 0	1.452 ± 0.101	0.0383 ± 0.0024	0.99906 ± 0.00001	1.861 ± 0.324
fold 1	1.463 ± 0.022	0.0420 ± 0.0024	0.99936 ± 0.00023	2.228 ± 0.358

Results

Global

- Metrics (per pixel, pixel-count weighted):
 - **RMSE**: per-pixel Root Mean SE accross the spectrum dimension
 - **Relative RMSE**: Normalized for radiance level to account for samples with low luminance
 - **GFC**: Goodness of fit coefficient
 - **CIE ΔE_{00}** : Perceptual color difference

Method	RMSE	RMSERel	GFC	ΔE_{00}
[Arad2016] reported	2.633	0.0756	-	-
[Arad2016] default params	2.440 ± 0.066	-40.3% 0.1003 ± 0.022	-60.0%	—
[Arad2016] optimized	2.184 ± 0.064	0.0872 ± 0.004	—	—
[Galliani2017] reported	1.980	0.0587	—	—
Ours	1.457 ± 0.040	0.0401 ± 0.0024	0.99921 ± 0.00012	2.044 ± 0.341
fold 0	1.452 ± 0.101	0.0383 ± 0.0024	0.99906 ± 0.00001	1.861 ± 0.324
fold 1	1.463 ± 0.022	0.0420 ± 0.0024	0.99936 ± 0.00023	2.228 ± 0.358

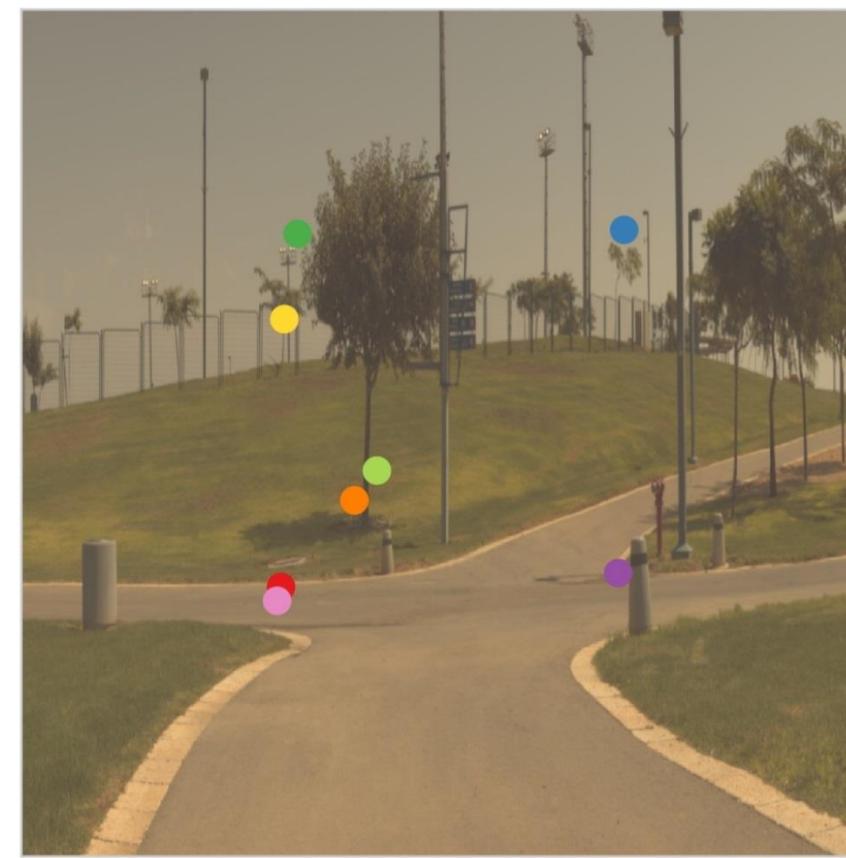
Results

Global

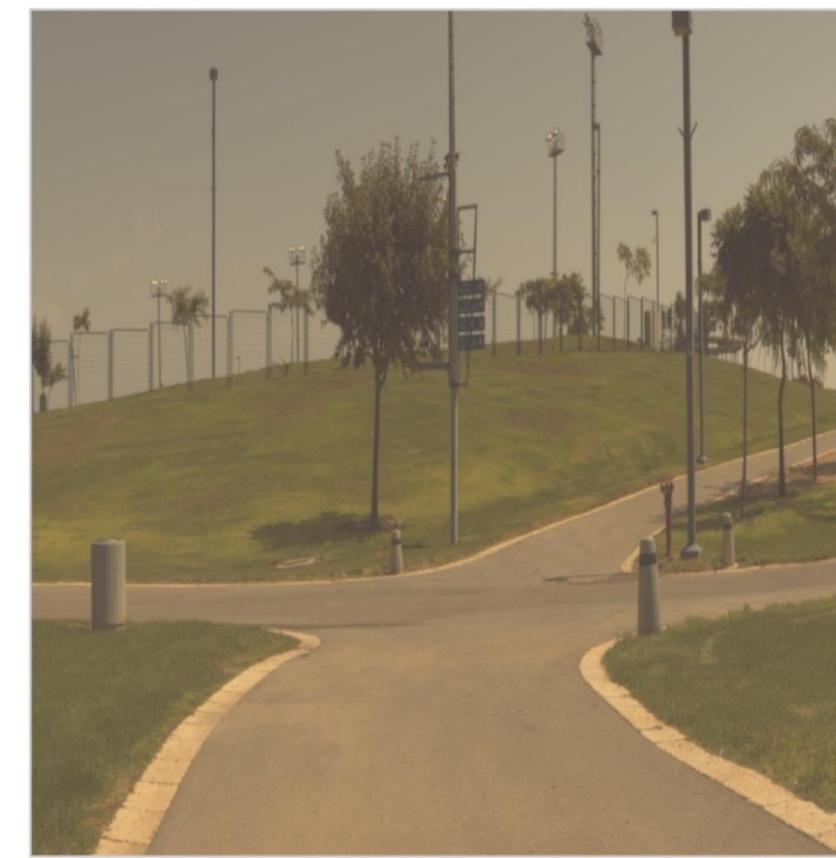
- Metrics (per pixel, pixel-count weighted):
 - **RMSE**: per-pixel Root Mean SE accross the spectrum dimension
 - **Relative RMSE**: Normalized for radiance level to account for samples with low luminance
 - **GFC**: Goodness of fit coefficient
 - **CIE ΔE_{00}** : Perceptual color difference

Method	RMSE	RMSERel	GFC	ΔE_{00}
[Arad2016] reported	2.633	0.0756	-	-
[Arad2016] default params	2.440 ± 0.066	-40.3% 0.1003 ± 0.022	-60.0%	—
[Arad2016] optimized	2.184 ± 0.064	-33.2% 0.0872 ± 0.004	-54.0%	—
[Galliani2017] reported	1.980	0.0587	—	—
Ours	1.457 ± 0.040	0.0401 ± 0.0024	0.99921 ± 0.00012	2.044 ± 0.341
fold 0	1.452 ± 0.101	0.0383 ± 0.0024	0.99906 ± 0.00001	1.861 ± 0.324
fold 1	1.463 ± 0.022	0.0420 ± 0.0024	0.99936 ± 0.00023	2.228 ± 0.358

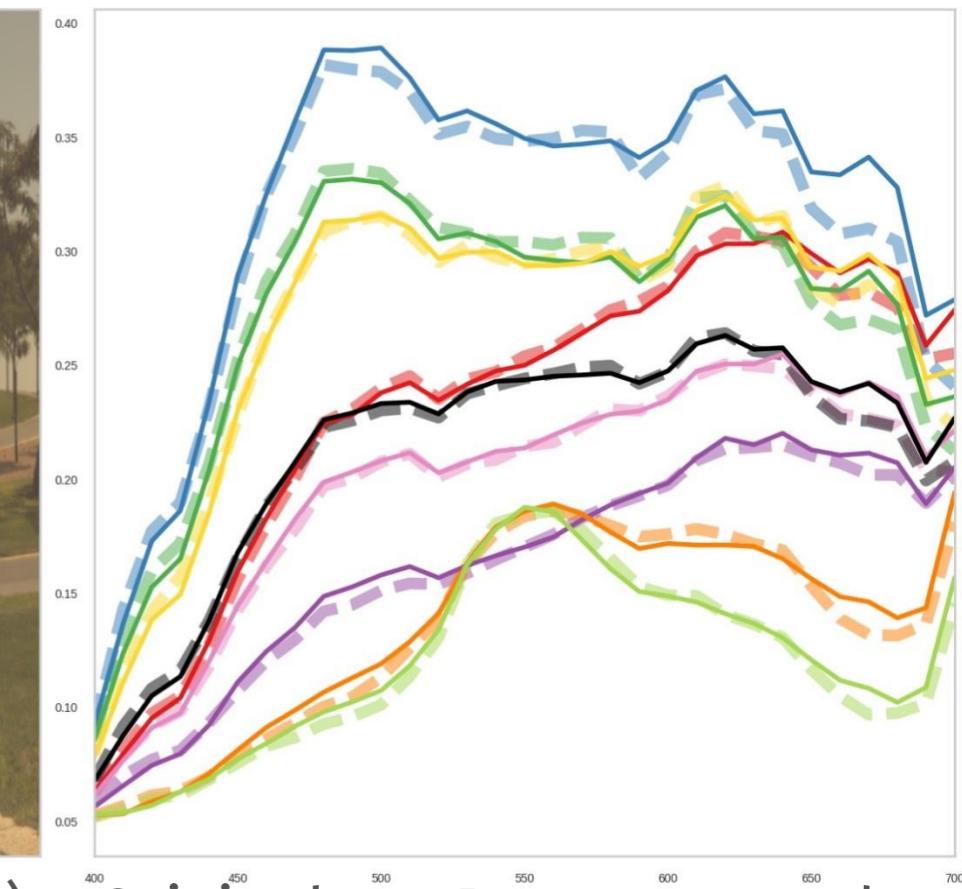
Results



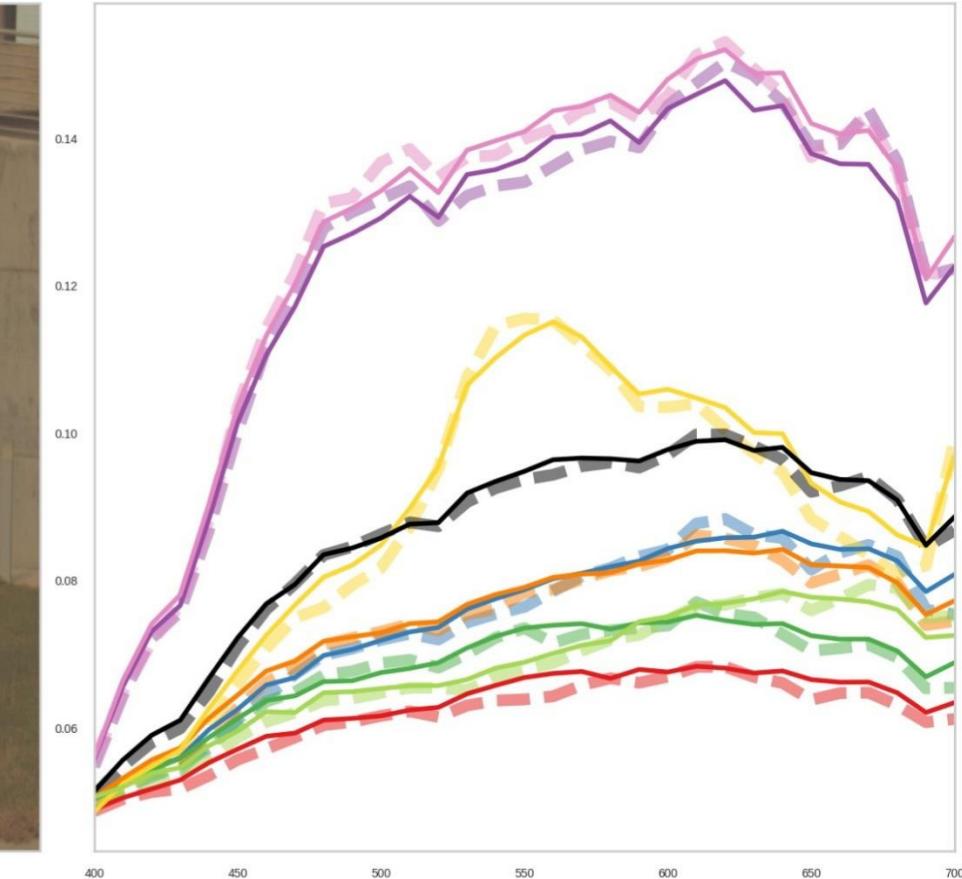
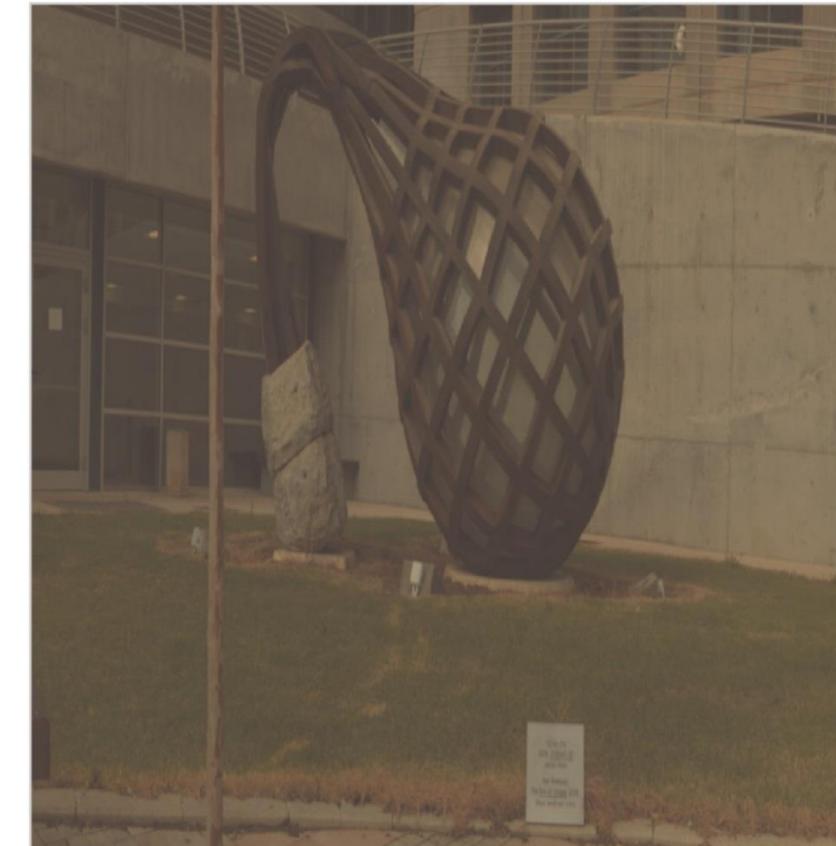
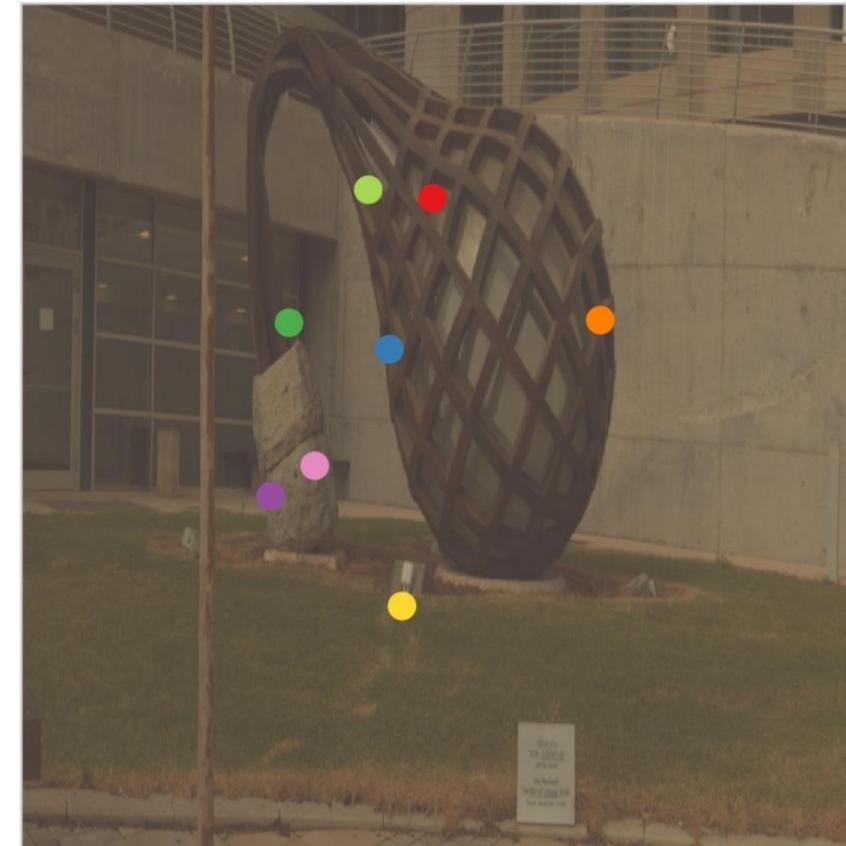
Original (sRGB render)



Reconstructed (sRGB render)



Original --- Reconstructed —



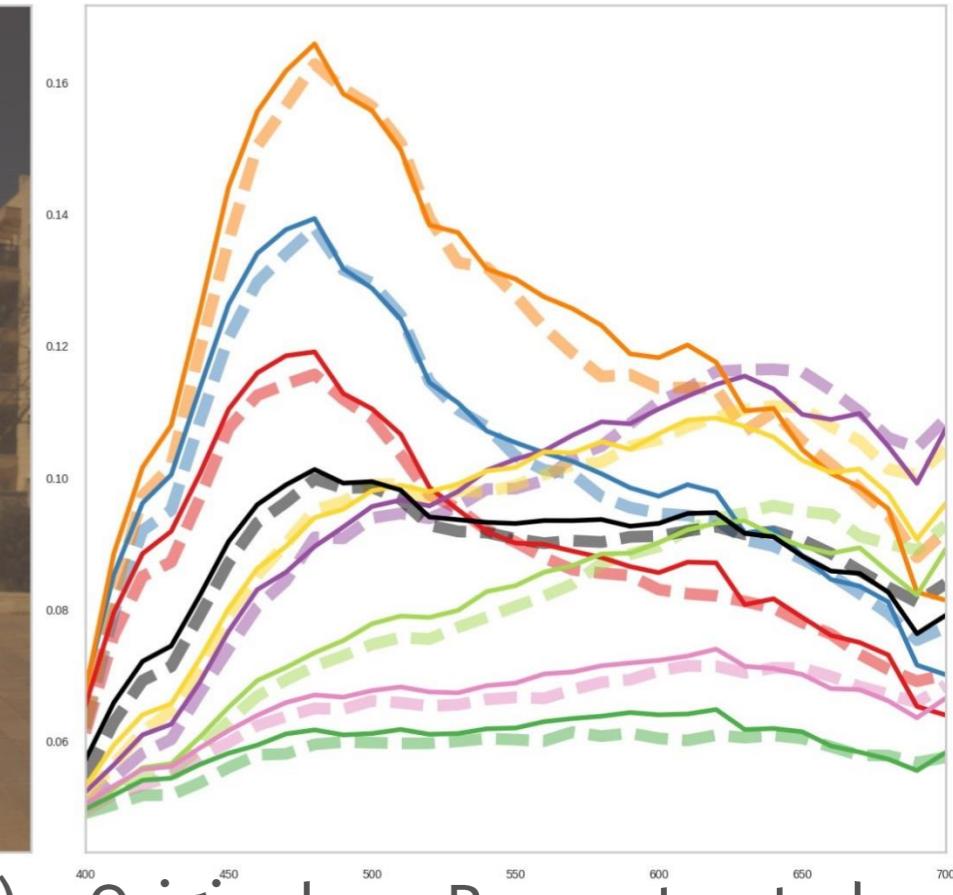
Results



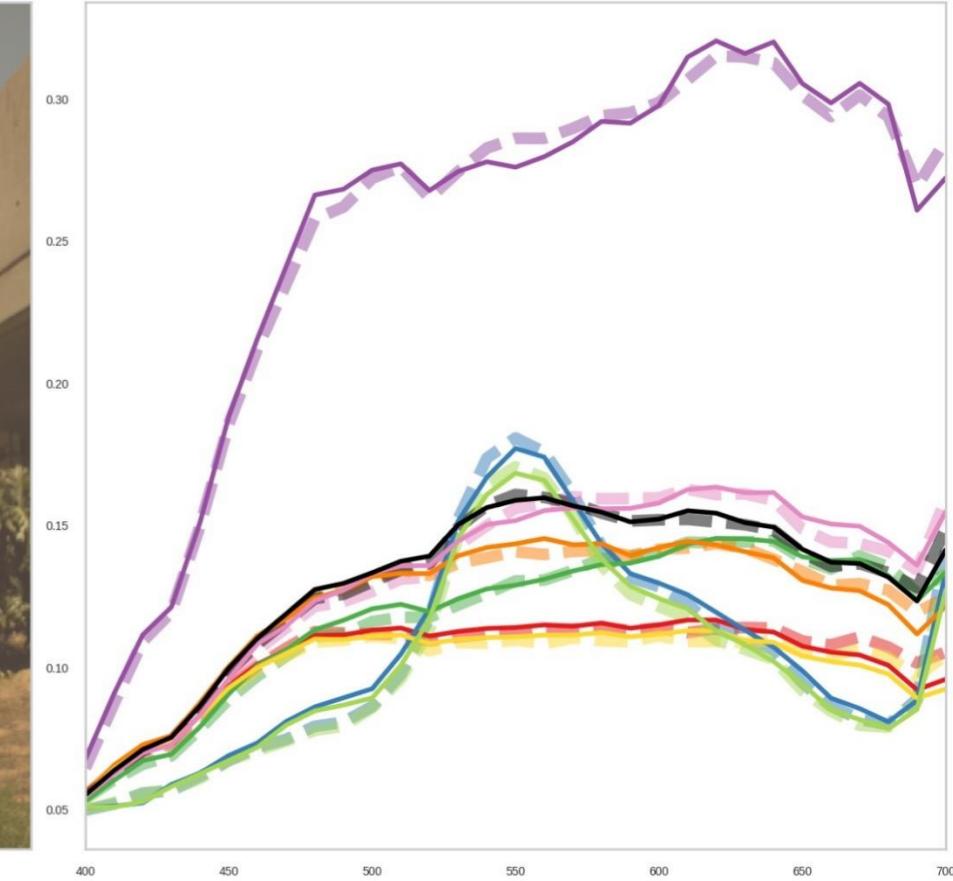
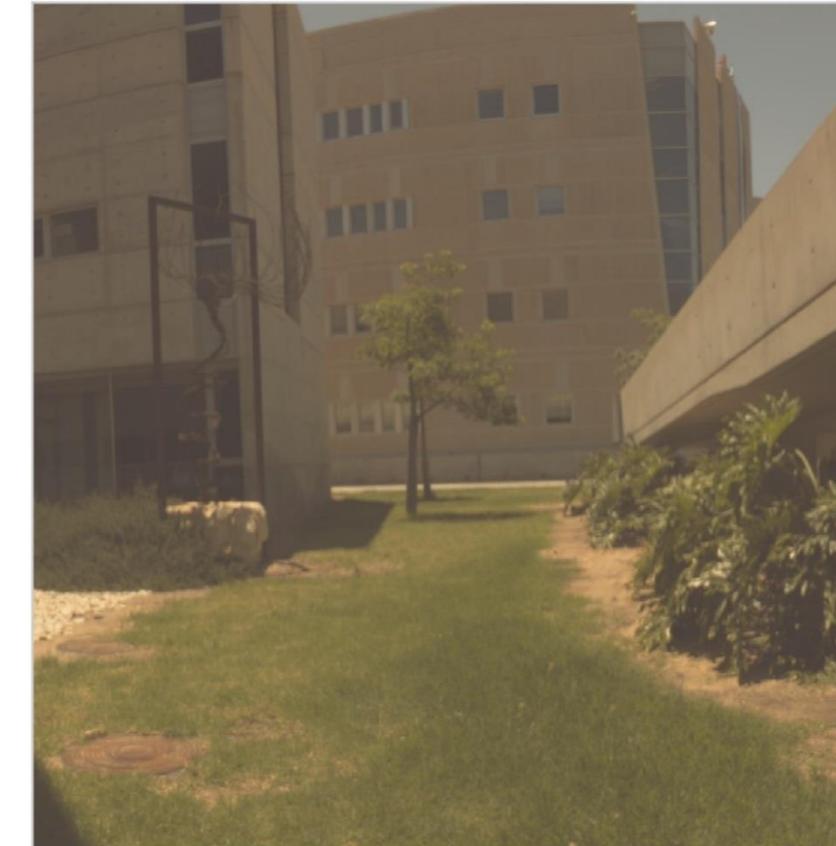
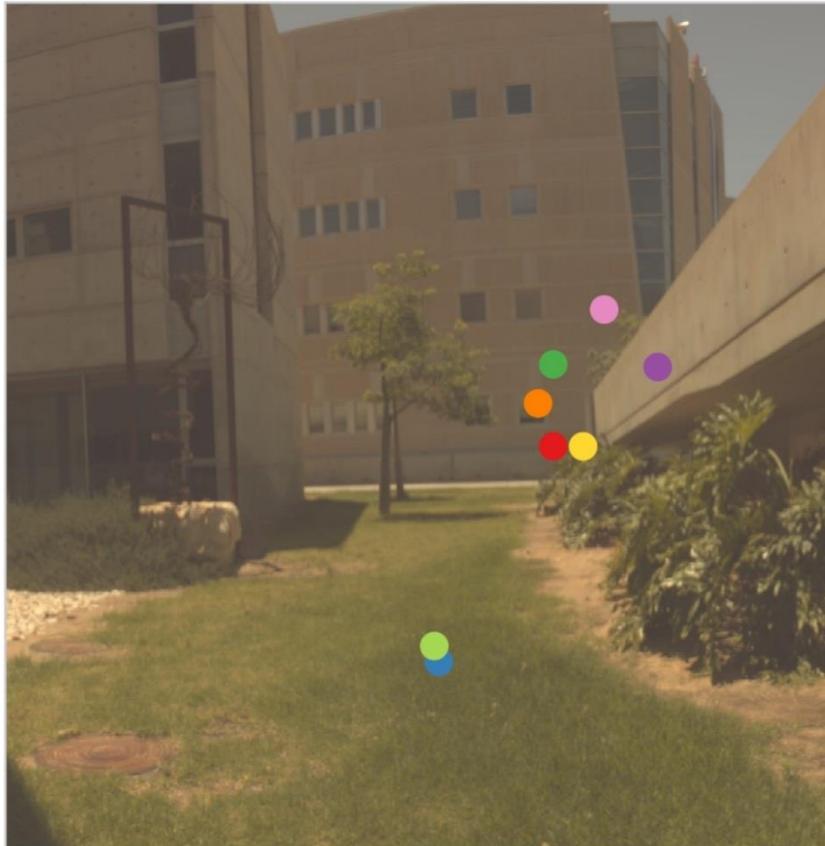
Original (sRGB render)



Reconstructed (sRGB render)



Original --- Reconstructed —



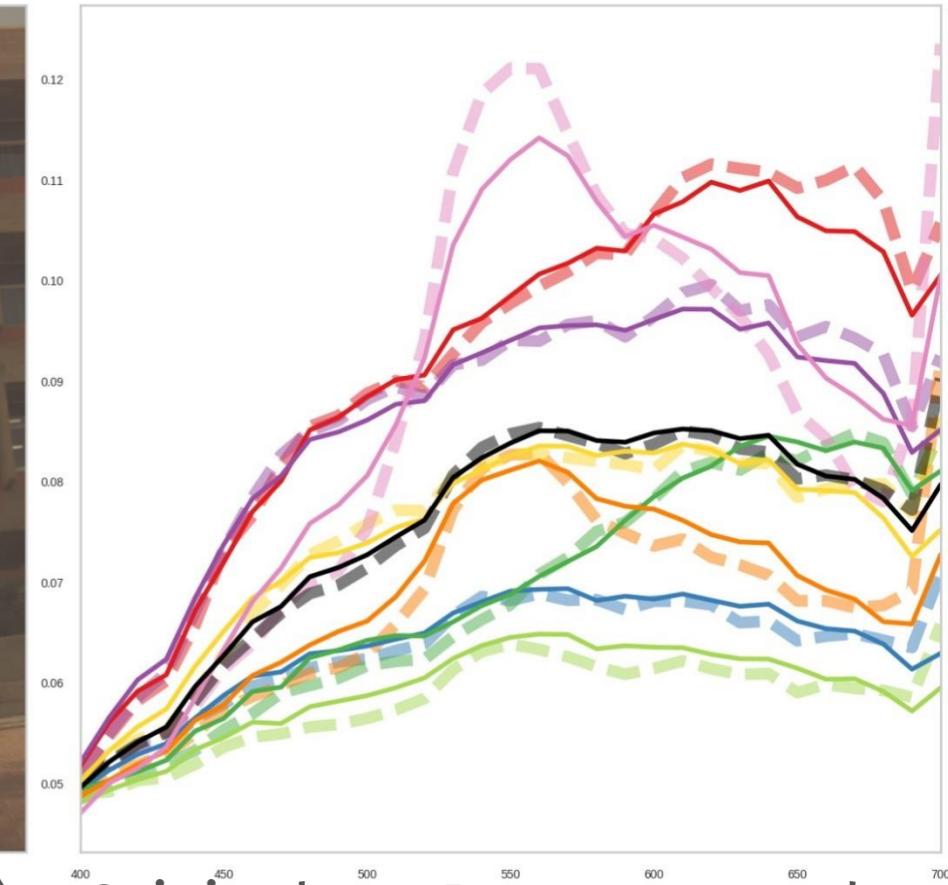
Results



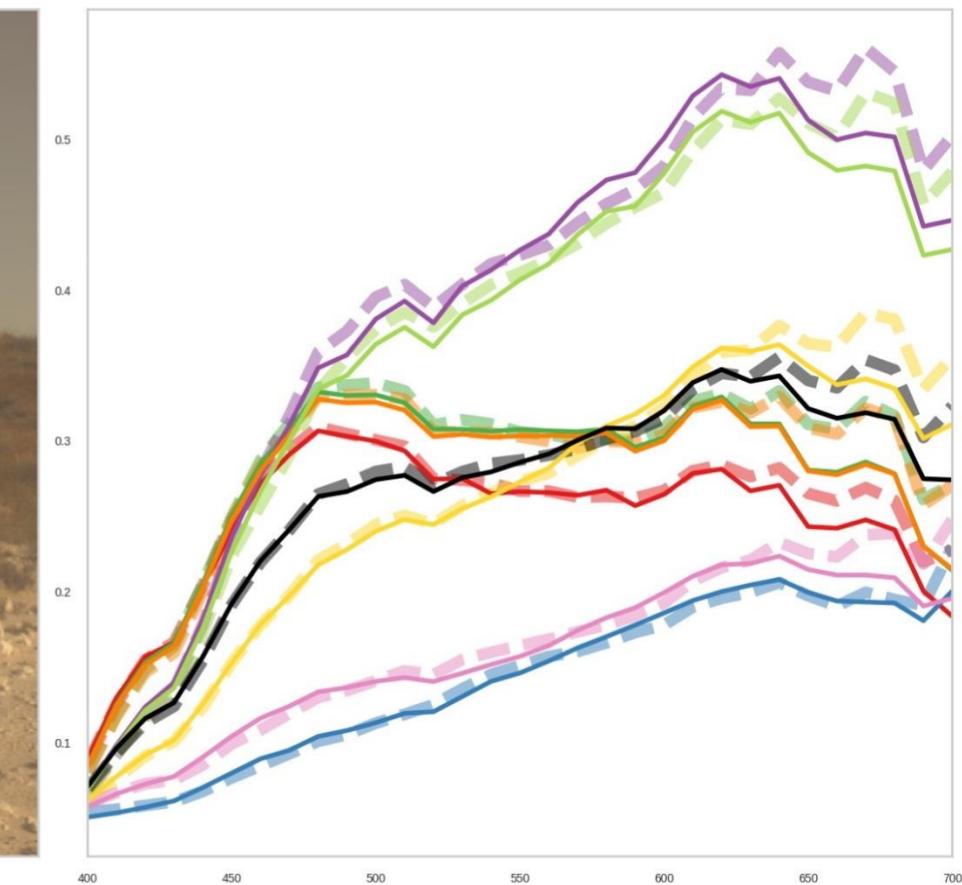
Original (sRGB render)



Reconstructed (sRGB render)



Original --- Reconstructed —



Results

Adversarial Networks for Spatial
Context-Aware Spectral Image
Reconstruction from RGB

Results

Adversarial Networks for Spatial
Context-Aware Spectral Image
Reconstruction from RGB

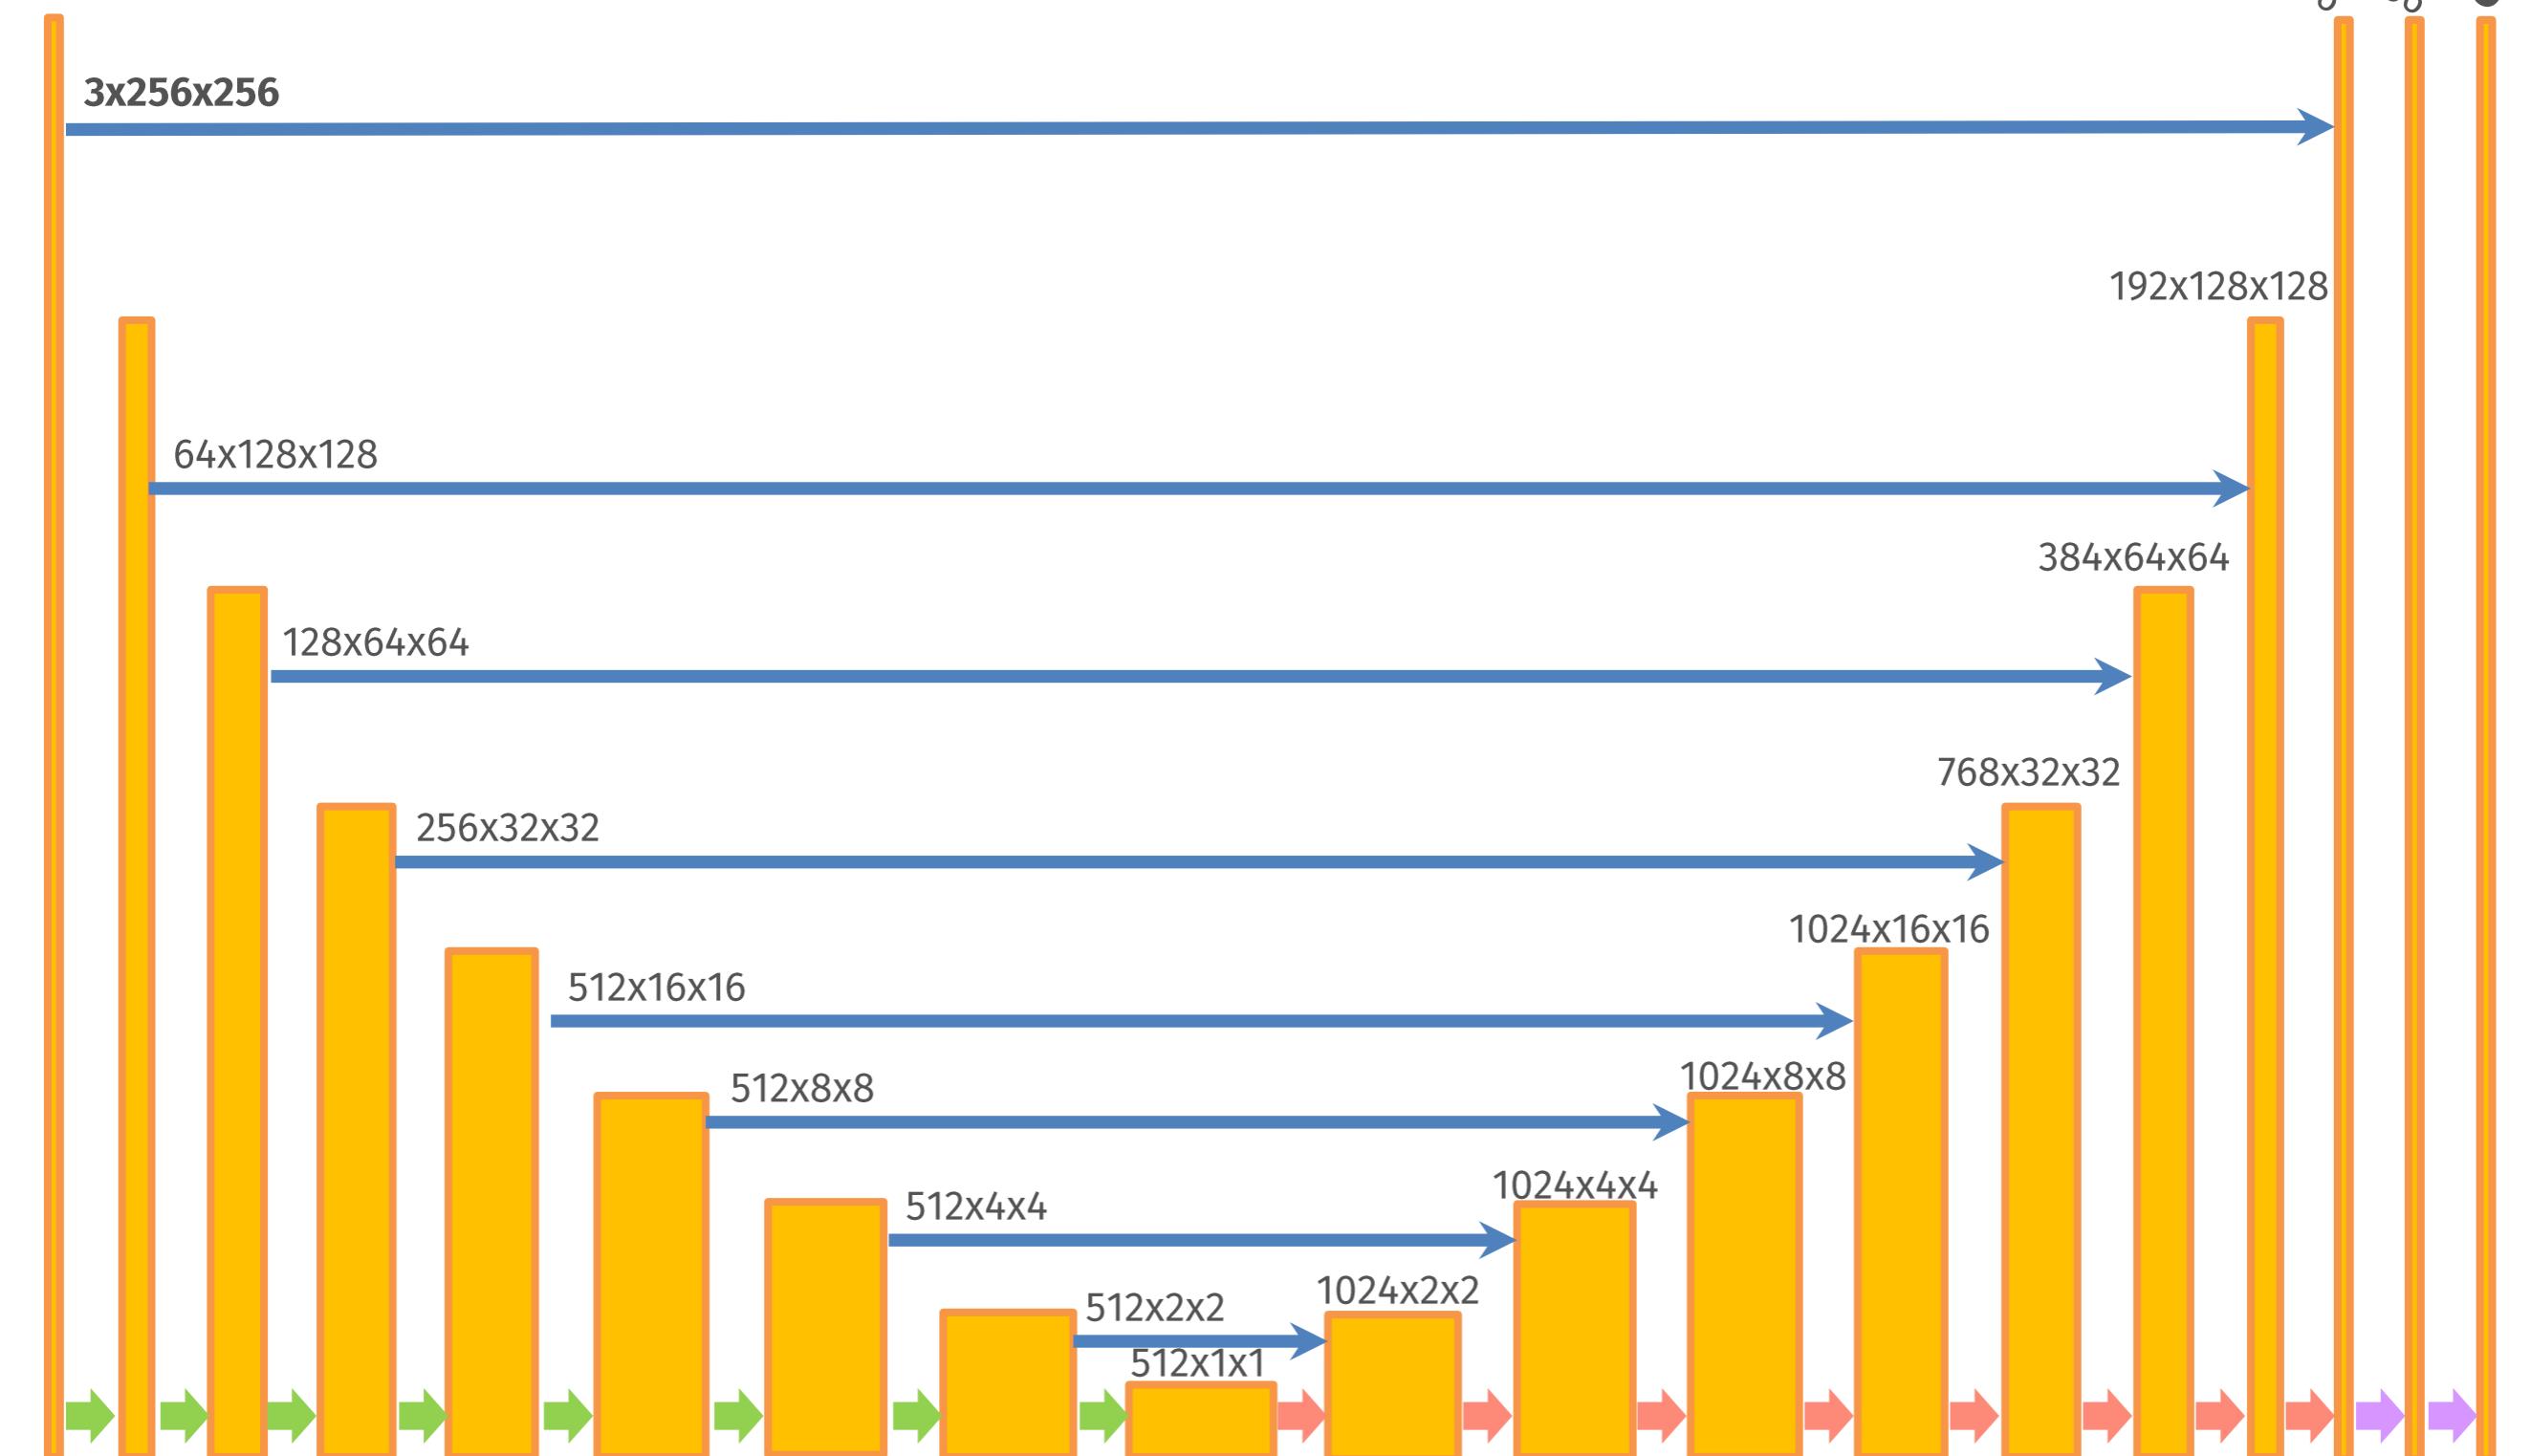
Results

Does spatial information really help?

Adversarial Networks for Spatial
Context-Aware Spectral Image
Reconstruction from RGB

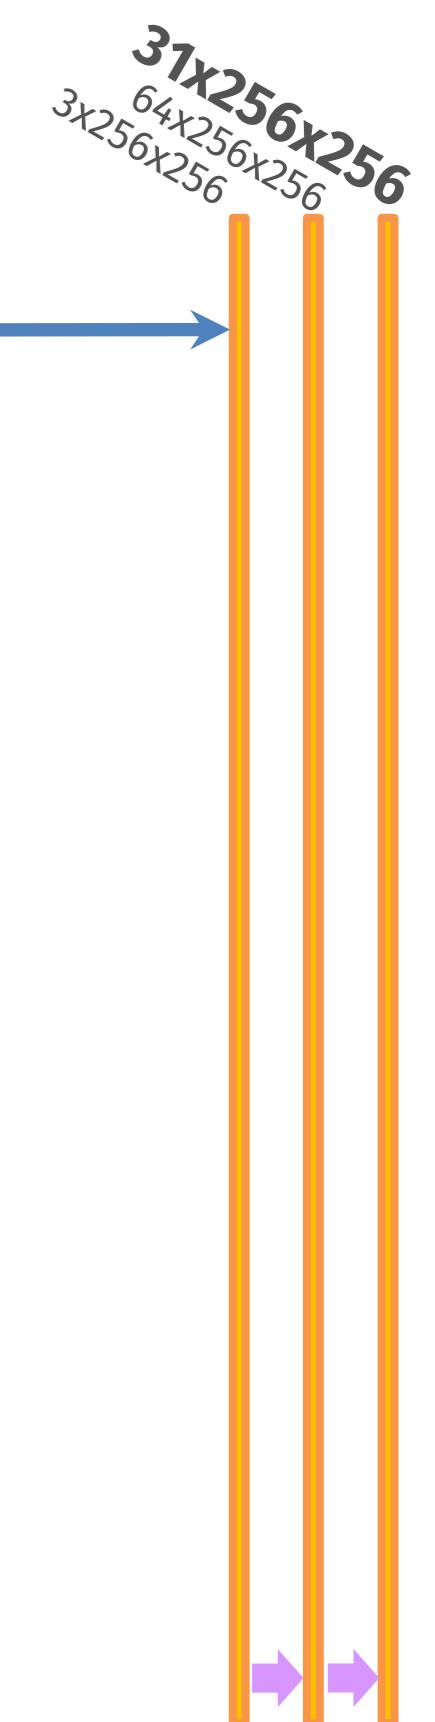
Does spatial information really help?

Full G Net



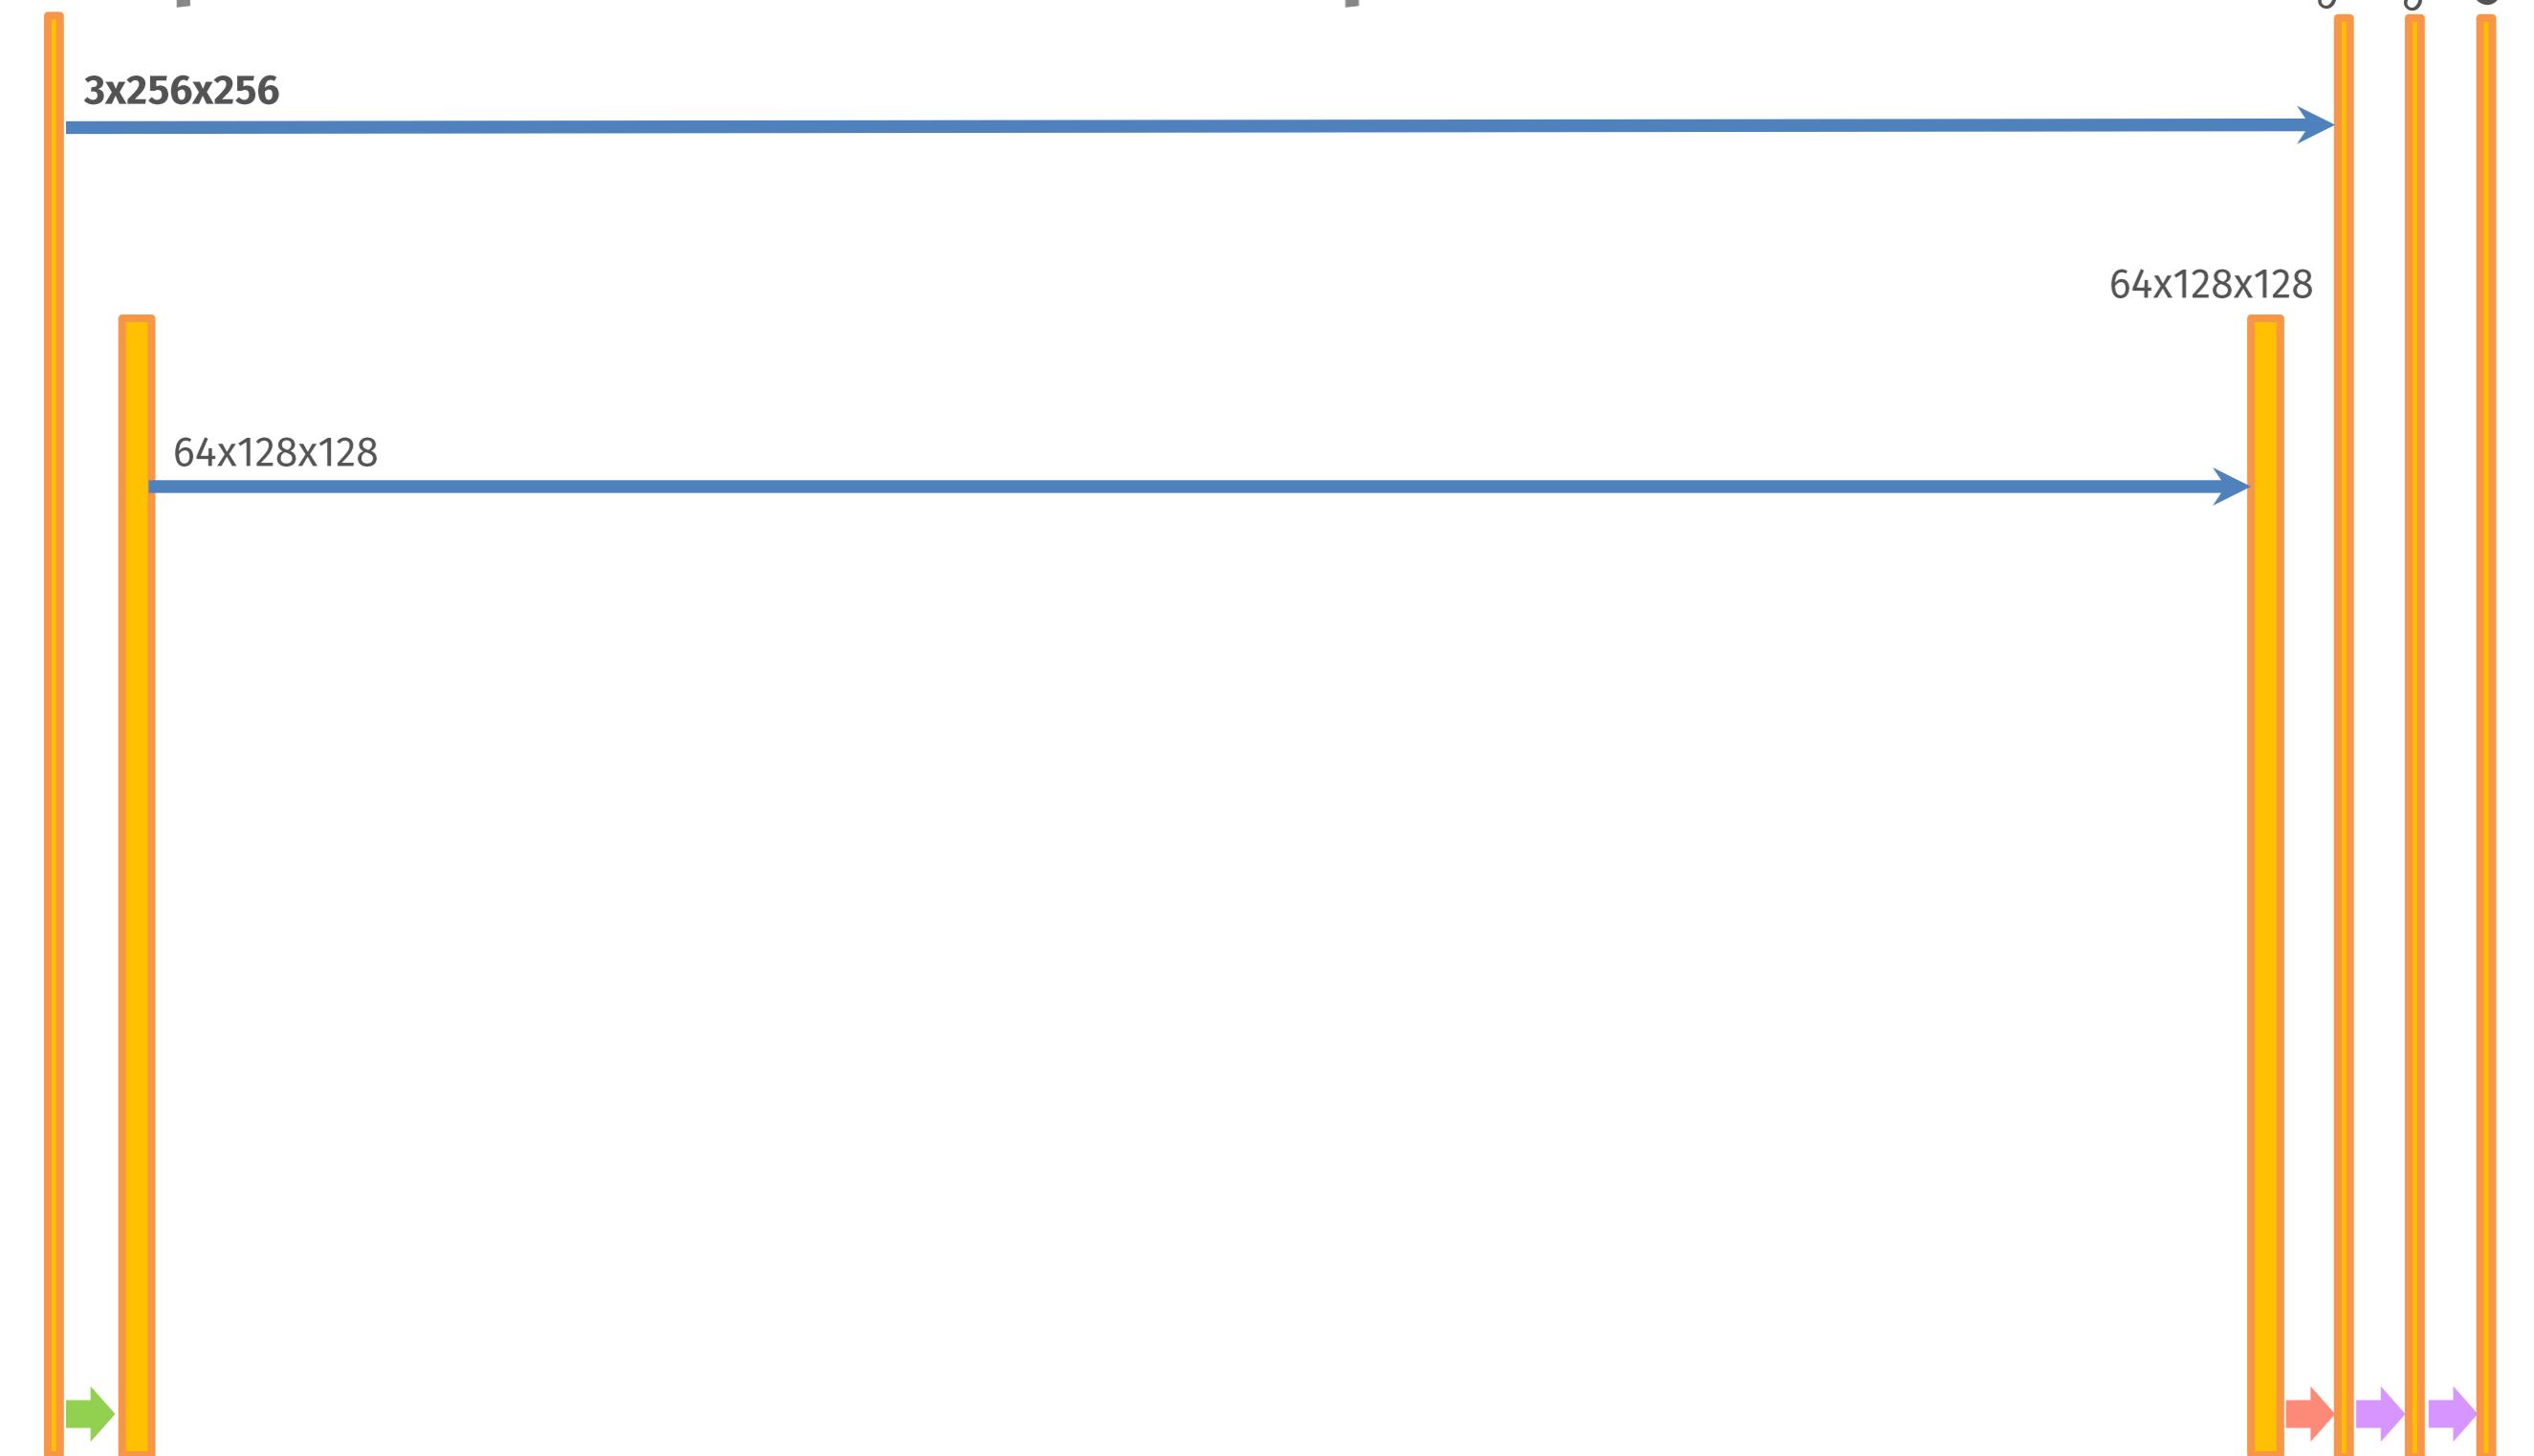
Does spatial information really help?

G pruned to 1 branch. Receptive field: 1x1



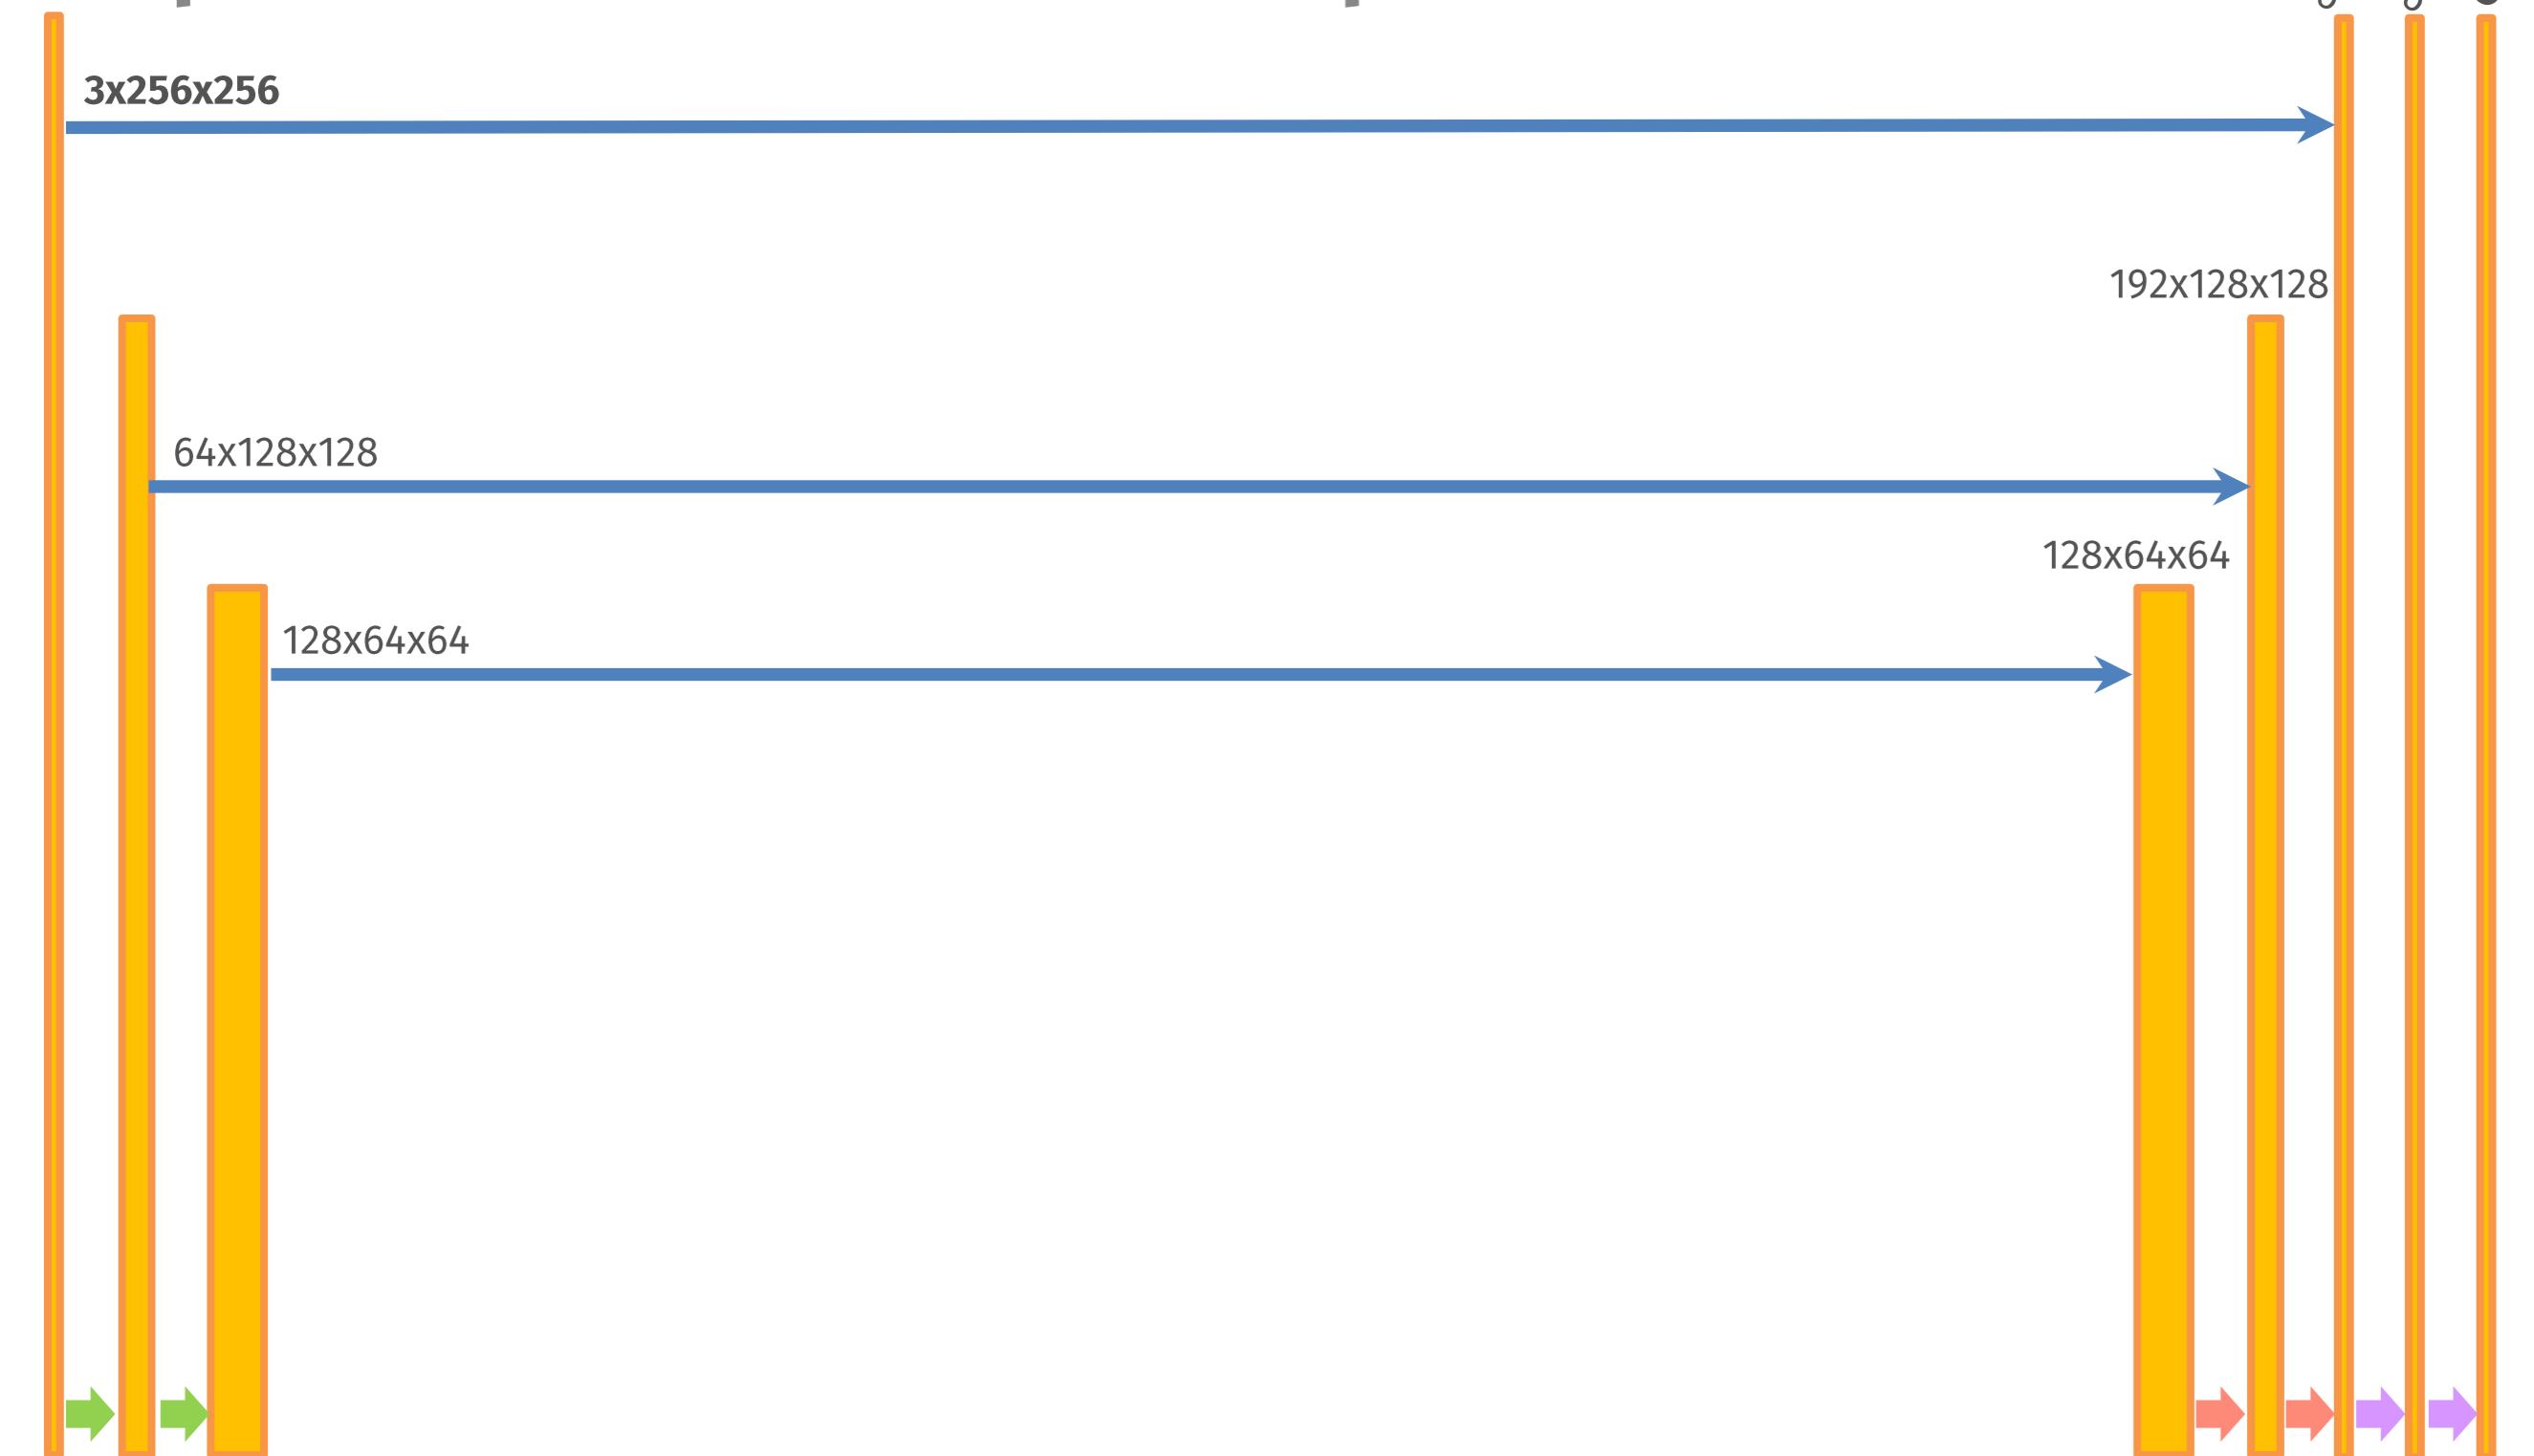
Does spatial information really help?

G pruned to 2 branches. Receptive field: 3x3



Does spatial information really help?

G pruned to 3 branches. Receptive field: 7x7



Does spatial information really help? G pruned to 4 branches. Receptive field: 15x15



Does spatial information really help? G pruned to 5 branches. Receptive field: 31x31

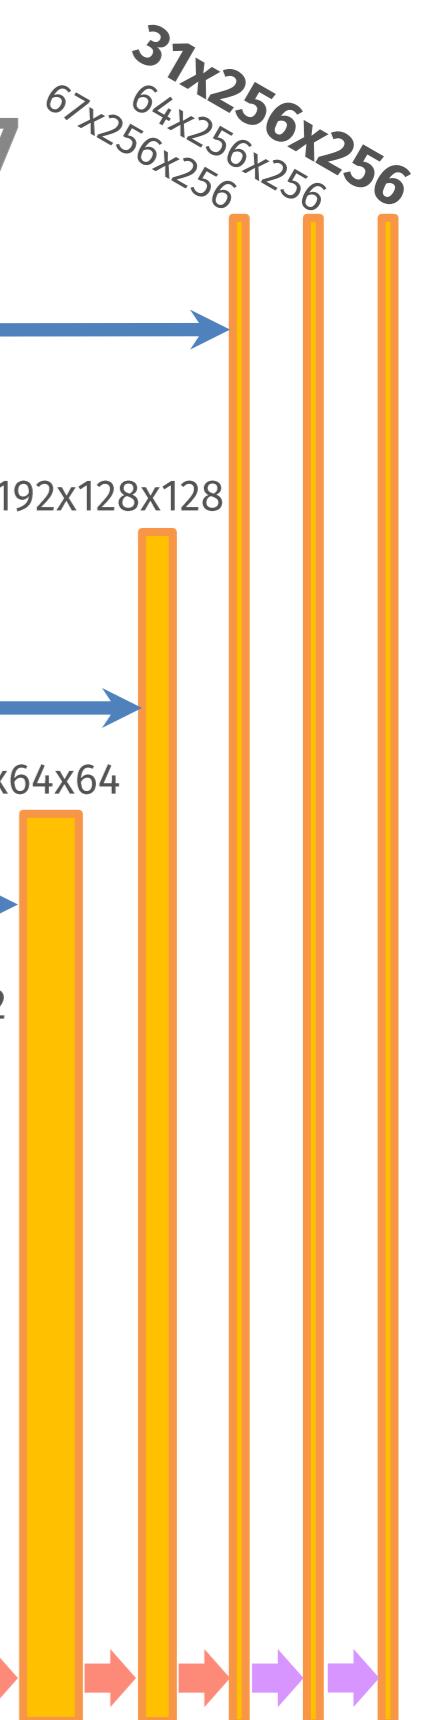


Does spatial information really help? G pruned to 6 branches. Receptive field: 63x63



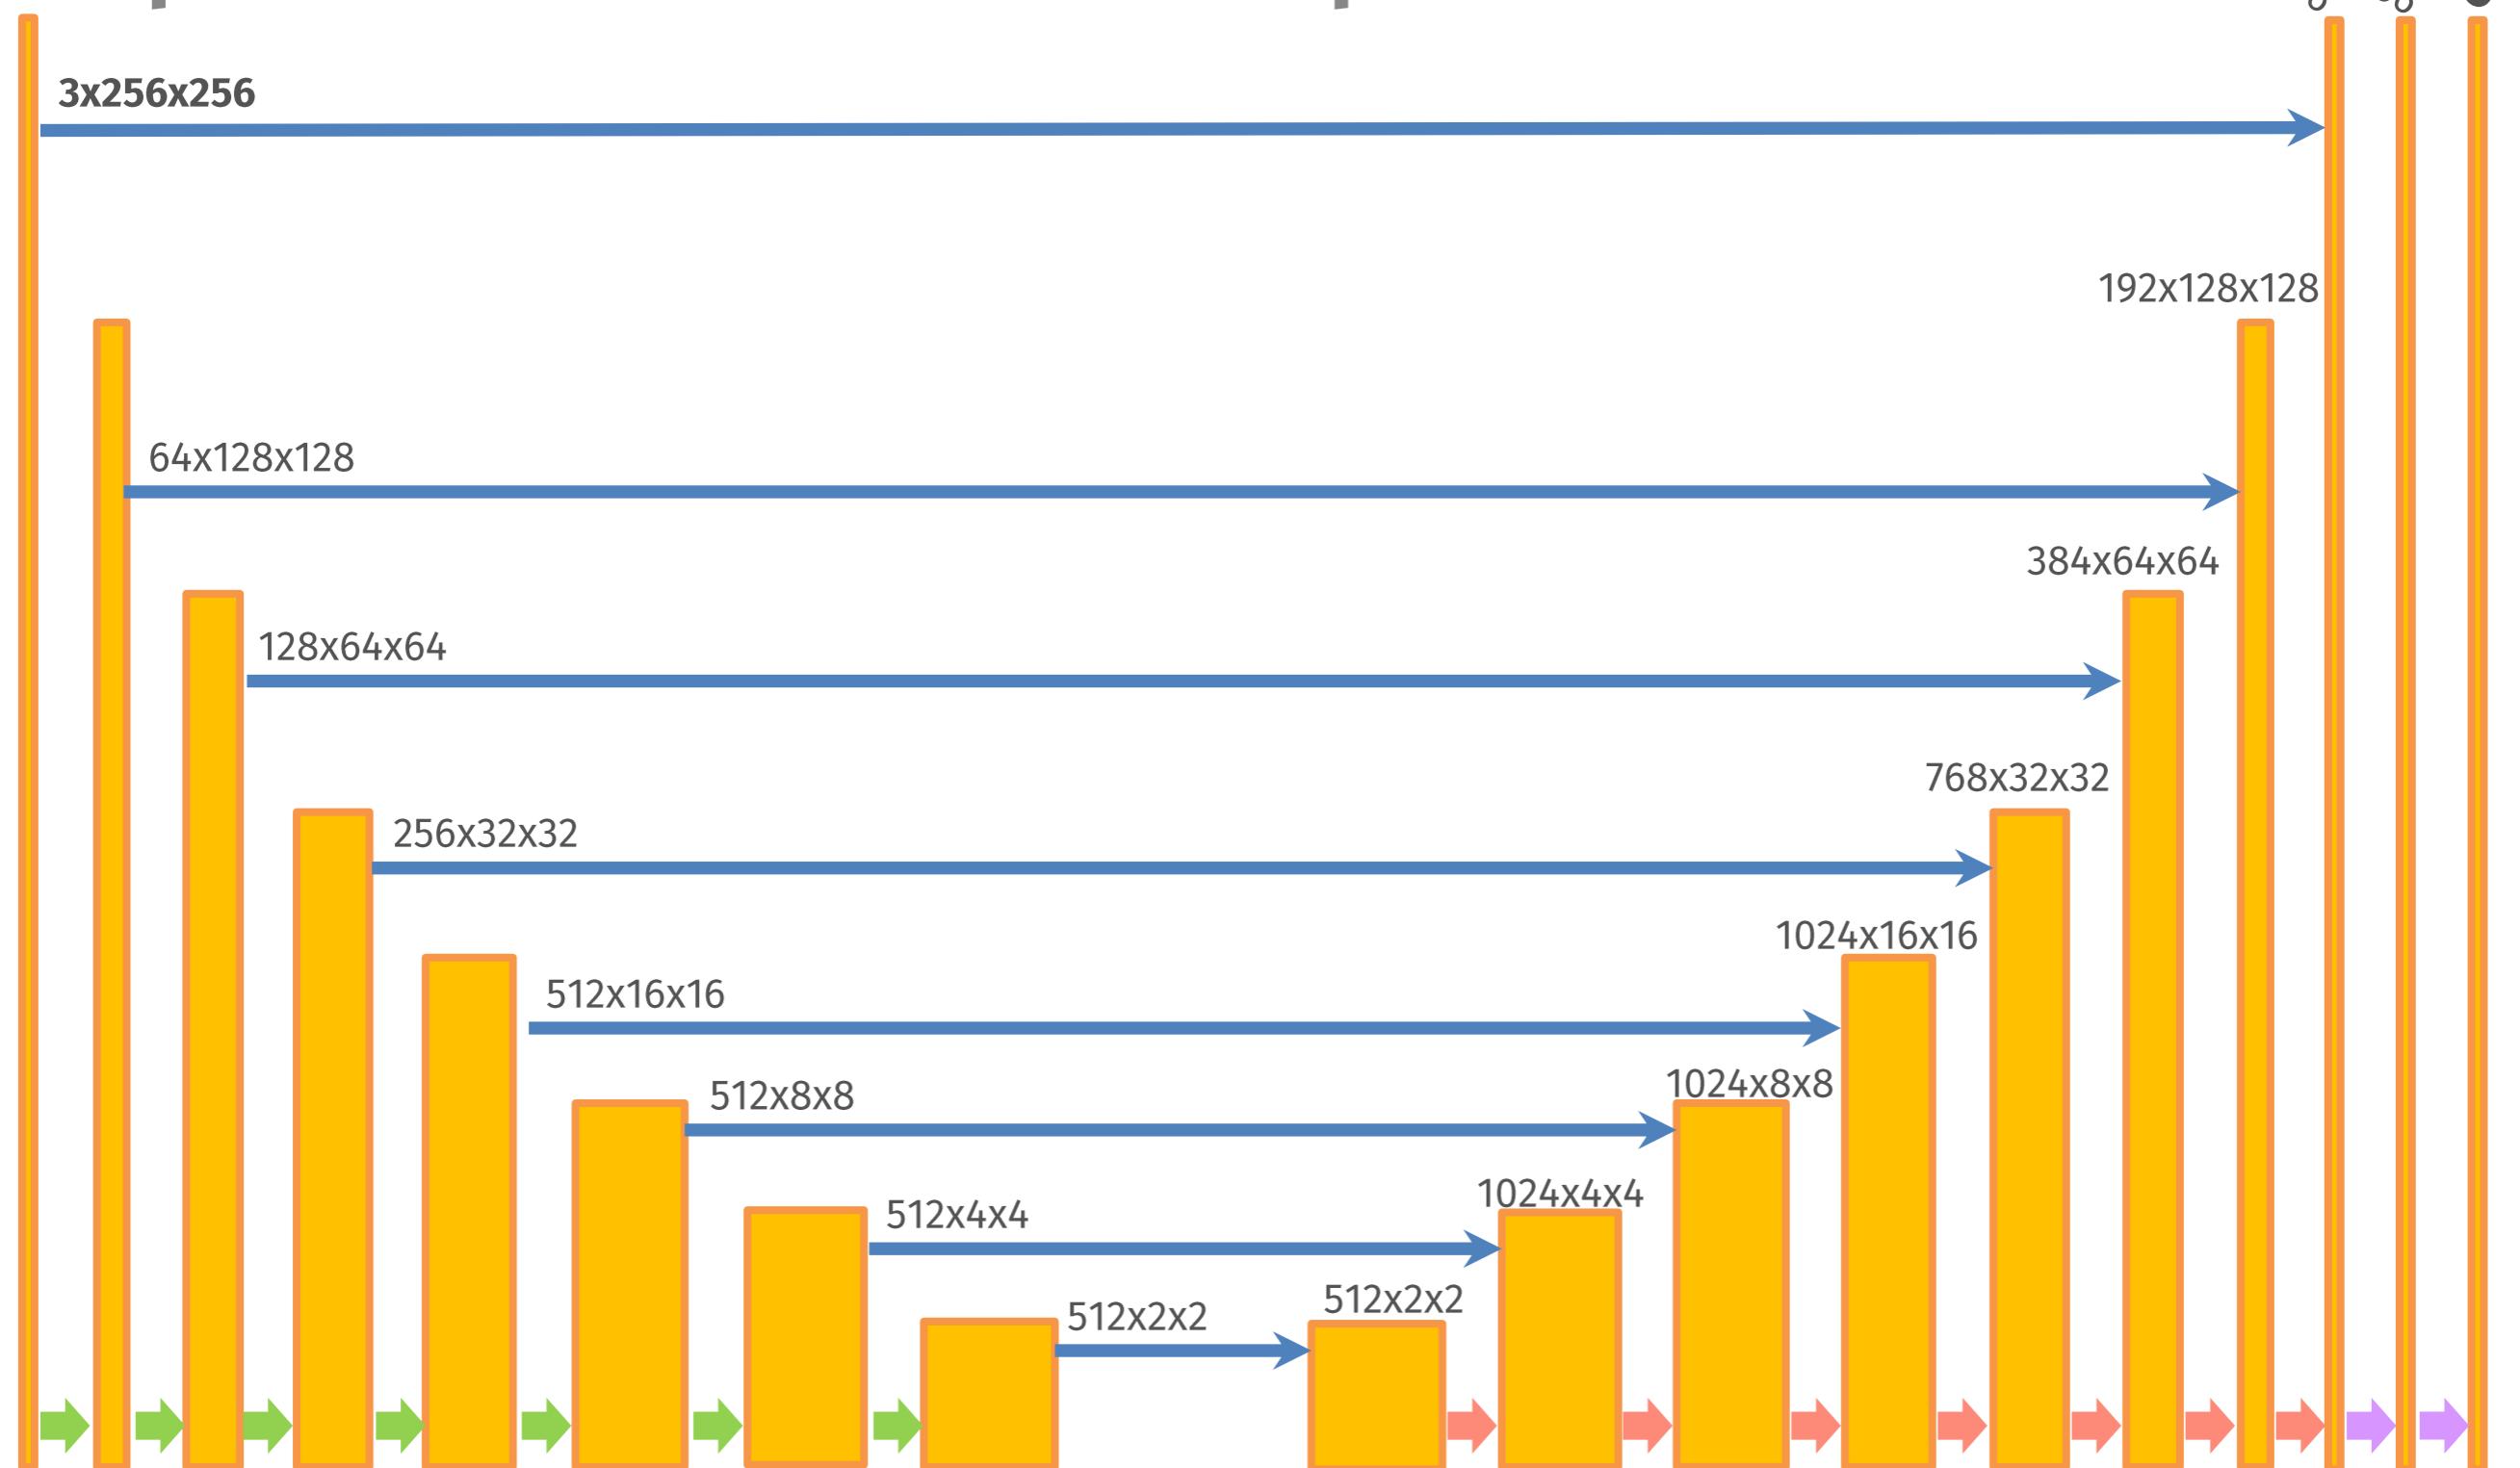
Does spatial information really help?

G pruned to 7 branches. Receptive field: 127x127



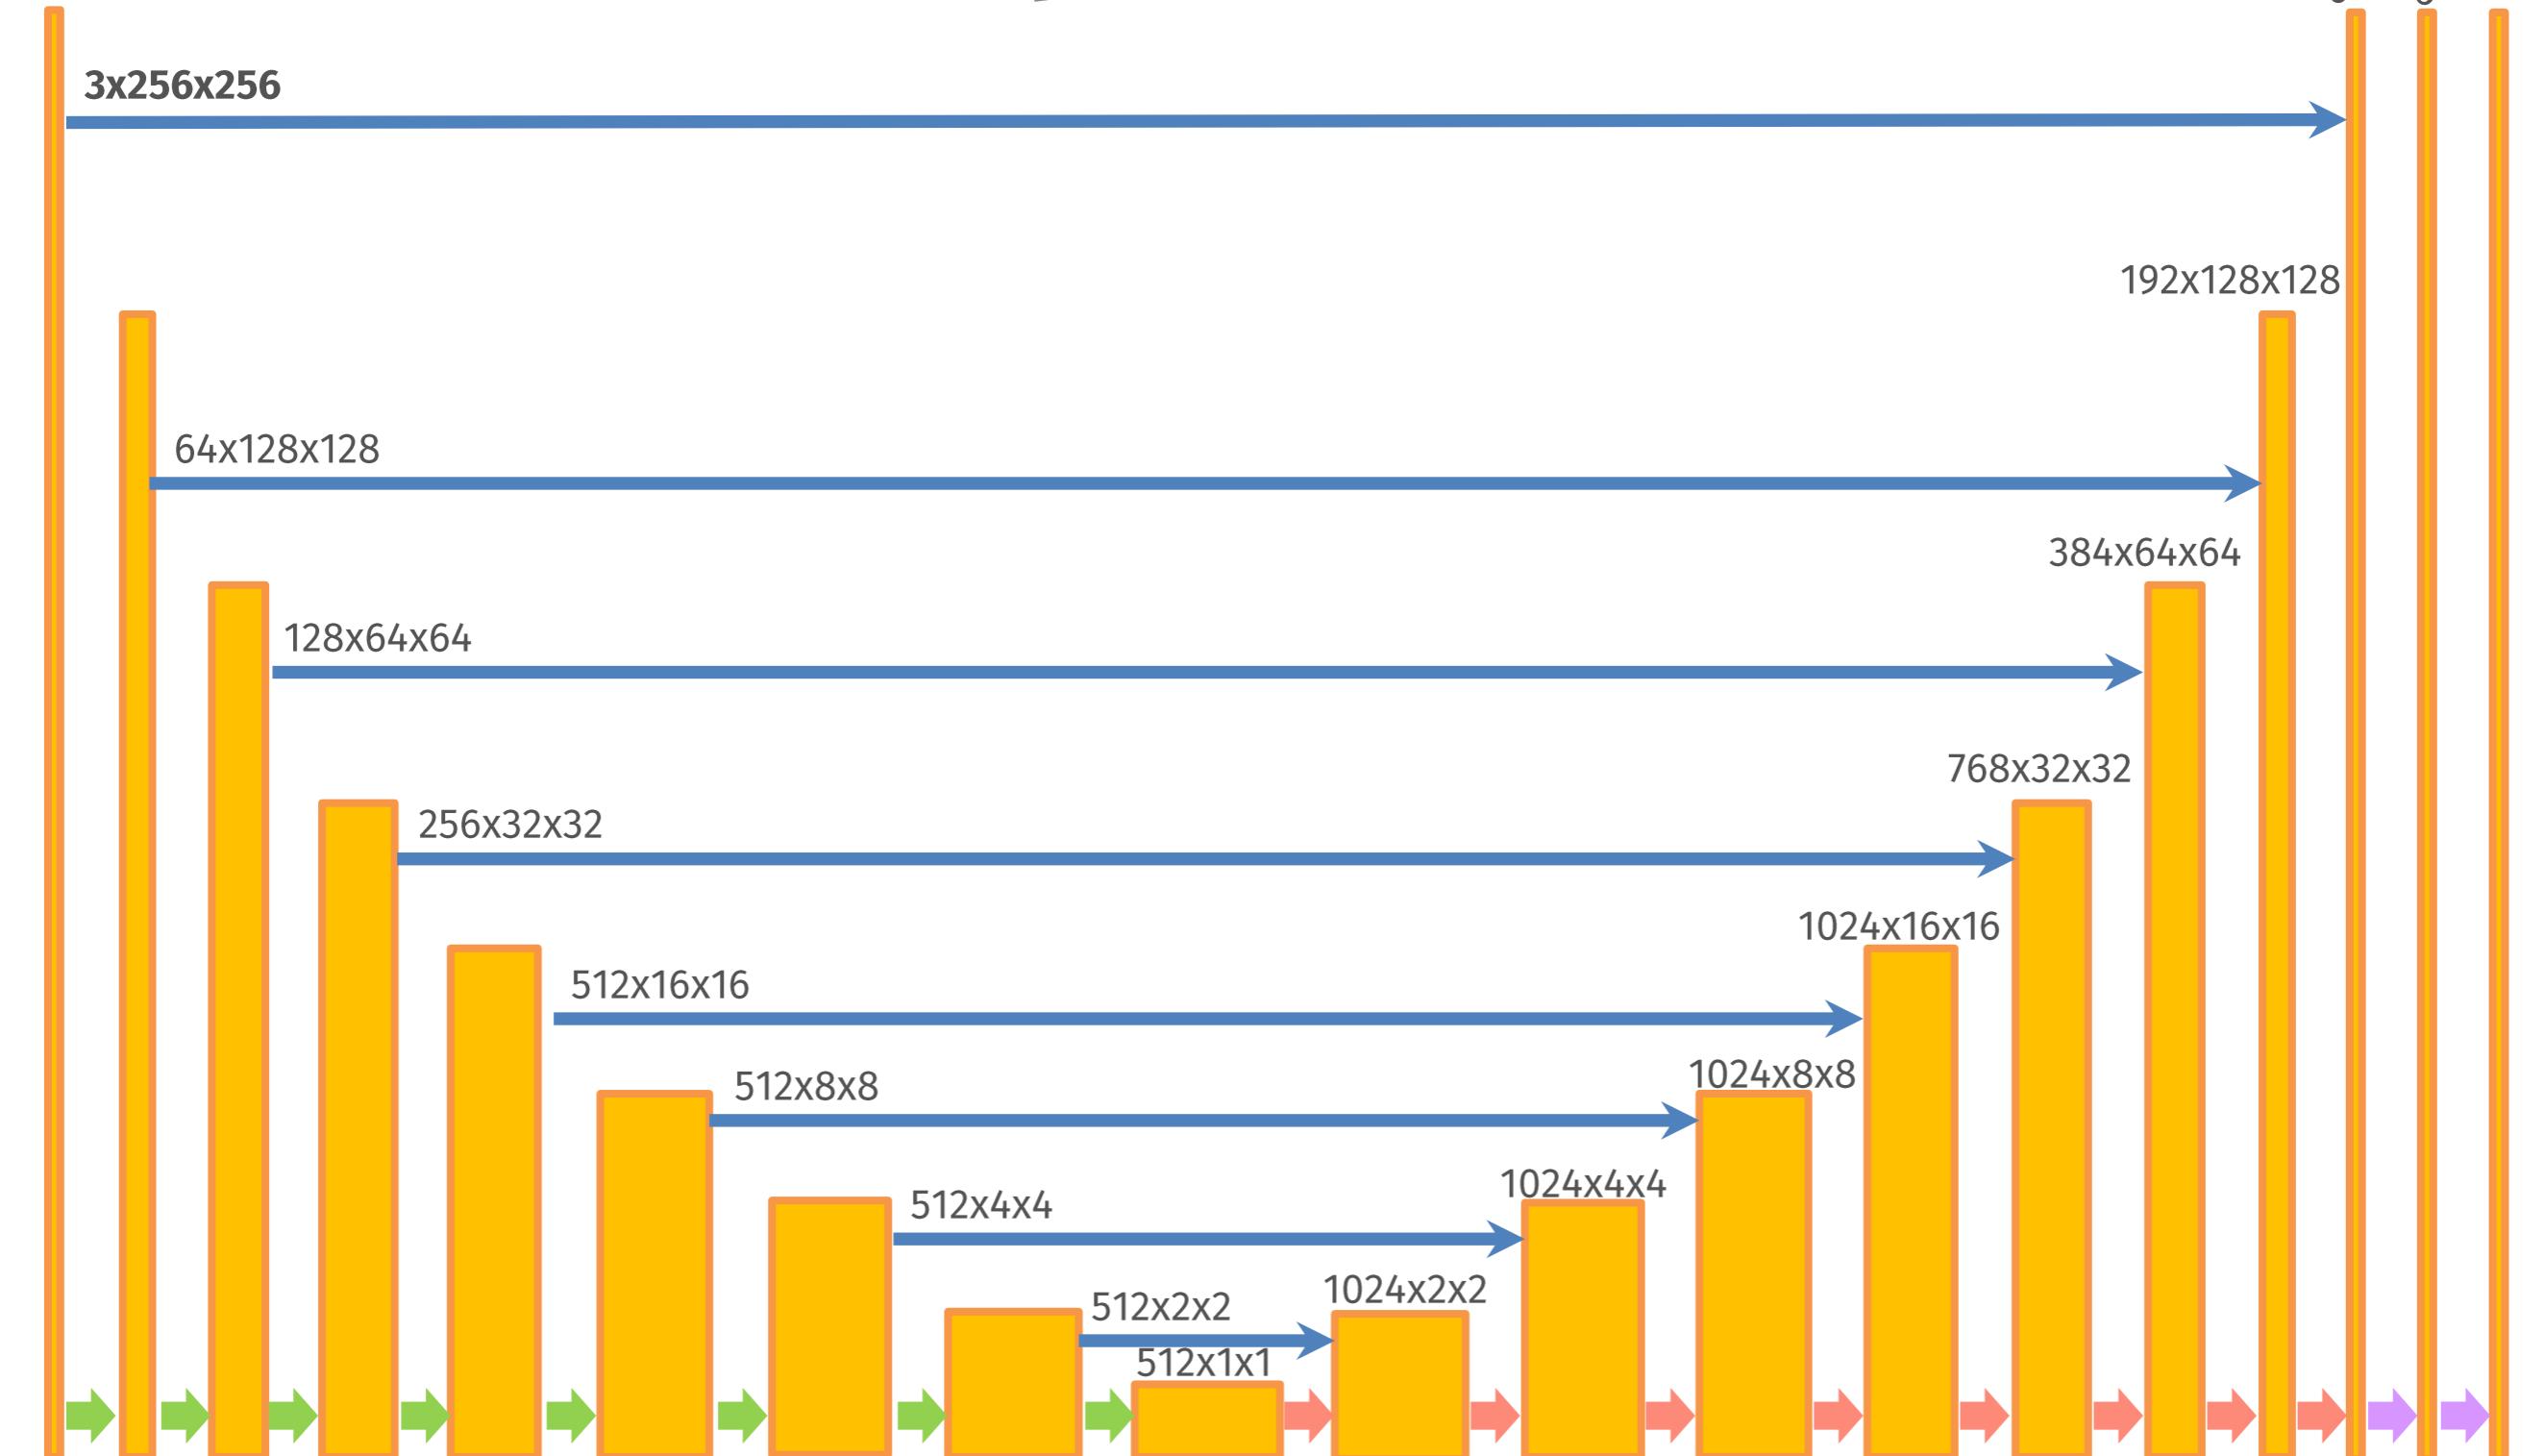
Does spatial information really help? G pruned to 8 branches. Receptive field: 255x255

3x256x256
64x256x256
67x256x256
37x256x256



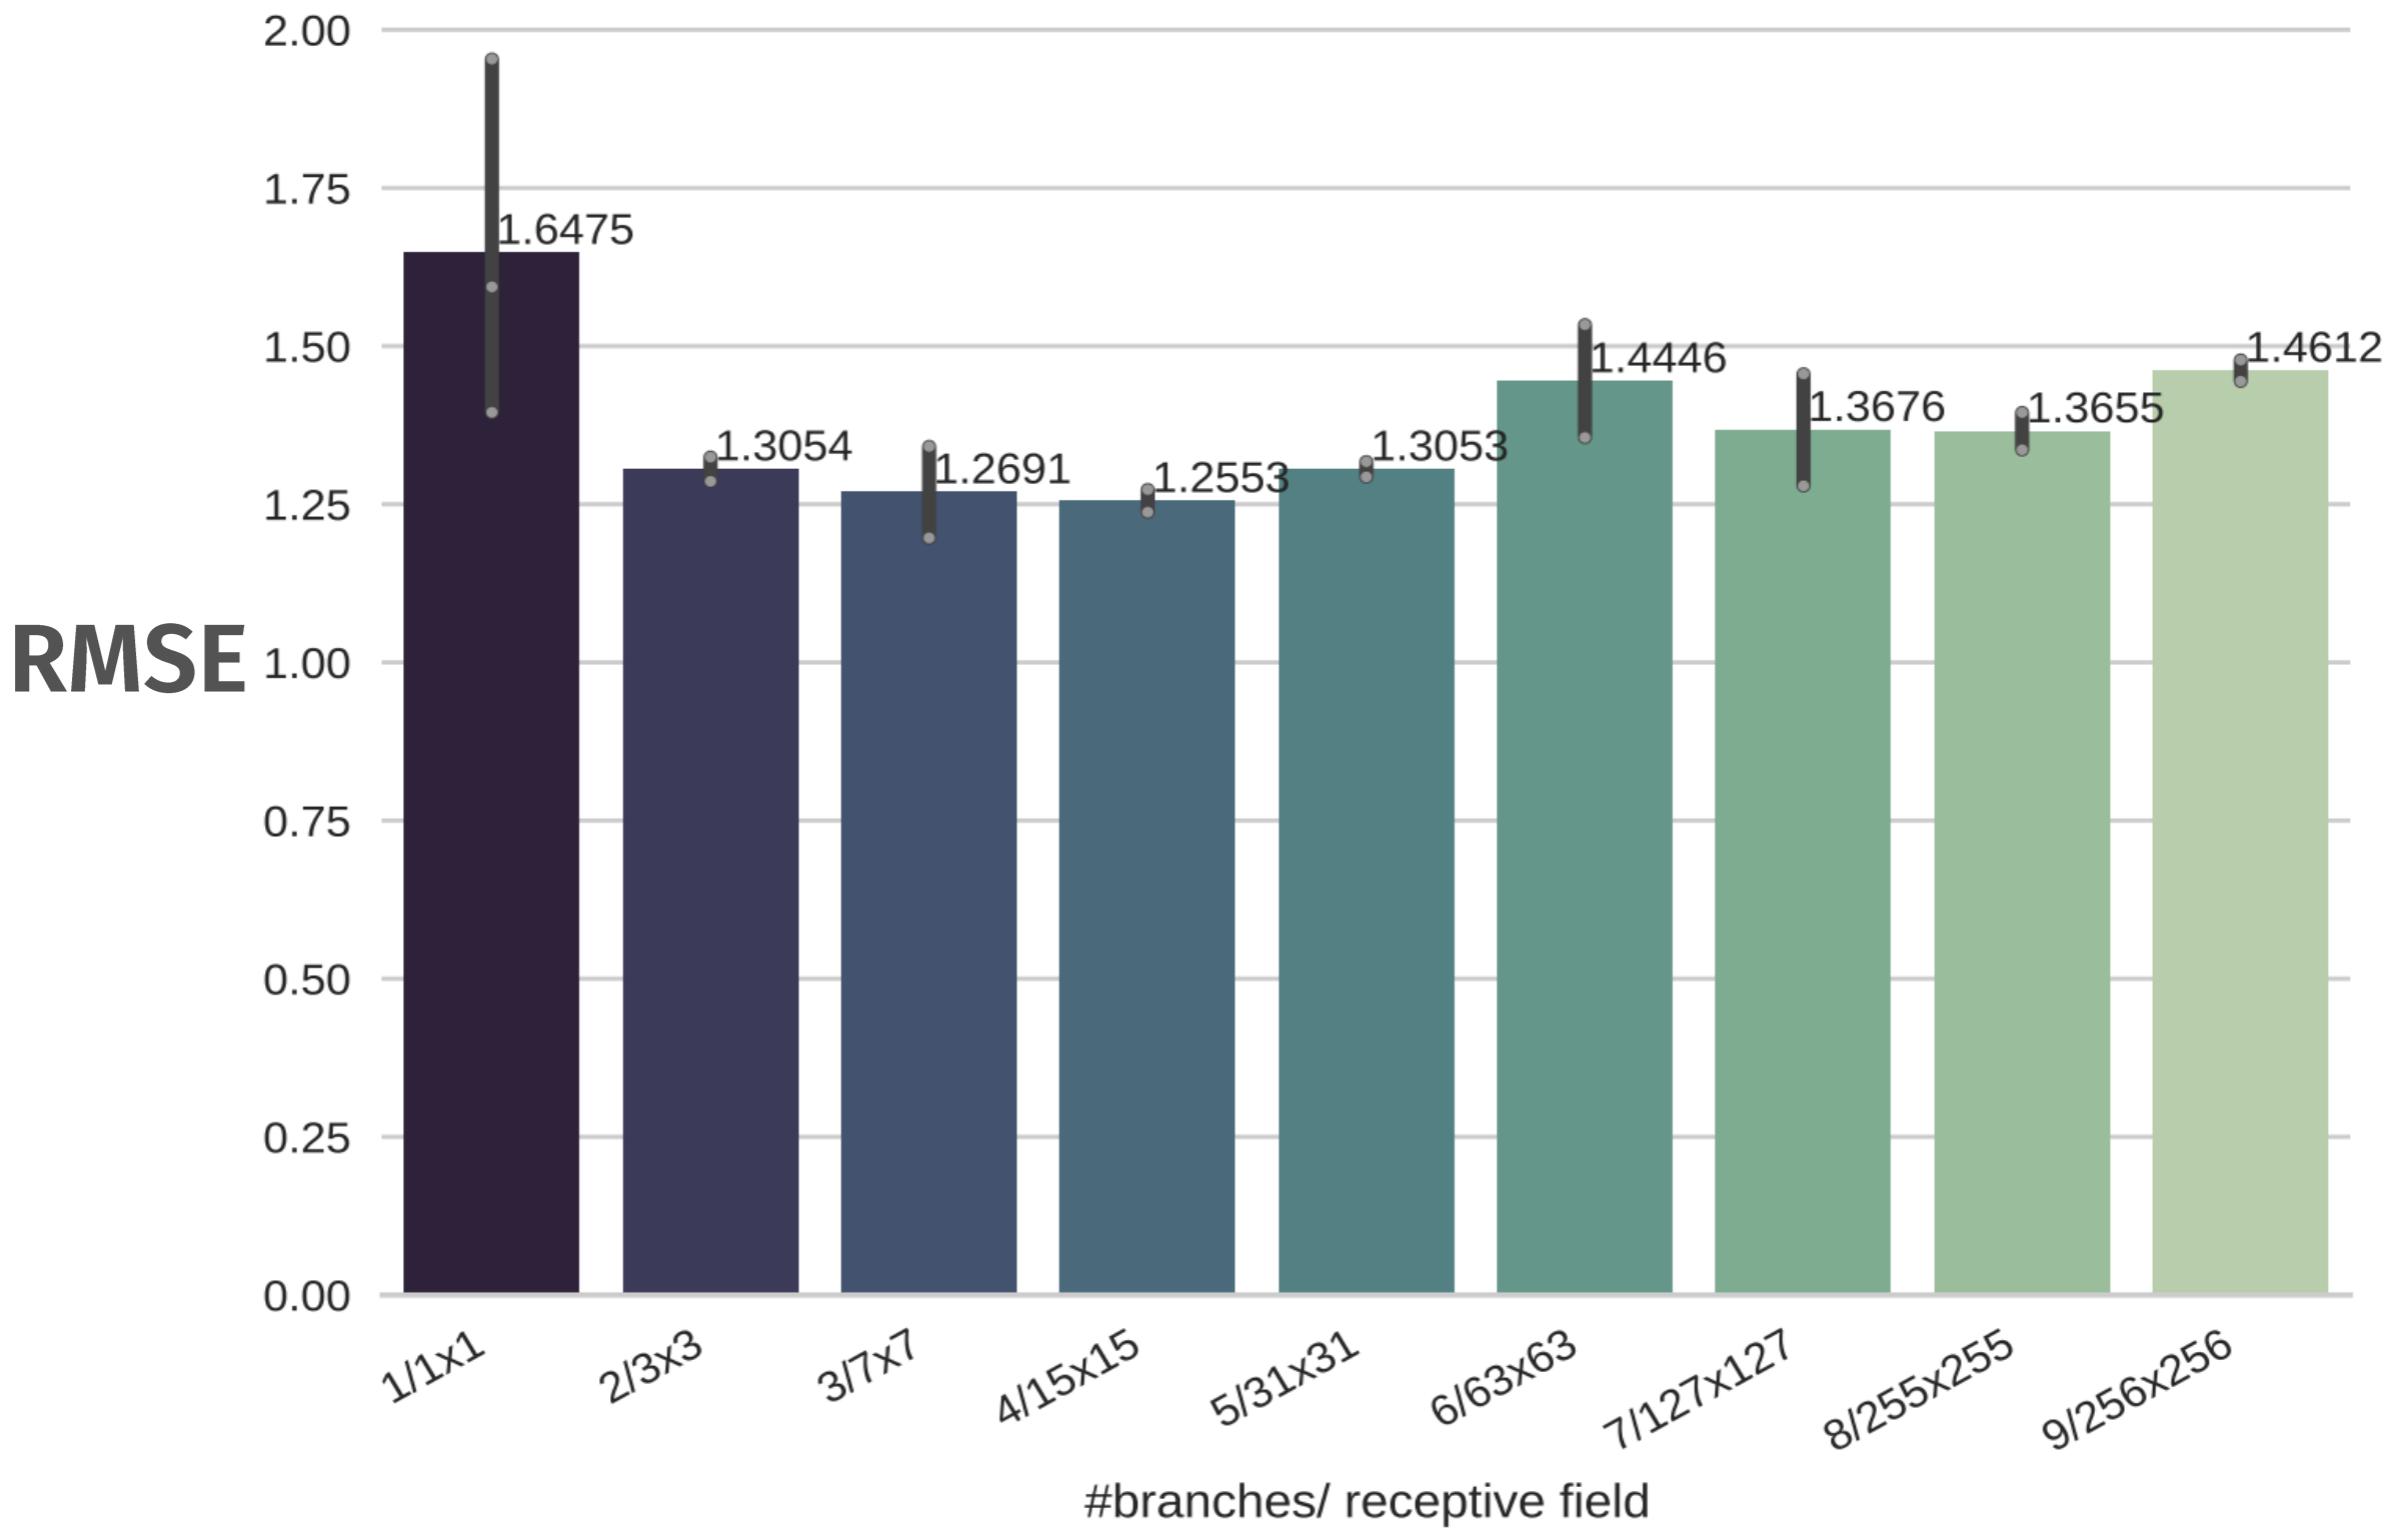
Does spatial information really help?

G is a U-Net Full. Receptive field: 256x256



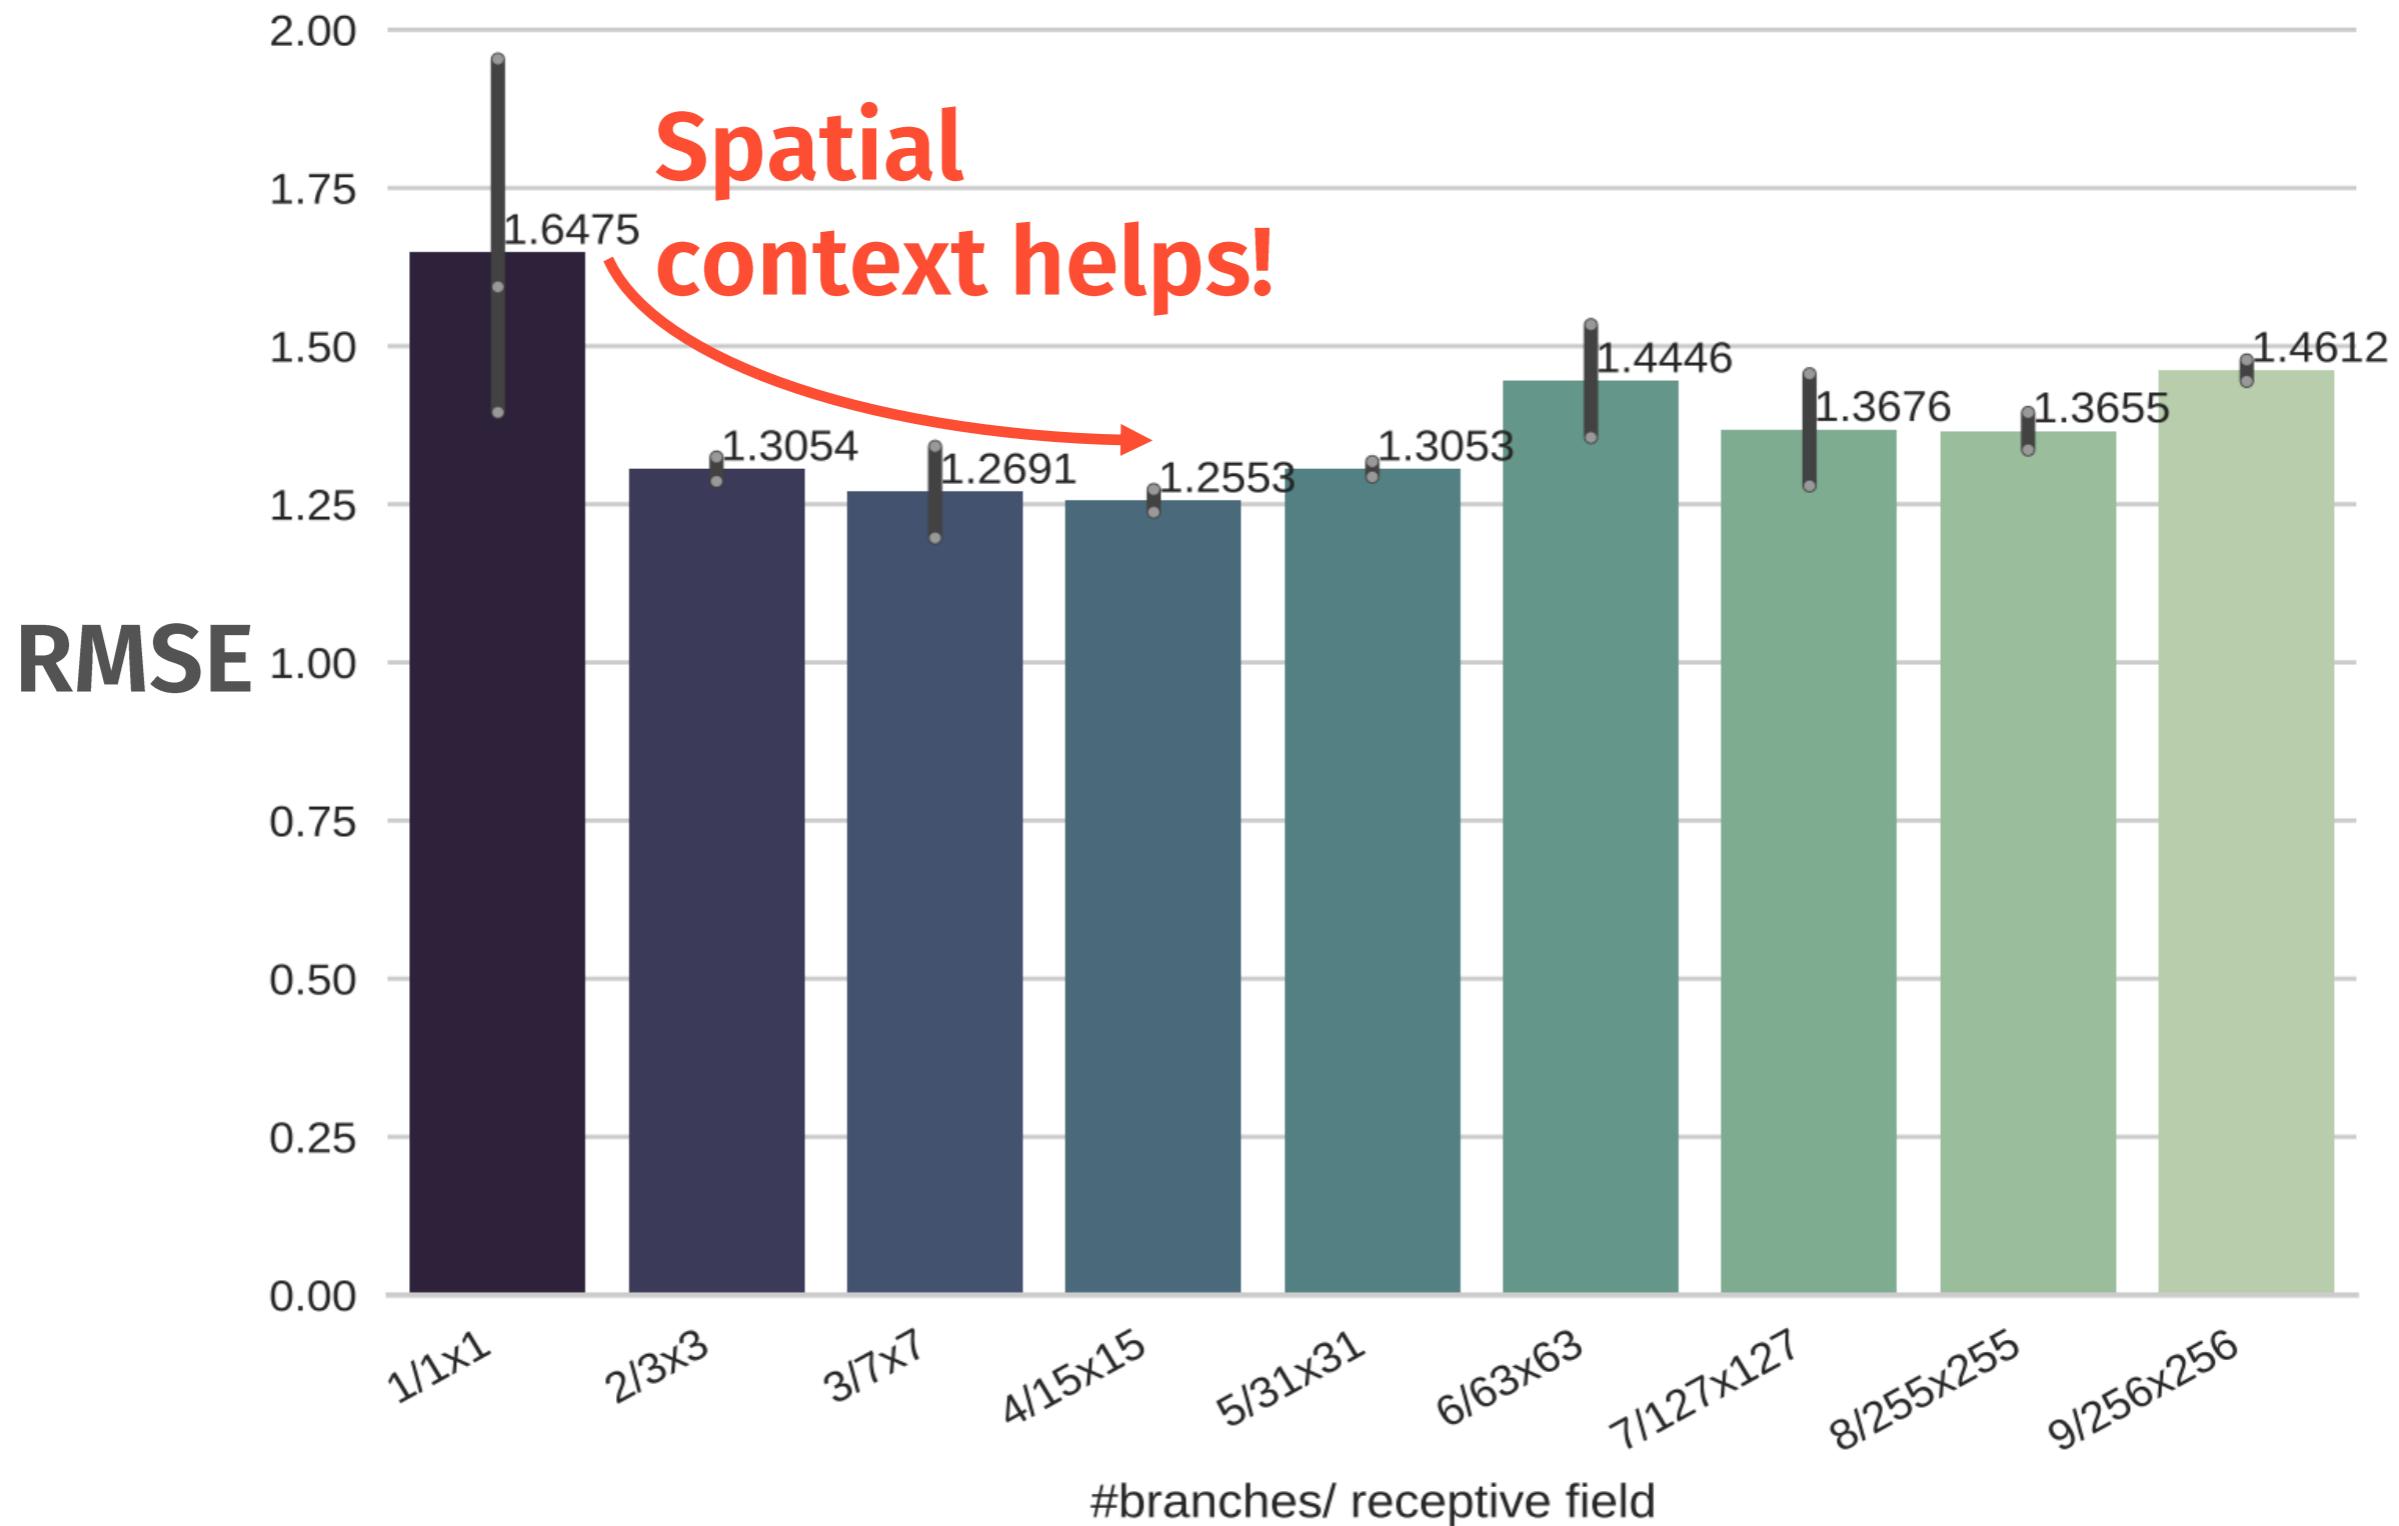
Results

Does spatial information really help?



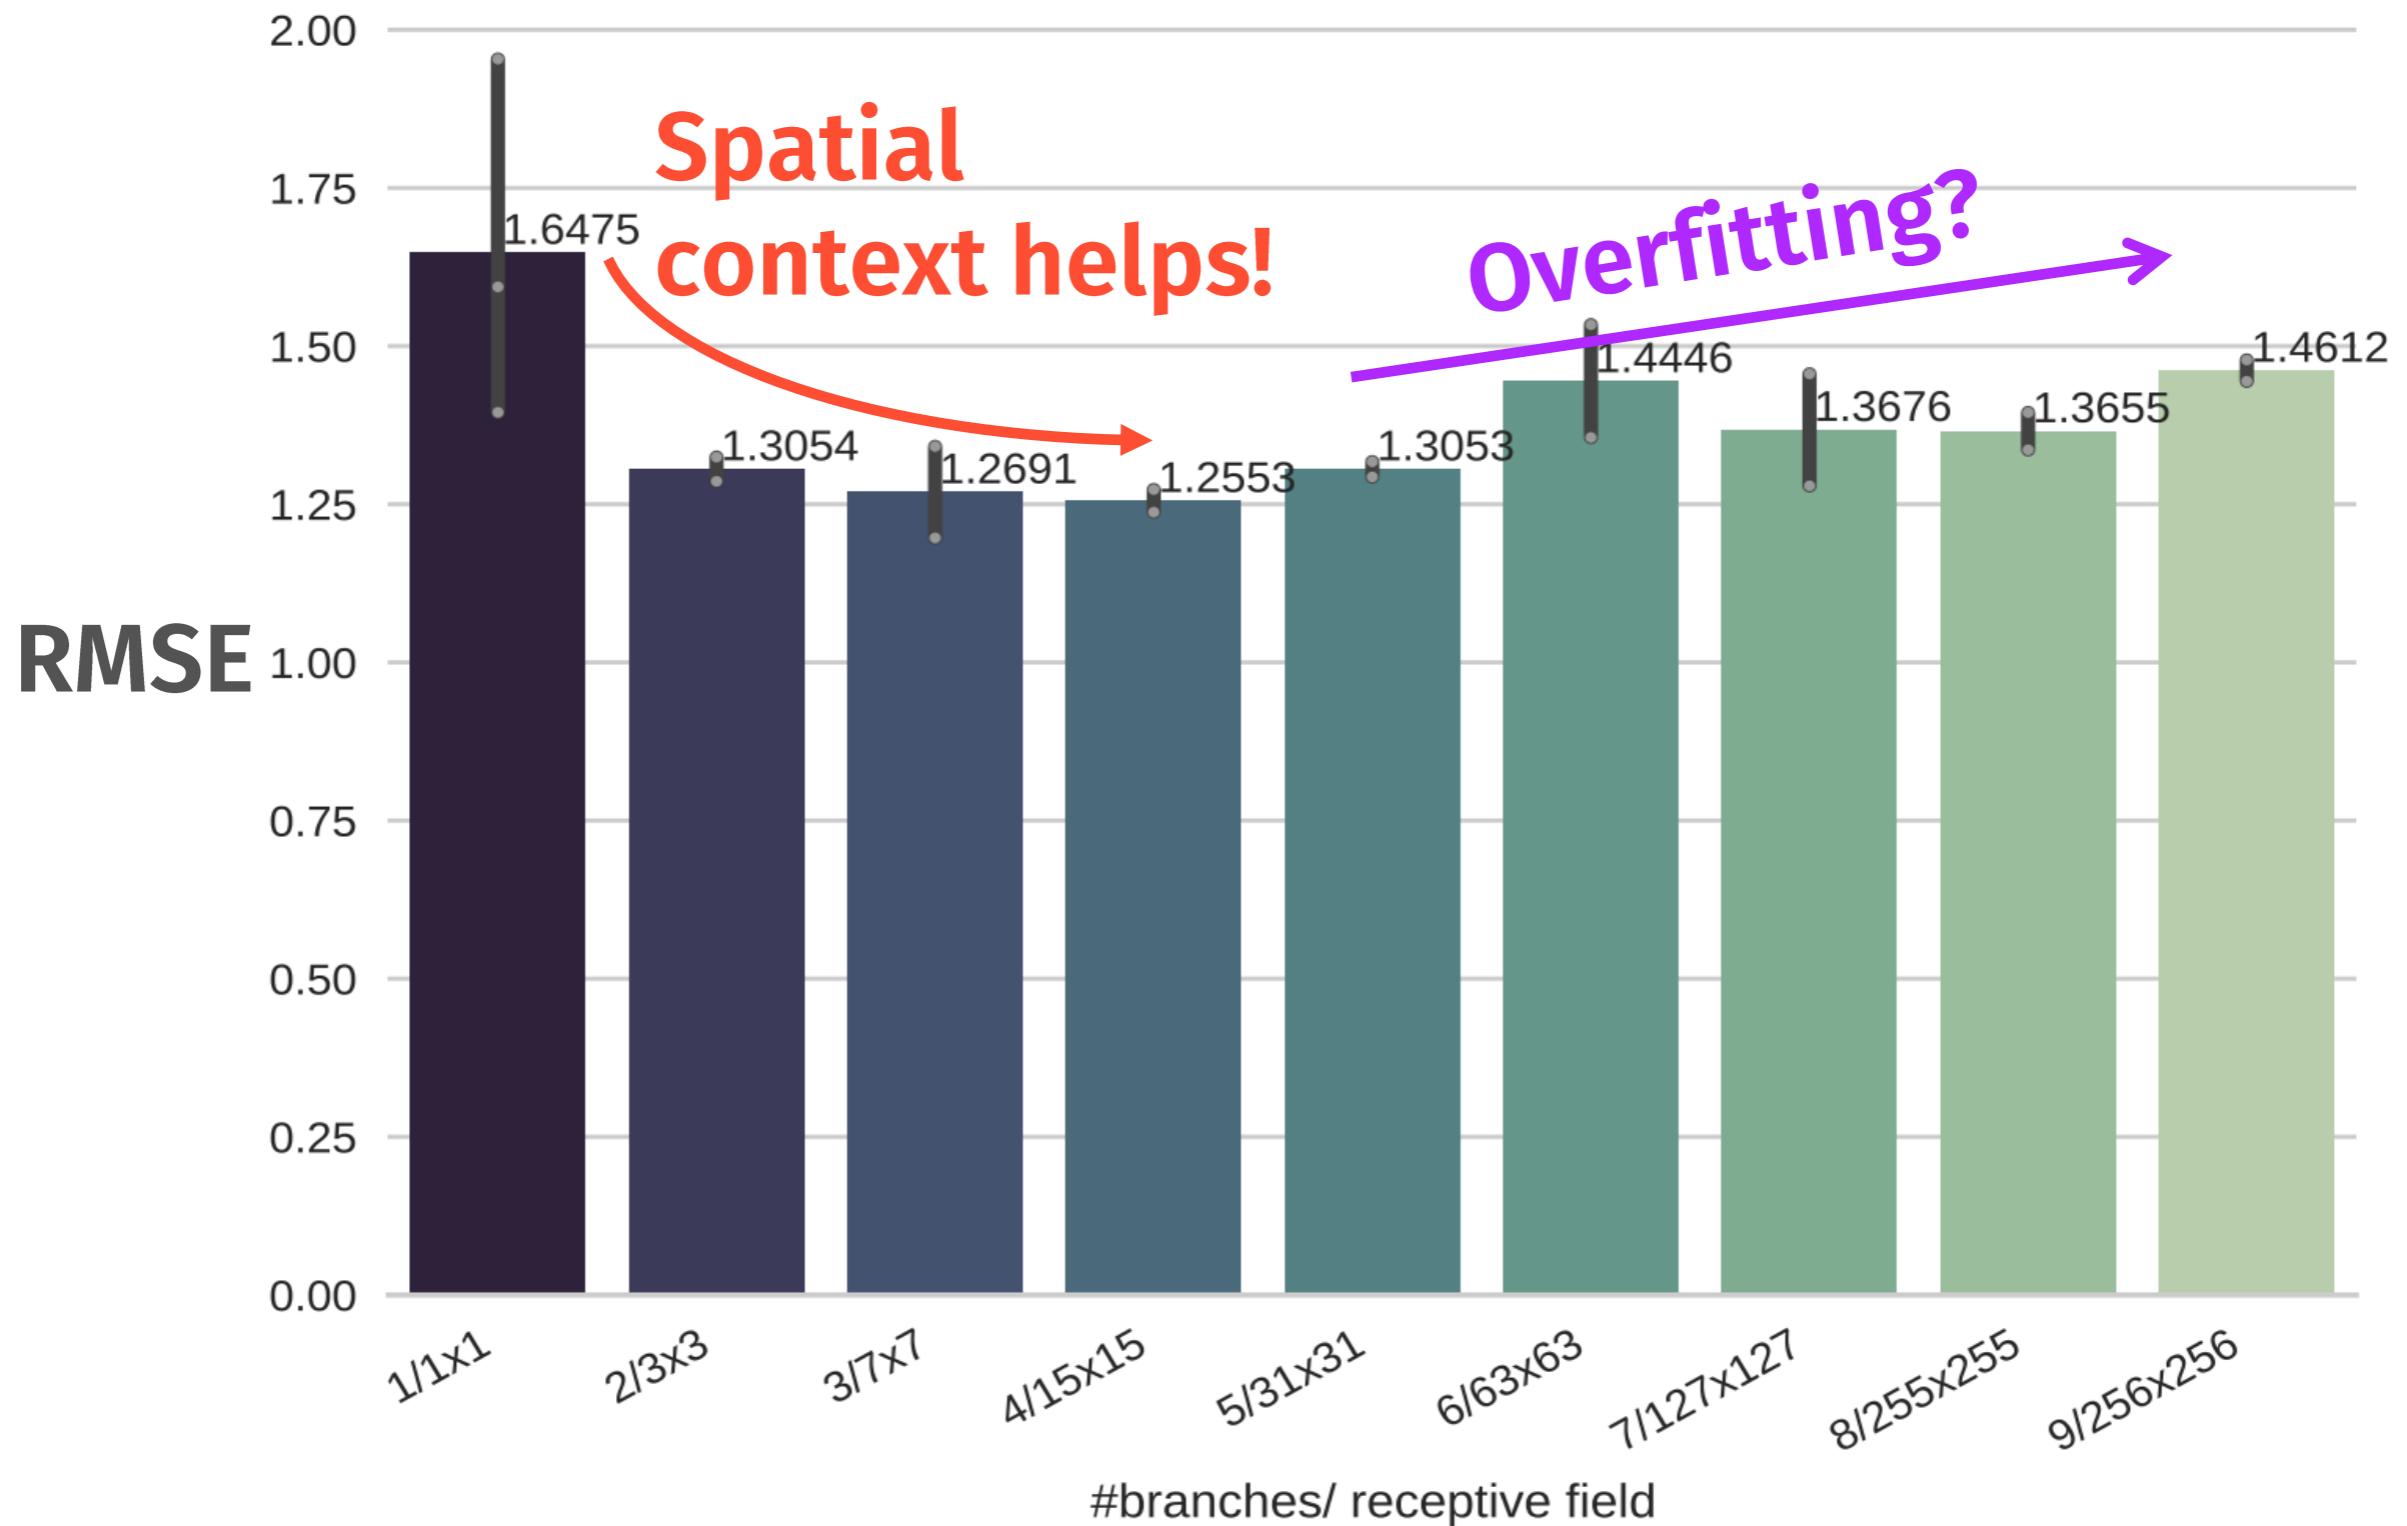
Results

Does spatial information really help?



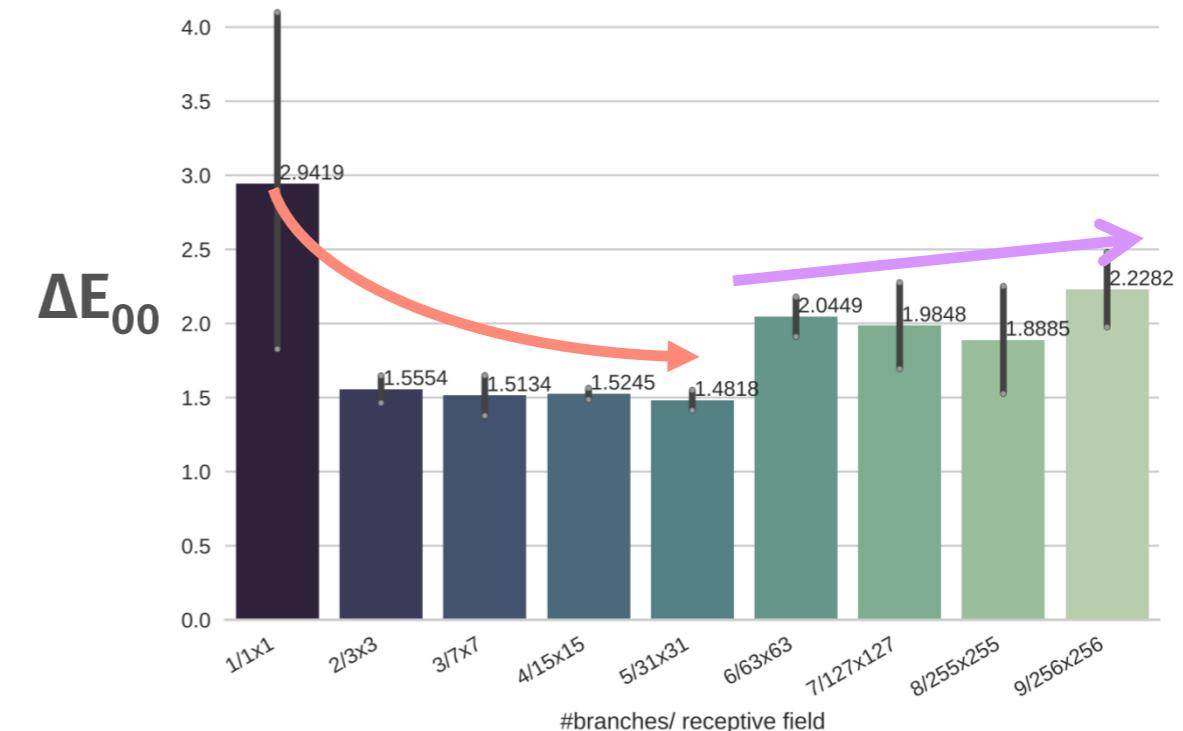
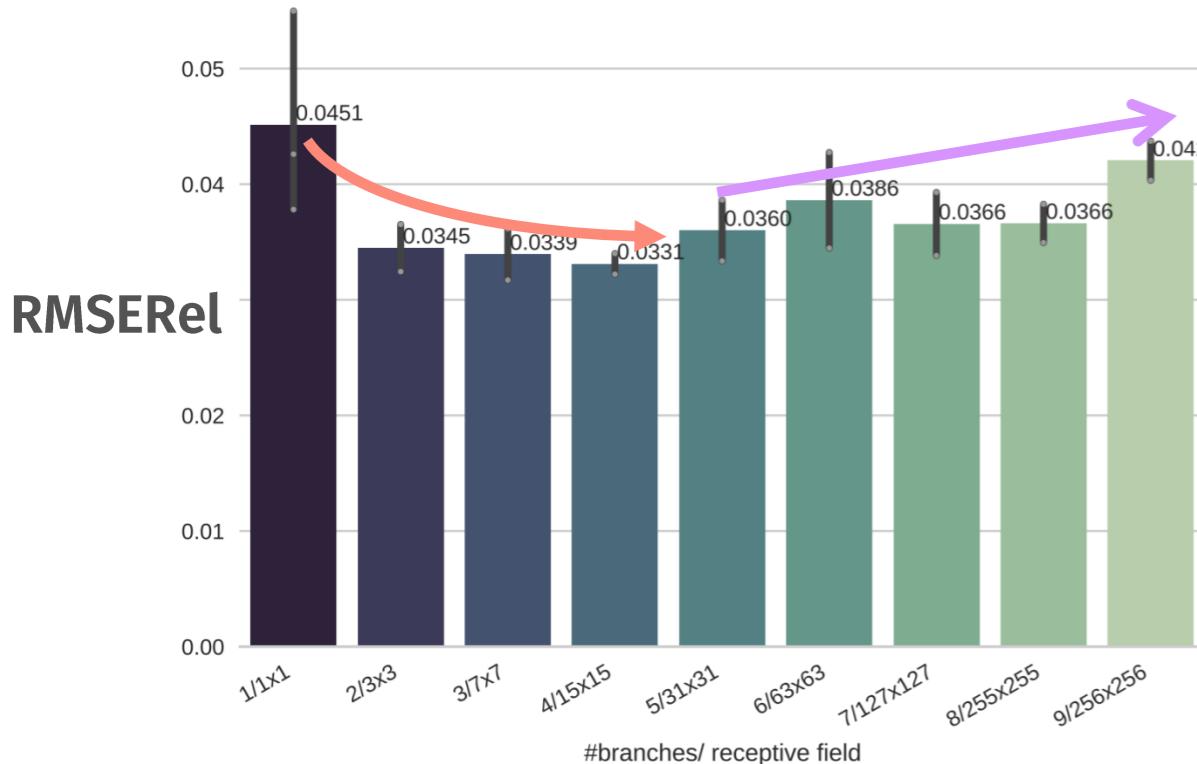
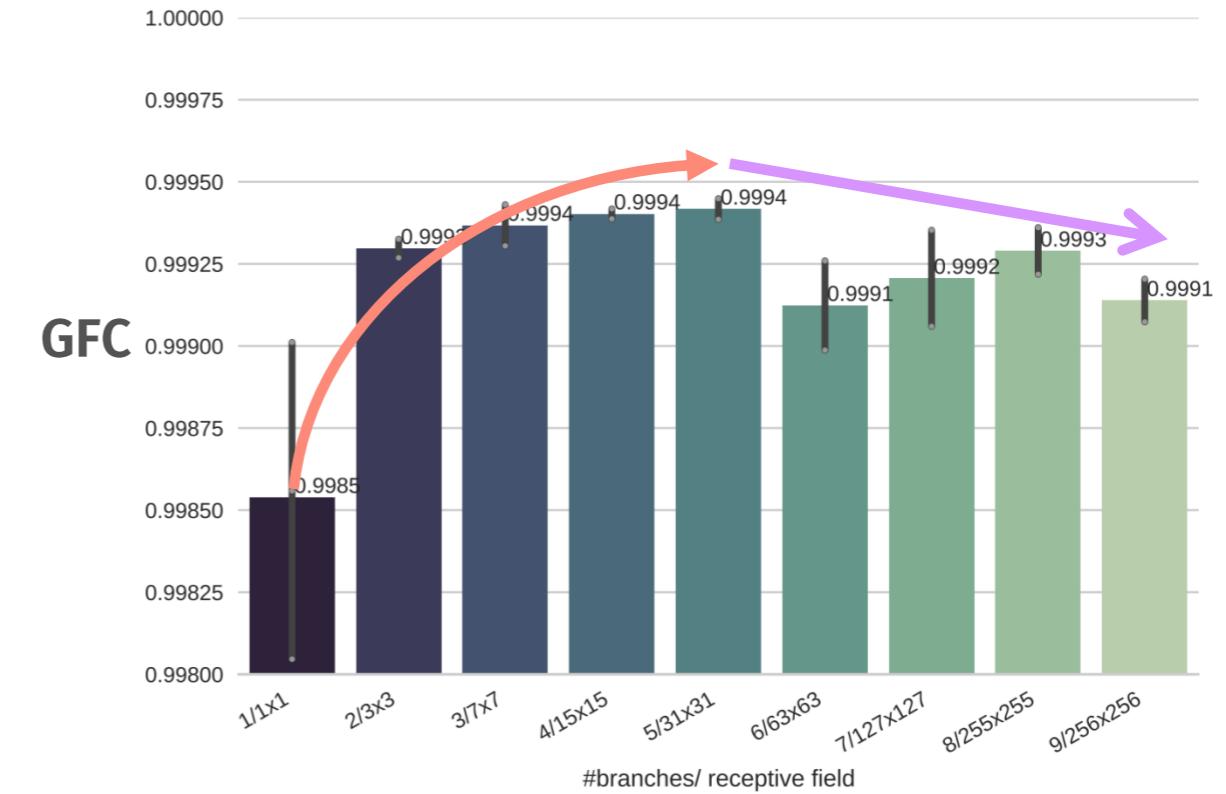
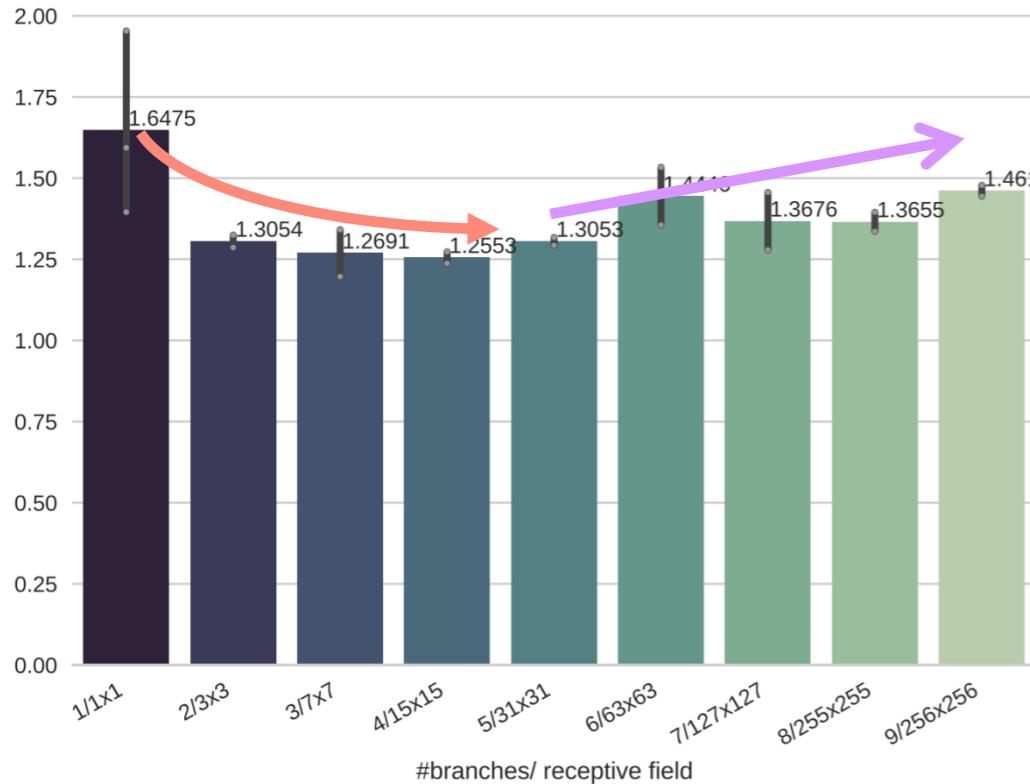
Results

Does spatial information really help?



Results

Does spatial information really help?



Conclusions...

- GANs applied to spectral image reconstruction for the 1st time
- State of the art over ICVL dataset
- Spatial context information helps

...and future work

- Evaluate on metamers
- Extend to a variety of spectral sensitivity curves
- Simultaneous {illuminant, spectral reflectance} estimation

Adversarial Networks for Spatial Context-Aware Spectral Image Reconstruction from RGB

Aitor Alvarez-Gila, Joost van de Weijer, Estibaliz Garrote

Personal:

<https://aitorshuffle.github.io>

Tecnalia Computer Vision group:

www.computervisionbytecnalia.com

CVC Learning And Machine Perception team:

<http://www.cvc.uab.es/LAMP/>